

ADOPT-99: Best Practice Summary

Series: ADOPT — Observability Adoption & Maturity | **Notebook:** 99 | **Created:** March 2026 | **Last Updated:** 08/04/2026

Overview

This notebook consolidates every actionable best practice from the ADOPT series (notebooks 01 through 06) into definitive, categorized guidance. Each practice specifies exactly what to set, the priority level, and the source notebook. Use this as a checklist for platform maturity and adoption readiness.

Table of Contents

1. [Agent Deployment and Host Monitoring](#)
 2. [Data Ingestion and Grail Configuration](#)
 3. [Alerting and Dynatrace Intelligence](#)
 4. [SLOs and Success Metrics](#)
 5. [Cost Optimization](#)
 6. [Automation and Workflows](#)
 7. [IAM and Access Control](#)
 8. [Team Enablement and Organizational Readiness](#)
 9. [Dashboards and Reporting](#)
 10. [Configuration Management](#)
 11. [Distributed Tracing and Spans](#)
 12. [Business Observability](#)
 13. [Coverage Audit and Staged Enablement](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS or Managed with Grail enabled
Permissions	<code>storage:logs:read</code> , <code>storage:metrics:read</code> , <code>storage:entities:read</code> , <code>storage:events:read</code> , <code>storage:buckets:read</code>
Context	Familiarity with ADOPT-01 through ADOPT-06
Audience	Platform engineers, SREs, engineering managers, leadership

1. Agent Deployment and Host Monitoring

#	Best Practice	Recommended Setting/Value	Priority	Source
1.1	OneAgent host coverage	> 95% of all known hosts monitored (compare entity count against CMDB)	Critical	ADOPT-02
1.2	Monitoring mode for application hosts	Set to <code>FULL_STACK</code> for all hosts running application workloads	Critical	ADOPT-01, ADOPT-02
1.3	Monitoring mode for non-application hosts	Set to <code>INFRASTRUCTURE</code> for hosts that only need OS-level metrics (reduces host unit cost)	Recommended	ADOPT-05
1.4	Agent version currency	Maintain all agents within 1 major version; maximum 2 distinct major versions across the fleet	Critical	ADOPT-02, ADOPT-05
1.5	Auto-update for non-production	Enable OneAgent auto-update on all non-production hosts	Recommended	ADOPT-05
1.6	Automated agent deployment	Use cloud-init, Ansible, or equivalent to deploy OneAgent on every new host automatically	Recommended	ADOPT-05

#	Best Practice	Recommended Setting/Value	Priority	Source
1.7	Host CPU/memory health check	Run a weekly health query confirming all hosts report CPU and memory metrics; any host with no data = agent issue	Recommended	ADOPT-02

2. Data Ingestion and Grail Configuration

#	Best Practice	Recommended Setting/Value	Priority	Source
2.1	Multi-signal ingestion	Ingest all 5 telemetry types: metrics, logs, spans, events, and business events	Critical	ADOPT-01
2.2	Log ingestion stability	Daily variance < 10%; monitor with a 7-day rolling trend query	Recommended	ADOPT-02
2.3	Span ingestion active	Confirm > 0 spans per hour at all times; a zero reading means broken instrumentation	Critical	ADOPT-02
2.4	Grail bucket naming convention	<code><team>_<datatype>_<env></code> (e.g., <code>checkout_logs_prod</code>)	Recommended	ADOPT-05
2.5	Bucket-specific retention	High-value data: 35 days. Medium-value: 7 days. Debug/verbose: 7-14 days	Recommended	ADOPT-05
2.6	OpenPipeline log routing	Route logs to the correct bucket by source using OpenPipeline rules; never leave all data in <code>default_logs</code>	Recommended	ADOPT-05
2.7	Drop DEBUG/TRACE in production	Configure OpenPipeline to discard <code>DEBUG</code> and <code>TRACE</code> log levels in production environments	Recommended	ADOPT-05
2.8	Use bucket targeting in queries	Always specify <code>bucket: {"name"}</code> in DQL fetch commands to reduce scan cost	Optional	ADOPT-05

3. Alerting and Dynatrace Intelligence

#	Best Practice	Recommended Setting/Value	Priority	Source
3.1	Alert noise ratio	Keep below 20% (frequent + duplicate problems / total problems); above 50% = overhaul needed	Critical	ADOPT-01, ADOPT-03
3.2	Dynatrace Intelligence enabled	Dynatrace Intelligence problem detection must be active; verify by confirming problems are detected in the last 7 days	Critical	ADOPT-01, ADOPT-02
3.3	Suppress auto-resolving problems	Problems that consistently resolve within 5 minutes: configure delayed notification (5-minute wait) or workflow-based triage	Recommended	ADOPT-05
3.4	Maintenance windows	Configure maintenance windows for all planned change periods to suppress known-noisy alerts	Recommended	ADOPT-03, ADOPT-05
3.5	Alert ownership	Assign every alerting profile to a specific team; no unowned alerts	Recommended	ADOPT-05
3.6	Monthly alert review	Review the top 10 noisiest problem types monthly; adjust sensitivity or suppress non-actionable alerts	Recommended	ADOPT-05
3.7	Filter frequent/duplicate events from metrics	Always exclude <code>dt.davis.is_frequent_event == true and dt.davis.is_duplicate == true</code> when calculating MTTD/MTTR	Critical	ADOPT-03
3.8	Exclude maintenance from MTTR	Always exclude <code>maintenance.is_under_maintenance == true</code> when calculating MTTR	Recommended	ADOPT-03

4. SLOs and Success Metrics

#	Best Practice	Recommended Setting/Value	Priority	Source
4.1	Define SLOs for top 10 services	Set availability and latency SLOs for the 10 most business-critical services	Critical	ADOPT-05
4.2	SLO burn-rate alerting	Enable burn-rate alerting on every SLO; static thresholds are insufficient	Critical	ADOPT-01, ADOPT-05
4.3	Track MTTD	Target: < 5 minutes average. Measure weekly using detected problem <code>event.start</code> vs detection timestamp	Critical	ADOPT-03
4.4	Track MTTR	Target: < 1 hour average (DORA Elite). Measure weekly; trend by problem category	Critical	ADOPT-03
4.5	Track weekly problem count	Trend must be flat or declining; an increasing trend signals environmental degradation	Recommended	ADOPT-03
4.6	Track change failure rate	Target: < 5% (DORA Elite). Correlate deployment events with post-deployment detected problems	Recommended	ADOPT-03
4.7	Establish 30-day baselines	Record current MTTD, MTTR, noise ratio, problem count, and CFR using a 30-day window before setting targets	Critical	ADOPT-03
4.8	Re-baseline quarterly	Recalculate baselines every 90 days to account for growth and environmental changes	Recommended	ADOPT-03
4.9	Set 3-month and 6-month targets	Define explicit improvement targets for each metric at 3-month and 6-month horizons	Recommended	ADOPT-03

5. Cost Optimization

#	Best Practice	Recommended Setting/Value	Priority	Source
5.1	Audit high-volume log sources	Identify top 15 log sources by volume weekly; evaluate each for filtering, sampling, or reduced retention	Recommended	ADOPT-05
5.2	Log retention tiers	High-value logs: 35 days. Standard logs: 14 days. Debug/verbose: 7 days	Recommended	ADOPT-05
5.3	Span sampling for high-volume services	Low-value spans: sample at 10:1. Medium-value: 2:1. High-value: full retention	Recommended	ADOPT-05
5.4	Use <code>samplingRatio</code> for exploratory queries	Set <code>samplingRatio:10</code> or <code>samplingRatio:100</code> on ad-hoc DQL queries against large datasets; multiply results back	Optional	ADOPT-05
5.5	Use <code>scanLimitGBytes</code> on expensive queries	Set <code>scanLimitGBytes:100</code> (or lower) to cap data scanning cost	Optional	ADOPT-05
5.6	Replace raw log storage with metric extraction	For log patterns you only need to count (not inspect), extract as metrics via OpenPipeline instead of storing raw records	Optional	ADOPT-05
5.7	Review monitoring mode quarterly	Verify that <code>FULL_STACK</code> hosts genuinely need full-stack; downgrade infrastructure-only hosts to <code>INFRASTRUCTURE</code> mode	Recommended	ADOPT-05
5.8	Track daily log ingestion over 7 days	Run a 7-day daily log count query monthly; flag any day exceeding 2x the baseline as a cost anomaly	Recommended	ADOPT-02, ADOPT-05

6. Automation and Workflows

#	Best Practice	Recommended Setting/Value	Priority	Source
6.1	Automate alert triage and routing	Create a Dynatrace Workflow that routes detected problems to the correct team channel; run daily	Critical	ADOPT-05
6.2	Automate top 5 frequent problem responses	Build Workflows for the 5 most frequent, well-understood problem types; start with triage, progress to remediation	Recommended	ADOPT-05
6.3	Delayed notification for transient problems	Configure a 5-minute delay before alerting on problem types that auto-resolve within 5 minutes	Recommended	ADOPT-05
6.4	Scheduled DQL reports for leadership	Create a weekly scheduled notebook that runs MTTR, MTTD, noise ratio, and problem count queries automatically	Recommended	ADOPT-05
6.5	Automated postmortem data collection	Build a Workflow triggered by problem closure that collects timeline, root cause, affected entities, and duration into a structured record	Optional	ADOPT-05
6.6	Dynatrace Intelligence forecasting for capacity	Enable Dynatrace Intelligence metric forecasting for CPU, memory, and disk on all production hosts; review monthly	Optional	ADOPT-01, ADOPT-05

7. IAM and Access Control

#	Best Practice	Recommended Setting/Value	Priority	Source
7.1	Eliminate shared admin accounts	Replace shared admin credentials with individual user accounts bound to IAM groups	Critical	ADOPT-05
7.2	Least-privilege IAM policies	Define policies per team granting only the permissions required for their role	Critical	ADOPT-05

#	Best Practice	Recommended Setting/Value	Priority	Source
7.3	Data-scoped access	Use segments (or management zones for Gen2) to restrict data visibility by team	Recommended	ADOPT-05
7.4	IAM group per team	Create an IAM group for every team that uses Dynatrace (e.g., sre-team , platform-team , app-dev)	Recommended	ADOPT-05

8. Team Enablement and Organizational Readiness

#	Best Practice	Recommended Setting/Value	Priority	Source
8.1	1 Dynatrace champion per team	Identify and train at least 1 internal champion per engineering team; complete the platform engineer track within 6 weeks	Critical	ADOPT-04
8.2	2 DQL-literate members per team	Every team must have at least 2 people who can write fetch , filter , summarize queries independently	Critical	ADOPT-04
8.3	100% teams with self-service dashboards	Every team creates and maintains their own service dashboard within 6 months	Recommended	ADOPT-04
8.4	12-week enablement program	Run a structured 12-week program: Weeks 1-2 orientation, 3-4 DQL, 5-6 role-specific, 7-8 dashboards, 9-10 incident practice, 11-12 champion certification	Recommended	ADOPT-04
8.5	Reduce escalation tickets by 50%	Track escalation tickets to the platform team; target 50% reduction within 3 months of champion program launch	Recommended	ADOPT-04
8.6	Monthly champion sync	Hold a monthly meeting for all Dynatrace champions to share learnings and review platform changes	Recommended	ADOPT-04
8.7	Quarterly skills refresh	Run a quarterly training session for champions covering new features and advanced patterns	Optional	ADOPT-04

#	Best Practice	Recommended Setting/Value	Priority	Source
8.8	Role-based learning paths	Assign the correct notebook study sequence per role: SRE (ONBRD, OPLOGS, SPANS, WFLOW, SYNTH), Developer (ONBRD 1-5, SPANS, OTEL, ORGNZ 1-3), Platform (ONBRD, K8S, IAM, AUTOM, ORGNZ, OPMIG), Manager (ADOPT, ONBRD 1-3, ORGNZ 1)	Recommended	ADOPT-04

9. Dashboards and Reporting

#	Best Practice	Recommended Setting/Value	Priority	Source
9.1	3 standard dashboard templates	Create exactly 3 organizational templates: Infrastructure Health, Application Health, Business KPIs	Recommended	ADOPT-05
9.2	Retire unused dashboards	Audit dashboards quarterly; delete any dashboard with no views in the last 90 days	Optional	ADOPT-05
9.3	Weekly platform health scorecard	Review 8 metrics weekly: host coverage, agent version currency, service discovery, log ingestion stability, span ingestion, ActiveGate health, Dynatrace Intelligence activity, alert noise ratio	Critical	ADOPT-02
9.4	Scorecard thresholds	8/8 green = Healthy. 6-7 = Minor gaps (fix this sprint). 4-5 = Significant (prioritize). < 4 = Critical (escalate)	Critical	ADOPT-02
9.5	Monthly ROI report	Combine MTTR/MTTD trends, problem count trends, data volume, enablement progress, and estimated cost avoidance into a single monthly report for leadership	Recommended	ADOPT-05
9.6	SLO dashboard for leadership	Create a dedicated SLO dashboard showing burn rate and compliance for the top 10 services	Recommended	ADOPT-05

10. Configuration Management

#	Best Practice	Recommended Setting/Value	Priority	Source
10.1	Configuration-as-code with Monaco	Export all Dynatrace configuration to Monaco format and store in a Git repository	Recommended	ADOPT-01, ADOPT-05
10.2	CI/CD pipeline for config deployment	Deploy Dynatrace configuration changes through a CI/CD pipeline; never apply changes manually in production	Recommended	ADOPT-05
10.3	Drift detection	Enable Monaco drift detection to catch manual configuration changes; run weekly	Recommended	ADOPT-05
10.4	Dashboard-as-code templates	Store dashboard JSON templates in the config repo; provision new service dashboards from templates via Monaco	Optional	ADOPT-05

11. Distributed Tracing and Spans

#	Best Practice	Recommended Setting/Value	Priority	Source
11.1	OpenTelemetry for uninstrumented services	Implement OpenTelemetry SDKs for any service not covered by OneAgent auto-instrumentation	Recommended	ADOPT-05
11.2	OTLP endpoint routing	Send all OTel data through the Dynatrace OTLP endpoint; do not use a separate collector unless required	Recommended	ADOPT-05
11.3	Validate trace context propagation	Verify end-to-end trace context propagation across all service boundaries; broken propagation = blind spots	Critical	ADOPT-05
11.4	Span volume monitoring	Query top 10 services by span count weekly; apply head-based or tail-based sampling to high-volume, low-value services	Recommended	ADOPT-05

12. Business Observability

#	Best Practice	Recommended Setting/Value	Priority	Source
12.1	Instrument 3-5 key business transactions	Identify the 3-5 highest-value business transactions (e.g., checkout, login, search) and emit business events for each	Recommended	ADOPT-05
12.2	Business-context alerting	Create alerts that trigger on business metric thresholds (e.g., revenue drop, conversion rate decline), not just technical metrics	Recommended	ADOPT-05
12.3	Correlate technical health with business KPIs	Build dashboards that show service latency/error rate alongside business metrics (conversion, revenue) on the same view	Optional	ADOPT-05
12.4	Entity ownership assignment	Assign an owning team to every monitored service and application entity in Dynatrace	Recommended	ADOPT-01, ADOPT-04

13. Coverage Audit and Staged Enablement

#	Best Practice	Recommended Setting/Value	Priority	Source
13.1	Run a Wave-0 coverage audit	Measure hosts by monitoring mode, log coverage, consumption by mode, and deep-monitoring exceptions; record as the enablement baseline	Critical	ADOPT-06
13.2	Segment the estate before enabling	Match mode to segment: revenue-path apps → Full-Stack, backing services → Infrastructure, long-tail fleet → Discovery	Critical	ADOPT-06
13.3	Enable in connected segments, not layers	Bring one application's full path (Full-Stack + logs + RUM) online before spreading a single layer estate-wide	Recommended	ADOPT-06

#	Best Practice	Recommended Setting/Value	Priority	Source
13.4	Sequence Waves 1–4 in order	Core apps to Full-Stack → logs → RUM/synthetic → AppSec and optimization; each wave rides the modules the prior wave already enabled	Recommended	ADOPT-06
13.5	Gate each wave on a moved metric	Use ADOPT-03's success metrics as value gates (MTTR delta, log-evidence rate, user-impact rate, vulnerability-backlog trend) before funding the next wave	Critical	ADOPT-06
13.6	Re-run the Wave-0 audit at every gate	Coverage decays between waves as new hosts land in default modes and exclusion rules outlive their incidents; re-audit quarterly	Recommended	ADOPT-06

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.