

AIOPS-01: Dynatrace Intelligence

Overview

Series: AIOPS — Dynatrace Intelligence | **Notebook:** 1 of 8 | **Created:** May 2026 |
Last Updated: 09/24/2026

Overview

Dynatrace Intelligence is the umbrella over the platform's AI capabilities. Three categories work together — Causal AI grounds every analysis in topology, Predictive AI projects metrics forward, Generative AI translates between natural language and DQL.

This notebook orients you to the surface area: what the categories are, where they appear in the product, what data each one operates on, and where to dive in next.

Audience: Platform admin, SRE, observability lead.

Outcome: A mental map of Dynatrace Intelligence — and a starting place for any AIOps initiative.




Table of Contents

- 1. [The Three AI Categories](#)
 - 2. [Where AI Surfaces in the Product](#)
 - 3. [The Data Foundation: Smartscape + Grail](#)
 - 4. [Quick Reality Check: What Davis Is Doing Right Now](#)
 - 5. [The Series Map](#)
 - 6. [Next Steps](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Grail (Dynatrace Platform / Gen3)
Permissions	<code>events:read</code> and <code>storage:events:read</code> for problem queries; <code>davis:analyzers:execute</code> for analyzer notebooks (introduced later in series)
Apps	Problems app, Anomaly Detection app, Notebooks app
Optional	Dynatrace Assist enabled (most modern tenants); MCP server for AIOps-06

1. The Three AI Categories

Dynatrace Intelligence does not pile every AI buzzword into one model. It separates by **what kind of question you're asking**.

Category	Question it answers	Examples in product
Causal AI	<i>Why did this happen?</i>	Davis Problems, automatic root cause analysis, dependency-aware event correlation

Category	Question it answers	Examples in product
Predictive AI	<i>What will happen, or what's outside normal?</i>	Auto-adaptive baselines, seasonal baselines, forecasts, novelty detection
Generative AI	<i>Translate between English and DQL — explain, summarize, generate</i>	Dynatrace Assist (chat), DQL2NL skill, Problems app summaries

Causal ≠ Correlation. Causal AI reasons over the Smartscape topology graph: it knows that service A depends on service B which runs on host C, so when an anomaly fires on C it can attribute the symptom on A. Generic ML correlation engines see the symptoms but can't reliably point to the cause.

Predictive ≠ Generative. A seasonal baseline is a statistical model — it has nothing in common with an LLM. Both are AI; they sit in different boxes for different jobs.

2. Where AI Surfaces in the Product

Surface	Powered by	What you see
Problems app	Causal AI + Generative summaries	Grouped problems, ranked root cause, narrative summary
Anomaly Detection app	Predictive AI (analyzers)	Detector configurations, custom alerts, baseline reviews
Smartscape Davis	Causal AI	Dependency-aware problem grouping happens implicitly
Dynatrace Assist (chat)	Generative AI	Natural-language DQL, query explanation, conversation starters
Kubernetes app	Generative AI	Warning-signal explanations
Databases app	Generative AI	Execution plan clarification
Workflow tasks	All three categories	Schedule analyzer runs; summarize problems; notify
Notebooks (DQL cells)	All three categories	Davis CoPilot side panel; DQL2NL on selected query

Surface	Powered by	What you see
AI Observability	Your GenAI app, not Dynatrace's AI	GenAI / LLM span analysis, token cost, prompt evaluations — plus a dedicated Smartscape view and a dedicated Evaluations screen (SaaS 1.344)

Most surfaces compose categories. The Problems app is the clearest example — Causal AI groups the events into a problem, and a Generative AI summary explains it in English.

The last row is the inverse of the other eight. Rows 1–8 are Dynatrace applying AI to *your* data. AI Observability is *your team observing your own GenAI applications* — Dynatrace is the observability tool, not the intelligence. It earns a place in this table only because readers asking "where does AI show up in the product?" land here first and would otherwise miss it entirely. For ingest mechanics, the `gen_ai.*` attribute namespace, and the conversation-tracking pattern, see **AIOPS-06 § 1**.

SaaS 1.344 (rollout from 07/29/2026) added the dedicated Smartscape view and the standalone Evaluations screen. On a tenant still below 1.344, prompt evaluations surface inside the general app views.

3. The Data Foundation: Smartscape + Grail

Dynatrace Intelligence sits on two distinct foundations:

Smartscape is the real-time dependency graph — every entity (host, process, service, K8s object, AWS resource) and every observed relationship between them. Causal AI walks this graph to localize root cause.

Grail is the data lakehouse — logs, events, spans, business events, problems, all in one queryable store. Predictive analyzers run over Grail timeseries; the Problems app reads from `dt.davis.problems`; Dynatrace Assist generates DQL that targets Grail.

Without Smartscape, Dynatrace would have correlation. With Smartscape, it has causation. Without Grail, the AI has no consistent corpus to reason over. Both are prerequisites — neither is replaceable.

4. Quick Reality Check: What Davis Is Doing Right Now

Before we go deeper, let's see Davis at work in your tenant. The next two queries answer:

1. *What problems is Causal AI surfacing right now?*
2. *How busy is the underlying signal stream that gets grouped into those problems?*

If problem counts are low, that's a healthy environment — not a quiet one. If raw signal counts are very high but problem counts are low, Causal AI is doing the right thing: collapsing noise into a small number of actionable problems.

```
// Active problems right now, by category
fetch dt.davis.problems, from:-2h
| filter event.status == "ACTIVE"
| summarize active_problems = count(), by:{event.category}
| sort active_problems desc
```

```
// Raw Davis signal volume in the last hour
// (all Davis events – including INFO/WARNING observations that never open a
// problem)
fetch dt.davis.events, from:-1h
| summarize signal_count = count(), by:{event.category}
| sort signal_count desc
```

Reading the result: Compare the two queries. The signal-to-problem ratio is the Causal AI compression factor — the higher the ratio, the more noise the platform is keeping off your alerting plate. Compute the compression factor from the non-INFO/WARNING categories only; the full stream on a tenant with observation-mode detectors is dominated by events that never become problems.

Note on data objects: `dt.davis.problems` carries grouped problems (`event.kind == "DAVIS_PROBLEM"`). `dt.davis.events` carries raw signals (`event.kind == "DAVIS_EVENT"`). They are sibling streams. `fetch dt.davis.events | filter event.kind == "DAVIS_PROBLEM"` returns zero rows — **not because the form was deprecated, but because it filters the wrong table:** `dt.davis.events` has only ever carried `DAVIS_EVENT` . It is a category error, so there is no migration to look for; use `fetch dt.davis.problems` for problem-shaped queries. Verified 08/27/2026:

over 7 days `dt.davis.events` returned 206,023 rows, all `DAVIS_EVENT`, and 0 when filtered to `DAVIS_PROBLEM`, while `dt.davis.problems` returned 2,774.

5. The Series Map

Notebook	Topic	When to read
AIOPS-01	Overview (<i>you are here</i>)	Orientation
AIOPS-02	Anomaly Detection	When tuning detectors or reasoning about baseline strategy
AIOPS-03	Davis AI — Problems & RCA	When investigating a specific incident or designing problem-driven workflows
AIOPS-04	Davis CoPilot — Dynatrace Assist	When integrating natural-language interaction into your team's workflow
AIOPS-05	AI Models	When you need to understand which model is doing what under the hood
AIOPS-06	AI Integrations & Agentic Workflows	When extending AI into automation, MCP servers, or external IDEs
AIOPS-07	Putting It Together	End-to-end Detect → Investigate → Remediate patterns
AIOPS-99	Series Summary	Cross-reference index, next steps, where AIOPS sits in the broader maturity model

Cross-references to other series (because AIOPS does not stand alone):

- **WFLOW** (Workflows & Alert Notifications) — once you have problems detected here, route them through WFLOW
- **AUTOM** (Configuration Automation) — anomaly detection settings as code via Monaco / Terraform
- **ADOPT** (Adoption & Maturity) — AIOps maturity model and metrics
- **DASH** (Dashboards) — problem-driven SLO and executive reporting

6. Next Steps

1. **Open the Problems app** and scan the last 24 hours. Are problems grouped by Causal AI in a way that matches your operational mental model?
2. **Open Anomaly Detection app** to see what detectors are active. (We go deep in AIOPS-02.)
3. **Open a notebook and try Dynatrace Assist** — ask in plain English: *"show me the slowest services in the last hour."* (We go deep in AIOPS-04.)
4. **Read AIOPS-02** if you want to start tuning detectors. **Read AIOPS-03** if a problem in production has your attention right now.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.