

ALERT-99: Best-Practice Summary and Setup Checklist

Series: ALERT — Alerting Strategy and Design | **Notebook:** 99 of 05 | **Created:** June 2026 | **Last Updated:** 06/25/2026

Overview

The series on one page: the principles, a complete end-to-end setup checklist, and the cross-series map of where each piece is built. Use this as the standing reference once you have read ALERT-01 through ALERT-04.

Table of Contents

- 1. [Principles](#)
- 2. [Complete Setup Checklist](#)
- 3. [The Ongoing Audit Loop](#)
- 4. [Cross-Series Map](#)

Prerequisites

Requirement	Details
Prior reading	ALERT-01 through ALERT-04
Companion series	AIOPS, SLO, WFLOW, OPIPE

1. Principles

1. **OOTB Davis first.** Tune before you build; descend the funnel only when the layer above cannot express the condition.
2. **One enriched problem is the hub.** Every mechanism converges on a Davis problem; everything downstream operates on it.
3. **Enrich upstream or you cannot route downstream.** Team/zone/severity belong on the problem before a workflow ever sees it.
4. **Burn-rate beats threshold-on-SLI.** Alert on how fast the budget drains, with multiwindow confirmation.
5. **Prefer simple workflows.** Pay for multi-step only when you need conditional/multi-team logic.
6. **Match destination to urgency.** Fast-burn → page; slow-burn → ticket.
7. **Validate every query before it becomes an alert.** A detector on a silent query is worse than no alert.
8. **Audit on a cadence.** Building well is half the job. A standing review — classify every problem, trace each false positive to the config behind it, and track the trend per area — keeps the funnel honest as services change.

2. Complete Setup Checklist

Detection

- [] Reviewed OOTB Davis anomaly detection settings and tuned sensitivity
- [] Custom anomaly detectors only where Davis does not cover the condition (auto-adaptive/seasonal, not static-on-traffic)
- [] Log/span signals extracted to OpenPipeline metrics before alerting on them
- [] SLOs defined for critical user journeys with burn-rate alerts

Enrichment

- [] Detector event templates carry team / zone / service properties
- [] Entity tags and Smartscape ownership populated for OOTB problems
- [] `event.severity` used to drive priority

Routing

- [] Problem-trigger workflows filter on enriched metadata
- [] Simple workflows used wherever conditional logic is not needed

- [] Fast-burn → page channel; slow-burn → ticket channel
- [] One notification per problem (no duplicate alerts on raw metrics)

ITSM / destinations

- [] ServiceNow incidents carry assignment group, priority, and a de-dup correlation key
- [] Legacy alerting-profile integrations identified; migration to workflows considered

Governance

- [] Production detectors/SLOs promoted to config-as-code
- [] Alerts that fire without action reviewed and tuned

3. The Ongoing Audit Loop

The setup checklist above is build-time. Noise creeps back as services, traffic, and ownership change, so a standing audit is what keeps the funnel honest. The point is not to *count* false positives — it is to find the repeating configuration patterns behind them, so the same noise is not regenerated next period.

Classify every problem. Each problem in the audit window lands in exactly one bucket:

Class	Meaning	Action
True Positive (TP)	A real issue that required action	None — the detector is working
False Positive (FP)	Fired but needed no action — a threshold, scope, detection-method, or sliding-window issue	Trace it to the config and fix the source
Under Review (UR)	Not yet conclusive	Carry to the next period; gather more data points

Every FP must trace back to a concrete cause — the analyzer type, the entity scope, the sliding window, or the event template. An FP with no configuration change behind it is just a number; an FP with an identified root cause is progress.

Track one KPI: the false-positive rate over time, per area, and whether it is trending down. That single trend is the primary output of every cycle.

Signals to scan each cycle:

Signal	What it flags
Fire rate	A detector firing many times a day is a tuning or de-dup candidate
Silence rate	Never fired in 90+ days → redundant, or watching a dead path
Immediate-silence rate	Acknowledged and closed with no action = noise
Routing coverage	Every alert should route to exactly one owner; an unrouted alert is invisible
Runbook link	An alert with no runbook attached has no possible action
Owner / area property	Present, so the alert maps to a team — also catches copy-paste clones that never set it

Cadence:

Stage	When	Focus
Pre-go-live	Before an area or service reaches production	Every new or migrated detector fires correctly in non-prod; nothing fires on no-data
Early (first ~3 months)	Monthly	Top-noise alerts reviewed with the owning team
Steady state	Quarterly	Cross-area review; retire detectors silent 90+ days or at a 100% silence rate

Express these as recurring queries over the problem feed (`fetch events | filter event.kind == "DAVIS_PROBLEM"`); AIOPS-03 and ORGNZ-10 carry validated problem-feed query patterns to build on.

4. Cross-Series Map

Piece	Built in
Anomaly detector mechanics (analyzer, tuning, event template)	AIOPS-02

Piece	Built in
Metric events / detector schemas as code	AIOPS-02 §6, AUTOM-05/06
OpenPipeline metric extraction	OPIPE, FAQ-09
SLIs, error budgets, burn-rate alerting	SLO-02 / 03 / 04
SLOs as code	SLO-05
Routing, conditions, escalation	WFLOW-03 / 04
HTTP / webhook actions	WFLOW-08
ServiceNow incident integration	ALERT-04
Cost / DPS framing	FINOPS, FAQ-09

Sources: [Alerting and notifications \(DT docs\)](#), [Workflows \(DT docs\)](#). **Derived:** the checklist and funnel synthesise the ALERT, AIOPS, SLO, and WFLOW series into one operational sequence.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.