

APPSEC-06: Kubernetes and Container Security

Series: APPSEC — Application Security | **Notebook:** 6 of 10 | **Created:** June 2026 |
Last Updated: 09/24/2026

Overview

Kubernetes and container workloads produce findings across all three AppSec pillars at once — RVA on the running containerized application, RAP on its request traffic, and SPM on the cluster + image configuration. This notebook covers how the pieces fit together and the K8s-specific levers.

For K8s rollout patterns themselves (DynaKube CR, operator install, namespace selectors), see the K8S series. This notebook focuses on the *security* dimension specifically.



Table of Contents

1. [1. DynaKube and applicationMonitoring](#)
 2. [2. Container Image Vulnerabilities](#)
 3. [3. Cluster SPM Findings](#)
 4. [4. DQL: K8s-Scoped Findings](#)
 5. [5. Namespace Scoping for IAM](#)
 6. [6. Next Steps](#)
 7. [References](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	Gen3 SaaS with Grail; AppSec entitlement enabled
OneAgent	Full-Stack mode (or code-module attached) on monitored hosts
Read access	At minimum <code>environment:roles:view-security-problems</code> and <code>storage:security.events:read</code> — see APPSEC-09 for the full model
Background	APPSEC-01 (fundamentals + three-pillar framing)

1. DynaKube and applicationMonitoring

The DynaKube custom resource is the entry point for AppSec on Kubernetes. The relevant fields:

```
apiVersion: dynatrace.com/v1beta5
kind: DynaKube
metadata:
  name: dynakube
  namespace: dynatrace
spec:
  apiUrl: https://&lt;tenant&gt;.live.dynatrace.com/api
  oneAgent:
    applicationMonitoring: {}
```

`useCSIDriver` is not a `v1beta5` field: the DynaKube parameters reference lists it only for the retired `v1beta1` / `v1beta2` APIs. Whether code modules come from the CSI driver is decided when the Operator is installed (CSI or *Without CSI driver* variant).

Sources: [DynaKube parameters \(DT docs\)](#) — *"DynaKube API version v1beta2 is no longer available with Dynatrace Operator version 1.7.0"*; [Application observability setup \(DT docs\)](#) — *"CSI driver is optional (see step 2). If enabled, it gets deployed as DaemonSet and results in a CSI driver pod on each node."*

`applicationMonitoring` enables the code-module injection that RVA and RAP rely on inside workload containers. Monitoring mode then decides how well findings are assessed: per the Application Security docs, Infrastructure and Discovery modes still provide third-party and code-level detection (limited) and RAP — Discovery only once code-module injection is enabled — but without the Full-Stack topology that adjusts the Dynatrace Security Score, so DSS stays at the CVSS base score (APPSEC-01 § 3). SPM findings on the cluster itself do not depend on monitoring mode.

For cluster rollout the K8S series covers per-namespace targeting, monitoring modes, and operator versioning. APPSEC concerns are downstream of that.

Sources: [Application Security \(DT docs\)](#) — the *Monitoring modes coverage* table and *"For Application Security to work in Discovery mode, after enabling Discovery mode, you also need to enable code-module injection."* (re-read 09/18/2026). **Softened:** the exact DynaKube schema may shift per operator release — verify against the current K8S series and the operator CRD.

2. Container Image Vulnerabilities

Image-tier findings are a subset of RVA: Dynatrace inventories the libraries in running container images and reports CVEs against them. The result lands in `security.events` with the same vulnerability event-type as host-level RVA but with K8s-specific context fields (cluster, namespace, workload, image name + digest).

Practical pattern: report image-tier vulnerabilities **by namespace** rather than by individual workload — namespaces are usually a stable team/service boundary, while workload names churn with deployments.

Sources: [Application Security \(DT docs\)](#) for the K8s + image scope. **Softened:** the namespace-as-stable-boundary recommendation is community practice — verify the namespace-naming discipline in your tenant before relying on it for reports.

3. Cluster SPM Findings

SPM (APPSEC-05) covers Kubernetes cluster configuration: RBAC bindings, pod security, network policies, secret management. Common high-impact findings:

- Cluster-admin-bound users that shouldn't be cluster-admin
- Pods running as root without explicit need
- Network policies absent on namespaces that should be isolated
- Secrets stored as ConfigMaps instead of Secrets (or unencrypted at rest)
- Public LoadBalancer services exposing internal-only workloads

These findings are durable — they reflect the cluster's declared state and don't churn the way runtime events do.

Sources: [Application Security \(DT docs\)](#) for SPM scope. **Softened:** the high-impact finding list reflects the standard CIS-K8s control set; the exact Dynatrace SPM catalog should be confirmed against current docs.

4. DQL: K8s-Scoped Findings

Filter AppSec events to a specific cluster + namespace to slice findings by team boundary.

```
// AppSec findings in a specific cluster + namespace, last 24h
fetch security.events, from:-24h
| filter k8s.cluster.name == "prod-cluster-01"
| filter k8s.namespace.name == "team-payments"
| summarize count = count(), by:{event.type, vulnerability.risk.level}
| sort count desc
```

Sources: [IAM policy statements reference \(DT docs\)](#) lists `storage:k8s.cluster.name` and `storage:k8s.namespace.name` as conditions on `storage:security.events:read`.

Dictionary: `k8s.namespace.name` (stable), `vulnerability.risk.level` (stable), read 09/18/2026. **Live-verified 09/18/2026:** `k8s.namespace.name` is populated on `COMPLIANCE_FINDING` records on the validation tenant; the cluster and namespace names in the query are placeholders, so replace them with your own. `vulnerability.risk.level` is populated only on vulnerability events — for compliance records group by `compliance.rule.severity.level` instead.

5. Namespace Scoping for IAM

K8s namespace is the natural IAM boundary for AppSec findings: AppDev team X should see findings in its own namespaces, not in team Y's. The IAM permission model supports this via boundary conditions:

```
ALLOW storage:buckets:read WHERE storage:table-name = "security.events";
ALLOW storage:security.events:read
  WHERE storage:k8s.namespace.name IN ("team-payments-prod", "team-payments-
  staging");
```

Two details make or break this policy. IAM lists use `IN ("...", "...")` with parentheses — DQL's `{...}` array syntax does not validate here. And `storage:buckets:read` is required in addition to the table permission; without it the namespace-scoped grant reads nothing. `vulnerability-service:vulnerabilities:read` takes no conditions, so it cannot be

namespace-scoped — grant it separately only if the team needs the API (APPSEC-09 § 3).

See APPSEC-09 for the full pattern including the policy-vs-managed-policy decision and the privacy carve-out for `view-sensitive-request-data`.

Sources: [IAM policy statements reference \(DT docs\)](#) confirms K8s namespace as an available boundary condition on `storage:security.events:read` and says of `storage:buckets:read` : *"Grants permission to read records from Grail buckets. Required additionally to a table permission."*; [IAM policy statement syntax \(DT docs\)](#) — its `IN` example is `WHERE settings:schemaId IN ("builtin:container.monitoring-rule", "builtin:container.built-in-monitoring-rule")` (both re-read 09/18/2026). The full IAM model is in APPSEC-09.

6. Next Steps

1. Confirm DynaKube `applicationMonitoring` is enabled on clusters where AppDev-tier vulnerability data matters.
2. Run the namespace-scoped DQL above against a known-active namespace to verify field names match.
3. Read the **K8S series** for the rollout mechanics — this notebook intentionally defers to it.
4. Read **APPSEC-09** for the full K8s namespace IAM scoping pattern.

References

Source	Coverage
Application Security (DT docs)	K8s + container security framing
IAM policy statements reference (DT docs)	k8s.namespace.name boundary condition
IAM policy statement syntax (DT docs)	Condition operators (<code>=</code> , <code>IN (...)</code> , <code>startswith</code> , <code>MATCH</code>)

 **DISCLAIMER:** This information was AI generated and is provided "as-is" without warranty. It was produced as an independent, community-driven project and **not supported by Dynatrace**. Always refer to official [Dynatrace documentation](#) for the most current information.