

APPSEC-06: Kubernetes and Container Security

Series: APPSEC — Application Security | **Notebook:** 6 of 10 | **Created:** June 2026 |
Last Updated: 10/02/2026

Overview

Kubernetes and container workloads produce findings across all three AppSec pillars at once — RVA on the running containerized application, RAP on its request traffic, and SPM on the cluster + image configuration. This notebook covers how the pieces fit together and the K8s-specific levers.

For K8s rollout patterns themselves (DynaKube CR, operator install, namespace selectors), see the K8S series. This notebook focuses on the *security* dimension specifically.



Table of Contents

1. [1. DynaKube and applicationMonitoring](#)
 2. [2. Container Image Vulnerabilities](#)
 3. [3. Cluster SPM Findings](#)
 4. [4. DQL: K8s-Scoped Findings](#)
 5. [5. Namespace Scoping for IAM](#)
 6. [6. Next Steps](#)
 7. [References](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	Gen3 SaaS with Grail; AppSec entitlement enabled
OneAgent	Full-Stack mode (or code-module attached) on monitored hosts
Read access	To run the DQL: <code>storage:security.events:read</code> plus <code>storage:buckets:read</code> (a table permission alone reads nothing). The Vulnerabilities and Threats & Exploits apps have their own requirements — see APPSEC-09 for the full model
Background	APPSEC-01 (fundamentals + three-pillar framing)

1. DynaKube and applicationMonitoring

The DynaKube custom resource is the entry point for AppSec on Kubernetes. The relevant fields:

```

apiVersion: dynatrace.com/v1beta5
kind: DynaKube
metadata:
  name: dynakube
  namespace: dynatrace
spec:
  apiUrl: https://<tenant>.live.dynatrace.com/api
  oneAgent:
    applicationMonitoring: {}

```

`useCSIDriver` is not a `v1beta5` field: the DynaKube parameters reference lists it only for the retired `v1beta1` / `v1beta2` APIs. Through Operator 1.10.x, whether code modules come from the CSI driver is decided when the Operator is installed (CSI or *Without CSI driver* variant), and that remains the working path on those versions.

Operator 1.11.0+ (released 10/01/2026): image volumes. Dynatrace now recommends image volume-based code-module injection, which *"replaces the CSI driver as the recommended approach"*: each node pulls the code-modules image once and shares it with every instrumented pod, with no CSI DaemonSet. Turn it on per DynaKube with the annotation `feature.dynatrace.com/mount-code-modules-via-image-volume: "true"` (mutually exclusive with `feature.dynatrace.com/node-image-pull`). It needs **Kubernetes 1.35+** and **containerd 2.2+ or CRI-O 1.33+**, and *"Image volume injection is not compatible with the Dynatrace built-in tenant registry."* Try it on one workload first with the pod annotation `oneagent.dynatrace.com/volume-type: "image"`; *"A full migration requires a rolling restart of all injected workloads."* Clusters that do not meet those requirements stay on the CSI driver or ephemeral volumes. The delivery mode changes how code modules reach the pod; RVA and RAP still depend only on code-module injection being enabled.

Sources:

- [DynaKube parameters \(DT docs\)](#) — *"DynaKube API version v1beta2 is no longer available with Dynatrace Operator version 1.7.0"*
- [Application observability setup \(DT docs\)](#) — *"CSI driver is optional (see step 2). If enabled, it gets deployed as DaemonSet and results in a CSI driver pod on each node."*
- [Operator 1.11.0 release notes \(DT docs\)](#)
- [Use image volumes for code modules injection \(DT docs\)](#) — requirements, both annotations, and the registry limit
- [Migrate to image volumes \(DT docs\)](#)

`applicationMonitoring` enables the code-module injection that RVA and RAP rely on inside workload containers. Monitoring mode then decides how well findings are assessed: per the Application Security docs, Infrastructure and Discovery modes still provide third-party and code-level detection (limited) and RAP — Discovery only once code-module injection is enabled — but without the Full-Stack topology that adjusts the Dynatrace Security Score, so DSS stays at the CVSS base score (APPSEC-01 § 3). SPM findings on the cluster itself do not depend on monitoring mode.

For cluster rollout the K8S series covers per-namespace targeting, monitoring modes, and operator versioning. APPSEC concerns are downstream of that.

Sources: [Application Security \(DT docs\)](#) — the *Monitoring modes coverage* table and *"For Application Security to work in Discovery mode, after enabling Discovery mode, you also need to enable code-module injection."* (re-read 09/18/2026). **Softened:** the exact DynaKube schema may shift per operator release — verify against the current K8S series and the operator CRD.

2. Container Image Vulnerabilities

RVA reports vulnerable libraries loaded by processes running in containers, plus vulnerabilities on Kubernetes nodes. A container image is a *related* entity of those findings, not the thing RVA scans. RVA vulnerability events identify the cluster and workload through the arrays `related_entities.kubernetes_clusters.names` and `related_entities.kubernetes_workloads.names`; they carry no `k8s.namespace.name`.

Container-*image* findings are a different record: `VULNERABILITY_FINDING` events with `object.type == "CONTAINER_IMAGE"`, `container_image.registry` / `container_image.repository`, and the Kubernetes resource fields (`k8s.cluster.name`, `k8s.namespace.name`). These come from vulnerability findings, including ingested third-party image scanners.

Practical pattern: report image findings **by namespace** rather than by individual workload — namespaces are usually a stable team/service boundary, while workload names churn with deployments.

Sources: [Vulnerabilities concepts \(DT docs\)](#) — *"Affected entities Entities (process groups, processes, and Kubernetes nodes) for which a vulnerability was detected"* and *"Related container image In Kubernetes environments, the container image used by the affected processes."*; [Vulnerability events](#)

([DT semantic dictionary](#)) — the Kubernetes resource fields appear under the vulnerability finding and scan events, whose example query filters `object.type == "CONTAINER_IMAGE"` (re-read 10/02/2026). **Softened:** the namespace-as-stable-boundary recommendation is community practice — verify the namespace-naming discipline in your tenant before relying on it for reports.

3. Cluster SPM Findings

SPM (APPSEC-05) covers Kubernetes cluster configuration: RBAC bindings, pod security, network policies, secret management. Common high-impact findings:

- Cluster-admin-bound users that shouldn't be cluster-admin
- Pods running as root without explicit need
- Network policies absent on namespaces that should be isolated
- Secrets stored as ConfigMaps instead of Secrets (or unencrypted at rest)
- Public LoadBalancer services exposing internal-only workloads

These findings are durable — they reflect the cluster's declared state and don't churn the way runtime events do.

Sources: [Application Security \(DT docs\)](#) for SPM scope. **Softened:** the high-impact finding list reflects the standard CIS-K8s control set; the exact Dynatrace SPM catalog should be confirmed against current docs.

4. DQL: K8s-Scoped Findings

Filter AppSec events to a specific cluster + namespace to slice findings by team boundary.

```
// Currently failing compliance rules in one cluster + namespace (latest
result per object + rule)
// Counting every record counts scan results, most of them PASSED or
NOT_RELEVANT (APPSEC-05 § 3).
fetch security.events, from:-7d
| filter event.type == "COMPLIANCE_FINDING"
| filter k8s.cluster.name == "prod-cluster-01" and k8s.namespace.name ==
"team-payments"
| dedup {object.id, compliance.rule.id}, sort:{timestamp desc}
| filter compliance.result.status.level == "FAILED"
| summarize {failing = count()}, by:{compliance.standard.short_name,
compliance.rule.severity.level}
| sort failing desc

// Container-image vulnerability findings in the same namespace (docs pattern;
run separately):
// fetch security.events, from:-7d
// | filter event.type == "VULNERABILITY_FINDING" and object.type ==
"CONTAINER_IMAGE"
// | filter k8s.namespace.name == "team-payments"
// | dedup {object.id, vulnerability.id, component.name, component.version},
sort:{timestamp desc}
// | summarize {findings = count()}, by:{dt.security.risk.level}
```

Sources: [IAM policy statements reference \(DT docs\)](#) lists `storage:k8s.cluster.name` and `storage:k8s.namespace.name` as conditions on `storage:security.events:read` ; [Vulnerability events \(DT semantic dictionary\)](#) for the `VULNERABILITY_FINDING` dedup pattern (re-read 10/02/2026). **Dictionary:** `k8s.namespace.name` (`stable`), read 09/18/2026. **Live-verified 10/02/2026:** on one validation namespace the previous count-everything form returned 31,400 `COMPLIANCE_FINDING` records; the deduplicated, `FAILED` -only form returns 16. The cluster and namespace names above are placeholders — replace them with your own. The image query is the docs' pattern; the validation tenant has no vulnerability findings to run it against.

5. Namespace Scoping for IAM

K8s namespace is the natural IAM boundary for AppSec findings: AppDev team X should see findings in its own namespaces, not in team Y's. The IAM permission model supports this via boundary conditions:

```
ALLOW storage:buckets:read WHERE storage:table-name = "security.events";
ALLOW storage:security.events:read
  WHERE storage:k8s.namespace.name IN ("team-payments-prod", "team-payments-
  staging");
```

Two details make or break this policy. IAM lists use `IN ("...", "...")` with parentheses — DQL's `{...}` array syntax does not validate here. And `storage:buckets:read` is required in addition to the table permission; without it the namespace-scoped grant reads nothing. Two limits follow from the data. RVA vulnerability state and change events carry no `k8s.namespace.name`, so this condition cannot match them — the team sees compliance and vulnerability *findings* only; scope RVA visibility by `dt.security_context` instead (APPSEC-09 § 3). And `vulnerability-service:vulnerabilities:read` takes no conditions, so granting it shows every vulnerability in the tenant in the Vulnerabilities app.

See APPSEC-09 for the full pattern including the policy-vs-managed-policy decision and the privacy carve-out for `view-sensitive-request-data`.

Sources: [IAM policy statements reference \(DT docs\)](#) confirms K8s namespace as an available boundary condition on `storage:security.events:read` and says of `storage:buckets:read`: *"Grants permission to read records from Grail buckets. Required additionally to a table permission."*; [IAM policy statement syntax \(DT docs\)](#) — its `IN` example is `WHERE settings:schemaId IN ("builtin:container.monitoring-rule", "builtin:container.built-in-monitoring-rule")` (both re-read 09/18/2026). The full IAM model is in APPSEC-09.

6. Next Steps

1. Confirm DynaKube `applicationMonitoring` is enabled on clusters where AppDev-tier vulnerability data matters.
2. Run the namespace-scoped DQL above against a known-active namespace to verify field names match.
3. Read the **K8S series** for the rollout mechanics — this notebook intentionally defers to it.
4. Read **APPSEC-09** for the full K8s namespace IAM scoping pattern.

References

Source	Coverage
Application Security (DT docs)	K8s + container security framing
IAM policy statements reference (DT docs)	k8s.namespace.name boundary condition
IAM policy statement syntax (DT docs)	Condition operators (= , IN (...) , startswith , MATCH)

 **DISCLAIMER:** This information was AI generated and is provided "as-is" without warranty. It was produced as an independent, community-driven project and **not supported by Dynatrace**. Always refer to official [Dynatrace documentation](#) for the most current information.