

APPSEC-09: IAM and Gen3 Permissions for AppSec

Series: APPSEC — Application Security | **Notebook:** 9 of 10 | **Created:** June 2026 | **Last Updated:** 10/02/2026

Overview

AppSec data is among the most sensitive material in Dynatrace — vulnerability inventory, captured attack payloads, compliance findings. Granting too much access exposes private payloads; granting too little blocks the SOC from doing its job. This notebook is the IAM-design notebook for AppSec: the permission catalog (verified verbatim against the IAM policy statements reference), three persona policies, boundary patterns, and the privacy carve-out for `view-sensitive-request-data`.

This is the most-grounded notebook in the series — the [IAM policy statements reference \(DT docs\)](#) was resolvable at series-creation time and the tokens below were verified verbatim on 06/04/2026.




Table of Contents

- 1. 1. AppSec IAM in Gen3 — Same DSL, Specific Tokens
- 2. 2. The AppSec Permission Catalog (Verified)
- 3. 3. Persona Policies (Worked Examples)
- 4. 4. Boundary Patterns for AppSec Data
- 5. 5. Privacy: view-sensitive-request-data
- 6. 6. Reach for Managed Policies First
- 7. 7. OAuth Client vs Platform Token Scope Routing
- 8. 8. Audit and Periodic Review
- 9. 9. What's Still Evolving
- 10. 10. Next Steps
- 11. References

Prerequisites

Requirement	Details
Dynatrace Environment	Gen3 SaaS with Grail; AppSec entitlement enabled
OneAgent	Full-Stack mode (or code-module attached) on monitored hosts
Read access	To run the DQL: <code>storage:security.events:read</code> plus <code>storage:buckets:read</code> (a table permission alone reads nothing). The Vulnerabilities and Threats & Exploits apps have their own requirements — see APPSEC-09 for the full model
Background	APPSEC-01 (fundamentals + three-pillar framing)

1. AppSec IAM in Gen3 — Same DSL, Specific Tokens

AppSec data in Gen3 lives in Grail (`security.events`) and the `vulnerability-service` API. Both are governed by the same `ALLOW <service>:<resource>:<action> WHERE <conditions>` policy DSL as the rest of Gen3 — not by classic environment roles.

The `environment:roles:*` permissions — `view-security-problems`, `manage-security-problems`, `view-sensitive-request-data` — belong to the **deprecated classic apps** (Third-Party Vulnerabilities, Security Overview). In Latest Dynatrace the surfaces are:

- Grail data → `storage:security.events:read` **plus** `storage:buckets:read` (and `storage:fieldsets:read` for sensitive fields — § 5)
- The Vulnerabilities app and API → `vulnerability-service:vulnerabilities:read` / `:write`, plus the app's documented default policies: *Read Entities*, *Read Security Events*, and one of *Admin User*, *Pro User* or *Standard User*
- Classic apps only → `environment:roles:*`

A complete AppSec persona policy needs the Grail grants and the vulnerability-service grants; the classic roles matter only while you still use the classic apps. See IAM-04 § Policy Authoring for the underlying DSL grammar.

Sources: [IAM policy statements reference \(DT docs\)](#) — `storage:security.events:read` : "Read security.events from grail"; `vulnerability-service:vulnerabilities:read` : "Allows viewing vulnerabilities"; [Vulnerability Analytics \(DT docs\)](#) — "This permissions section refers to the classic Third-Party Vulnerabilities and Security Overview apps, which are deprecated."; [Vulnerabilities \(DT docs\)](#) — "Read Entities Read Security Events One of the following user policies: Admin User , Pro User , Standard User" (all re-read 10/02/2026).

2. The AppSec Permission Catalog (Verified)

Tokens below were verified verbatim on 10/02/2026 against the IAM policy statements reference. Re-verify before pinning these into a production policy.

AppSec-specific tokens

Token	Purpose	Optional conditions
<code>environment:roles:view-security-problems</code>	Classic apps only: view security problems	Management zone scope
<code>environment:roles:manage-security-problems</code>	Classic apps only: manage security problems	Management zone scope

Token	Purpose	Optional conditions
<code>environment:roles:view-sensitive-request-data</code>	View sensitive request data (classic; the <i>Read Sensitive Data</i> default policy pairs it with <code>storage:fieldsets:read</code>)	Management zone scope
<code>environment:roles:configure-request-capture-data</code>	Configure capture of sensitive data	—
<code>storage:security.events:read</code>	Read security events from Grail	Bucket, event-type, K8s namespace, cluster, host, cloud account
<code>storage:buckets:read</code>	Read records from Grail buckets — required in addition to <code>storage:security.events:read</code> ; scope it with <code>WHERE storage:table-name = "security.events"</code>	Table name, bucket name
<code>storage:fieldsets:read</code>	Read data from fieldsets — the Grail half of sensitive-data access (§ 5)	Table, bucket, fieldset name
<code>openpipeline:security.events:ingest</code>	Ingest security events into OpenPipeline	—
<code>vulnerability-service:vulnerabilities:read</code>	View vulnerabilities (programmatic)	—
<code>vulnerability-service:vulnerabilities:write</code>	Modify vulnerability information — in the app: mute status and ticket links of affected entities	—
<code>security-intelligence:enrichments:run</code>	Run enrichments (IP-address attribution, integration-app discovery) — relevant for Threats & Exploits / RAP investigation	—

Platform-collateral tokens AppSec apps load on top of the above

The AppSec product surfaces (Vulnerabilities app, Threats & Exploits app, Security Posture Management app) also need general Gen3 platform reads to load and operate — these are **not AppSec-specific**, they're what every Gen3 app needs:

```
storage:entities:read , storage:logs:read , storage:filter-segments:read ,
document:documents:read , settings:objects:read , hub:catalog:read , state:user-app-
states:read , state:user-app-states:write
```

These are usually already granted by the managed user policies covered in § 6 — call them out explicitly in audit conversations so reviewers know they're collateral, not AppSec-elevated access. `storage:buckets:read` is deliberately **not** in this list: every `fetch security.events` needs it, so the persona policies in § 3 grant it explicitly rather than assume it.

Critical corrections (verified)

1. **There is no `storage:security_problems:read` token.** Vulnerability read access in Latest Dynatrace is `vulnerability-service:vulnerabilities:read` ; `environment:roles:view-security-problems` only serves the classic apps. Common mistake when copying patterns from other Grail tables.
2. **The storage token is `security.events` with a dot, not `security_events` with an underscore.** The ingest token lives in the `openpipeline` service — `openpipeline:security.events:ingest` — not in `storage` .
3. **`environment-api:security-problems:read` is NOT in the IAM policy reference** as of 10/02/2026. If you see this token cited (some docs assistants surface it), it's not a current IAM permission token — use `vulnerability-service:vulnerabilities:read` for programmatic vuln access instead.

Sources: [IAM policy statements reference \(DT docs\)](#) — every token re-read 10/02/2026, including `storage:buckets:read` ("Grants permission to read records from Grail buckets. Required additionally to a table permission.") and `openpipeline:security.events:ingest` ("Grants permission to ingest security events into OpenPipeline").

3. Persona Policies (Worked Examples)

Three personas, three policies. Copy and adapt to your tenant's group structure. Each assumes the group also holds the Vulnerabilities app's documented default policies (*Read*

Entities, Read Security Events, one of *Admin/Pro/Standard User*) — the statements below are the AppSec-specific part.

Persona A — Security Analyst (read-only)

```
ALLOW storage:buckets:read WHERE storage:table-name = "security.events";
ALLOW storage:security.events:read;
ALLOW vulnerability-service:vulnerabilities:read;
```

This is the SOC's daily-driver policy. Read everything, change nothing. Does **not** include sensitive-data access — see § 5.

Persona B — Security Engineer (manage)

```
ALLOW storage:buckets:read WHERE storage:table-name = "security.events";
ALLOW storage:security.events:read;
ALLOW vulnerability-service:vulnerabilities:read;
ALLOW vulnerability-service:vulnerabilities:write;
```

`vulnerability-service:vulnerabilities:write` takes no conditions, so in Latest Dynatrace it **cannot be limited to production**: whoever holds it can change mute status and ticket links on any vulnerability in the environment. Separate production from non-production by using separate environments, or by keeping the group that holds write small. A `WHERE environment:management-zone = ...` condition on `environment:roles:manage-security-problems` only affects the deprecated classic apps — and management zones are not available in Latest Dynatrace.

Persona C — AppDev (namespace-scoped Grail read)

```
ALLOW storage:buckets:read WHERE storage:table-name = "security.events";
ALLOW storage:security.events:read
    WHERE storage:k8s.namespace.name IN ("team-payments-prod", "team-payments-
    staging");
```

This scopes **Grail reads of records that carry** `k8s.namespace.name` — compliance findings and vulnerability findings (APPSEC-06 § 2). It does not give the team its own view of RVA vulnerabilities:

- **RVA vulnerability events carry no** `k8s.namespace.name`, so the namespace condition cannot match them. To scope them, set `dt.security_context` on security events with OpenPipeline enrichment and condition on that instead — and test with a scoped user before rollout.

- **vulnerability-service:vulnerabilities:read** is **tenant-wide**. Adding it would show the team every vulnerability in the Vulnerabilities app. Leave it out unless tenant-wide visibility is acceptable.

Two syntax details: IAM lists use `IN ("...", "...")` with parentheses — DQL's `{...}` array syntax fails policy validation — and every persona needs `storage:buckets:read` alongside the table permission, or `fetch security.events` returns nothing.

Sources: [IAM policy statements reference \(DT docs\)](#) for the conditions and operators on each token (`storage:k8s.namespace.name` accepts `=`, `IN`, `startsWith`, `MATCH`; `vulnerability-service:vulnerabilities:write` — “Allows modifying vulnerability related information” — lists none); [IAM policy statement syntax \(DT docs\)](#) for the `IN ("...", "...")` list form; [Upgrade security notifications \(DT docs\)](#) — “Management zones are not available in Latest Dynatrace. Replace management zone scoping with Grail record-based field filters.”; [Vulnerability events \(DT semantic dictionary\)](#) — RVA vulnerability events name clusters and workloads in `related_entities.*` arrays, and the Kubernetes resource fields appear only on finding and scan events (all re-read 10/02/2026). **Derived:** a namespace condition cannot match records that lack the field — not tested with a scoped user on RVA data, which the validation tenant does not have.

4. Boundary Patterns for AppSec Data

Beyond the three personas above, in community practice three boundary patterns recur in AppSec IAM:

1. **Small write group** — Persona B above. Write cannot be environment-scoped inside one environment, so keep the group that holds it small, or split production into its own environment.
2. **Namespace-scoped read** — Persona C above. AppDev team sees only its own namespaces.
3. **Sensitive-payload carve-out** — sensitive-data access (`storage:fieldsets:read`, plus `view-sensitive-request-data` for the classic apps) granted separately to a small named group (SOC tier 2 / incident responders), not to the general read-only group.

For workflows that change vulnerabilities (ticket links, mute status), grant `vulnerability-service:vulnerabilities:write` to the dedicated service user the workflow runs as (§ 7), not to a shared OAuth client.

Sources: [IAM policy statements reference \(DT docs\)](#) for boundary syntax.

5. Privacy: view-sensitive-request-data

RAP attack events can include the actual request payload that triggered the detection — the SQL injection string, the JNDI lookup URL, the command-injection input. These payloads frequently contain PII or other sensitive material that came through the application boundary.

Dynatrace's *Read Sensitive Data* default policy grants two statements:

`storage:fieldsets:read` (Grail) and `environment:roles:view-sensitive-request-data` (classic). Track both — an audit that lists only the role misses the Grail half. In community practice in regulated environments, three operating rules are common:

1. **Don't grant it broadly.** Grant to incident responders + SOC tier 2, not to general read-only roles. The default-deny posture matters.
2. **Pair with** `configure-request-capture-data` for the small group responsible for tuning what gets captured.
3. **Audit access regularly.** Who has either permission, when did they last use it, against which events? Captured-payload access is exactly the surface auditors scrutinize.

This is the carve-out where AppSec IAM most often gets too generous — make it explicit who has it and why.

Sources: [Default policies \(DT docs\)](#) — *"Read Sensitive Data Grants unconditional permissions to read sensitive data in Dynatrace."*, whose policy is `ALLOW storage:fieldsets:read;` plus `ALLOW environment:roles:view-sensitive-request-data;` under `//Classic`; [IAM policy statements reference \(DT docs\)](#) — `environment:roles:configure-request-capture-data` : *"Grants user the Configure capture of sensitive data permission."* (re-read 10/02/2026). **Softened:** the three operating rules are community practice.

6. Reach for Managed Policies First

Following IAM-04 § 6 guidance, prefer Dynatrace's default policies over hand-rolled custom policies wherever they cover the use case. The ones that matter for AppSec are documented by name:

- **Read Security Events** — security events from the `events` and `security.events` tables and the default security-event buckets (includes `storage:buckets:read` for `security.events`)
- **Read Entities** — required by the Vulnerabilities app
- **Admin User / Pro User / Standard User** — one of these is required by the Vulnerabilities app
- **Read Sensitive Data** — `storage:fieldsets:read` + the classic sensitive-data role (§ 5)

If a default policy covers 80% of a persona's need, use it and add a small custom policy for the deltas, rather than rebuilding the whole grant set from scratch. This shrinks the audit surface and reduces drift across persona definitions.

Sources: [Default policies \(DT docs\)](#) — "Read Security Events Grants unconditional access to the security events from both, `events` and `security.events` tables and to the default security event buckets."; [Vulnerabilities \(DT docs\)](#) — "Read Entities Read Security Events One of the following user policies: Admin User , Pro User , Standard User" (re-read 10/02/2026).

7. OAuth Client vs Platform Token Scope Routing

Workflows and external automation authenticate differently:

- **Workflows** run as their *actor* — by default the user who created the workflow, or a non-interactive service user you select. No token is involved: the workflow gets that identity's policies. For security workflows, give **both** the creating user and the service user `storage:security.events:read` (plus `storage:buckets:read`).
- **External automation** (scripts, CI, Terraform) calling Grail or the vulnerability-service API: a platform token on a service user. AUTOM-04 § 3, *Provider Configuration*, covers the three-things-align model.
- **For classic config APIs** that AppSec workflows occasionally touch (rare in v1 AppSec, more common as workflows chain into broader automation): classic API Token may still be required for some endpoints.

In community practice, teams avoid granting `vulnerability-service:vulnerabilities:write` on a long-lived shared OAuth client — too much blast radius — and bind it to a dedicated Service User that holds only this permission. The policy-statement reference lists no conditions for `vulnerability-service:vulnerabilities:write` , so the permission itself cannot be narrowed — and management zones are not available in Latest Dynatrace to narrow it with. Scope *which* findings the workflow acts on in its trigger filter instead, using `dt.security_context` , a primary Grail field, or `primary_tags.*` .

Sources: [Workflow security \(DT docs\)](#) — "By default, the workflow actor is the user who created the workflow. However, there is the option to select a non-interactive service user as the actor of a workflow."; [Upgrade security notifications \(DT docs\)](#) — "Management zones are not available in Latest Dynatrace. Replace management zone scoping with Grail record-based field filters." and "Use custom metadata enrichment to set `dt.security_context` on security events via `OpenPipeline`, then filter by it in the workflow trigger." and "Assign this permission to both the human user creating the Workflow and the service user the Workflow runs as."; [IAM policy statements \(DT docs\)](#) — `vulnerability-service:vulnerabilities:write` is listed with no conditions.

8. Audit and Periodic Review

Audit IAM bindings for AppSec the same way you would audit any sensitive-data access. In community practice a quarterly minimum cadence is common — adapt it to your audit regime:

1. **Who has `storage:fieldsets:read` or `view-sensitive-request-data` ?** Pull the group memberships; remove anyone who hasn't used it in 90 days.
2. **Who has `vulnerability-service:vulnerabilities:write`** (and, while classic apps remain, `manage-security-problems`)? Write is a privileged, tenant-wide grant.
3. **Which service users act as workflow actors** with AppSec permissions? Confirm each one is still in active workflow service.
4. **Are namespace-scoped policies still aligned with team ownership?** Team boundaries shift; policies don't auto-update.

The query for "who can read `security.events` in production?" can be answered via the IAM API (see IAM-04 § 7 Policy Testing and Validation).

Sources: [IAM policy statements reference \(DT docs\)](#).

9. What's Still Evolving

Three areas where the AppSec IAM model is less settled and verification-at-policy-time is warranted:

1. **Boundary conditions on `vulnerability-service:vulnerabilities:read`** — at 10/02/2026 the policy reference shows no condition options for this token. If this changes,

narrower programmatic AppDev access becomes possible.

2. **Default-policy catalog** — coverage changes per release. Re-read the default-policies page before relying on one in a persona design.
3. **Cross-tenant AppSec access patterns** — multi-tenant SOC operations are still operator-discipline rather than first-class IAM features.

Softened: none of the three are official Dynatrace roadmap items — they're observations of where the model is thinner than the rest of Gen3 IAM as of 10/02/2026. Re-verify at policy-author time.

10. Next Steps

1. Adopt the three persona policies as starting points. Adapt the namespace names to your tenant, and decide how you will scope RVA visibility (`dt.security_context`).
2. List who has `storage:fieldsets:read` or `view-sensitive-request-data` today. If the list is longer than your incident-responder roster, narrow it.
3. Read **APPSEC-08** for workflow actors and ticket linkage.
4. Read **IAM-04** and **IAM-05** for the underlying DSL grammar and review patterns — this notebook only covers the AppSec-specific overlay.

References

Source	Coverage
IAM policy statements reference (DT docs)	Every permission token verified verbatim
Default policies (DT docs)	Read Security Events, Read Sensitive Data, user policies
Vulnerabilities (DT docs)	Vulnerabilities app permission requirements
Application Security (DT docs)	Three-pillar framing

 **DISCLAIMER:** This information was AI generated and is provided "as-is" without warranty. It was produced as an independent, community-driven project and **not supported by Dynatrace**. Always refer to

official [Dynatrace documentation](#) for the most current information.