

AUTOM-09: Terraform GitOps Setup

Recipe

Series: AUTOM — Dynatrace Automation | **Notebook:** 9 of 9 | **Created:** May 2026 |
Last Updated: 07/24/2026

A practical, opinionated recipe for standing up a Terraform GitOps shop for Dynatrace from scratch. This notebook covers what AUTOM-04 (Terraform resources) and AUTOM-07 (CI/CD integration) deliberately don't — repo layout, state backend choices, multi-environment promotion, lifecycle protections, secrets handling end-to-end, team onboarding, and operational realities. Use it as the bootstrap reference; consult AUTOM-04 for resource-level patterns and AUTOM-07 for CI/CD pipeline specifics.

Scope: how to *set up* a Terraform GitOps shop. Not "Terraform basics" (use HashiCorp Learn) and not "what Dynatrace resources exist" (use AUTOM-04). This is the operational architecture between those two layers.

Table of Contents

- [1. Introduction — How This Recipe Fits](#)
- [2. Opinionated Repo Layout](#)
- [3. State Backend Setup](#)
- [4. Provider Configuration and Version Discipline](#)
- [5. Module Strategy](#)
- [6. Multi-Environment Promotion](#)
- [7. CI/CD Pipeline \(cross-reference\)](#)
- [8. Secrets Handling End-to-End](#)
- [9. Lifecycle Protections](#)
- [10. Onboarding New App Teams](#)

- 11. [Cost and DPS Considerations](#)
- 12. [Operational Realities — State Lock, Break-Glass, DR](#)
- 13. [Next Steps](#)

Prerequisites

Requirement	Details
Audience	Platform engineering team standing up a Terraform GitOps shop for Dynatrace
Format	Opinionated recipe — concrete examples + explicit defaults; not a buffet of options
Assumes	Working knowledge of Terraform (HCL, modules, state, providers); familiarity with one Git host (GitHub / GitLab / Bitbucket Cloud)
Companion notebooks	AUTOM-04 for resource-level patterns; AUTOM-07 for CI/CD pipeline details; AUTOM-01 §5 for when-to-add-Monaco framing; FAQ-02 for tagging strategy
Out of scope	Terraform fundamentals (HashiCorp Learn covers this); Monaco workflows (AUTOM-03); Workflows-based automation (AUTOM-05)

1. Introduction — How This Recipe Fits

AUTOM-04 documents what Terraform resources exist for Dynatrace. AUTOM-07 documents how to wire those resources to a CI/CD pipeline. This notebook fills the operational gap between them: the repo layout, backend setup, lifecycle protections, secrets handling, and team-onboarding mechanics that determine whether your Terraform GitOps shop will be cleanly operable in 12 months — or a state-management swamp.

What this notebook is not

- **Not a Terraform tutorial.** Use the [HashiCorp Learn](#) tracks if you need provider basics.
- **Not a substitute for AUTOM-04 + AUTOM-07.** Those notebooks remain the canonical references for resource catalog and pipeline patterns; this one orchestrates the

surrounding operational architecture.

- **Not a Monaco-vs-Terraform debate.** That question is addressed in **AUTOM-01 §5: When a Terraform Shop Should Add Monaco**. This notebook assumes you've already committed to Terraform as the primary tool.

The opinionated stance

There are many "valid" Terraform layouts. This recipe picks one — separate-modules-and-environments — because in practice that's the layout most production Terraform shops converge to after a year of operating. Picking it on day one saves you the inevitable refactor.

Sources: [Terraform configuration language \(HashiCorp\)](#) — module and root configuration concepts.

Derived: the "separate-modules-and-environments" recommendation is community / engagement guidance grounded in the modules-as-reusable-units pattern documented by HashiCorp, applied to Dynatrace's per-environment tenant model.

2. Opinionated Repo Layout

Opinionated Repo Layout — Terraform GitOps for Dynatrace

Modules separated from environments; one root per environment; shared CI definitions at the top

dynatrace-terraform/

├─ .github/workflows/

CI/CD pipelines (or .gitlab-ci.yml, bitbucket-pipelines.yml)

├─ modules/

Reusable modules — versioned, env-agnostic

│ └─ management-zone/

│ └─ alerting-profile/

│ └─ slo/

│ └─ synthetic-http/

├─ envs/

One root per environment — references modules + per-env tfvars

│ └─ dev/

│ │ └─ main.tf

│ │ └─ backend.tf

│ │ └─ terraform.tfvars

│ └─ staging/

│ └─ production/

├─ policy/

Sentinel / OPA rules enforced in CI

├─ DECISIONS.md

Architectural decisions log

└─ README.md

Onboarding entry point + repo conventions

Why this layout

Layer	Purpose	What changes when
<code>modules/<resource-family>/</code>	Reusable building blocks; environment-agnostic; versioned independently	A pattern is added or refined (new field, new validation)
<code>envs/<env>/</code>	Per-environment composition; references modules with <code>source = "../../modules/<...>"</code> or by Git tag	An environment-specific value changes (new MZ name, threshold, location)
<code>policy/</code>	Pre-apply guardrails: Sentinel for HCP Terraform, OPA/Conftest for self-hosted CI	A new governance rule is added (mandatory tag, naming convention)

The line between `modules/` and `envs/` is the **environment-agnostic vs environment-specific** line. If a value differs between staging and production, it lives in `envs/<env>/terraform.tfvars`, not in the module.

What goes where — concrete examples

Concern	Module input or root variable?
Management Zone naming pattern (<code>{env}-{app}-{role}</code>)	Module input — the pattern itself is shared
The specific env / app / role values	Root tfvars per env
Severity thresholds	Module input with sensible defaults; override in root tfvars per env
Synthetic location IDs	Root tfvars per env — the IDs differ by tenant
Sentinel policies	Top-level <code>policy/</code> directory — not per-env
Provider version constraint	Top-level <code>versions.tf</code> shared by every root via symlink, or duplicated identically

Two-repo variant — modules separate from consumer (community-validated pattern)

A common variant promoted in [a May 2026 Dynatrace Guild thread](#) splits modules and consumer into **two repositories** rather than the single-repo single-tree above:

Repo	Role
<code>dynatrace-modules</code> (or per-resource-family: <code>dynatrace-dashboards</code> , <code>dynatrace-alerting</code> , <code>dynatrace-anomalies</code> , <code>dynatrace-global-settings</code>)	Reusable modules. Versioned with Git tags (<code>v1.4.0</code>); released as immutable artifacts.
<code>dynatrace-config-platform</code> (consumer)	Environment-specific configuration only. References modules via pinned Git tags.

The consumer repo organizes by ownership rather than by environment alone:

```
dynatrace-config-platform/
├─ org/
│  └─ _account_template/          # ← Shared root used by every environment
│     below
│     │ └─ backend.tf             # Remote state backend config (S3 +
│     │ DynamoDB locking)
│     │ └─ modules.tf             # Calls reusable modules with pinned Git-
│     │ tag versions
│     │ └─ variables.tf           # Input variable definitions
│     │ └─ locals.tf             # Derived values (env flags, naming
│     │ conventions)
│     └─ provider.tf             # Provider auth/region setup
│        └─ spoke/
│           └─ team-a/
│              └─ dev/
│                 └─ terraform.tfvars # Dev-only values
│                 └─ management_zones/ # JSON configs consumed by modules
│                    └─ alerting_profiles/
│                       └─ prod/
│                          └─ terraform.tfvars # Prod-only values
│                          └─ dashboards/
└─ .github/workflows/            # (or .azure-pipelines/, .buildkite/,
etc.)
```

The `_account_template/` **shared root** is the key element. Rather than duplicating `backend.tf`, `provider.tf`, `versions.tf` into each env directory (the single-repo pattern

above), each team / env directory **symlinks or templates from** `_account_template/` to inherit the shared Terraform boilerplate. Changes to the backend or provider config land in one place; every env picks them up.

Module enablement via `locals` **+ conditional** `for_each` / `count` is the per-env-customization knob inside the shared root. Pattern:

```
# _account_template/locals.tf
locals {
  env_flags = {
    dev = { enable_synthetic = false, enable_slos = false }
    prod = { enable_synthetic = true, enable_slos = true }
  }
  flags = local.env_flags[var.env]
}

# _account_template/modules.tf
module "synthetic_http" {
  source = "git::https://example.com/dynatrace-modules.git//synthetic-http?ref=v1.4.0"
  for_each = local.flags.enable_synthetic ? var.synthetic_monitors : {}
  # ...
}
```

The flag pattern lets staging skip expensive resources (synthetic monitors consume DPS — see §11) while production runs the full set, all from the same root configuration.

When to pick the two-repo variant vs the single-repo layout above

Single-repo (modules/ + envs/)	Two-repo (modules separate)
Small platform team (1-3 people); modules co-developed with envs	Multiple platform teams; modules consumed by many product repos
Module changes and env changes go in the same PR routinely	Module releases are deliberate events (Git tag → consumer bumps version)
Lower ceremony; faster iteration on module shape	Stronger versioning discipline; module authors and consumers have separate review chains

The single-repo layout is the right starting point. **Migrate to the two-repo variant when:** (a) more than ~3 product teams are consuming the modules, OR (b) module changes start blocking env changes (or vice versa) due to PR-queue contention, OR (c) audit /

compliance requires that the modules code-base is reviewed separately from environment configs.

When to deviate further

- **Monorepo with many products.** Sub-divide `envs/` further (e.g., `envs/<product>/<env>/`) and consider git submodules for modules consumed by multiple products.
- **Truly separate platform teams.** Put `modules/` in its own repo, version with Git tags or a private Terraform registry, consume from product-team repos by reference — the two-repo variant above is the canonical shape of this.

Sources:

- [Module composition \(HashiCorp\)](#) — the "root module + child modules" model.
- [Standard module structure \(HashiCorp\)](#) — the `main.tf / variables.tf / outputs.tf` convention used in `modules/<resource-family>/` .

3. State Backend Setup

State is the durable record of what Terraform manages. Remote backends matter for three reasons: concurrent-apply safety (state locking), audit (state history), and shared visibility (multiple people can see the same state). Local-file state is fine for solo experimentation; it is not fine for a team.

Choosing a backend

Backend	Choose when	Trade-off
HCP Terraform (<code>cloud {}</code>)	You want managed state + run-on-HCP + Sentinel policy enforcement out-of-box	Paid SaaS; ties you to HashiCorp; runs originate on HCP unless you configure agents
S3 + DynamoDB lock	You're an AWS shop already; want self-hosted; have ops capacity to maintain it	Two AWS resources to keep healthy; bucket + DynamoDB table both need encryption + versioning

Backend	Choose when	Trade-off
GCS (with built-in locking)	You're a GCP shop; want self-hosted	Built-in locking simpler than the S3+DynamoDB pair
Azure Storage (with built-in lease-based locking)	You're an Azure shop	Built-in locking via blob lease

The four examples below show one `backend.tf` per environment. Each environment has its own state file / workspace — the line between envs is the line between state files. Cross-env access happens via `terraform_remote_state` data sources (used sparingly).

HCP Terraform backend

```
# envs/production/backend.tf
terraform {
  cloud {
    organization = "my-org"
    workspaces {
      name = "dynatrace-production"
    }
  }
}
```

Workspace per environment; HCP manages state, locking, and run history. Sentinel policies attach to the workspace.

S3 + DynamoDB lock backend

```
# envs/production/backend.tf
terraform {
  backend "s3" {
    bucket      = "my-org-terraform-state-prod"
    key         = "dynatrace/terraform.tfstate"
    region      = "us-east-1"
    encrypt     = true
    dynamodb_table = "terraform-state-lock"
    kms_key_id   = "arn:aws:kms:us-east-1:123456789012:key/abcd-..."
  }
}
```

Bucket setup checklist:

- Versioning enabled (recovery from accidental state deletion)
- Server-side encryption with KMS CMK (not just SSE-S3)
- Public-access block on
- Lifecycle policy keeping at least 90 days of versions
- Bucket policy restricting access to the CI/CD service principal + the platform team

State-key naming — derive from repo + env path: for shops with multiple Dynatrace-Terraform repos sharing one state bucket, derive the S3 key from `<repo-name>/<env-path>/terraform.tfstate` rather than a flat `dynatrace/<env>.tfstate`. Examples: `dynatrace-config-platform/team-a/prod/terraform.tfstate`, `dynatrace-tenant-bootstrap/sandbox/terraform.tfstate`. This isolation lets multiple repos coexist in one bucket without key collisions; the bucket-policy + KMS-key boundary still applies at the prefix level. A simple `backend.tf` template that interpolates the key from repo + env:

```
terraform {
  backend "s3" {
    bucket = "my-org-terraform-state-prod"
    # key set per-env in backend config file or via -backend-config flag:
    # key = "<repo-name>/<env-path>/terraform.tfstate"
    region      = "us-east-1"
    encrypt     = true
    dynamodb_table = "terraform-state-lock"
    kms_key_id   = "arn:aws:kms:us-east-1:123456789012:key/abcd-..."
  }
}
```

Initialize per-env: `terraform init -backend-config="key=dynatrace-config-platform/team-a/prod/terraform.tfstate"`.

DynamoDB lock table:

- Primary key: `LockID` (string)
- On-demand billing (locks are short-lived; provisioned is wasteful)
- Point-in-time recovery enabled

GCS backend

```
# envs/production/backend.tf
terraform {
  backend "gcs" {
    bucket = "my-org-terraform-state-prod"
    prefix = "dynatrace/"
  }
}
```

Built-in locking via GCS object generation. Less infrastructure to maintain than S3+DynamoDB. Enable bucket versioning + CMEK for parity.

Azure Storage backend

```
# envs/production/backend.tf
terraform {
  backend "azurerm" {
    resource_group_name = "rg-terraform-state-prod"
    storage_account_name = "myorgtfstateprod"
    container_name       = "tfstate"
    key                  = "dynatrace.tfstate"
  }
}
```

Lease-based locking via Azure Blob. Enable soft-delete + versioning on the container; configure customer-managed keys if compliance requires.

State backup and DR

Regardless of backend, treat the state file like a tier-0 asset:

1. **Versioning enabled** on the storage layer (S3 versions / GCS object versions / Azure soft-delete). Not optional.
2. **Periodic snapshots** to a separate cold storage location for true DR (cross-region copy, separate account). Quarterly minimum; monthly if you make frequent destructive

changes.

3. **Documented restoration procedure.** Test it. A state file you can't restore from is a state file you don't have.
4. **Never edit state by hand in production.** Use `terraform state` subcommands; record every state-surgery operation in a runbook.

Sources:

- [Backend types \(HashiCorp\)](#) — backend configuration reference.
- [S3 backend \(HashiCorp\)](#), [GCS backend \(HashiCorp\)](#), [Azure backend \(HashiCorp\)](#) — per-backend config details.
- **Derived:** the bucket-setup checklist (versioning + CMK + public-access-block + lifecycle + bucket policy) is industry-standard cloud-storage hardening applied to state buckets specifically; not vendor-documented as a single checklist.

4. Provider Configuration and Version Discipline

Provider version drift bites Terraform shops harder than almost anything else. A simple `~>` constraint discipline avoids most of the pain.

Top-level `versions.tf` per environment

```
# envs/production/versions.tf
terraform {
  required_version = "~> 1.15"

  required_providers {
    dynatrace = {
      source = "dynatrace-oss/dynatrace"
      version = "~> 1.96"
    }
    aws = {
      source = "hashicorp/aws"
      version = "~> 6.0"
    }
  }
}
```

Constraint discipline

Constraint	Meaning	Use when
<code>version = "1.96.3"</code>	Exact pin	Never (loses bug fixes)
<code>version = "~> 1.96"</code>	Allows 1.96.x patch versions only (<code>>= 1.96</code> , <code>< 1.97</code>)	Production — get patches, not feature changes
<code>version = "~> 1.96, < 2.0"</code>	Allows $1.x \geq 1.96$	Most environments — get minor and patch versions, gate at major
<code>version = ">= 1.96"</code>	Any ≥ 1.96	Avoid — no upper bound, breaks unpredictably

The middle two are the practical defaults. Use the **tighter** form (`~> 1.96`) in production once you've validated a specific minor; use the **looser** form (`~> 1.96, < 2.0`) in dev/staging to catch minor-version changes early.

Combined auth in the provider config

```
# envs/production/main.tf
provider "dynatrace" {
  dt_env_url    = var.dt_env_url
  dt_api_token  = var.dt_api_token
  # Platform Token is read from DT_PLATFORM_TOKEN env var
  # HTTP_OAUTH_PREFERENCE=true should be set in the CI/CD env
}
```

See **AUTOM-07 §3.1 Terraform Workflow with Combined Auth** for the full pattern — Platform Token + API Token together give full Dynatrace resource coverage. Don't pick one; use both.

Upgrade cadence

Cadence	What to do
Every Dynatrace provider release	Skim the release notes; flag breaking changes; defer minor bumps until staging-validated

Cadence	What to do
Quarterly	Bump the production version constraint after staging has been on the newer version for 2–4 weeks without issue
On a known breaking change	Coordinate the bump across all envs in one PR; bring a rollback plan

The Dynatrace provider has been on a steady cadence — 13 minor versions between November 2025 and July 2026, v1.88 through v1.100. Staying current is straightforward if you treat it as a recurring chore.

Sources:

- [Provider requirements \(HashiCorp\)](#) — `required_providers` syntax.
- [Version constraints \(HashiCorp\)](#) — `~>`, `>=`, exact-pin semantics.
- [dynatrace-oss/terraform-provider-dynatrace releases \(GitHub\)](#) — release cadence reference (current: v1.100.0, released 07/02/2026).

5. Module Strategy

Modules are how you keep the codebase from sprawling as you add resources. The decisions: what to modularize, how to version, and where to put modules consumed by multiple products.

What to modularize

Modularize when:

- The same resource shape will be created multiple times with different inputs (e.g., 20 management zones, each following the same naming + tagging pattern → one MZ module instantiated 20 times)
- A resource has many supporting resources that always go together (e.g., a synthetic monitor + alerting profile + notification channel → one bundle module)
- A pattern encodes governance rules (e.g., a module that *a/ways* sets `dt.security_context` and a mandatory tag set)

Don't modularize when:

- A resource exists exactly once (most platform-wide configs — wrapping a single management-zone definition in a module just adds indirection)
- The "module" would have one input and one output (premature abstraction)

Versioning modules

Approach	Use when
Local-path references (<code>source = "../../modules/management-zone"</code>)	Same repo as the consumer; fastest iteration
Git tag references (<code>source = "git::https://gitlab.example.com/platform/terraform-modules.git//management-zone?ref=v1.4.0"</code>)	Modules in a separate repo, consumed by multiple product repos; clear version semantics
Terraform registry (private or public)	You publish to a registry (HCP Terraform private registry, Artifactory) and want UI browsing + version metadata

Pick one per repo; mixing leads to confusion about where modules live.

Semantic versioning for modules

- **Patch (1.4.0 → 1.4.1)** — bug fix, no input or output change
- **Minor (1.4.0 → 1.5.0)** — new optional input with a sensible default; new output; no breaking change
- **Major (1.x → 2.0.0)** — required input added, input removed, output removed, resource address changed

Treat module versions as a contract with consumers. Renaming a resource inside a module without a major version bump breaks every consumer's plan on next apply.

A representative module

```
# modules/management-zone/main.tf
variable "env"      { type = string }
variable "app"      { type = string }
variable "role"     { type = string }
variable "owner_tag" { type = string }

resource "dynatrace_management_zone_v2" "this" {
  name = "${var.env}-${var.app}-${var.role}"

  rules {
    type          = "ME"
    enabled       = true
    propagation_type = "HOST_TO_PROCESS_GROUP_INSTANCE"

    conditions {
      key {
        attribute = "HOST_GROUP_NAME"
      }
      string_conditions {
        operator      = "EQUALS"
        value         = "${var.env}-${var.app}"
        case_sensitive = false
      }
    }
  }
}

output "id" { value = dynatrace_management_zone_v2.this.id }
output "name" { value = dynatrace_management_zone_v2.this.name }
```


The module encodes the naming pattern (`{env}-{app}-{role}`) and the rule structure. Consumers can't accidentally name a management zone differently — the convention is enforced at the module boundary.

Sources:

- [Module development \(HashiCorp\)](#) — module-authoring guidance.
- [Module sources \(HashiCorp\)](#) — local paths, Git, registry references.
- [Semantic Versioning 2.0.0](#) — `MAJOR.MINOR.PATCH` semantics applied to modules.

6. Multi-Environment Promotion

The whole point of the modules + envs split is enabling promotion: the same module code runs against three different tenants with three different `tfvars` and three different state files.



The promotion workflow

1. **Change goes through `envs/dev/` first.** A PR adds or modifies a resource in `envs/dev/main.tf` (or bumps a module-source ref). CI runs `terraform plan` against the dev tenant; on PR approval and merge, `terraform apply` against dev.
2. **Validation in dev.** The team confirms the change behaves correctly in dev — alerts fire, dashboards render, synthetics pass.
3. **Promotion PR to `envs/staging/`.** Same module source ref + corresponding `tfvars` change. CI plans against staging. On merge → apply.
4. **Production promotion gated.** A PR to `envs/production/` requires a separate review (CODEOWNERS) and a manual approval step in the CI pipeline. The plan is reviewed by a human before apply.

What gets copied vs what diverges

Element	Same across envs?
Module source ref (<code>source = "../../modules/<name>"</code> or <code>?ref=v1.4.0</code>)	Yes — that's the whole point
Resource-shape inputs (which fields the module accepts)	Yes — schema is shared
Resource-value inputs (the values themselves)	No — these live in per-env <code>tfvars</code>
Backend block (where state lives)	No — each env has its own state
Provider config (tenant URL, credentials)	No — each env points at its own tenant
Sentinel / OPA policies	Yes — same governance rules everywhere

Catching env divergence early

It's common for `envs/production/main.tf` to drift from `envs/staging/main.tf` in subtle ways — someone added a resource to prod directly without backporting to staging. Two guardrails:

- **Periodic diff check.** A CI job that runs `diff envs/staging/main.tf envs/production/main.tf` (filtering out expected differences like backend) and flags divergence above a tolerance.
- **Required PR template** that asks "Did you make the same change in staging?" before allowing a prod PR.

When to use Terraform workspaces instead

Terraform workspaces (`terraform workspace new staging`) are an alternative to the `envs/<env>/` directory pattern: same files, different state per workspace. **Don't.**

Workspaces share `main.tf` between envs, which means production changes get applied to dev (and vice versa) the moment you forget to switch workspaces. Directory-per-env is mechanically harder to misuse — if you're editing `envs/production/main.tf`, you know what you're working on.

Sources:

- [Terraform workspaces \(HashiCorp\)](#) — workspace mechanics (which we recommend *against* in favor of directory-per-env).
- **Derived:** the directory-per-env-vs-workspaces recommendation is community / engagement guidance grounded in the failure mode of workspace switching across envs; HashiCorp does not take a position either way in its docs.

7. CI/CD Pipeline (cross-reference)

This recipe assumes a Terraform plan-and-apply pipeline. The full implementation details for GitHub Actions, GitLab CI, Bitbucket Pipelines, ArgoCD, and FluxCD are in **AUTOM-07**. Highlights relevant to the recipe:

AUTOM-07 section	What it covers
§3 GitHub Actions	Combined-auth Terraform workflow, Vault integration, drift detection, reusable workflows
§3.1 Terraform Workflow with Combined Auth	Platform Token + API Token together — the canonical pattern
§3.2 Vault Integration	Pulling short-lived credentials at pipeline start
§3.3 Policy-as-Code Gates	Sentinel / OPA / Conftest integration
§3.4 Drift Detection	Scheduled <code>terraform plan</code> runs that report diffs to a channel
§4 GitLab CI/CD	GitLab equivalent of the above
§5 Bitbucket Pipelines	Bitbucket Cloud Pipelines + FabianSchurig/bitbucket Terraform provider for the combined-workspace pattern
§9 Best Practices	Branch protection, required reviewers, plan-output retention
§10 Governance Architecture — Single SA Writer	Operational model for who can apply

Recipe-specific CI/CD recommendations

Recommendation	Why
Plan on every PR; apply only on merge to a long-lived env branch	Plan is read-only; apply is the privileged operation
Plan output retained as an artifact	Required for audit; required for "what was about to happen?" forensics
Plan summary posted to the PR	Reviewers see the diff without leaving the PR
Required reviewers per environment (CODEOWNERS)	Production requires platform-team review; dev can be self-merge
Apply uses <code>--auto-approve</code> only after gate passes	Manual approval gate (GH environments / GitLab environments / Bitbucket deployments) before production apply
Sentinel / OPA runs before apply	Policy gate fails the apply, not the plan — keep plan output unblocked for review

Sources: AUTOM-07: CI/CD Integration (this repo) covers each pipeline platform in detail.

8. Secrets Handling End-to-End

Three places secrets live in a Terraform GitOps shop: **in pipeline storage**, **in flight**, and **at rest in state**. Each needs explicit handling.

Pipeline-stored secrets

Where the Dynatrace tokens, cloud provider credentials, and Vault tokens come from.

Options:

Mechanism	Trade-off
GitHub Actions encrypted secrets	Easy; tied to GitHub; secrets exposed only to workflow runs

Mechanism	Trade-off
GitLab CI/CD variables	Easy; tied to GitLab; per-env masking + protection flags
Bitbucket repository / deployment variables	Easy; tied to Bitbucket; secured variable becomes write-only
Vault (or AWS Secrets Manager / Azure Key Vault)	Best — short-lived credentials, audit trail, decoupled from CI tool
OIDC federation to cloud → cloud secrets manager → Dynatrace token	Most secure but most plumbing; appropriate for high-stakes envs

The pragmatic baseline: pipeline-native secrets for the bootstrap; migrate to Vault once you have multiple envs and need rotation discipline.

In-flight secrets — Terraform variables

Mark every sensitive input:

```
variable "dt_api_token" {
  type      = string
  sensitive = true
}

variable "dt_platform_token" {
  type      = string
  sensitive = true
}
```

`sensitive = true` does three things:

1. **Plan output masks the value** — `terraform plan` shows `(sensitive value)` instead of the token
2. **Apply output masks the value** — same masking in the success log
3. **Outputs that pass sensitive values through inherit the marking** — propagation prevents accidental exposure downstream

It does **not** encrypt the value in state. State is in plaintext; the masking is for *log* output only.

At-rest secrets — state encryption

Terraform state contains every value the provider returned, including some sensitive ones (refresh tokens, generated secrets). Protect state at rest:

Backend	Mechanism
HCP Terraform	Encrypted by default; access via workspace permissions
S3	KMS CMK (not just SSE-S3) — KMS gives you key rotation + audit
GCS	CMEK (customer-managed encryption key)
Azure Storage	Customer-managed key in Key Vault

Additionally for self-hosted backends: bucket policy restricting access to the CI service principal + the platform team — and **nobody else**. State files exfiltrated from a misconfigured bucket give the attacker every secret Terraform managed.

Plan output discipline

`terraform plan -out=tfplan` writes the plan to a binary file that contains the proposed state — including pending secret values. Two rules:

1. **Don't retain the binary tfplan artifact** beyond the apply step. Encrypt it in transit between plan and apply jobs.
2. `terraform show -no-color tfplan > plan.txt` **for human review** — `-no-color` keeps it diffable, and the `sensitive` markings hide values in the rendered text.

Sources:

- [Sensitive input variables \(HashiCorp\)](#) — `sensitive` semantics.
- [Sensitive output values \(HashiCorp\)](#) — propagation through outputs.
- [Backend encryption \(HashiCorp\)](#) — state-at-rest concerns.

9. Lifecycle Protections

Lifecycle blocks are Terraform's safety net for production. Four meta-arguments matter; each has a specific use case.

Lifecycle Protections — Decision Guide

When to use `prevent_destroy`, `ignore_changes`, `create_before_destroy`, and `replace_triggered_by`

`prevent_destroy = true`

Refuses `terraform destroy` and any plan that would delete this resource

Use for:

- Production Management Zones
- Production IAM policies / group bindings
- Production Grail buckets (lose data if destroyed)
- ActiveGate token resources backing oncall flows

Caveat: cannot terraform-apply a deliberate rename — remove the lifecycle block first

`ignore_changes = [field, ...]`

Stops Terraform from "fixing" drift on the listed fields after the initial apply

Use for:

- Fields a human team owns (description, custom tags)
- Threshold values tuned via the UI by domain SMEs
- Auto-populated metadata (last_modified, version)
- Schema fields that drift via Dynatrace platform updates

Caveat: ignored fields drift silently — confirm you're OK with that before adding

`create_before_destroy = true`

Stand up the replacement before tearing down the original — no observability gap

Use for:

- Synthetic monitors whose recreation = data gap in SLO
- Alerting profiles tied to active oncall rotations
- Notification channels with active webhooks

Caveat: temporarily doubles the resource — check unique-name constraints.

`replace_triggered_by = [...]`

Force-recreate this resource when a referenced resource changes

Use for:

- Webhook URLs that need to refresh when secret rotates
- Alerting profiles tied to a renamed Management Zone
- Dashboards whose data sources reference a recreated entity

Caveat: chained recreations can have non-obvious blast radius.

`prevent_destroy` — for resources whose loss is catastrophic

```
resource "dynatrace_management_zone_v2" "platform_baseline" {
  name = "platform-baseline"
  # ...

  lifecycle {
    prevent_destroy = true
  }
}
```

Catches `terraform destroy` and any plan that would `-/+` the resource. The plan fails with an error and the apply doesn't proceed.

Use for: production Management Zones (lose IAM scope + alerting routing if destroyed), production IAM policies / group bindings (lose access), production Grail buckets (lose data), ActiveGate token resources backing on-call notifications.

Caveat: you can't terraform-apply a deliberate rename. To rename, remove the lifecycle block in PR 1, apply, then re-add it in PR 2.

`ignore_changes` — for fields owned by humans

```
resource "dynatrace_calculated_service_metric" "request_count" {
  name = "request-count"
  # ... most fields managed by Terraform ...

  lifecycle {
    ignore_changes = [
      description,    # SMEs edit this in the UI
      conditions[0].condition_key.attribute, # Auto-populated by Dynatrace
    ]
  }
}
```

Use for: description fields, custom tags, threshold values tuned via the UI by domain experts, auto-populated metadata (`last_modified`, `version`), schema fields that drift via Dynatrace platform updates.

Caveat: ignored fields drift silently. Confirm you're OK with that before adding — the field becomes invisible to drift detection.

`create_before_destroy` — for resources where recreation = outage

```
resource "dynatrace_http_monitor" "checkout_api" {
  name = "checkout-api-availability"
  # ...

  lifecycle {
    create_before_destroy = true
  }
}
```

Use for: synthetic monitors (recreation = data gap in SLO), alerting profiles tied to active on-call rotations, notification channels with active webhooks.

Caveat: temporarily doubles the resource. Check that unique-name constraints in Dynatrace allow the new resource to exist alongside the old — most do; some don't (e.g., some named-singleton settings).

`replace_triggered_by` — for chained recreations

```
resource "dynatrace_alerting_profile" "payments_alerts" {
  # ...

  lifecycle {
    replace_triggered_by = [
      dynatrace_management_zone_v2.payments.id,
    ]
  }
}
```

Use for: webhook URLs that need to refresh when a secret rotates, alerting profiles tied to a renamed Management Zone, dashboards whose data sources reference a recreated entity.

Caveat: chained recreations have non-obvious blast radius — adding `replace_triggered_by` to a hub resource can cascade to many dependents. Add deliberately.

Apply protections at the module boundary, not per-resource

If your module-management-zone always wants `prevent_destroy` for production, encode it in the module — don't depend on every consumer to remember.

```
# modules/management-zone/main.tf
variable "protect_from_destroy" {
  type    = bool
  default = false
}

resource "dynatrace_management_zone_v2" "this" {
  # ...

  lifecycle {
    prevent_destroy = var.protect_from_destroy
  }
}
```

Then in `envs/production/main.tf`:

```

module "platform_baseline" {
  source      = "../../modules/management-zone"
  env         = "prod"
  app         = "platform"
  role        = "baseline"
  protect_from_destroy = true # ← only prod sets this
}

```

Sources: [The lifecycle meta-argument \(HashiCorp\)](#) — all four meta-arguments documented.

10. Onboarding New App Teams

Once the platform team's Terraform GitOps shop is running, the next question is how product teams join. The architecture decision is the **brokered self-service** pattern from AUTOM-04 §7 — but the **onboarding mechanics** are what determine whether new teams ramp up in days or in weeks.

The onboarding workflow

1. **Team requests onboarding.** Lightweight intake form: team name, product, expected resource types (which modules will they use), expected scale (handful of MZs? hundreds of synthetics?), preferred Git host.
2. **Platform team provisions the team's namespace.**
 - A subdirectory under `envs/<env>/` per product (e.g., `envs/production/payments/`)
 - A Bitbucket / GitLab / GitHub team with write access to that subdirectory only (CODEOWNERS enforces)
 - An OPA / Sentinel policy that restricts that team's resources to a specific management zone scope (the "team-fencing" pattern)
3. **Team gets a template PR.**
 - `envs/<env>/<product>/main.tf` with a starter module instantiation (one MZ following the team's naming pattern)
 - `envs/<env>/<product>/terraform.tfvars` with placeholder values
 - A README pointing at the company's Terraform conventions
4. **Team's first PR validates the path.** Real change, real review, real apply against dev. Platform team co-reviews. After this PR merges, the team has a working pipeline + the muscle memory.

5. **Promotion drill.** Within the first 2 weeks, the team promotes the same change to staging and production with platform-team coaching. Now they own the operational rhythm.

Repo template

Maintain a `templates/team-onboarding/` directory containing:

```
templates/team-onboarding/  
├─ main.tf           # Minimal example with one module  
├─ terraform.tfvars  # Placeholder values  
├─ README.md         # Team-specific conventions  
└─ CODEOWNERS.snippet # CODEOWNERS lines to append
```

The first PR copies this directory into `envs/<env>/<product>/` and customizes. The template is itself versioned — when the platform team learns something new about onboarding, the template improves and the next team benefits.

Boundary enforcement

A new team should never be able to:

- Modify modules they don't own (CODEOWNERS on `modules/` requires platform-team approval)
- Modify another team's directory (CODEOWNERS on `envs/<env>/<other-product>/`)
- Bypass the Sentinel / OPA policies (CI cannot be merged-around)
- Get direct API access to Dynatrace (per AUTOM-04 §7 — the pipeline is the only writer)

These are enforced at three layers: Git host permissions, CODEOWNERS, and CI gates. All three need to be set for new teams during onboarding; one missing layer is the gap a future incident exploits.

Off-boarding teams

Less glamorous but equally important. When a team's product is decommissioned:

1. **Identify resources owned by the team** via the `owner` tag (see FAQ-02 for tagging strategy)
2. **PR removing the team's** `envs/<env>/<product>/` **subdirectory**
3. **Run** `terraform plan` **to surface what gets destroyed** — review carefully

4. **Apply destruction in dev first**, validate clean removal, then staging, then production
5. **Off-board the Git team and remove CODEOWNERS** entries in a follow-up PR

The `prevent_destroy` lifecycle on critical resources may need to be temporarily removed for deliberate off-boarding — do this in a dedicated PR, not bundled with destruction.

Sources:

- **AUTOM-04 §7 Brokered Self-Service** (this repo) — the architectural pattern this section operationalizes.
- **FAQ-02: Tagging — Sources, Standards, and Strategy** (this repo) — `owner` tag convention referenced in off-boarding.
- **Derived:** the onboarding-workflow steps (intake → namespace → template → first PR → promotion drill) and off-boarding sequence are community / engagement guidance — not a vendor-documented procedure but consistent practice across platform teams running Terraform GitOps shops.

11. Cost and DPS Considerations at Apply Time

Two cost surfaces that surprise teams new to Terraform GitOps for Dynatrace.

Dynatrace Platform Subscription (DPS) impact at apply time

Some Dynatrace resources consume DPS / DDU at creation or use:

Resource	DPS / DDU impact
<code>dynatrace_calculated_service_metric</code>	Generates a custom metric — counts against custom-metric DDU allocation
<code>dynatrace_http_monitor</code> (synthetic)	Each monitor execution consumes Synthetic actions
<code>dynatrace_browser_monitor</code>	Each monitor execution consumes Synthetic actions (higher rate than HTTP)
<code>dynatrace_application_anomalies</code> (custom thresholds)	No direct DPS cost; can cause more events → more retention DDU
Mass-creation via a module loop (<code>count = 100</code>)	Multiplies the above by 100

The pattern that bites: someone runs `terraform apply` against a module configured with `count = var.synthetic_count`, the variable defaults to a number that looked sensible in dev but is too large for production, and the apply provisions hundreds of synthetics — each consuming actions.

Guardrails:

- **Mandatory `count` / `for_each` review** for resources that create per-instance Dynatrace billable units. CODEOWNERS or a Sentinel policy flagging these.
- **Pre-apply cost check** for environments above a size threshold — script that counts billable resources in the plan and posts a comment on the PR.
- **Quota alerts in Dynatrace** at the tenant level so unexpected scale surfaces before the bill does.

Terraform plan / apply runtime cost

The Dynatrace API has rate limits. Plans against a tenant with hundreds of resources can take 5-20 minutes — apply, longer. Costs:

- **CI minutes** — runner time × number of plans per day. Drift-detection runs every 15 minutes against a 500-resource state file are not free.
- **Throttling failures** — large applies that hit API rate limits return errors mid-apply, leaving state partially updated. Recovery requires manual `terraform refresh` or partial apply.

Guardrails:

- `terraform apply -parallelism=10` instead of the default 10 — actually 10 is the default; consider lowering to 5 for very large tenants to avoid throttling.
- **Split large states.** If a single state file is over ~500 resources, consider splitting by product or by resource family. The boundary is the line of "this group rarely changes when that group changes."
- **Drift detection cadence.** Hourly is plenty for most environments; every-15-minutes is rarely justified by the value it adds.

Sources:

~ **Derived:** the DPS-impact list combines documented Dynatrace billing categories with how each maps to Terraform resource types; not a single vendor-published checklist. Verify specific resource

cost behavior against current [Dynatrace pricing documentation](#) at use time.

- Terraform CLI command — apply (HashiCorp) — `-parallelism` flag.

12. Operational Realities — State Lock, Break-Glass, DR

Three operational scenarios every Terraform GitOps shop hits eventually.

Stuck state lock

Symptom: `terraform plan` or `apply` hangs with *"Error acquiring the state lock"*.

Causes:

- Previous CI run killed mid-apply; lock not released
- Network partition between the CI runner and the lock store
- Lock-table outage (DynamoDB)
- Someone ran `terraform apply` from their laptop and Ctrl-C'd

Recovery:

1. **Confirm the lock is genuinely abandoned.** Check the lock-store record — the lock includes `who`, `operation`, `created` timestamps. If `created` is recent (< 10 min) and someone owns it, wait or message them.
2. `terraform force-unlock <LOCK_ID>` if the lock is abandoned. This removes the lock without doing anything to state.
3. **For HCP Terraform**, use the workspace UI's "Force cancel" — same idea, different UX.

Document the procedure in your team's runbook. Force-unlock is safe **when the lock is genuinely abandoned**; it is dangerous if another apply is mid-flight, because two applies running simultaneously corrupt state. Verify, don't guess.

Break-glass UI changes

Sometimes someone has to change Dynatrace config in production at 3am via the UI — production is broken and waiting for a PR + review + apply takes too long. This is fine **as long as the change gets back into Terraform afterward**.

Recovery workflow:

1. **Document the change** in the incident channel at the time it's made (resource type, name, what changed).
2. **Within 24 hours**, open a PR that brings the change into Terraform. Two approaches:
 - **Edit the existing resource** in `envs/production/main.tf` to match what the UI now has.
 - **Run `terraform import`** if it's a brand-new resource created via UI.
3. **Run `terraform plan`** — it should now show **no changes** (Terraform sees the same state the UI created). If the plan shows changes, the resource has drifted; reconcile.
4. **Apply**. This is a no-op apply that ratifies the UI change as Terraform-managed.

The danger pattern: break-glass change happens, nobody backports it, two weeks later a routine apply *reverts the emergency fix*. Audit periodically that no production resources have drifted from the Terraform definition.

Disaster recovery — corrupted or lost state

Scenarios:

- State bucket accidentally deleted
- `terraform state rm` removed a resource and the next apply destroyed it
- A `terraform refresh` corrupted state due to a provider bug

Recovery sequence:

1. **Stop all CI pipelines** targeting the affected env. Prevent concurrent applies during recovery.
2. **Restore from state versioning**. S3 / GCS / Azure all have versioning if enabled (per §3). Restore the most recent known-good version. For HCP Terraform, restore from the workspace's state versions.
3. `terraform plan` against the restored state. The plan tells you what Terraform thinks is missing or extra compared to actual cloud / Dynatrace state.
4. **Reconcile manually**. If state shows a resource missing that actually exists in Dynatrace, `terraform import` it. If state shows a resource present that's been removed in Dynatrace, `terraform state rm` it.
5. **Resume CI**. With reconciled state, the next normal apply should be a no-op or only the intended pending changes.

If state was **never versioned** (a setup mistake — fix immediately), recovery is much harder: you essentially have to import every resource by hand. This is why §3's bucket-

setup checklist insists on versioning.

Practice the DR procedure

Run a state-loss DR drill once a quarter against a non-production environment:

1. Snapshot the dev state file
2. Simulate loss (rename the state file in storage)
3. Walk through the recovery procedure end-to-end
4. Time the recovery; identify the slowest step; document a fix

The drill is not optional — discovering that your "documented" DR procedure has a missing step during an actual incident is the worst time to find out.

Sources:

- [Force-unlock state \(HashiCorp\)](#) — `terraform force-unlock` semantics and safety caveats.
- [Terraform import \(HashiCorp\)](#) — bringing UI-created resources into Terraform state.
- **Derived:** the break-glass and DR workflows are community / engagement guidance — HashiCorp documents the CLI commands but not the operational discipline around using them; the discipline is what separates a working Terraform GitOps shop from a broken one.

13. Next Steps

You now have the recipe to stand up a Terraform GitOps shop for Dynatrace. Practical sequence to follow:

1. **Pick a state backend** (§3) — usually whichever cloud your team already runs in
2. **Set up the repo layout** (§2) with `modules/`, `envs/`, `policy/`
3. **Provision the dev environment first** — backend, providers, a single management-zone module, one root config
4. **Wire up CI/CD for dev** following AUTOM-07 §3 (or §4/§5 for GitLab/Bitbucket)
5. **Validate the round-trip** — PR, plan, merge, apply, observe in Dynatrace
6. **Add staging**, repeat steps 3-5
7. **Add production** with the manual approval gate
8. **Implement Sentinel / OPA policies** in `policy/` and wire them into CI

9. **Add lifecycle protections** (§9) on critical production resources
10. **Run a state-loss DR drill** (§12) before claiming the setup is production-ready
11. **Onboard the first product team** following §10 — use the experience to refine the template
12. **Document everything** in the repo README and your team's runbooks

Cross-references

- **AUTOM-01 §5** — When a Terraform Shop Should Add Monaco (positioning question this recipe assumes you've already answered)
- **AUTOM-04** — Terraform Provider resource catalog and patterns (the *what*; this notebook is the *how*)
- **AUTOM-07** — CI/CD Integration deep-dive (pipelines for each Git host)
- **FAQ-02** — Tagging — Sources, Standards, and Strategy (the tag conventions referenced in §10 onboarding and §11 cost)

Summary

Standing up a Terraform GitOps shop for Dynatrace is more than wiring a CI pipeline to a Dynatrace provider. The operational architecture between those two layers — repo layout, state backend, multi-env promotion, lifecycle protections, secrets handling, team onboarding, operational realities — determines whether the shop is operable in a year. This recipe codifies the opinionated defaults that production Terraform shops converge to in practice, applied to Dynatrace's per-tenant model and the Dynatrace Terraform provider's combined-auth requirements.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official [Dynatrace documentation](#).