

AUTOM-96 LAB: GitHub Actions CI/CD for Dynatrace Terraform

Series: AUTOM — Dynatrace Automation | **Reference:** 96 — GitHub Actions CI/CD LAB |
Created: May 2026 | **Last Updated:** 10/02/2026

Overview

Hands-on lab that takes the GitHub Actions material from **AUTOM-07: CI/CD Integration** §3 and the credential-handling pattern from **AUTOM-04: Terraform Provider** §3 and walks through them end-to-end. By the end of this lab you have a working pipeline that:

- Runs `terraform fmt -check`, `terraform validate`, and `terraform plan` on every pull request, posting the plan summary back as a PR comment.
- Runs `terraform apply` on merge to `main`, gated by a GitHub Environment with required reviewers (on a private repo that gate needs GitHub Enterprise — see Prerequisites).
- Uses one GitHub OIDC token for **two parallel trusts**: an **AWS IAM role** for the S3 state backend, and a **HashiCorp Vault `jwt` role** for the runtime fetch of a Dynatrace Platform Token. **No static Dynatrace token lives in GitHub Secrets.**
- Stores Terraform state in S3 with native S3 lockfile locking, encrypted at rest.
- Manages a single Settings 2.0 resource (an ownership team, `builtin:ownership.teams`) end-to-end so the credential plumbing is the focus, not the resource model.

This LAB is the hands-on companion to AUTOM-07 §3 (which describes the patterns) and AUTOM-04 §3 (which describes the token model). It does **not** repeat material from those notebooks — it walks through the steps.

Out of scope for this LAB (cross-references provided):

Topic	See
Other CI platforms (GitLab, Bitbucket, Bamboo, Azure DevOps)	AUTOM-07 §4–§7
State backend hardening (KMS, IAM scopes, multi-account)	AUTOM-09 §3
Multi-environment promotion patterns	AUTOM-09 §6
Lifecycle protections (<code>prevent_destroy</code> , <code>replace_triggered_by</code>)	AUTOM-09 §9
Token-model deep dive (Platform vs classic API)	AUTOM-04 §3
<code>dynatrace_api_token</code> / state-leakage concerns	AUTOM-04 §3

Table of Contents

1. [Prerequisites](#)
 2. [Architecture](#)
 3. [Step 1 — Bootstrap the Terraform repo layout](#)
 4. [Step 2 — Mint the Platform Token \(out-of-band\)](#)
 5. [Step 3 — Wire OIDC trust](#)
 6. [Step 4 — Author the workflow file](#)
 7. [Step 5 — Run a plan-only test on a PR](#)
 8. [Step 6 — Promote to apply on merge](#)
 9. [Step 7 — Validate](#)
 10. [Common Pitfalls](#)
 11. [Where to go next](#)
 12. [Summary Checklist](#)
 13. [Sources and References](#)
-

1. Prerequisites


Requirement	Details
Completed	AUTOM-04 (Terraform Provider) and AUTOM-07 (CI/CD Integration) §3 read at least once
Dynatrace Environment	Gen3 SaaS tenant; account-admin access to mint a Platform Token on a service user
GitHub Account	Repository with admin access (to add Environments, configure OIDC permissions, set repo secrets)
GitHub plan	The apply gate uses an environment with required reviewers . On GitHub Free, Pro and Team those are available only for public repositories — use a public repo for this LAB (nothing secret lives in it) or GitHub Enterprise. See Step 4.
AWS Account	Account where you can create an IAM OIDC identity provider, an IAM role, an S3 bucket, and (optionally) a KMS key. A Free Tier account works.
HashiCorp Vault	An instance reachable from GitHub Actions, with a KV v2 engine at <code>secret/</code> : a self-hosted Vault behind a public endpoint or a runner-reachable VPN endpoint, or HCP Vault Dedicated (check its current tiers and trial). On HCP Vault Dedicated the top-level namespace is <code>admin</code> , which the commands below must target — see Step 3.

Requirement	Details
Local tools	terraform >= 1.11 (S3 native locking is GA from 1.11), aws CLI, vault CLI, gh CLI (GitHub CLI), git
Time budget	Roughly 1.5–2 hours end-to-end, the first time you run it.

Tenant URL placeholder. Throughout this LAB, replace `https://<your-tenant>.apps.dynatrace.com` (UI host) and `https://<your-tenant>.live.dynatrace.com` (API host) with your actual tenant identifier.

2. Architecture

The pipeline has three zones: **GitHub** (source + runner + state), a **trust boundary** crossed by short-lived federated identity (no static secrets), and **Dynatrace** (the Platform Token mint side and the eventual API target).



Why this shape

- **Mint out-of-band, deliver in-band** (the pattern from AUTOM-04 §3): the long-lived Dynatrace credential is generated once by an admin, deposited in Vault, and never touched by Terraform or by GitHub Secrets.
- **OIDC federation** removes the chicken-and-egg of "what authenticates the credential fetch?" — GitHub's OIDC identity authenticates to AWS IAM and to Vault directly, so there is no static AWS access key or static Vault token in GitHub Secrets either.

- **Single** `id-token: write` permission on the workflow is the only GitHub-side configuration that matters; the rest is set up once on the AWS and Vault sides.

3. Step 1 — Bootstrap the Terraform repo layout

The LAB uses the simplest layout that exercises the credential plumbing — one directory, three `.tf` files, one Settings 2.0 resource. For multi-environment / module-based layouts see **AUTOM-09 §2**.

Create the directory and files

```
mkdir -p dynatrace-terraform/.github/workflows
cd dynatrace-terraform
git init
touch versions.tf main.tf backend.tf .gitignore
```

`.gitignore` (mandatory before first commit)

```
# Terraform
.terraform/
*.tfstate
*.tfstate.*
*.tfplan
crash.log

# Local credentials
*.tfvars
!example.tfvars
.env
```

Why this matters. State files and `.tfvars` files routinely contain sensitive values — even with `sensitive = true` they are still plain-text on disk (see AUTOM-04 §3 *Operational Safety*). Committing them to Git is the most common credential-leak path in real Terraform shops.

Do commit `.terraform.lock.hcl`. It records the exact provider build `terraform init` selected, so the PR plan and the merge apply run the same provider. HashiCorp: *"You should include this file in your version control repository"* ([Dependency lock file \(Terraform docs\)](#)). Run `terraform init` once locally and commit the file it creates.

versions.tf — pin everything

```
terraform {
  required_version = ">= 1.11.0" # use_lockfile (backend.tf) is GA from 1.11

  required_providers {
    dynatrace = {
      source = "dynatrace-oss/dynatrace"
      version = "~> 1.93"
    }
  }
}
```

backend.tf — S3 state, native S3 lockfile

```
terraform {
  backend "s3" {
    bucket = "my-org-terraform-state"
    key    = "dynatrace/lab/terraform.tfstate"
    region = "us-east-1"

    # Native S3 lockfile – added in Terraform 1.10, GA in 1.11; replaces the
    # DynamoDB lock table, which HashiCorp has deprecated.
    # See https://developer.hashicorp.com/terraform/language/backend/s3 for the
    # current locking model.
    use_lockfile = true

    encrypt = true
  }
}
```

Bucket setup is out of scope for this LAB. Create the bucket once via console / CloudFormation / a separate Terraform configuration; enable bucket versioning; consider an SSE-KMS key. See **AUTOM-09 §3** for the hardening checklist.

main.tf — one Settings 2.0 resource

We pick an **ownership team** (`builtin:ownership.teams`) because it (a) is a Settings 2.0 schema, so it works with a Platform Token alone, (b) is a **multi-object** schema (up to 1,000 teams), so creating one never overwrites an existing object, (c) is not deprecated and is not on Dynatrace's removed-schemas list, and (d) is harmless to create and delete in a non-prod tenant.

```

provider "dynatrace" {
  # Picked up from environment variables at apply time:
  #   DYNATRACE_ENV_URL
  #   DYNATRACE_PLATFORM_TOKEN
}

resource "dynatrace_generic_setting" "lab_team" {
  schema = "builtin:ownership.teams"
  scope  = "environment"

  value = jsonencode({
    name          = "LAB Team"
    identifier     = "lab-team"
    description    = "Created by the AUTOM CI/CD LAB - safe to delete"
    responsibilities = {
      development      = false
      infrastructure    = false
      lineOfBusiness    = false
      operations        = true
      security          = false
    }
    supplementaryIdentifiers = []
    contactDetails           = []
    links                    = []
    additionalInformation    = []
  })
}

```

`schema`, **not** `schema_id`. `dynatrace_generic_setting` takes `schema`, `scope` and a JSON value; `schema_id` fails `terraform validate`. The value above passes the Settings API's own validation for `builtin:ownership.teams` (checked with a `validate-only` request, 09/18/2026). Earlier versions of this LAB used `builtin:kubernetes.generic.metadata.enrichment`, which is a deprecated **singleton** (`maxObjects: 1`) — creating it would have overwritten a tenant's existing object — and whose rule fields were misnamed.

No `dt_platform_token` argument in the provider block. The Dynatrace provider reads `DYNATRACE_ENV_URL` and `DYNATRACE_PLATFORM_TOKEN` from the environment when the corresponding HCL arguments are absent. This is what lets the `OIDC -> Vault -> env-vars` chain work without HCL-level coupling.

4. Step 2 — Mint the Platform Token (out-of-band)

This step happens **once**, in the Dynatrace UI, by a human admin. It is deliberately not part of the pipeline. The full rationale (mint-out-of-band-vs-in-pipeline trade-offs) is in **AUTOM-04 §3 Operational Safety — State File Leakage** — read that subsection if you have not already.

UI steps

1. Sign in to `https://<your-tenant>.apps.dynatrace.com` as an account administrator.
2. Open the **Account Management** console (the user-icon menu in the top right) -> **Identity & access management** -> **Service users**.
3. Create or select the service user that will own this token (e.g. `svc-terraform-lab`).
4. Confirm the service user is in a group with the IAM permissions needed for Settings 2.0 management — at minimum the equivalent of **"Change monitoring settings"** (`environment:roles:manage-settings`) on the tenant.
5. Switch to **Platform tokens** (`myaccount.dynatrace.com` -> Platform tokens) and click **Generate new token**.
6. **On behalf of:** select the `svc-terraform-lab` service user.
7. **Scopes:** select `settings:objects:read` and `settings:objects:write`. (Add `settings:schemas:read` if your tenant requires it for schema discovery.)
8. Set an **expiration** appropriate to your rotation policy (the LAB recommendation is 90 days — short enough to force you through the rotation flow at least once during the LAB lifetime).
9. Copy the token value (`dt0s16...`). It is shown **once**.

Deposit the token into Vault

```
# Authenticate to Vault as a human admin (any auth method that works for you)
vault login -method=oidc

# On HCP Vault Dedicated: export VAULT_NAMESPACE=admin (see Step 3)
# If no KV v2 engine is mounted at secret/ yet:
#   vault secrets enable -path=secret kv-v2

# Write the token to the path the workflow will read from
vault kv put secret/dynatrace/lab \
  platform_token="dt0s16.XXXX..." \
  env_url="https://&lt;your-tenant&gt;.live.dynatrace.com"
```

Three things must align for the Platform Token to actually work — service user IAM permissions, the creator's `iam:service-users:use` permission, and the scopes selected at token creation. AUTOM-04 §3 *Service User Credentials for Terraform* covers the model in full. If `terraform apply` later fails with a 403, this is the first place to look.

5. Step 3 — Wire OIDC trust

Two trust relationships need to exist before the workflow can run. Both rest on the same GitHub OIDC token, and they are **independent** of each other:

1. **GitHub -> AWS** — the workflow assumes an AWS IAM role to reach S3 (state) without a static AWS access key. AWS is used for nothing else.
2. **GitHub -> Vault** — Vault's `jwt` auth method trusts GitHub's OIDC issuer directly, so the workflow pulls the Dynatrace Platform Token without a static Vault token. No AWS credential is involved.

The [GitHub OIDC docs \(GitHub docs\)](#) name both kinds of target: "access a cloud provider (such as AWS, Azure, GCP, HashiCorp Vault, and others)".

The `sub` claim — check its format first

Both trusts match on the token's `sub` claim, and its format depends on the repository and the job:

Case	<code>sub</code> looks like
Repo created on GitHub.com after July 15, 2026 (or opted in) — includes any repo you create for this LAB	<code>repo:<org>@<org-id>/<repo>@<repo-id>:ref:refs/heads/main</code>
Older repo	<code>repo:<org>/<repo>:ref:refs/heads/main</code>
A job that references an environment (the <code>apply</code> job uses <code>production</code>)	<code>...:environment:production</code> — the ref is not included

Get the two IDs with `gh api repos/<your-org>/<your-repo> --jq '.owner.id, .id'`. If in doubt, run [github/actions-oidc-debugger \(GitHub\)](#) once and copy the `sub` it prints.

Sources: [OIDC reference \(GitHub docs\)](#) — "repositories created after July 15, 2026 now use an immutable default subject format that includes both the owner ID and repository ID"; "The subject claim includes the branch name of the workflow, but only if the job doesn't reference an environment".

Preview — skipping the stored Dynatrace token entirely. Dynatrace now documents [Workload identity federation \(DT docs\)](#) as a **Preview** feature: "Workload identity federation (WIF) lets an external workload call the Dynatrace API with a token issued by its own identity provider." The page names a GitHub Actions job as a fitting workload, so the GitHub OIDC token could authenticate to Dynatrace directly, mapped to a service user — no Platform Token in Vault at all. Verify that WIF is available in your account before relying on it; until then, the Vault-held Platform Token in this LAB remains the working path.

5.1 — Create the GitHub OIDC identity provider in AWS

This is a one-time, account-wide setup. If your AWS account already has the GitHub provider registered, skip this.

```
# Use the AWS account where the S3 state bucket lives
aws iam create-open-id-connect-provider \
  --url https://token.actions.githubusercontent.com \
  --client-id-list sts.amazonaws.com
```

No thumbprint needed. `--thumbprint-list` is optional: *"If it is not included, IAM will retrieve and use the top intermediate certificate authority (CA) thumbprint of the OpenID Connect identity provider server certificate."* ([CreateOpenIDConnectProvider \(AWS docs\)](#))

5.2 — Create the AWS IAM role the workflow assumes

Create `trust-policy.json`:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Federated": "arn:aws:iam::<your-aws-account-id>;oidc-provider/token.actions.githubusercontent.com"
      },
      "Action": "sts:AssumeRoleWithWebIdentity",
      "Condition": {
        "StringEquals": {
          "token.actions.githubusercontent.com:aud": "sts.amazonaws.com"
        },
        "StringLike": {
          "token.actions.githubusercontent.com:sub": "repo:<your-org>:@<org-id>/<your-repo>@<repo-id>:*"
        }
      }
    }
  ]
}
```

The `sub` claim is the **load-bearing scoping mechanism**. The policy uses the immutable format of a repo created after July 15, 2026 (see the table above); for an older repo use `repo:<your-org>/<your-repo>:*`. The `:*` suffix admits any branch, pull request and environment in the repo — right for a shared plan/apply role like this LAB's. To tighten a separate **apply** role, match the environment, not the branch: `repo:<your-org>@<org-id>/<your-repo>@<repo-`

`id>:environment:production .A :ref:refs/heads/main` condition never matches the apply job, because a job that references an environment presents `:environment:...` instead of the ref. A mismatched `sub` is a frequent cause of a failed role assumption — see [Common Pitfalls](#).

Create the role and attach a least-privilege policy that allows S3 state access only:

```
aws iam create-role \
  --role-name github-actions-dynatrace-terraform \
  --assume-role-policy-document file://trust-policy.json

aws iam put-role-policy \
  --role-name github-actions-dynatrace-terraform \
  --policy-name terraform-state-access \
  --policy-document '{
    "Version": "2012-10-17",
    "Statement": [
      {
        "Effect": "Allow",
        "Action": ["s3:GetObject", "s3:PutObject", "s3:DeleteObject"],
        "Resource": "arn:aws:s3:::my-org-terraform-state/dynatrace/lab/*"
      },
      {
        "Effect": "Allow",
        "Action": ["s3:ListBucket"],
        "Resource": "arn:aws:s3:::my-org-terraform-state"
      }
    ]
  }'
```

5.3 — Configure Vault to trust GitHub OIDC

Vault accepts the GitHub OIDC JWT directly through its `jwt` auth method — AWS plays no part in this trust. You only need this once per Vault instance.

On HCP Vault Dedicated, the top-level namespace is `admin`, not `root`. Run `export VAULT_NAMESPACE=admin` before these commands (and before the `vault kv put` in Step 2), and uncomment `namespace: admin` on both `vault-action` steps in Step 4. Without it, the commands and the workflow reach the wrong namespace. ([What is HCP Vault Dedicated](#) (HashiCorp), [Vault CLI](#) `VAULT_NAMESPACE` (Vault docs))

```
# Enable the jwt auth method (idempotent; skip if already enabled)
vault auth enable jwt

# Configure it to trust GitHub's OIDC issuer
vault write auth/jwt/config \
  oidc_discovery_url="https://token.actions.githubusercontent.com" \
  bound_issuer="https://token.actions.githubusercontent.com"

# Create a policy that grants read on the secret we wrote in Step 2
vault policy write dynatrace-lab-read - &&&'EOF'
path "secret/data/dynatrace/lab" {
  capabilities = ["read"]
}
EOF

# Create a role that is assumable by the workflow
vault write auth/jwt/role/dynatrace-terraform-lab \
  role_type="jwt" \
  user_claim="actor" \
  bound_audiences="https://github.com/&lt;your-org&gt;" \
  bound_claims_type="glob" \
  bound_claims='{ "sub": "repo:&lt;your-org&gt;:@&lt;org-id&gt;/&lt;your-repo&gt;:@&lt;repo-id&gt;:*"}' \
  policies="dynatrace-lab-read" \
  ttl="15m"
```

`bound_claims` uses the same `sub` format as the AWS trust policy — for a repo created before July 15, 2026, use `repo:<your-org>/<your-repo>:*` .

`bound_audiences` must match the audience the workflow requests. If `jwtGithubAudience` is omitted, the token carries GitHub's default audience — *"the URL of the repository owner"*, i.e. `https://github.com/<your-org>` (OIDC reference (GitHub docs)). The workflow below sets it explicitly so the value visibly matches `bound_audiences` .

See the [Vault JWT auth method \(Vault docs\)](#) for the full claim-binding reference.

6. Step 4 — Author the workflow file

Two jobs in one file: `plan` runs on every pull request; `apply` runs only on push to `main` and is gated by a GitHub Environment. Both jobs share a composite step that performs the OIDC -> Vault -> Dynatrace-token fetch, so the credential plumbing exists once.

Create `.github/workflows/terraform.yml` :

```

name: Terraform – Dynatrace Lab

on:
  pull_request:
    branches: [main]
    paths:
      - "**/*.tf"
      - ".github/workflows/terraform.yml"
  push:
    branches: [main]
    paths:
      - "**/*.tf"
      - ".github/workflows/terraform.yml"

# Required for OIDC. Without id-token:write the workflow has no OIDC token to present.
permissions:
  contents: read
  id-token: write
  pull-requests: write

env:
  AWS_REGION: us-east-1
  TF_IN_AUTOMATION: "true"
  TF_INPUT: "false"

# Action majors checked 10/2026 – all run on Node 24 (GitHub removed Node 20 from the
# runners on 09/23/2026). Terraform pinned to an exact release.
jobs:
  plan:
    name: terraform plan
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v7

      - name: Setup Terraform
        uses: hashicorp/setup-terraform@v4
        with:
          terraform_version: "1.16.5"

      - name: Configure AWS credentials via OIDC
        uses: aws-actions/configure-aws-credentials@v6
        with:
          role-to-assume: arn:aws:iam::<your-aws-account-id>;role/github-actions-
dynatrace-terraform
          aws-region: ${ env.AWS_REGION }

      - name: Fetch Dynatrace Platform Token from Vault
        id: vault
        uses: hashicorp/vault-action@v4
        with:
          url: https://vault.example.com:8200
          # namespace: admin # HCP Vault Dedicated only
          method: jwt
          jwtGithubAudience: https://github.com/<your-org>;

```

```

    role: dynatrace-terraform-lab
    secrets: |
      secret/data/dynatrace/lab platform_token | DYNATRACE_PLATFORM_TOKEN ;
      secret/data/dynatrace/lab env_url         | DYNATRACE_ENV_URL

- name: terraform fmt -check
  run: terraform fmt -check -recursive

- name: terraform init
  run: terraform init -input=false

- name: terraform validate
  run: terraform validate -no-color

- name: terraform plan
  id: plan
  run: |
    terraform plan -no-color -out=tfplan
    terraform show -no-color tfplan > plan.txt
  env:
    DYNATRACE_PLATFORM_TOKEN: ${ env.DYNATRACE_PLATFORM_TOKEN }
    DYNATRACE_ENV_URL: ${ env.DYNATRACE_ENV_URL }

- name: Post plan to PR
  if: github.event_name == 'pull_request'
  uses: actions/github-script@v9
  with:
    script: |
      const fs = require('fs');
      const plan = fs.readFileSync('plan.txt', 'utf8').slice(0, 60000);
      const body = [
        '### Terraform Plan',
        '',
        '<details><summary>Click to expand</summary>',
        '',
        ``hcl`,
        plan,
        ```,
 '',
 '</details>',
].join('\n');
 await github.rest.issues.createComment({
 owner: context.repo.owner,
 repo: context.repo.repo,
 issue_number: context.issue.number,
 body,
 });

apply:
 name: terraform apply
 needs: plan
 if: github.ref == 'refs/heads/main' && github.event_name == 'push'
 runs-on: ubuntu-latest
 environment: production # gated by GitHub Environment Protection Rules

```

```

steps:
 - uses: actions/checkout@v7

 - name: Setup Terraform
 uses: hashicorp/setup-terraform@v4
 with:
 terraform_version: "1.16.5"

 - name: Configure AWS credentials via OIDC
 uses: aws-actions/configure-aws-credentials@v6
 with:
 role-to-assume: arn:aws:iam::<your-aws-account-id>;role/github-actions-
dynatrace-terraform
 aws-region: ${ env.AWS_REGION }

 - name: Fetch Dynatrace Platform Token from Vault
 uses: hashicorp/vault-action@v4
 with:
 url: https://vault.example.com:8200
 # namespace: admin # HCP Vault Dedicated only
 method: jwt
 jwtGithubAudience: https://github.com/<your-org>
 role: dynatrace-terraform-lab
 secrets: |
 secret/data/dynatrace/lab platform_token | DYNATRACE_PLATFORM_TOKEN ;
 secret/data/dynatrace/lab env_url | DYNATRACE_ENV_URL

 - name: terraform init
 run: terraform init -input=false

 - name: terraform apply
 run: terraform apply -auto-approve -input=false

```

## What the workflow does, in order

1. **Checkout the repo** — standard.
2. **Setup Terraform** — `hashicorp/setup-terraform@v4` pins to a known version. Pin a real version, not `latest`, so plan output is reproducible across reruns.
3. **AWS credentials via OIDC** — the workflow's OIDC token is exchanged for short-lived AWS credentials by assuming the IAM role created in Step 3.2. **No static AWS access key in GitHub Secrets.**
4. **Vault token fetch** — the same OIDC token (re-issued for the Vault audience) is exchanged for a short-lived Vault token bound to `dynatrace-terraform-lab` role. **No static Vault token in GitHub Secrets.** `vault-action` automatically masks the secrets in subsequent log lines.
5. **terraform fmt -check** — fail fast on style.
6. **terraform init** — downloads the Dynatrace provider; AWS credentials from step 3 authorize state read/write.
7. **terraform validate** — schema/HCL validation.

8. **terraform plan** — generates the plan; environment variables for the Dynatrace provider come from the Vault step.
9. **PR-comment step** — only runs on `pull_request` events; posts the plan back as a comment using `actions/github-script`.
10. **Apply job** — only runs on push-to-main, only after `plan` succeeds, and only after the GitHub Environment's required reviewers approve — if your plan supports required reviewers on this repo (below).

`workflow_dispatch` is intentionally omitted. Adding manual-trigger support is a one-line addition (`on.workflow_dispatch:` block) but it widens the surface — anyone with `workflow-write` permission can apply on demand. Add it explicitly if your team needs it; defaulting it off is safer.

## Configure the GitHub Environment

In the GitHub UI: **Settings -> Environments -> New environment -> production**. Then:

- **Required reviewers:** add at least one team member who is not the author. This is the apply gate.

**Check your GitHub plan before relying on the gate.** *"If you are on a GitHub Free, GitHub Pro, or GitHub Team plan, other deployment protection rules, such as a wait timer or required reviewers, are only available for public repositories."* On GitHub Free, a private repository gets no environments at all. Without the gate, every merge to `main` runs `terraform apply -auto-approve` with no approval and no warning. Use a public repository for this LAB (the token lives in Vault, not the repo) or GitHub Enterprise. ([Managing environments \(GitHub docs\)](#))

- **Wait timer:** optional 5–15 minute delay before the gate can be approved (useful for last-call review).
- **Deployment branches:** restrict to `main` only.
- **Environment secrets:** none needed for this LAB — secrets come from Vault at runtime.

---

## 7. Step 5 — Run a plan-only test on a PR

---

```
From your local clone of the repo
git checkout -b lab-add-ownership-team
git add .
git commit -m "feat: add Settings 2.0 ownership team"
git push -u origin lab-add-ownership-team

Open the PR
gh pr create --base main --title "LAB: add ownership team" --body "Lab change."
```

## What you should see, in order

1. The `plan` job starts immediately on PR open. Expect ~60–90 seconds total wall-clock.
2. The first run downloads the Dynatrace provider — adds about 20 seconds you won't pay on subsequent runs.
3. The Vault step prints `Reading secret from secret/data/dynatrace/lab` and the masked secret values are not echoed. If you see the actual token value in the log, **stop**: the masking is broken or you typed it into the workflow YAML by mistake.
4. The `terraform plan` step prints something like:

```
Plan: 1 to add, 0 to change, 0 to destroy.
```

5. Within 5–10 seconds of the run completing, a new comment appears on the PR titled **Terraform Plan**, containing the collapsed plan output.

## What you should NOT see

- A comment with raw token values. If the plan or apply ever leaks a token into logs, treat it as a credential incident: rotate the Platform Token in DT, rotate any downstream tokens, audit recent activity.
- A failed apply on a PR. Apply should be skipped entirely on PRs (the `if:` condition gates it). If apply runs on a PR, the `if:` is wrong.

---

## 8. Step 6 — Promote to apply on merge

```
After PR review and approval
gh pr merge --squash --delete-branch
```

## What you should see, in order

1. A new workflow run starts on `main`, triggered by the merge commit.
2. The `plan` job runs again (and should produce the same plan as the PR-time run — see Common Pitfalls if it doesn't).
3. The `apply` job appears as **Waiting** with a *Review deployments* button — this is the GitHub Environment gate.
4. An approver clicks *Review deployments*, optionally adds a comment, and approves.
5. The `apply` job runs through the same checkout / Terraform / OIDC / Vault steps, then runs `terraform apply -auto-approve`. Output ends with:

```
Apply complete! Resources: 1 added, 0 changed, 0 destroyed.
```

## Manual approval is the load-bearing safety net

Because `apply` is `-auto-approve`, the GitHub Environment gate is the **only** human checkpoint — and it exists only where your plan provides required reviewers for this repository (Step 4). Without it, a merge applies unattended. Treat the required-reviewers list with the same care as production deploy access for any other system. Two-person rule (author cannot self-approve) is the standard configuration in **Settings -> Environments -> production -> Prevent self-review**.

---

## 9. Step 7 — Validate

---

Three places to check, in order:

### 9.1 — Dynatrace tenant

1. Sign in to `https://<your-tenant>.apps.dynatrace.com`.
2. Open the Settings app and search for **Teams** (schema `builtin:ownership.teams`) — or list the objects with `GET /api/v2/settings/objects?schemaIds=builtin:ownership.teams`.
3. Confirm a team exists with name `LAB Team` and identifier `lab-team`.
4. Check the object's modification info — it should show the service user (e.g. `svc-terraform-lab`), not a human.

### 9.2 — GitHub

1. **Actions** tab -> the merge-commit run -> confirm both jobs are green.
2. **Settings -> Environments -> production -> Deployment history** — confirm one deployment, with the approver's name.
3. The closed PR shows the *Terraform Plan* comment from the plan run.

### 9.3 — Vault audit log

Vault starts with auditing **disabled**, and `vault audit list` only lists the audit devices that are enabled — it shows no requests. To check the reads, enable a file audit device **before** the merge run (on a Vault where you can manage audit devices), then search the log on the Vault server:

```
Before the run
vault audit enable file file_path=/var/log/vault_audit.log

After the run, on the Vault server
grep 'secret/data/dynatrace/lab' /var/log/vault_audit.log
Expect two reads from the merge run (plan + apply), each authenticated
via auth/jwt with role dynatrace-terraform-lab.
```

Sources: [Audit logging \(Vault docs\)](#) — "When you initialize a new Vault cluster, auditing is disabled."; [audit list \(Vault docs\)](#) — "List information about active audit devices."

If all three places show what you expect, the pipeline is wired correctly end-to-end.

## 10. Common Pitfalls

#	Symptom	Cause	Fix
1	Error: Could not satisfy plugin requirements on terraform init after a state-backend change	Stale .terraform/ cache	<code>rm -rf .terraform &amp;&amp; terraform init</code>
2	Error: state lock could not be acquired	A previous run died holding the lock; with <code>use_lockfile = true</code> the lock object is <code>key.tflock</code> in the same prefix	Verify no run is actually in progress, then <code>terraform force-unlock &lt;LOCK_ID&gt;</code> (the ID is in the error message)
3	Not authorized to perform <code>sts:AssumeRoleWithWebIdentity</code> at the <i>Configure AWS credentials</i> step ( <code>terraform init</code> never runs)	The trust policy's <code>sub</code> (or <code>aud</code> ) condition does not match the token — commonly the old <code>repo:</code> <code>&lt;org&gt;/&lt;repo&gt;:*</code> form on a repo that issues the immutable <code>repo:</code> <code>&lt;org&gt;@&lt;id&gt;/&lt;repo&gt;@&lt;id&gt;:...</code> format, or a <code>:ref:refs/heads/main</code> condition on the apply job, which presents <code>:environment:production</code>	Print the token's claims once with <code>github/actions-oidc-debugger</code> and match the condition to the <code>sub</code> it shows (Step 3 table)
4	permission denied: cannot retrieve credential from Vault	Workflow lacks <code>id-token: write</code> permission, the Vault role's <code>bound_audiences</code> does not match <code>jwtGithubAudience</code> , or its <code>bound_claims</code> <code>sub</code> glob uses the wrong <code>sub</code> format (Step 3 table)	Add <code>id-token: write</code> to the workflow's <code>permissions:</code> block; align the audience strings exactly — Vault is whitespace and case sensitive on this field

#	Symptom	Cause	Fix
5	Plan posts the Dynatrace token in cleartext into the PR comment	<code>terraform plan</code> rendered an attribute that contains the token (e.g. an output, or a debug <code>local</code> )	Mark any <code>output</code> referencing the token as <code>sensitive = true</code> and, more importantly, <b>never echo <code>TF_VAR_*</code> or <code>DYNATRACE_PLATFORM_TOKEN</code> from a <code>run: step</code> —</b> <code>vault-action</code> 's masking only catches values it injected, not values you re-stringify
6	<code>Error: 403 Forbidden</code> from the Dynatrace API on apply	Three-things-align failure: service user IAM permissions, creator's <code>iam:service-users:use</code> at mint time, or scope selection	Re-read AUTOM-04 §3 <i>Service User Credentials</i> ; the most common variant is the service user lacking the <b>Change monitoring settings</b> permission
7	Plan on PR succeeds; same plan on merge-to-main fails	Drift between PR-time and merge-time tenant state, or the apply role can read state but the plan role cannot (or vice versa)	If you split the IAM roles per environment, both must permit S3 state access on the same prefix; for this LAB they share a role, so look for tenant drift first
8	<code>AccessDenied</code> on <code>s3:GetObject</code> / <code>s3:PutObject</code> during <code>terraform init</code>	The role was assumed, but its <b>permission</b> policy does not cover the bucket, the <code>dynatrace/lab/</code> prefix or the <code>.tflock</code> lock object	Fix the resource ARNs in the role's <code>terraform-state-access</code> policy (Step 3.2)

## 11. Where to go next

Topic	Notebook
Other CI platforms (GitLab, Bitbucket, Bamboo, Azure DevOps) following the same OIDC + Vault shape	<b>AUTOM-07</b> §4–§7
State backend hardening — KMS keys, IAM scopes per environment, multi-account state, audit logging	<b>AUTOM-09</b> §3
Multi-environment promotion — <code>envs/dev/</code> , <code>envs/staging/</code> , <code>envs/prod/</code> directories vs Terraform workspaces	<b>AUTOM-09</b> §6

Topic	Notebook
Lifecycle protections — <code>prevent_destroy</code> , <code>ignore_changes</code> , <code>replace_triggered_by</code> for critical resources	<b>AUTOM-09</b> §9
Token-model deep dive — Platform Token vs classic API Token vs OAuth Client; when each applies	<b>AUTOM-04</b> §3
<code>dynatrace_api_token</code> / state-leakage operational safety	<b>AUTOM-04</b> §3 <i>Operational Safety</i>
Drift detection (scheduled <code>terraform plan -detailed-exitcode</code> workflow)	<b>AUTOM-07</b> §3 <i>Drift Detection</i>
Reusable workflows for multi-team / multi-tenant orgs	<b>AUTOM-07</b> §3 <i>Reusable Workflows</i>
Policy-as-code gates (OPA/Conftest, Sentinel)	<b>AUTOM-07</b> §3 <i>Policy-as-Code Gates</i>

## 12. Summary Checklist

Tick off each item to confirm a working LAB:

- ☐ Repo created with `versions.tf` , `main.tf` , `backend.tf` , `.gitignore` ; `.terraform.lock.hcl` committed
- ☐ S3 state bucket created, versioning enabled, encryption enabled
- ☐ Platform Token minted in DT UI on a service user; deposited in Vault at `secret/dynatrace/lab`
- ☐ AWS OIDC identity provider for `token.actions.githubusercontent.com` exists
- ☐ AWS IAM role `github-actions-dynatrace-terraform` with trust policy scoped to your repo, in the `sub` format your repo issues
- ☐ Vault `jwt` auth method enabled and bound to GitHub's OIDC issuer
- ☐ Vault role `dynatrace-terraform-lab` with policy granting read on `secret/data/dynatrace/lab`
- ☐ `.github/workflows/terraform.yml` committed; permissions block contains `id-token: write`
- ☐ PR run shows green `plan` job; PR comment with plan summary posted
- ☐ No raw token values appear in any log line
- ☐ `production` GitHub Environment created; required reviewers configured (public repo or GitHub Enterprise); deployment branches restricted to `main`
- ☐ Merge-to-main run shows `apply` job waiting on approval; approval triggers apply
- ☐ Validation: DT tenant shows the new ownership team; modified-by = service user
- ☐ Validation: GitHub Environments page shows the deployment with approver's name

- [ ] Validation: Vault audit log (device enabled before the run) shows two reads from the workflow's OIDC identity
- 

## 13. Sources and References

---

### Dynatrace

- [Platform Tokens \(DT docs\)](#)
- [Terraform for Dynatrace \(DT docs\)](#)
- [Dynatrace Terraform Provider \(Terraform Registry\)](#)
- [Dynatrace Terraform Provider \(Dynatrace GitHub\)](#)

### GitHub Actions

- [About security hardening with OpenID Connect \(GitHub docs\)](#)
- [OpenID Connect reference — subject claims and immutable format \(GitHub docs\)](#)
- [Deprecation of Node 20 on GitHub Actions runners \(GitHub changelog\)](#) — *"remove Node20 on September 23rd, 2026"*
- [Configuring OpenID Connect in Amazon Web Services \(GitHub docs\)](#)
- [Using secrets in GitHub Actions \(GitHub docs\)](#)
- [Using environments for deployment \(GitHub docs\)](#)

### Actions used in the workflow

- [hashicorp/setup-terraform \(HashiCorp GitHub\)](#)
- [hashicorp/vault-action \(HashiCorp GitHub\)](#)
- [aws-actions/configure-aws-credentials \(AWS GitHub\)](#)
- [actions/github-script \(GitHub\)](#)
- [github/actions-oidc-debugger \(GitHub\)](#)

### Terraform / Vault

- [S3 backend \(Terraform docs\)](#) — current S3-lockfile locking model
- [Sensitive data in state \(Terraform docs\)](#)
- [JWT auth method \(Vault docs\)](#)
- [Dependency lock file \(Terraform docs\)](#)
- [Terraform v1.11.0 release notes \(HashiCorp GitHub\)](#) — *"S3 native state locking is now generally available."*

- [What is HCP Vault Dedicated \(HashiCorp\)](#)
- [Audit logging \(Vault docs\)](#)

## AWS

- [CreateOpenIDConnectProvider \(AWS docs\)](#)
- 
- 

*This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.*