

CLOUD-04: AWS Lambda & Serverless Monitoring

Series: CLOUD — Cloud Provider Integrations | **Notebook:** 4 of 8 | **Created:** March 2026 | **Last Updated:** 08/12/2026

Overview

This notebook covers serverless monitoring with Dynatrace, focusing on AWS Lambda. You will learn how to monitor Lambda function performance (cold starts, duration, errors, throttles), integrate API Gateway tracing, analyze Step Functions workflows, assess DynamoDB performance, and build end-to-end serverless application tracing.

Sprint 1.337 (April 2026): Service Detection v2 for Lambda

Sprint 1.337 SaaS introduced **Service Detection v2 (SDv2) for AWS Lambda** in Early Access — a major upgrade for serverless monitoring:

- Unified rules for OTel and OneAgent.** SDv2 detects Lambda functions whether they emit traces via OneAgent (Lambda layer) or via the OpenTelemetry Lambda extension. The same service entity is enriched from both paths — no more parallel services for the same function.
- Three FaaS-specific metrics:**
 - **Invocation/failure counts** — `dt.faas.invocations` and `dt.faas.failures`
 - **Duration** — `dt.faas.duration` (cold-start vs warm split)
 - **Trigger type breakdown** — `dt.faas.trigger.type` dimension (HTTP, S3, EventBridge, SQS, etc.)
- OTel `service.name` enrichment** — when Lambda functions emit OTel spans with `service.name`, Dynatrace uses that value to enrich the SDv2-detected service rather than creating a parallel one.

Sample DQL — Lambda failure rate by trigger type:

```
// dt.faas.failures / dt.faas.invocations / dt.faas.trigger.type do not exist
// (corrected
// 08/12/2026). The OneAgent FaaS series is dt.service.faas_invoke.count,
split by faas.trigger
// with a boolean `failed` dimension – so the failure/total split is a
dimension filter, not two
// separate metrics.
timeseries invocations = sum(dt.service.faas_invoke.count), from:-1h, by:
{faas.trigger, failed}
| fieldsAdd total = arraySum(invocations)
| summarize {
    failures = sum(if(failed == true, total, else: 0)),
    invocations = sum(total)
}, by:{faas.trigger}
| fieldsAdd failure_rate_pct = round(failures * 100.0 / invocations, decimals:
2)
| sort failure_rate_pct desc
```

Status: Early Access — confirm tenant availability before production rollout. Existing Lambda monitoring via OneAgent layer or OTel extension keeps working.

Table of Contents

1. [Lambda Monitoring Fundamentals](#)
 2. [Cold Start Detection](#)
 3. [Lambda Error Analysis](#)
 4. [Throttle Analysis](#)
 5. [API Gateway Integration](#)
 6. [Step Functions Tracing](#)
 7. [DynamoDB Performance](#)
 8. [End-to-End Serverless Tracing](#)
 9. [Summary and Next Steps](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Grail enabled
Permissions	<code>metrics.read</code> , <code>entities.read</code> , <code>spans.read</code>
AWS Integration	AWS cloud integration configured (CLOUD-02)
Lambda Functions	At least one Lambda function with Dynatrace layer or OneAgent extension
Prior Knowledge	CLOUD-01 fundamentals, CLOUD-02 AWS integration

1. Lambda Monitoring Fundamentals

Dynatrace monitors Lambda functions through two complementary approaches:

Approach	Data Source	What It Provides
Cloud Integration	CloudWatch metrics via ActiveGate	Invocations, duration, errors, throttles, concurrent executions
Dynatrace Lambda Layer	OneAgent in Lambda runtime	Distributed traces, code-level visibility, custom metrics

Key Lambda Metrics

Metric	Description	Healthy Range
<code>dt.cloud.aws.lambda.invocations</code>	Total invocations	Application-dependent
<code>dt.cloud.aws.lambda.duration</code>	Execution time (ms)	< function timeout
<code>dt.cloud.aws.lambda.errors</code>	Execution errors	0 (ideally)
<code>dt.cloud.aws.lambda.throttlers</code>	Throttled invocations	0

Metric	Description	Healthy Range
<code>dt.cloud.aws.lambda.conc_executions</code>	Concurrent runs	< reserved concurrency

List Monitored Lambda Functions

```
// List all monitored Lambda functions with runtime and code size
//
// Corrected 08/12/2026: the runtime attribute is awsRuntime, not
awsLambdaFunctionRuntime – the
// latter does not exist and failed the whole cell with FIELD_DOES_NOT_EXIST.
Time range added
// because dt.entity.* returns only entities seen in the query window, not the
standing inventory.
fetch dt.entity.aws_lambda_function, from:-7d
| fieldsKeep id, entity.name, awsRuntime, awsCodeSize, awsMemorySize,
awsTimeout, tags
| sort entity.name asc
| limit 25

// Smartscape note (dt.entity.* is deprecated but still functional): AWS
Lambda functions ARE
// modeled on Smartscape as smartscapeNodes "AWS_LAMBDA_FUNCTION", but the
classic aws* attribute
// fields differ there – inspect smartscapeNodes "AWS_LAMBDA_FUNCTION" | limit
1 for the node
// field names. Keep the classic query above until the fields are mapped.
```

Lambda Execution Time Over Time

```
// Lambda average execution duration over the last 6 hours
//
// Metric keys corrected 08/12/2026. The AWS CloudWatch Lambda metrics are
named
// cloud.aws.lambda.<CloudWatchName>.By.FunctionName - Invocations /
Errors / Duration /
// Throttles / ConcurrentExecutions / IteratorAge. The `dt.cloud.aws.lambda.*`
names used before
// do not exist in any spelling, and a timeseries against a missing key
returns an EMPTY result
// instead of an error, so every one of these tiles silently drew nothing. The
split dimension is
// FunctionName - dt.entity.aws_lambda_function is null on these metrics.
Enumerate with:
//  metrics | filter startsWith(metric.key, "cloud.aws.lambda") | fields
metric.key | sort metric.key asc
timeseries avgDuration = avg(cloud.aws.lambda.Duration.By.FunctionName),
from:-6h, by:{FunctionName}
| fieldsAdd avgDurationValue = arrayAvg(avgDuration)
| sort avgDurationValue desc
| limit 10
```

2. Cold Start Detection

Cold starts occur when Lambda creates a new execution environment. They add latency (100ms to several seconds depending on runtime and package size).

Understanding Cold Starts

Factor	Impact on Cold Start
Runtime	Java/C# > Python/Node.js
Package size	Larger = slower initialization
VPC attachment	Adds ENI creation time (~1-2s)
Provisioned concurrency	Eliminates cold starts (at cost)
SnapStart	Reduces Java cold starts to ~200ms

Detecting Cold Starts via Init Duration

CloudWatch provides an `InitDuration` metric for cold starts. With Dynatrace Lambda Layer, you can detect cold starts from spans.

```
// Detect Lambda cold starts from spans (init phase)
//
// Corrected 08/12/2026. The old cell filtered `contains(span.name,
"lambda")`, but a Lambda span
// is named for the FUNCTION, not the word "lambda" – so the filter discarded
every matching span
// and the cell reported cold_start_count = 0 while 14 real cold starts (avg
2.45 s) sat in the
// window. A zero inside a summarize row is the easiest kind of wrong answer
to miss: the cell
// returns a row, it just returns a false one. faas.coldstart alone is the
correct discriminator.
// Also dropped faas.invocation_id, which is not a field (faas.* has no
invocation_id).
fetch spans, from:-6h
| filter faas.coldstart == true
| summarize {cold_start_count = count(), avg_duration_ms = avg(duration) /
1ms}, by:{faas.name}
| sort cold_start_count desc
| limit 20
```

Lambda Invocations with Duration Percentiles

```
// Lambda duration percentiles over the last 6 hours – computed from spans,
// not the metric.
//
// Corrected 08/12/2026 (two defects). The old cell called
// percentile(dt.cloud.aws.lambda.duration, 50): the key does not exist, AND
// percentile() does not
// work on the real CloudWatch Duration metric either –
cloud.aws.lambda.Duration.By.FunctionName
// is a pre-aggregated gauge, so percentile() over it returns an empty result
// even when the key is
// spelled correctly. Only avg/min/max/sum are meaningful there. Real
// percentiles need the
// underlying distribution, which lives on the spans.
fetch spans, from:-6h
| filter isNotNull(faas.name)
| summarize {
    p50 = percentile(duration, 50) / 1ms,
    p90 = percentile(duration, 90) / 1ms,
    p99 = percentile(duration, 99) / 1ms,
    invocations = count()
}, by:{faas.name}
| sort p99 desc
| limit 10
```

3. Lambda Error Analysis

Lambda errors fall into two categories:

Error Type	Description	Metric
Function errors	Unhandled exceptions, timeout, OOM	dt.cloud.aws.lambda.errors
Invocation errors	Permission issues, throttling, invalid payload	dt.cloud.aws.lambda.throttlers

Error Rate by Function

```
// Lambda error count by function over the last 24 hours
//
// Metric keys corrected 08/12/2026. The AWS CloudWatch Lambda metrics are
// named
// cloud.aws.lambda.<CloudWatchName>.By.FunctionName – Invocations /
// Errors / Duration /
// Throttles / ConcurrentExecutions / IteratorAge. The `dt.cloud.aws.lambda.*`
// names used before
// do not exist in any spelling, and a timeseries against a missing key
// returns an EMPTY result
// instead of an error, so every one of these tiles silently drew nothing. The
// split dimension is
// FunctionName – dt.entity.aws_lambda_function is null on these metrics.
Enumerate with:
// metrics | filter startsWith(metric.key, "cloud.aws.lambda") | fields
// metric.key | sort metric.key asc
timeseries errors = sum(cloud.aws.lambda.Errors.By.FunctionName), from:-24h,
by:{FunctionName}
| fieldsAdd totalErrors = arraySum(errors)
| filter totalErrors > 0
| sort totalErrors desc
| limit 10
```


Error-to-Invocation Ratio

```
// Error rate percentage by function over the last 24 hours
//
// Corrected 08/12/2026 (two defects). Besides the non-existent metric keys,
the old cell used
// `append [ ... ]` to bring invocations alongside errors. append STACKS ROWS
– it does not join
// columns – so totalErrors and totalInvocations never landed on the same row
and error_pct was
// null for every row even with valid keys. Request both series inside ONE
timeseries block, which
// is what actually aligns them on a shared timeline and grouping.
//
// Metric keys corrected 08/12/2026. The AWS CloudWatch Lambda metrics are
named
// cloud.aws.lambda.<CloudWatchName>.By.FunctionName – Invocations /
Errors / Duration /
// Throttles / ConcurrentExecutions / IteratorAge. The `dt.cloud.aws.lambda.*`
names used before
// do not exist in any spelling, and a timeseries against a missing key
returns an EMPTY result
// instead of an error, so every one of these tiles silently drew nothing. The
split dimension is
// FunctionName – dt.entity.aws_lambda_function is null on these metrics.
Enumerate with:
// metrics | filter startsWith(metric.key, "cloud.aws.lambda") | fields
metric.key | sort metric.key asc
timeseries {
  invocations = sum(cloud.aws.lambda.Invocations.By.FunctionName),
  errors = sum(cloud.aws.lambda.Errors.By.FunctionName)
}, from:-24h, by:{FunctionName}
| fieldsAdd totalInvocations = arraySum(invocations), totalErrors =
arraySum(errors)
| filter totalInvocations > 0
| fieldsAdd error_pct = round((totalErrors / totalInvocations) * 100,
decimals: 2)
| sort error_pct desc
| limit 10
```

4. Throttle Analysis

Throttling occurs when Lambda cannot allocate an execution environment, typically due to:

- **Account-level concurrency limit** (default 1,000 per region)
- **Reserved concurrency** on the function being exhausted
- **Burst limit** exceeded (3,000 in US regions, varies by region)

Throttled Invocations Over Time

```
// Lambda throttles over the last 24 hours by function
//
// Corrected 08/12/2026: the key is Throttles, not `throttlers`.
// The `filter totalThrottles > 0` guard was also dropped – on a healthy
estate every function
// throttles zero times, so the guard emptied the table and made a correct
result look like a
// broken query. Seeing the zeros is the point.
//
// Metric keys corrected 08/12/2026. The AWS CloudWatch Lambda metrics are
named
// cloud.aws.lambda.<CloudWatchName>.By.FunctionName – Invocations /
Errors / Duration /
// Throttles / ConcurrentExecutions / IteratorAge. The `dt.cloud.aws.lambda.*`
names used before
// do not exist in any spelling, and a timeseries against a missing key
returns an EMPTY result
// instead of an error, so every one of these tiles silently drew nothing. The
split dimension is
// FunctionName – dt.entity.aws_lambda_function is null on these metrics.
Enumerate with:
// metrics | filter startsWith(metric.key, "cloud.aws.lambda") | fields
metric.key | sort metric.key asc
timeseries throttles = sum(cloud.aws.lambda.Throttles.By.FunctionName),
from:-24h, by:{FunctionName}
| fieldsAdd totalThrottles = arraySum(throttles)
| fields FunctionName, totalThrottles
| sort totalThrottles desc
| limit 20
```

Concurrent Executions Trend

```
// Concurrent Lambda executions over the last 6 hours
//
// Corrected 08/12/2026: the key is ConcurrentExecutions, not
`conc_executions`.
//
// Metric keys corrected 08/12/2026. The AWS CloudWatch Lambda metrics are
named
// cloud.aws.lambda.<CloudWatchName>.By.FunctionName — Invocations /
Errors / Duration /
// Throttles / ConcurrentExecutions / IteratorAge. The `dt.cloud.aws.lambda.*`
names used before
// do not exist in any spelling, and a timeseries against a missing key
returns an EMPTY result
// instead of an error, so every one of these tiles silently drew nothing. The
split dimension is
// FunctionName — dt.entity.aws_lambda_function is null on these metrics.
Enumerate with:
//  metrics | filter startsWith(metric.key, "cloud.aws.lambda") | fields
metric.key | sort metric.key asc
timeseries concurrency =
max(cloud.aws.lambda.ConcurrentExecutions.By.FunctionName), from:-6h, by:
{FunctionName}
| fieldsAdd peakConcurrency = arrayMax(concurrency)
| sort peakConcurrency desc
| limit 10
```

5. API Gateway Integration

AWS API Gateway is a common front-end for Lambda functions. Dynatrace monitors the API Gateway-to-Lambda path for end-to-end latency.

API Gateway Monitoring Points

Metric	What It Measures
Count	Total API requests
Latency	End-to-end latency (API GW → Lambda → response)
IntegrationLatency	Lambda execution time only

Metric	What It Measures
4XXError	Client-side errors
5XXError	Server-side errors (Lambda failures)

Tracing API Gateway Requests

With the Dynatrace Lambda Layer, traces propagate through API Gateway into Lambda:

```
Client → API Gateway → Lambda Function → DynamoDB / SQS / etc.
|         |           |           |
| span    | span      | span      | span
|         |           |           |
\         /           \         /
      Complete distributed trace
```

```
// Trace spans from API Gateway and Lambda in the last hour
//
// Corrected 08/12/2026: the old `contains(span.name, "lambda") or
contains(span.name,
// "api-gateway")` name-guessing filter matched nothing, because spans are
named for the function
// or route. Select on the semantic field instead – faas.name is present on
every FaaS span.
fetch spans, from:-1h
| filter span.kind == "server" and isNotNull(faas.name)
| fieldsKeep timestamp, trace.id, span.name, duration, span.status_code,
faas.name
| sort timestamp desc
| limit 20
```

6. Step Functions Tracing

AWS Step Functions orchestrate Lambda functions into workflows. Dynatrace can trace the entire state machine execution.

Step Functions Monitoring

What to Monitor	Why
Execution duration	Detect slow workflows

What to Monitor	Why
Failed executions	Identify error patterns
State transitions	Find bottleneck states
Lambda cold starts within workflows	Impact on workflow latency

Step Function Execution Tracking via Logs

```
// Step Function execution logs (if forwarded to Dynatrace)
fetch logs, from:-24h
| filter contains(content, "StepFunction") or contains(content,
"StateMachine")
| fieldsKeep timestamp, content, log.source
| sort timestamp desc
| limit 20
```

7. DynamoDB Performance

DynamoDB is frequently used with Lambda for serverless data storage. Key monitoring areas:

Metric	Concern
Read/Write capacity	Throttling if exceeded
Latency	Slow queries impacting Lambda duration
Consumed vs provisioned	Over/under-provisioning
Error count	Conditional check failures, throttles

DynamoDB Span Analysis

```
// Database call duration from Lambda spans in the last 6 hours
//
// Corrected 08/12/2026: `db.operation` is not a field – the stable name is
`db.operation.name`.
// It went unnoticed because the preceding `db.system == "dynamodb"` filter
already emptied the
// stream, so the bad field in the summarize was never evaluated. The hard-
coded dynamodb filter
// is also removed: it returns nothing unless that engine is present. Group by
db.system and read
// what your estate actually calls.
fetch spans, from:-6h
| filter span.kind == "client" and isNotNull(db.system)
| summarize {avg_duration_ms = avg(duration) / 1ms, call_count = count()}, by:
{db.system, db.namespace, db.operation.name}
| sort avg_duration_ms desc
| limit 15
```

8. End-to-End Serverless Tracing

A complete serverless application might follow this pattern:

```
API Gateway → Lambda (auth) → Lambda (business logic) → DynamoDB
                                     → SQS → Lambda (async processor)
                                     → SNS → Email notification
```

Dynatrace traces the entire flow as a single distributed trace when the Lambda Layer propagates context headers.

Trace Analysis Across Services

```
// Slowest traces involving Lambda functions in the last hour
fetch spans, from:-1h
| filter span.kind == "server"
| summarize {trace_duration = max(duration), span_count = count()}, by:
{trace.id}
| sort trace_duration desc
| limit 10
```

Serverless Monitoring Best Practices

Practice	Description
Always install Dynatrace Lambda Layer	Enables distributed tracing and code-level visibility
Set meaningful function names	Avoid auto-generated names for better identification
Monitor concurrency headroom	Alert before hitting account limits
Track cold start ratio	High ratio indicates scaling or provisioned concurrency needs
Correlate API Gateway with Lambda	End-to-end latency matters more than Lambda duration alone

9. Summary and Next Steps

Key Takeaways

- Lambda monitoring combines **CloudWatch metrics** (invocations, errors, throttles) with **Dynatrace Lambda Layer** (traces, code-level visibility)
- **Cold starts** can be detected via spans and mitigated with provisioned concurrency or SnapStart
- **Error-to-invocation ratio** is a more meaningful health metric than raw error count
- **End-to-end tracing** through API Gateway, Lambda, and downstream services requires context propagation

Next Steps

- **CLOUD-05: Azure Integration** — Azure-specific monitoring patterns
- **CLOUD-07: CloudWatch Log Ingestion** — Forward Lambda logs to Dynatrace
- **CLOUD-08: Multi-Cloud Patterns** — Unified serverless monitoring across providers

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace

documentation.