FAQ-27: How Do I Get Microsoft 365 Audit Logs Into Dynatrace?

> **Series:** FAQ — Frequently Asked Questions | **Reference:** 27 — Microsoft 365 Audit Logs Into Dynatrace | **Created:** October 2026 | **Last Updated:** 10/29/2026

Overview

The request usually arrives as one line — "get the O365 audit logs into Dynatrace" — and the design that comes back is usually one pipeline: a collector polls the Office 365 Management Activity API and posts everything to the Dynatrace log ingest API. That design is half right. Three things about this integration change it:

- **There is a Microsoft 365 extension on the Dynatrace Hub, and it is not an audit feed.** It runs on an ActiveGate and produces service-health, usage and license metrics. Its documentation says it reads audit logs only to count users and that no change logs are recorded or exported. Installing it is worthwhile for those metrics; it does not answer the audit question.
- **Microsoft Entra ID sign-ins already have a documented Dynatrace route.** Entra ID diagnostic settings stream to an Event Hub, Azure log forwarding sends them to Dynatrace, and OpenPipeline has a built-in bundle for them. Collecting the same sign-ins a second time through the Management Activity API is duplication, not coverage.
- **The workload audit — Exchange, SharePoint and OneDrive, Teams and the rest, DLP — has no documented Dynatrace route.** A ready-made OpenTelemetry Collector receiver exists, but its current version misses rules this entry depends on (§ 8.1). Either way, you run and own a collector. It has to cope with a feed that Microsoft delivers late, out of order and with duplicates, against an ingest API that drops records stamped more than 24 hours in the past.

That last pairing is the expensive one. A collector that stamps each record with its event time works on the first day and then loses records quietly — after any outage longer than a day, and whenever Microsoft is late. Nothing errors: the ingest API answers `200` instead of `204`, and the records are gone.

This entry is the third **worked example of FAQ-19**, after Zscaler (**FAQ-20**) and Box (**FAQ-23**). The generic decisions — signal classes, route choice, extract-then-discard, bucket isolation — are made there; here they are made for Microsoft 365.

> **A note on the code and the DQL in this entry.** Both collectors (§ 9 and § 10) were checked against Microsoft's Management Activity API reference and exercised against mocked Microsoft and Dynatrace endpoints on 10/27/2026 — pagination, duplicate removal, the timestamp rule, the partial-ingest stop and the 7-day gap check. **They have not been run against a live Microsoft 365 tenant.** Run them against a test tenant before production. The assessment of the `m365` Collector receiver in § 8.1 comes from reading its v1.2.0 source on 10/29/2026; it was not run either. The parse and duplicate queries (§ 13.1, § 13.2) run against synthetic records with `data record(...)` at zero bytes scanned; the rest read real logs and are billed as queries — each scanned up to about a gigabyte on the validation tenant with no Microsoft 365 data present, so keep their timeframes short and scope them to your bucket once it exists. All queries were executed against a live tenant on 10/27/2026.

Table of Contents

1. Short Answer
2. What Exists Today — and What It Does Not Do
3. Two Routes, Split by Signal
4. The Management Activity API — Constraints That Shape the Collector
5. The 24-Hour Timestamp Trap
6. Duplicates and the Checkpoint
7. App Registration and Credentials
8. Choosing the Collector — Python, Workflow or the Collector Receiver
9. Implementation A — Python
10. Implementation B — Dynatrace Workflow
11. Ingestion, Processing, and Storage
12. Field Mapping Starting Point
13. Querying Audit Records in Grail
14. Operating the Collector

15. Implementation Sequence and Checklist
16. Common Gotchas
17. Recommended Approach

Prerequisites

Requirement	Details
Microsoft 365 auditing	On by default for most organizations, but **not** for small and medium business licenses or free enterprise trials. Verify in Exchange Online PowerShell with `Get-AdminAuditLogConfig | Format-List UnifiedAuditLogIngestionEnabled` — run it there, not in Security & Compliance PowerShell, where the property always reads `False` (§ 7)
Entra ID app registration	An app with the **Office 365 Management APIs** application permission *Read activity data for your organization*, and admin consent granted. Add the DLP permission only if you collect `DLP.All` (§ 7)
Credential	A certificate for the Python collector; a client secret in the Credential Vault for the Workflow collector (§ 7)
Dynatrace environment	SaaS with **Grail** and **OpenPipeline**
Ingest credential	Platform token with `openpipeline:logs:ingest` (Python). On a classic or hybrid environment a classic access token with `logs.ingest` also works
Outbound connectivity	Workflow path: `login.microsoftonline.com` and `manage.office.com` under **Settings > General > External requests**
Security sign-off	This feed names people, the files they opened and the mail they read. Decide retention, masking and who can read it before switching anything on (§ 11)
Related reading	**FAQ-19** (the pattern), **FAQ-23** (Box — the closest sibling), **FAQ-24** (data trust), **ORGNZ** (buckets), **OPIPE** / **OPLOGS** (OpenPipeline), **CLOUD** (Azure), **WFLOW** / **FAQ-21** (alert routing)

---

1. Short Answer

Use two routes. Send **Entra ID sign-ins and directory audit** through the documented Azure log forwarding route. Build a scheduled collector on the **Office 365 Management Activity API** for everything else. Do not expect the Microsoft 365 Hub extension to deliver audit events — it does not.

The decisions, made:

Decision	Answer	Why
Entra ID sign-ins and audit	Diagnostic settings → Event Hub → Azure log forwarding	Documented by Dynatrace, with a built-in OpenPipeline bundle for sign-ins (§ 2.2)
Workload audit	Pull from the Management Activity API	No documented Dynatrace route; Microsoft is de-emphasizing webhooks (§ 4)
Content types	`Audit.Exchange`, `Audit.SharePoint`, `Audit.General`, plus `DLP.All` if needed	Leave out `Audit.AzureActiveDirectory` while Entra arrives by the other route (§ 3.1)
Collector	**Python** by default; a **Workflow** for small tenants; the **`m365` Collector receiver** only where its v1.2.0 gaps are acceptable	Microsoft's volume is set by your tenant, not by you; a Workflow run stops at 120 s; the receiver reads only the first page of each content listing (§ 8)
Timestamp	Event time only when it is under 23 h old; always keep it in `m365.creation_time`	Ingest drops older timestamps — and Microsoft keeps content for 7 days (§ 5)
Checkpoint	Per content type, on `contentCreated`, written only after a successful ingest	The API lists content by when it was published, not when the event happened (§ 6)
Duplicates	Remove within a run; count once with `dedup` at query time	Microsoft says duplicates are expected and removing them is the consumer's job (§ 6)

Decision	Answer	Why
Destination	Log Ingestion API v2, `log.source` `microsoft365.audit`, a dedicated bucket	Separate retention and a separate access boundary (§ 11)

---

2. What Exists Today — and What It Does Not Do

Microsoft 365 is a service you do not run, so everything you learn about it comes through Microsoft's APIs — the same position as Zscaler (**FAQ-20**) and Box (**FAQ-23**). The difference here is that two pieces already exist, and both are easy to misread.

2.1 The Hub extension is not an audit feed

The **Microsoft 365, Office 365** extension on the Dynatrace Hub is real and current. It runs on an ActiveGate, calls Microsoft Graph, the Office 365 Management API and the reporting web services, and produces metrics: service health, Outlook, OneDrive, Teams, SharePoint, licenses, and optional per-user usage reports.

What it does with audit logs is narrow. Its documentation says it uses them to count users, and that no change logs are recorded or exported. Its per-user reports are counts, not events — no timestamps, no individual actions.

Install it for what it does. Service health and adoption metrics are useful, and they cost no collector. Then build the audit route separately. Treating the extension as the audit integration is the Microsoft 365 version of FAQ-23's webhook trap: something real, correctly configured, that never carries the events you needed.

2.2 Entra ID has a documented route

For identity, Dynatrace documents the route end to end. In the Azure portal, **Microsoft Entra ID > Diagnostic settings** streams the categories you choose — `AuditLogs`, `SignInLogs`, `NonInteractiveUserSignInLogs`, `ServicePrincipalSignInLogs`, `ManagedIdentitySignInLogs`, `ProvisioningLogs` — to an Event Hub, and Azure log forwarding sends them to Dynatrace. The setting is tenant-wide, so one setting covers every user and application. For sign-ins, OpenPipeline has a built-in technology bundle,

Azure Entra ID Audit Logs, that maps them onto the semantic dictionary's security-event model.

2.3 Nothing documented covers the workload audit

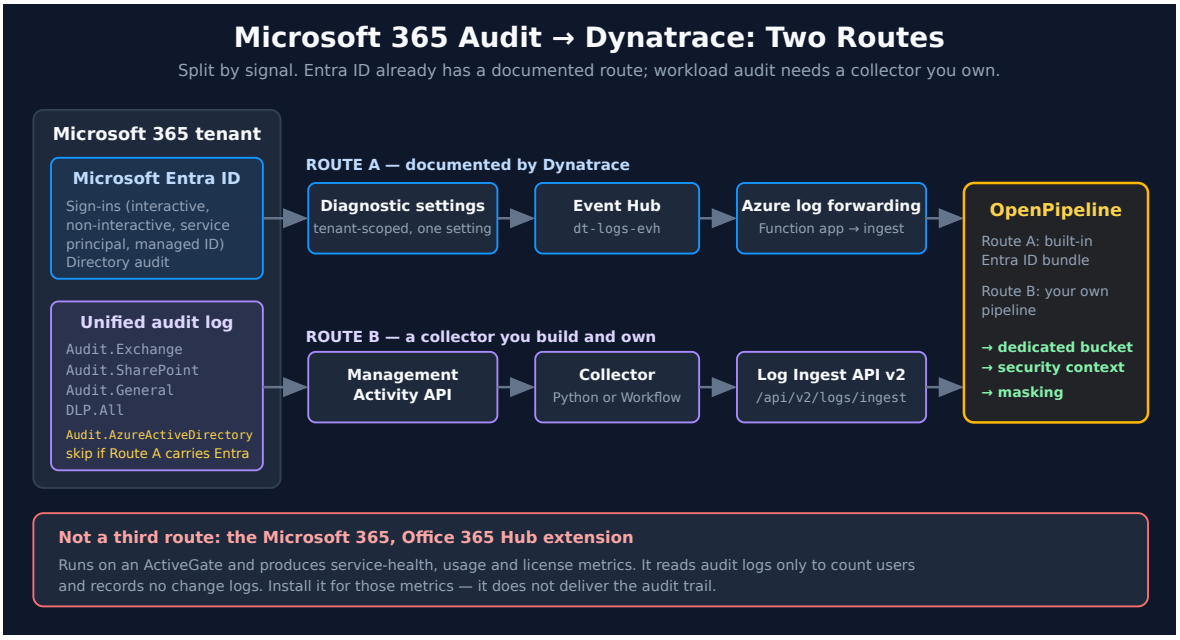
For Exchange, SharePoint and OneDrive, Teams and the other workloads, and for DLP events, there is no Dynatrace extension or documented route as of 10/27/2026.

There is one ready-made component. The `m365` receiver for the OpenTelemetry Collector polls the same Management Activity API. It is maintained in a Dynatrace GitHub repository and built into the Dynatrace Bindplane Distribution of the Collector. In v1.2.0 it reads only the first page of each content listing and shares one checkpoint across all content types; § 8.1 sets out where that matters and how to configure it safely. Whether you adopt it or build your own, you run and own the collector, and the rest of this entry is about doing that well.

> **Sources:**
>
> - Microsoft 365, Office 365 extension (DT docs) — *"The extension uses audit logs to count O365 service users."*; *"No change logs are recorded or exported by the extension."*; *"The extension runs on Dynatrace ActiveGate."*
> - Forward Azure logs (DT docs) — *"Entra ID diagnostic settings are tenant-scoped, not subscription-scoped."*
> - Ingest Microsoft Entra ID sign-in logs (DT docs)
> - m365receiver README (Dynatrace GitHub)
> - **Catalog observation (10/27/2026):** § 2.3's no-route claim rests on the Dynatrace documentation assistant, a web search of dynatrace.com and the extension page above, which covers metrics only — re-check the Hub before quoting it, because listings appear over time

---

3. Two Routes, Split by Signal



FAQ-19 § 2 says to choose the signals before the transport. For Microsoft 365 the split falls out of what already exists: identity has a route, the workloads do not, so each gets the route that fits it. Both land in the same bucket with the same retention and the same access boundary, and query side by side.

3.1 Do not collect Entra ID twice

The Management Activity API has an `Audit.AzureActiveDirectory` content type, and it overlaps Route A. Microsoft's own troubleshooting page says that Entra ID sign-in activity reaches the unified audit log *from* Entra ID's Event Hubs — so Route A reads the same events closer to their source, with categories (non-interactive, service-principal and managed-identity sign-ins) that you choose individually.

While Route A is running, leave `Audit.AzureActiveDirectory` out of the collector. If you cannot use Route A — no Azure subscription for the Event Hub, say — add it to the collector's content types instead. Do not run both and plan to deduplicate later: you pay for every sign-in twice, and the two routes produce records of different shapes, so a `dedup` on one field will not merge them.

3.2 Is Route A already running?

Many organizations already forward Azure logs. Check before building anything:

```
// Is Route A (Entra ID via Azure log forwarding) already delivering sign-ins?
// Matches the routing condition of the built-in Entra ID bundle.
// Control on the validation tenant (10/27/2026): cloud.provider == "azure"
alone
// returned 35,750 records in the same hour, so an empty result here means no
// sign-in forwarding — not a broken query.
fetch logs, from:-1h
| filter cloud.provider == "azure" and contains(content, "SignInLogs")
| summarize {records = count(), last_seen = takeMax(timestamp)}
```

A non-zero count means sign-ins are arriving, and § 3.1 applies. Zero means one of three things — Route A is not set up, Entra ID's diagnostic setting does not include `SignInLogs`, or nothing is forwarding Azure logs at all. Run the query without the `contains` clause to tell the last one apart.

> **Sources:** Troubleshooting the Management Activity API (Microsoft Learn) — *"audit logs for Microsoft Entra sign-in activities are ingested into the unified audit log from Microsoft Entra Event Hubs."*; Ingest Microsoft Entra ID sign-in logs (DT docs). **Derived:** § 3.1's leave-it-out rule combines Microsoft's statement that the API's sign-ins come from Entra ID's Event Hubs with the Dynatrace route that reads those Event Hubs directly.

---

4. The Management Activity API — Constraints That Shape the Collector

The API does not hand you events. You **subscribe** to a content type, **list** the content blobs Microsoft has published for it, and **fetch** each blob, which holds one or more audit records. Every constraint below comes from that shape.

Constraint	Value	Design consequence
Content types	`Audit.AzureActiveDirectory`, `Audit.Exchange`, `Audit.SharePoint`, `Audit.General`, `DLP.All`	One subscription and one checkpoint per type. `Audit.General` is every workload not in the other types

Constraint	Value	Design consequence
Listing window	`startTime` / `endTime` at most **24 h** apart; start at most **7 days** back	Catch-up after an outage walks forward one day at a time
What the window filters on	`contentCreated` — when Microsoft published the blob	You cannot ask for events by when they happened (§ 6)
Pagination	`NextPageUri` response header	Keep following it until it is absent. Stopping at page one loses data with no error
Retention	Content fetchable for **7 days** after it is published	The longest outage you can recover from (§ 14.3)
Latency	No guaranteed maximum. Core services typically 60–90 minutes; others longer. First content on a new subscription can take up to 12 hours	Measure it (§ 13.4); it is what makes § 5 matter
Ordering	Records inside blobs are not in event order	Never assume the newest record ends a run
Duplicates	Expected; Microsoft does not remove them in the API	Remove them yourself (§ 6)
Throttling	Baseline 2,000 requests/minute per tenant; `AF429` errors, which can arrive as HTTP 403	Send `PublisherIdentifier` on every call — including the blob GET — and back off on `AF429`
Subscriptions	Wait 15 minutes between start requests; a stopped subscription cannot recover content published while it was stopped	Never stop a subscription to "reset" a collector

4.1 Poll, do not use the webhook

The API can push a notification to a webhook when content is published. Microsoft's own troubleshooting page says webhooks are being de-emphasized, and that webhook notifications have recently been *slower* than listing content directly. A webhook also needs a public HTTPS listener that answers Microsoft's validation request before the subscription is created. Poll every 5–15 minutes instead.

4.2 Government clouds

GCC, GCC High and DoD tenants use different API hosts (`manage-gcc.office.com`, `manage.office365.us`, `manage.protection.apps.mil`). Sovereign-cloud tenants may also sign in at a different Entra ID host, so both collectors take the API host and the sign-in host as settings.

> **Sources:**
>
> - Management Activity API reference (Microsoft Learn) — *"they must be no more than 24 hours apart, with the start time no more than 7 days in the past."*; *"Make sure this parameter is specified in all issued requests to get a dedicated quota."*; *"Once you submit a subscription request, a waiting period of 15 minutes is required before another request can be made."*; *"You will not be able to retrieve content that was available between the time the subscription was stopped and restarted."*
> - Troubleshooting the Management Activity API (Microsoft Learn) — *"There is no guaranteed maximum latency for notification delivery"*; *"For core services (such as Exchange, SharePoint, OneDrive, and Teams) event availability is typically 60 to 90 minutes."*; *"Webhooks are being de-emphasized by Microsoft because of the difficulty in debugging and troubleshooting."*; *"should also be included in the GET request to the contentUri URL when retrieving the content blob."*

---

5. The 24-Hour Timestamp Trap



Two documented limits disagree by a factor of seven, and the collector sits between them:

- **Microsoft keeps content for 7 days**, and delivers it late — typically an hour or so, with no guarantee.
- **Dynatrace drops a log record whose timestamp is more than 24 hours in the past.** If some records in a request are too old, the API answers `200` with a count of what it did not ingest. Only if *every* record is too old does it answer `400`.

So the natural mapping — `timestamp = CreationTime` — loses data in exactly the situations a collector exists for: catching up after it was down for more than a day, a slow workload, and the first run of a new subscription. Every one of those is a `200`, which most HTTP code treats as success.

The fix has two parts, and both collectors in this entry do both:

1. **Send the event time as its own field**, `m365.creation_time`, on every record. Set `timestamp` from it only when it is under 23 hours old; otherwise leave `timestamp` unset so the record takes its ingest time. Query event time on `m365.creation_time` — recovered records then sit at the moment they arrived on any chart built on `timestamp`, but their real time is still there.
2. **Treat `200` from the ingest API as a failure.** It means records were dropped. Stop, and do not advance the checkpoint.

> **SaaS 1.348 widens the window.** SaaS 1.348 (staged tenant rollout from 10/12/2026) raises the limit to 72 hours. Verify it has reached your tenant before relying on it; until it does, 24 hours is the limit. Even at 72 hours, Microsoft's 7-day retention is longer, so the fix above still applies — change the 23-hour constant only after the new limit is live on your tenant.

`CreationTime` has no time-zone suffix (Microsoft's schema says it is UTC). Append `Z` before parsing, or a parser that assumes local time shifts every record by your UTC offset.

> **Sources:**
>
> - Log ingestion limits (DT docs) — *"The event is dropped if the log event contains a timestamp before the current time minus 24 hours."*
> - POST ingest logs (DT docs) — *"Only a part of input events were ingested due to event invalidity."*
> - What's new in Dynatrace SaaS 1.348 (DT docs) — *"up to 72 hours in the past"*; rollout start 10/12/2026, re-read 10/27/2026
> - Troubleshooting the Management Activity API (Microsoft Learn) — *"Content is available to fetch via the API for 7 days after the notification of the content availability."*
> - Management Activity API schema (Microsoft Learn) — *"The date and time in Coordinated Universal Time (UTC) when the audit log record was generated."*

---

6. Duplicates and the Checkpoint

6.1 Two clocks

Every audit record carries two times, and the API only lets you ask about one of them:

Field	Means	Can you filter on it?
`contentCreated` (on the blob)	When Microsoft **published** the blob	**Yes** — it is what `startTime` / `endTime` filter on
`CreationTime` (on each record)	When the **event** happened	No — the API has no event-time query

The checkpoint therefore follows `contentCreated`. A checkpoint built on the newest `CreationTime` seen goes wrong in both directions: a late record from an early event moves it backwards, and out-of-order blobs move it past records that have not arrived yet.

6.2 The checkpoint, exactly

One checkpoint per content type, holding two values:

Value	Holds	Why
`since`	`contentCreated` of the last blob ingested, to the second	The start of the next listing window
`last_ids`	The `contentId`s already ingested *at* that second	`startTime` is inclusive, so the next listing returns them again

When a whole window has been ingested, `since` moves to the window's end — which is exclusive — and `last_ids` empties. The collector writes the checkpoint **only after the ingest succeeds**, so a failed run re-reads blobs instead of skipping them. Re-reading a blob produces duplicates, which the next section deals with; skipping one is a gap nobody can see afterwards.

6.3 Duplicates are part of the contract

Microsoft states that duplicates are expected: the audit service retries on errors, and the API does not remove the resulting copies, which Purview's export does. Removing them is your job. Both collectors drop repeated record `Id`s within a run, which is cheap. Copies that span runs — and re-reads after a failed run — are counted once at query time with `dedup audit.id` (§ 13.2). Watch the duplicate rate; a sudden rise usually means a checkpoint that is not being written.

> **Sources:** Troubleshooting the Management Activity API (Microsoft Learn) — *"It is the responsibility of the SIEM solution to implement logic to remove such duplicated events upon ingestion."*; *"This is the date the notification was created."* (of `contentCreated`); Management Activity API reference (Microsoft Learn). **Derived:** § 6.2's `since` / `last_ids` pair follows from the documented inclusive `startTime` and exclusive `endTime`.

---

7. App Registration and Credentials

7.1 Turn auditing on, and check it

Auditing is on by default for most Microsoft 365 organizations — but not for small and medium business licenses (Business Basic, Standard and Premium) or unmanaged free trials of enterprise licenses. Check it before debugging anything else:

```
# Exchange Online PowerShell — not Security &amp;amp; Compliance PowerShell,
# where this property always reads False even when auditing is on.
Get-AdminAuditLogConfig | Format-List UnifiedAuditLogIngestionEnabled
```

`True` means auditing is on. If you turn it on, allow up to 60 minutes for the change to take effect, and then up to 12 hours for the first content on a new subscription (§ 4).

7.2 The app registration

Step	What	Notes
1. Register	Entra ID > **App registrations** > **New registration**	Single-tenant is enough for your own organization
2. Permission	**API permissions** > **Office 365 Management APIs** (not Microsoft Graph) > **Application permissions** > *Read activity data for your organization*	Add *Read DLP policy events, including detected sensitive information* only if you collect `DLP.All`
3. Consent	**Grant admin consent**	Without it every call fails
4. Credential	Upload a certificate, or create a client secret (§ 7.3)	Record the tenant ID and client ID
5. Check the token	Decode one token and look for `ActivityFeed.Read` in its `roles` claim	The API checks for that claim and rejects the call with `AF10001` if it is missing

Permission names do not match across pages. Microsoft's API reference names the claim `ActivityFeed.Read`; the Dynatrace extension page lists `ActivityFeed.Read.All`. Pick the permission by its description in the portal, then confirm the token carries the claim the API checks — `ActivityFeed.Read`.

7.3 Certificate or secret

Microsoft's getting-started guide says to configure an X.509 certificate for app-only tokens, and its troubleshooting page shows the same flow with a client secret. Both work. Prefer the certificate where you can:

Collector	Credential	Why
Python (§ 9)	**Certificate** (`.pfx`) in Key Vault or your secret store	No shared secret on the wire; MSAL signs the assertion
Workflow (§ 10)	**Client secret** in the Credential Vault	A certificate would mean writing the signed-assertion code yourself inside the JS runtime; a secret needs one form post

A client secret expires. Put its expiry date in your calendar, and alert on token failures — an expired secret looks exactly like a quiet tenant until the staleness alert fires (§ 14).

> **Sources:**
>
> - Turn auditing on or off (Microsoft Learn) — *"Audit logging is on by default for Microsoft 365 organizations."*; *"Enabling auditing might take up to 60 minutes to take effect."*
> - Get started with Office 365 Management APIs (Microsoft Learn) — *"You must configure an X.509 certificate with your application to be used as client credentials when requesting app-only access tokens from Microsoft Entra ID."*
> - Management Activity API reference (Microsoft Learn) — *"the access token must contain the ActivityFeed.Read claim"*
> - Troubleshooting the Management Activity API (Microsoft Learn) — the client-secret token example
> - Microsoft 365, Office 365 extension (DT docs)

---

8. Choosing the Collector — Python, Workflow or the Collector Receiver

	Python	Workflow (JS runtime)
Infrastructure	You run it — Azure Functions, a container, a VM	None
Execution ceiling	None	**120 s** per run (hard)
Credential	Certificate	Client secret in the Credential Vault
Checkpoint	Your own durable store	`stateClient` app state — verify it works from a workflow first (§ 10.1)
Catch-up after an outage	Drains every missed day in one run	One day-window per content type per run; a long outage drains across many runs
Scheduling, retry, failure alerting	You build them	Built in

Default to Python. In FAQ-23 the volume was known and a Workflow cleared the ceiling with room to spare. Here the volume is set by your Microsoft 365 tenant: every mailbox access, file view and Teams action is a record, so the only reliable figure is the one you measure. A Python collector has no ceiling, and Microsoft's own sample for this API is an Azure Function.

Use a Workflow when the tenant is small enough that one run drains a content type in well under two minutes — measure it: the Workflow returns `finished: false` for any content type it could not drain. It keeps everything inside Dynatrace, and failure alerting comes free.

If you have a choice, a sensible split is FAQ-23's: Workflow for steady state, with the Python collector built and tested as the catch-up tool.

> **Sources:** JavaScript runtime (Dynatrace Developer) — *"Function execution times out after 120 seconds."*; Run JavaScript action (DT docs) — *"The task result size is limited to 6 MB."*; Office 365 Activity Feed Azure Function sample (Microsoft Learn).

8.1 The `m365` Collector receiver

A third option needs no code of your own. The `m365` receiver for the OpenTelemetry Collector subscribes to the Management Activity API, polls it every 5 minutes by default, and emits each audit record as an OTLP log record whose timestamp is the event's `CreationTime`. It also scrapes Microsoft Graph usage reports as metrics. It is maintained in the `dynatrace-bindplane-otel-contrib` repository and built into the full Dynatrace Bindplane Distribution of the Collector (DBDOT). It is not in the slim DBDOT build or in the Dynatrace OTel Collector distribution. Its README configures it in plain Collector YAML; Bindplane can manage the same receiver as its **Microsoft 365** source, and the README's examples need no Bindplane server. Its stability is **beta**.

Measured against the rules this entry depends on, v1.2.0 (released 10/28/2026) falls short in several places:

Rule from this entry	**`m365` receiver v1.2.0**	**Effect**
Follow every `NextPageUri` (§ 4)	Reads the first page of each content listing only	A busy tenant loses content blobs, with no error
Checkpoint per content type, written only after success (§ 6)	One checkpoint for all content types, advanced after every poll even when a type's listing failed	A failed poll loses that window for that content type
Catch up one day-window at a time (§ 4)	Lists from the checkpoint to now in one request	After an outage of more than 24 h the listing exceeds Microsoft's limit and fails, and the checkpoint moves past the gap
Survive a restart (§ 14)	The checkpoint persists only with a `storage` extension	Without one, a restart resumes at now minus the poll interval and the downtime is lost

Rule from this entry	`m365` receiver v1.2.0	Effect
`PublisherIdentifier` on every call (§ 4)	Not sent	No dedicated throttling quota
Leave out `Audit.AzureActiveDirectory` while Route A runs (§ 3.1)	`logs.azureAD` defaults to `true`	Every sign-in twice unless you turn it off
Government clouds (§ 4.2)	Commercial `manage.office.com` and `login.microsoftonline.com` hosts only	Not usable for GCC, GCC High or DoD
Certificate credentials (§ 7)	Client secret only	Track the secret's expiry and rotate it
Event time only under 23 h (§ 5)	`timestamp` is always `CreationTime`	OTLP log ingest discards records older than the Log Age limit — the same trap

If you use it, configure it like this:

```
extensions:
  file_storage:
    directory: /var/lib/otelcol/m365   # must survive restarts
receivers:
  m365:
    tenant_id: ${env:M365_TENANT_ID}
    client_id: ${env:M365_CLIENT_ID}
    client_secret: ${env:M365_CLIENT_SECRET}
    logs:
      azureAD: false   # Entra ID arrives by Route A (§ 3.1)
      dlp: false       # turn on only if you need DLP.All
    storage: file_storage
```

Add `file_storage` to the `service.extensions` list and the receiver to a logs pipeline that exports to your environment's OTLP endpoint. The queries in § 13 and § 14 read fields the Python and Workflow collectors write — `audit.id`, `m365.content_type`, `m365.collected_at`, `log.source == "microsoft365.audit"`. The receiver writes the record ID as `id` and the content type as the resource attribute `m365.audit`, with short

names such as `exchange` rather than `Audit.Exchange`, and sets none of the others, so map them in OpenPipeline or adapt the queries before you rely on them.

When it fits: a commercial-cloud tenant where you have confirmed that a listing for your busiest content type returns no `NextPageUri` header, with Route A carrying Entra ID, a `storage` extension, and a staleness alert that catches an outage within hours. Otherwise use the Python collector. The receiver is new and actively developed, so re-check this table against each release before relying on it.

> **Sources:**
>
> - m365receiver README (Dynatrace GitHub) — *"The storage extension prevents duplication of data after a agent restart by remembering which data were previously collected. No storage is used when omitted."*
> - m365receiver metadata (Dynatrace GitHub) — *"beta: [metrics, logs]"*
> - DBDOT Collector manifest (Dynatrace GitHub) — *"receiver/m365receiver v1.2.0"*
> - Microsoft 365 source (Bindplane docs)
> - Ingest OTLP logs (DT docs) — *"Log events older than the Log Age limit are discarded."*
> - Management Activity API reference (Microsoft Learn) — *"they must be no more than 24 hours apart, with the start time no more than 7 days in the past."*
> - **Code read (10/29/2026):** `client.go` and `logs.go` at tag `v1.2.0` contain no reference to `NextPageUri` or `PublisherIdentifier` (control: `contentUri`, `startTime` and `CreationTime` are found in the same files); one `NextStartTime` checkpoint is shared by all content types and advanced after every poll; the listing window runs from that checkpoint to now with no 24-hour split; every content-type flag defaults to `true`; the only credential field is `client_secret`; the token and API hosts are hard-coded. Not run against a Microsoft 365 tenant
> - **Catalog observation (10/29/2026):** no `m365receiver` entry in the Dynatrace OTel Collector manifest (0.59.0) or the slim DBDOT manifest

9. Implementation A — Python

9.1 Install

```
pip install "msal&gt;=1.29" requests
```

`msal` is Microsoft's authentication library for Python. Version 1.29 added the `.pfx` certificate form used below; the older PEM-plus-thumbprint form still works, but MSAL's own documentation marks it deprecated because it uses a SHA-1 thumbprint.

9.2 The collector

Checked against the Management Activity API reference and `msal` 1.39.0 on 10/27/2026, and run against mocked Microsoft and Dynatrace endpoints (pagination, duplicates, the timestamp rule, a partial ingest and a 7-day gap). **Not yet run against a live Microsoft 365 tenant.** It authenticates to Dynatrace with a **platform token** (scope `openpipeline:logs:ingest`); on a classic or hybrid environment a classic access token with `logs.ingest` also works — send `Api-Token <token>` instead of `Bearer <token>`.

```
"""Collect Microsoft 365 audit records from the Management Activity API
and ship them to the Dynatrace Log Ingestion API.

Run every 5-15 minutes. One checkpoint file per content type.
Requires Python 3.11+, msal&gt;=1.29, requests.
"""
import json
import os
import sys
import time
from datetime import datetime, timedelta, timezone
from pathlib import Path

import msal
import requests

TENANT_ID   = os.environ["M365_TENANT_ID"]
CLIENT_ID   = os.environ["M365_CLIENT_ID"]
CERT_PFX    = os.environ["M365_CERT_PFX_PATH"]          # certificate +
private key
CERT_PASS   = os.environ.get("M365_CERT_PASSPHRASE")
API_ROOT    = os.environ.get("M365_API_ROOT", "https://manage.office.com")  #
GCC etc. differ
LOGIN_HOST  = os.environ.get("M365_LOGIN_HOST",
"https://login.microsoftonline.com")
PUBLISHER   = os.environ.get("M365_PUBLISHER_ID", TENANT_ID)  # your own
tenant GUID

DT_ENV_URL        = os.environ["DT_ENV_URL"]          #
https://abc12345.live.dynatrace.com
DT_PLATFORM_TOKEN = os.environ["DT_PLATFORM_TOKEN"]   # scope:
openpipeline:logs:ingest

CHECKPOINT_DIR = Path(os.environ.get("M365_CHECKPOINT_DIR", "/var/lib/m365-
audit"))

# Audit.AzureActiveDirectory is left out on purpose when Entra ID logs arrive
# through Azure log forwarding (section 3). Add it if they do not.
CONTENT_TYPES = ["Audit.Exchange", "Audit.SharePoint", "Audit.General",
"DLP.All"]

LOG_SOURCE    = "microsoft365.audit"
INGEST_BATCH  = 1000                    # well under 50,000 records / 10 MB
MAX_EVENT_AGE = timedelta(hours=23)     # ingest drops timestamps &gt; 24 h
old (section 5)
MAX_LOOKBACK  = timedelta(days=7) - timedelta(minutes=10)  # content lives 7
days
```

```
SETTLE        = timedelta(minutes=5)    # do not close a window on the last
few minutes
FEED = f"{API_ROOT}/api/v1.0/{TENANT_ID}/activity/feed"

_app = msal.ConfidentialClientApplication(
    CLIENT_ID,
    authority=f"{LOGIN_HOST}/{TENANT_ID}",
    client_credential={"private_key_pfx_path": CERT_PFX, "passphrase":
CERT_PASS},
)


def token() -&gt; str:
    """App-only token; MSAL caches it and renews it before it expires."""
    result = _app.acquire_token_for_client(scopes=[f"{API_ROOT}/.default"])
    if "access_token" not in result:
        raise RuntimeError(f"Entra ID token request failed:
{result.get('error_description')}")
    return result["access_token"]


def m365_get(url: str, params: dict | None = None) -&gt; requests.Response:
    """GET with PublisherIdentifier on every call, retrying throttling and
5xx."""
    params = dict(params or {})
    if "PublisherIdentifier=" not in url:
        params["PublisherIdentifier"] = PUBLISHER
    for attempt in range(6):
        resp = requests.get(url, params=params, timeout=60,
                            headers={"Authorization": f"Bearer {token()}"})
        # Throttling arrives as AF429, sometimes with HTTP 403 rather than
429.
        throttled = resp.status_code == 429 or "AF429" in resp.text[:200]
        if resp.ok and not throttled:
            return resp
        if throttled or resp.status_code &gt;= 500:
            time.sleep(2 ** attempt)
            continue
        raise RuntimeError(f"Management API {resp.status_code}:
{resp.text[:500]}")
    raise RuntimeError(f"Management API still failing after retries: {url}")


def ensure_subscriptions() -&gt; None:
    """Start any content-type subscription that is not enabled yet."""
    resp = m365_get(f"{FEED}/subscriptions/list")
    enabled = {s["contentType"] for s in resp.json() if s.get("status") ==
"enabled"}
```

```
    for ct in CONTENT_TYPES:
        if ct not in enabled:
            r = requests.post(f"{FEED}/subscriptions/start", timeout=60,
                              params={"contentType": ct,
"PublisherIdentifier": PUBLISHER},
                              headers={"Authorization": f"Bearer {token()}"})
            r.raise_for_status()
            print(f"started subscription {ct}; first content can take up to 12
h")


def utc(ts: str) -&gt; datetime:
    """Parse Microsoft's UTC timestamps; CreationTime carries no offset."""
    dt = datetime.fromisoformat(ts.replace("Z", "+00:00"))
    return dt if dt.tzinfo else dt.replace(tzinfo=timezone.utc)


def fmt(dt: datetime) -&gt; str:
    """startTime/endTime format the API accepts: YYYY-MM-DDTHH:MM:SS (UTC)."""
    return dt.astimezone(timezone.utc).strftime("%Y-%m-%dT%H:%M:%S")


def list_blobs(ct: str, start: datetime, end: datetime) -&gt; list[dict]:
    """Every blob published in [start, end), following NextPageUri to the
end."""
    blobs: list[dict] = []
    url = f"{FEED}/subscriptions/content"
    params: dict | None = {"contentType": ct, "startTime": fmt(start),
"endTime": fmt(end)}
    while url:
        resp = m365_get(url, params)
        blobs.extend(resp.json())
        url, params = resp.headers.get("NextPageUri"), None
    return sorted(blobs, key=lambda b: b["contentCreated"])


def to_dt_record(rec: dict, ct: str, content_id: str, now: datetime) -&gt;
dict:
    """Map one audit record to a Dynatrace log record."""
    created = rec.get("CreationTime")
    record = {
        "content":            json.dumps(rec, separators=(",", ":")),
        "log.source":         LOG_SOURCE,
        "audit.id":           rec.get("Id"),
        "audit.action":       rec.get("Operation"),
        "audit.identity":     rec.get("UserId"),
        "audit.result":       rec.get("ResultStatus"),
        "m365.content_type":  ct,
```

```
        "m365.workload":      rec.get("Workload"),
        "m365.record_type":   rec.get("RecordType"),
        "m365.client_ip":     rec.get("ClientIP"),
        "m365.object_id":     rec.get("ObjectId"),
        "m365.content_id":    content_id,
        "m365.collected_at":  now.isoformat(),
    }
    if created:
        event_time = utc(created)
        record["m365.creation_time"] = event_time.isoformat()
        # Stamp event time only when ingest will accept it. Older records keep
        # their ingest time as timestamp; their real time is
m365.creation_time.
        if now - event_time &lt; MAX_EVENT_AGE:
            record["timestamp"] = event_time.isoformat()
    return record


def ingest(records: list[dict]) -&gt; None:
    """POST to the Dynatrace Log Ingestion API, batched."""
    url = f"{DT_ENV_URL}/api/v2/logs/ingest"
    headers = {"Authorization": f"Bearer {DT_PLATFORM_TOKEN}",
               "Content-Type": "application/json; charset=utf-8"}
    for i in range(0, len(records), INGEST_BATCH):
        batch = records[i:i + INGEST_BATCH]
        for attempt in range(5):
            resp = requests.post(url, headers=headers, data=json.dumps(batch),
timeout=60)
            if resp.status_code == 204:
                break
            # 200 is a PARTIAL success: some records were not ingested.
            # Fail so the checkpoint does not move past them.
            if resp.status_code == 200:
                raise RuntimeError(f"Dynatrace partial ingest: {resp.text}")
            if resp.status_code == 413 and len(batch) &gt; 1:
                mid = len(batch) // 2
                ingest(batch[:mid])
                ingest(batch[mid:])
                break
            if resp.status_code in (429, 500, 502, 503, 504):
                time.sleep(2 ** attempt)
                continue
            raise RuntimeError(f"Dynatrace ingest {resp.status_code}:
{resp.text[:500]}")
        else:
            raise RuntimeError("Dynatrace ingest failed after 5 attempts")
```

```
def read_checkpoint(ct: str, now: datetime) -&gt; dict:
    path = CHECKPOINT_DIR / f"{ct}.json"
    if path.exists():
        return json.loads(path.read_text())
    return {"since": fmt(now - timedelta(hours=24)), "last_ids": []}


def write_checkpoint(ct: str, since: datetime, last_ids: list[str]) -&gt;
None:
    """Write atomically — a torn checkpoint is worse than none."""
    CHECKPOINT_DIR.mkdir(parents=True, exist_ok=True)
    path = CHECKPOINT_DIR / f"{ct}.json"
    tmp = path.with_suffix(".tmp")
    tmp.write_text(json.dumps({"since": fmt(since), "last_ids": last_ids}))
    tmp.replace(path)


def collect(ct: str) -&gt; tuple[int, bool]:
    """Drain one content type. Returns (records ingested, gap detected)."""
    now = datetime.now(timezone.utc)
    cp = read_checkpoint(ct, now)
    since, last_ids = utc(cp["since"]), list(cp["last_ids"])
    gap = False
    if since &lt; now - MAX_LOOKBACK:
        # Content older than 7 days cannot be fetched. Say so instead of
        # pretending the collection is complete.
        print(f"GAP {ct}: checkpoint {fmt(since)} is beyond the 7-day window",
file=sys.stderr)
        since, last_ids, gap = now - MAX_LOOKBACK, [], True

    seen: set[str] = set()
    total = 0
    while since &lt; now - SETTLE:
        end = min(since + timedelta(hours=24), now - SETTLE)
        pending: list[dict] = []
        cursor, cursor_ids = since, last_ids
        for blob in list_blobs(ct, since, end):
            if blob["contentId"] in last_ids:
                continue                    # already ingested before a
restart
            for rec in m365_get(blob["contentUri"]).json():
                rid = rec.get("Id")
                if rid in seen:
                    continue                # duplicates are expected (section
6)
                seen.add(rid)
                pending.append(to_dt_record(rec, ct, blob["contentId"], now))
            created = utc(blob["contentCreated"]).replace(microsecond=0)
```

```
            cursor_ids = cursor_ids + [blob["contentId"]] if created == cursor
else [blob["contentId"]]
            cursor = created
            if len(pending) &gt;= INGEST_BATCH:
                ingest(pending)
                total += len(pending)
                pending = []
                write_checkpoint(ct, cursor, cursor_ids)   # only after
success
        if pending:
            ingest(pending)
            total += len(pending)
        # The whole window is done: the next one starts at its exclusive end.
        since, last_ids = end, []
        write_checkpoint(ct, since, last_ids)
    return total, gap


if __name__ == "__main__":
    ensure_subscriptions()
    any_gap = False
    for content_type in CONTENT_TYPES:
        count, had_gap = collect(content_type)
        any_gap |= had_gap
        print(f"{content_type}: ingested {count} records")
    sys.exit(2 if any_gap else 0)
```

9.3 What in that code is load-bearing

Detail	Why
`PublisherIdentifier` on every call, including the blob GET	Without it your calls share a pool with every other caller and are throttled
`AF429` checked in the body, not just HTTP 429	Microsoft's own example of throttling arrives as HTTP 403
`NextPageUri` followed to the end	A busy tenant's listing is paged; stopping at page one loses blobs silently
`timestamp` only under 23 h	§ 5 — the record is otherwise dropped with a `200`
`200` raises	It is a partial ingest. The checkpoint stays where it was

Detail	Why
Checkpoint written after ingest, atomically	A failed run re-reads; a torn file cannot point past unread content
Exit code 2 on a 7-day gap	Content older than 7 days cannot be fetched. Report it rather than clamp past it quietly
`SETTLE` of 5 minutes	A window is not closed on its newest few minutes, in case a just-published blob is not listed yet

9.4 Where to run it

Host	Fit
Azure Functions (timer trigger)	Natural fit — certificate in Key Vault, checkpoint in Table or Blob Storage rather than a file
Kubernetes `CronJob`	Good if you already run Kubernetes; persistent volume for the checkpoint
A VM with cron or a systemd timer	Works; keep the checkpoint on persistent disk

Whatever runs it needs monitoring of its own — see § 14.

> **Sources:** msal on PyPI; Management Activity API reference (Microsoft Learn); Troubleshooting the Management Activity API (Microsoft Learn) — *"Response Code 403"* in the throttling example; POST ingest logs (DT docs) — *"Only a part of input events were ingested due to event invalidity."*
> **Observation (10/27/2026):** the `private_key_pfx_path` form and the deprecation of the SHA-1 form were read from `msal` 1.39.0's own docstring.

---

10. Implementation B — Dynatrace Workflow

10.1 Set up first

Step	What	Why it matters
1. Start the subscriptions	Once per content type: `POST /subscriptions/start?contentType=…` — or run the Python collector once, which does it	The Workflow below lists and fetches; it does not start subscriptions
2. Credential Vault entry	Client ID as username, client secret as password; **AppEngine** scope, **Allow access without app context** on, workflow actor granted access	Workflow JavaScript runs without an app context; a credential missing any of the three cannot be read
3. Outbound allowlist	`login.microsoftonline.com` and `manage.office.com` under **Settings > General > External requests**	Every HTTP call from the task is checked against this list
4. Standard workflow	Not a simple workflow	Simple workflows have no **Run JavaScript** action
5. Test the checkpoint store	A two-task workflow: `setAppState` in the first, `getAppState` in the second	See the note below

> **Verify app state works from a workflow before you schedule this.** The collector keeps its checkpoint in app state through `stateClient`, as FAQ-23's does. Two Dynatrace pages leave open whether a Run JavaScript task can reach app state: the state service says states *"are only accessible from their respective apps"*, and the Run JavaScript page says *"JavaScript tasks run without an app context."* If the test value does not come back, keep the checkpoint in a store the workflow can read and write; nothing else in the design changes.

10.2 The collector

A single **Run JavaScript** action on a 5–15 minute schedule. Checked against the `@dynatrace-sdk/client-classic-environment-v2` 9.1.0 and `@dynatrace-sdk/client-`

`state` 1.11.0 typings and run against mocked SDK and Microsoft endpoints on 10/27/2026; not yet run in a workflow against a live tenant.

```
import { credentialVaultClient, logsClient } from "@dynatrace-sdk/client-classic-environment-v2";
import { stateClient, isNotFound } from "@dynatrace-sdk/client-state";

const M365_CREDENTIAL_ID = "CREDENTIALS_VAULT-XXXXXXXXXXXXXXXX"; // username = client ID, password = secret
const TENANT_ID     = "00000000-0000-0000-0000-000000000000";
const API_ROOT      = "https://manage.office.com";
const FEED          = `${API_ROOT}/api/v1.0/${TENANT_ID}/activity/feed`;
const PUBLISHER     = TENANT_ID;
const CONTENT_TYPES = ["Audit.Exchange", "Audit.SharePoint", "Audit.General", "DLP.All"];
const LOG_SOURCE    = "microsoft365.audit";
const TIME_BUDGET_MS = 90 * 1000;          // stop well before the 120 s ceiling
const INGEST_BATCH  = 1000;
const MAX_EVENT_AGE_MS = 23 * 3600 * 1000; // ingest drops timestamps &gt; 24 h old (§ 5)
const LOOKBACK_MS   = 7 * 24 * 3600 * 1000 - 10 * 60 * 1000;
const SETTLE_MS     = 5 * 60 * 1000;
const DAY_MS        = 24 * 3600 * 1000;

const fmt = (d) =&gt; d.toISOString().slice(0, 19);                // YYYY-MM-DDTHH:MM:SS
const utc = (s) =&gt; new Date(/Z|[+-]\d\d:\d\d$/.test(s) ? s : `${s}Z`);

export default async function () {
  const started = Date.now();
  const outOfTime = () =&gt; Date.now() - started &gt; TIME_BUDGET_MS;

  // --- 1. Token (client secret from the Credential Vault) -----------------
  const creds = await credentialVaultClient.getCredentialsDetails({ id: M365_CREDENTIAL_ID });
  const tokenRes = await fetch(`https://login.microsoftonline.com/${TENANT_ID}/oauth2/v2.0/token`, {
    method: "POST",
    headers: { "Content-Type": "application/x-www-form-urlencoded" },
    body: new URLSearchParams({
      grant_type: "client_credentials",
      client_id: creds.username,
      client_secret: creds.password,
      scope: `${API_ROOT}/.default`,
    }).toString(),
  });
  if (!tokenRes.ok) throw new Error(`Entra ID token failed: ${tokenRes.status} ${await tokenRes.text()}`);
  const { access_token } = await tokenRes.json();
```

```
  const m365 = async (url) =&gt; {
    const u = new URL(url);
    if (!u.searchParams.has("PublisherIdentifier"))
u.searchParams.set("PublisherIdentifier", PUBLISHER);
    const res = await fetch(u, { headers: { Authorization: `Bearer
${access_token}` } });
    const text = await res.text();
    // Throttling is AF429, sometimes on HTTP 403. Stop; the next run resumes.
    if (res.status === 429 || text.slice(0, 200).includes("AF429")) return {
throttled: true };
    if (!res.ok) throw new Error(`Management API ${res.status}:
${text.slice(0, 500)}`);
    return { body: JSON.parse(text), next: res.headers.get("NextPageUri") };
  };

  const ingest = async (records) =&gt; {
    for (let i = 0; i &lt; records.length; i += INGEST_BATCH) {
      const result = await logsClient.storeLog({
        type: "application/json; charset=utf-8",
        body: records.slice(i, i + INGEST_BATCH),
      });
      // storeLog resolves empty on 204 and returns a body on 200, which is a
      // PARTIAL ingest — some records were dropped. Fail before the
checkpoint moves.
      if (result) throw new Error(`Dynatrace partial ingest:
${JSON.stringify(result)}`);
    }
  };

  const summary = {};
  for (const ct of CONTENT_TYPES) {
    if (outOfTime()) break;
    const key = `m365-audit-${ct}`;
    const now = new Date();

    // --- 2. Checkpoint ----------------------------------------------------
-
    let cp = { since: fmt(new Date(now - DAY_MS)), lastIds: [] };
    try {
      const saved = await stateClient.getAppState({ key });
      if (saved?.value) cp = JSON.parse(saved.value);
    } catch (e) {
      if (!isNotFound(e)) throw e;     // first run has no checkpoint yet
    }
    let since = utc(cp.since);
    let lastIds = cp.lastIds;
    let gap = false;
```

```
    if (now - since &gt; LOOKBACK_MS) {   // older than Microsoft keeps
content
      since = new Date(now - LOOKBACK_MS); lastIds = []; gap = true;
    }
    const end = new Date(Math.min(since.getTime() + DAY_MS, now - SETTLE_MS));
    if (end &lt;= since) { summary[ct] = { ingested: 0, gap }; continue; }

    // --- 3. List every blob in [since, end), following NextPageUri ---------
--
    const blobs = [];
    let url = `${FEED}/subscriptions/content?
contentType=${ct}&amp;startTime=${fmt(since)}&amp;endTime=${fmt(end)}`;
    let throttled = false;
    while (url) {
      const page = await m365(url);
      if (page.throttled) { throttled = true; break; }
      blobs.push(...page.body);
      url = page.next;
    }
    if (throttled) { summary[ct] = { throttled: true, gap }; continue; }
    blobs.sort((a, b) =&gt; a.contentCreated.localeCompare(b.contentCreated));

    // --- 4. Fetch, transform, ingest, checkpoint ----------------------------
--
    const seen = new Set();
    let pending = [], total = 0, finished = true;
    let cursor = since, cursorIds = lastIds;
    for (const blob of blobs) {
      if (lastIds.includes(blob.contentId)) continue;
      if (outOfTime()) { finished = false; break; }
      const got = await m365(blob.contentUri);
      if (got.throttled) { finished = false; break; }
      const collectedAt = new Date().toISOString();
      for (const rec of got.body) {
        if (seen.has(rec.Id)) continue;    // duplicates are expected (§ 6)
        seen.add(rec.Id);
        const record = {
          content: JSON.stringify(rec),
          "log.source": LOG_SOURCE,
          "audit.id": rec.Id,
          "audit.action": rec.Operation,
          "audit.identity": rec.UserId,
          "audit.result": rec.ResultStatus,
          "m365.content_type": ct,
          "m365.workload": rec.Workload,
          "m365.record_type": rec.RecordType,
          "m365.client_ip": rec.ClientIP,
          "m365.object_id": rec.ObjectId,
```

```
          "m365.content_id": blob.contentId,
          "m365.collected_at": collectedAt,
        };
        if (rec.CreationTime) {
          const eventTime = utc(rec.CreationTime);
          record["m365.creation_time"] = eventTime.toISOString();
          if (Date.now() - eventTime &lt; MAX_EVENT_AGE_MS) record.timestamp =
eventTime.toISOString();
        }
        pending.push(record);
      }
      const created = new
Date(utc(blob.contentCreated).setUTCMilliseconds(0));
      cursorIds = created.getTime() === cursor.getTime() ? [...cursorIds,
blob.contentId] : [blob.contentId];
      cursor = created;
    }

    if (pending.length) { await ingest(pending); total += pending.length; }
    // A finished window advances to its exclusive end; an interrupted one
    // resumes after the last blob ingested.
    const next = finished ? { since: fmt(end), lastIds: [] } : { since:
fmt(cursor), lastIds: cursorIds };
    await stateClient.setAppState({ key, body: { value: JSON.stringify(next) }
});
    summary[ct] = { ingested: total, finished, gap };
  }

  // Keep the task result small (6 MB limit): counts only.
  return summary;
}
```

10.3 What in that code is load-bearing

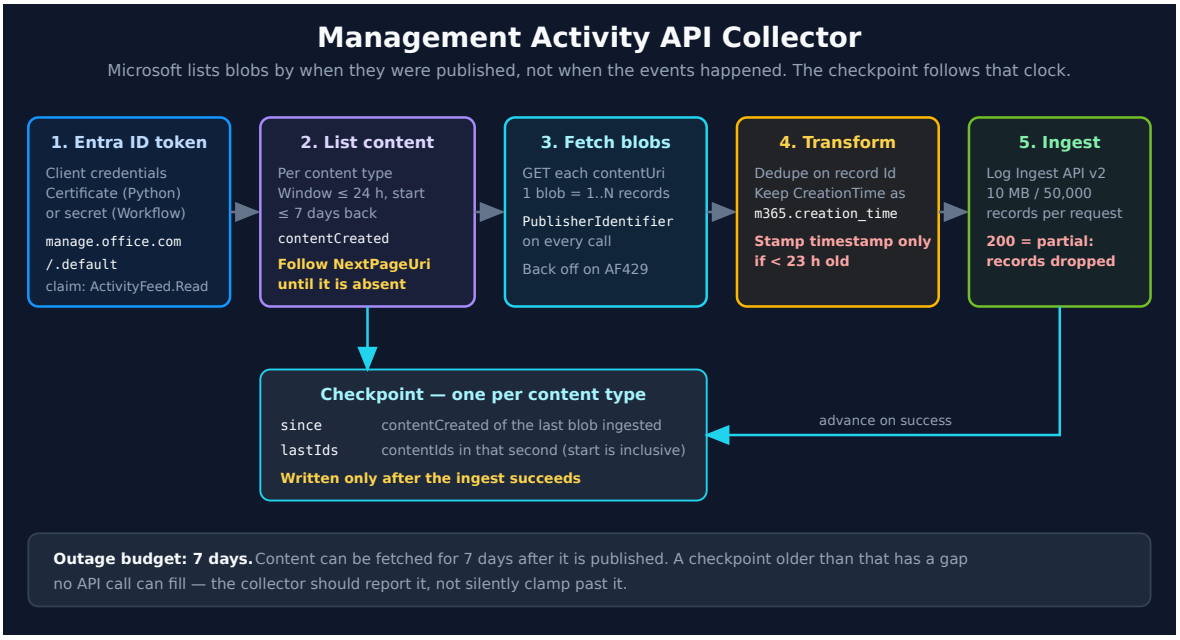
Detail	Why
`TIME_BUDGET_MS` of 90 s	Stops fetching before the 120 s ceiling and saves where it got to. The next run resumes
`storeLog` result checked	The SDK resolves with nothing on `204` and with a body on `200`. A body means records were dropped (§ 5)
`finished` in the result	`false` means the content type did not drain in one run. Persistently `false` means the tenant is too busy for a Workflow — move to Python

Detail	Why
One window per content type per run	Keeps every run inside the ceiling. A long outage drains across many runs
Counts only in the return value	The task result is limited to 6 MB
No TTL on the app state	A checkpoint that expires silently restarts the collector at "the last 24 hours"

> **Sources:**
>
> - Run JavaScript action (DT docs) — *"The AppEngine scope is selected."*; *"All HTTP calls are validated against the global allowlist."*; *"JavaScript tasks run without an app context."*
> - State service (Dynatrace Developer) — *"These states are only accessible from their respective apps and can't be accessed using a different app or ad-hoc functions."*
> - JavaScript runtime (Dynatrace Developer) — *"Function execution times out after 120 seconds."*
> - **Observation (10/27/2026):** `storeLog` returning `SuccessEnvelope | void` — a body on `200`, nothing on `204` — was read from `@dynatrace-sdk/client-classic-environment-v2` 9.1.0

---

11. Ingestion, Processing, and Storage



11.1 Ingest limits

Limit	Value	Consequence
Payload per request	**10 MB**	Both collectors send 1,000 records per request
Records per request	**50,000**	Rarely binding
After processing	Rejected above **16 MB**	A batch can grow during processing; halve on `413` (the Python collector does)
Log age	**24 h** — 72 h from SaaS 1.348 (staged rollout)	§ 5

11.2 Extract, then discard — selectively

FAQ-19 § 4 holds the mechanics and the stage-order trap: discard with `No storage assignment` in the Bucket assignment stage, never `Drop record` in Processing, or the metrics you meant to extract are never built.

Audit is the case where discarding needs the most care. Some of it exists to be produced on request, months later, and a counter cannot answer "who opened this file." Decide with

whoever owns compliance:

Records	Typical posture
Admin and configuration changes, permission and sharing changes, mailbox-permission changes, DLP matches	**Keep** for the full audit retention period
Sign-in failures and unusual sign-ins	**Keep** (they arrive by Route A)
High-volume reads — `FileAccessed`, `MailItemsAccessed`, page views	**Extract counts, then decide.** Keep raw only where a regulation or an investigation process requires it

11.3 Isolation and masking

Put both routes into a **dedicated bucket** (**ORGNZ**): audit retention is usually longer than application-log retention, and this data should be readable by security and compliance, not by everyone who can read logs.

Set `dt.security_context` in the pipeline so an IAM policy can scope it — and read **FAQ-24** first: where `dt.security_context` holds an array, `=` or `IN` against it is always false.

Mask in the Processing stage before storage: for SharePoint and OneDrive, `ObjectId` holds the full path of the file or folder accessed. Hash where you need correlation without identity — remembering that a hash of a corporate email address can be enumerated.

> **Sources:** Log ingestion limits (DT docs); Configure data storage and retention for logs (DT docs) — *"This means you can extract metrics and generate alerts from records that you won't store."*; Management Activity API schema (Microsoft Learn).

---

12. Field Mapping Starting Point

The collectors map Microsoft's common schema onto semantic-dictionary fields where one exists, and onto `m365.*` where none does. Mapping to `audit.*` is what lets a Microsoft 365 record sit next to records from other audit sources in one query.

Microsoft field	Dynatrace field	Notes
`Id`	`audit.id`	**Duplicate key** — unique per audit record
`Operation`	`audit.action`	`FileDownloaded`, `MailItemsAccessed`, `Set-Mailbox`, … — the primary dimension
`UserId`	`audit.identity`	The UPN. Includes system accounts and `app@sharepoint`
`ResultStatus`	`audit.result`	`Succeeded` / `PartiallySucceeded` / `Failed`; `True` / `False` for Exchange admin
`CreationTime`	`m365.creation_time`, and `timestamp` when under 23 h old	UTC with no suffix — append `Z` (§ 5)
`Workload`	`m365.workload`	`Exchange`, `SharePoint`, `OneDrive`, `MicrosoftTeams`, …
`RecordType`	`m365.record_type`	An integer enum; `6` is a SharePoint file operation
`ClientIP`	`m365.client_ip`	Kept as a string: some services send a trusted application's address, and Entra ID events send none. Convert with `toIp()` in the pipeline once you have seen your data
`ObjectId`	`m365.object_id`	File path or object name — mask (§ 11.3)
(blob) `contentId`	`m365.content_id`	Ties a record back to the blob it came from
(collector)	`m365.content_type`, `m365.collected_at`	Which subscription, and when it was collected (§ 14.1)

`audit.id` is `experimental` in the semantic dictionary and `client.ip` is too; `audit.action`, `audit.identity` and `audit.result` are `stable`. An experimental field can change name — check the dictionary before building a dashboard on one.

If you also forward Entra ID through Route A, map both routes' user field to the same name. Correlating a file download with the sign-in before it is one of the most useful things this data does, and it only works if the user is spelled the same way in both.

> **Sources:** Management Activity API schema (Microsoft Learn) — *"Unique identifier of an audit record."*; Global field reference (DT docs). **Dictionary:** `audit.action`, `audit.identity`, `audit.result` (`stable`); `audit.id`, `audit.time`, `client.ip` (`experimental`), read 10/27/2026.

---

13. Querying Audit Records in Grail

§ 13.1 and § 13.2 run against synthetic records at zero bytes scanned — swap the `data record(...)` lines for `fetch logs, from:-24h | filter log.source == "microsoft365.audit"` once data flows. § 13.3 and § 13.4 read real logs and are billed as queries.

13.1 Parse a raw record

Both collectors put the raw record in `content` and the main fields alongside it. Parsing `content` is useful while validating, and for fields the collectors do not lift out:

```
// Parse a raw Microsoft 365 audit record into semantic-dictionary fields.
// Runs standalone against a synthetic record — zero bytes scanned.
// CreationTime has no time-zone suffix; it is UTC, so append "Z".
data record(content = "{\"CreationTime\":\"2026-10-
27T14:03:19\",\"Id\":\"80c76bd2-9d81-4c57-a97a-
accfc3443dca\",\"Operation\":\"FileDownloaded\",\"OrganizationId\":\"41463f53-
8812-40f4-890f-
865bf6e35190\",\"RecordType\":6,\"UserKey\":\"i:0h.f|membership|1003bffd8ec47c
a6@live.com\",\"UserType\":0,\"Workload\":\"SharePoint\",\"ResultStatus\":\"Su
cceeded\",\"ObjectId\":\"https://contoso.sharepoint.com/sites/finance/Shared
Documents/q3-
forecast.xlsx\",\"UserId\":\"jdoe@contoso.com\",\"ClientIP\":\"203.0.113.42\"}
",
            m365.content_type = "Audit.SharePoint")
| parse content, "JSON:rec"
| fieldsAdd audit.id       = rec[Id],
            audit.action   = rec[Operation],
            audit.identity = rec[UserId],
            audit.result   = rec[ResultStatus],
            audit.time     = toTimestamp(concat(rec[CreationTime], "Z")),
            client.ip      = toIp(rec[ClientIP]),
            m365.workload  = rec[Workload],
            m365.record_type = toLong(rec[RecordType]),
            m365.object_id = rec[ObjectId]
| fields audit.time, m365.content_type, m365.workload, audit.action,
audit.identity, audit.result, client.ip, m365.object_id, audit.id
```

13.2 Duplicates — measure them, then count once

The first query is the weekly health check: the duplicate rate per content type. The second is the pattern for every report — `dedup audit.id` before counting, so a record Microsoft delivered twice counts once.

```
// Duplicate rate per content type. Microsoft states duplicates are expected
// and that removing them is the consumer's job.
// Runs standalone against synthetic records — zero bytes scanned.
data record(m365.content_type = "Audit.Exchange",   audit.id = "a1"),
     record(m365.content_type = "Audit.Exchange",   audit.id = "a1"),
     record(m365.content_type = "Audit.Exchange",   audit.id = "a2"),
     record(m365.content_type = "Audit.SharePoint", audit.id = "b1"),
     record(m365.content_type = "Audit.SharePoint", audit.id = "b2")
| summarize {ingested = count(), unique = countDistinctExact(audit.id)}, by:
{m365.content_type}
| fieldsAdd duplicates = ingested - unique
| fieldsAdd duplicate_pct = round(100.0 * duplicates / ingested, decimals: 1)
| sort duplicate_pct desc
```

```
// Count each audit record once, however many copies Microsoft delivered.
// Runs standalone against synthetic records — zero bytes scanned.
data record(audit.id = "a1", audit.action = "FileDownloaded"),
     record(audit.id = "a1", audit.action = "FileDownloaded"),
     record(audit.id = "a2", audit.action = "FileDownloaded"),
     record(audit.id = "a3", audit.action = "UserLoggedIn")
| dedup audit.id
| summarize records = count(), by:{audit.action}
| sort records desc
```

13.3 Confirm records landed in the right bucket

A record in the wrong bucket is still queryable — it just has the wrong retention and the wrong readers.

```
// Which bucket did the audit records land in? Expect your dedicated bucket,
// not default_logs. An empty result means no data yet — not success.
fetch logs, from:-24h
| filter log.source == "microsoft365.audit"
| summarize records = count(), by:{dt.system.bucket, m365.content_type}
| sort records desc
```

13.4 How late is Microsoft?

The number that decides whether § 5 bites. Lag is collection time minus event time; `over_23h` counts records that arrived with ingest time as `timestamp`. Checked on 10/27/2026 against two synthetic records 1 hour and 26 hours late, which returned 60 and 1,560 minutes and one record over 23 hours.

```
// How late does Microsoft deliver, per content type? Lag is collection time
// minus event time. Watch p99 against the 24 h ingest window.
fetch logs, from:-24h
| filter log.source == "microsoft365.audit"
| fieldsAdd lag = toTimestamp(m365.collected_at) -
toTimestamp(m365.creation_time)
| summarize {records = count(),
             p50_min = percentile(lag, 50) / 1m,
             p99_min = percentile(lag, 99) / 1m,
             over_23h = countIf(lag &gt; 23h)},
            by:{m365.content_type}
| sort p99_min desc
```

---

14. Operating the Collector

The failure mode: the collector stops — an expired client secret, a revoked consent, a full disk. It produces no records, and no records looks like a quiet tenant. Microsoft keeps content for 7 days, so an outage nobody notices for a week becomes a permanent gap in the audit trail.

14.1 The staleness check — on the right clock

Measure freshness on **when records were collected** (`m365.collected_at`), not on `timestamp`. Microsoft is normally an hour or more behind, so a check on event time reports an outage every morning. And keep the explicit `events == 0` branch: with no data, `null > 30` is false, and the check would report `OK` for a collector that has never run — the trap FAQ-23 § 13.1 reproduced.

```
// Collector health, measured on when records were COLLECTED, not when the
// events happened — Microsoft's own delay would otherwise read as an outage.
// The events == 0 branch matters: with no data, null &gt; 30 is false and the
// check would report OK.
fetch logs, from:-6h
| filter log.source == "microsoft365.audit"
| summarize {events = count(),
             last_collected = takeMax(toTimestamp(m365.collected_at)),
             content_types = countDistinctExact(m365.content_type)}
| fieldsAdd minutes_since_last = round((now() - last_collected) / 1m,
decimals: 1)
| fieldsAdd collector_status = if(events == 0, "NO DATA — collector never ran
or is misconfigured",
                                else: if(minutes_since_last &gt; 30, "STALLED",
                                else: "OK"))
```

On the validation tenant, with no collector, this returned `NO DATA` on 10/27/2026 — the branch working as intended. Set the threshold from what you see: a tenant whose busiest content type is quiet at night needs a longer one than 30 minutes, or a check per content type.

14.2 What to alert on

Condition	Severity	Why
No records collected for longer than your threshold	High	The collector is down; the 7-day clock is running
`NO DATA`	High	Usually auditing off, consent missing, or the token claim (§ 7.2)
Collector exit code 2, or a Workflow `gap: true`	High	Content has already been lost; record the window for the audit file
Token request failures	High	An expiring client secret announces itself here first
`over_23h` rising (§ 13.4)	Medium	Microsoft is running late; check the collector is not falling behind too
Duplicate rate rising (§ 13.2)	Medium	Usually a checkpoint that is not being written
Workflow `finished: false` on every run	Medium	The tenant has outgrown a Workflow (§ 10.3)

14.3 Recovering from an outage

Outage length	Recovery
Under 7 days	Restart the collector. It resumes from its checkpoint and walks forward a day at a time. Records more than 23 hours old land with ingest time as `timestamp` and their real time in `m365.creation_time`
7 days or more	The API cannot help: content older than 7 days cannot be fetched. Both collectors report the gap. Records in that window are still searchable in Purview Audit for your audit retention period, and exportable from there — outside this pipeline
Checkpoint lost	Both collectors start at "the last 24 hours". Set the checkpoint back by hand to as far as 7 days to replay; `dedup audit.id` absorbs the overlap

Never stop and restart a **subscription** as a recovery step. Microsoft says content published while a subscription was stopped cannot be retrieved afterwards.

> **Sources:** Troubleshooting the Management Activity API (Microsoft Learn) — *"Content is available to fetch via the API for 7 days after the notification of the content availability."*; Management Activity API reference (Microsoft Learn) — *"You will not be able to retrieve content that was available between the time the subscription was stopped and restarted."* **Tenant observation (10/27/2026):** the § 14.1 query returned `NO DATA` with zero records on the validation tenant.

---

15. Implementation Sequence and Checklist

15.1 Sequence

Phase	Do	Gate before moving on
1. Inventory	Run § 3.2; install the Hub extension if you want service-health metrics	Know whether Route A already exists

Phase	Do	Gate before moving on
2. Route A	Entra ID diagnostic settings → Event Hub → Azure log forwarding; enable the Entra ID bundle	Sign-ins arriving (§ 3.2)
3. Microsoft side	Verify auditing (§ 7.1); register the app, grant consent, check the token's claim (§ 7.2)	A manual `subscriptions/list` call succeeds
4. Dynatrace side	Bucket, security context, pipeline; token or Credential Vault entry; allowlist and checkpoint test for a Workflow	A test run reaches Microsoft and writes a checkpoint
5. Collector	Python (§ 9) or Workflow (§ 10)	Records in the right bucket (§ 13.3)
6. Prove the timestamp rule	Stop the collector for more than a day, then restart it	Records over 23 h old arrive with `m365.creation_time` set, and the ingest log shows no `200`
7. Operate	Staleness alert, duplicate check, lag check	Alerts proven by stopping the collector
8. Downstream	Dashboards, alert routing per FAQ-21	—

Phase 6 is the one teams skip, and it is the only way to find out whether the collector loses records before an auditor does.

15.2 Checklist

Microsoft side

- [ ] `UnifiedAuditLogIngestionEnabled` is `True` (checked in Exchange Online PowerShell)
- [ ] App registered with the **Office 365 Management APIs** application permission, admin consent granted
- [ ] Token's `roles` claim includes `ActivityFeed.Read`
- [ ] Certificate (Python) or secret (Workflow) stored, with its expiry date recorded
- [ ] Subscriptions started for each content type you collect — and not for `Audit.AzureActiveDirectory` while Route A runs

Dynatrace side

- [ ] Route A configured, or `Audit.AzureActiveDirectory` added to the collector
- [ ] Dedicated bucket with agreed retention; `dt.security_context` set
- [ ] Pipeline: parse → mask → extract → **`No storage assignment`** in the Bucket assignment stage, never `Drop record`
- [ ] Workflow path: Credential Vault settings, both hosts allowlisted, standard workflow, checkpoint store tested

Operations

- [ ] Checkpoint per content type, on `contentCreated`, written only after ingest succeeds
- [ ] `200` from the ingest API treated as a failure
- [ ] Staleness alert on `m365.collected_at`, tested by stopping the collector
- [ ] Gap reporting (exit code 2 / `gap: true`) routed to someone
- [ ] Retention, masking and readers signed off by whoever owns the data

---

16. Common Gotchas

#	Gotcha	Symptom	Fix
1	**Treating the Hub extension as the audit integration**	Metrics arrive; no audit events ever do	It counts users from audit logs and records no change logs. Build Route B (§ 2.1)
2	**`timestamp = CreationTime` on every record**	Works for a day; records vanish after any outage over 24 h	Event time only under 23 h; always send `m365.creation_time` (§ 5)
3	**Treating ingest `200` as success**	Dropped records with no error	`200` is a partial ingest; stop and keep the checkpoint (§ 5)
4	**Checkpoint on the newest `CreationTime`**	Records skipped or replayed at random	Checkpoint on `contentCreated` (§ 6.1)

#	Gotcha	Symptom	Fix
5	**Not following `NextPageUri`**	Fine on a quiet test tenant; gaps in production	Loop until the header is absent (§ 4)
6	**No `PublisherIdentifier` on the blob GET**	`AF429` throttling	Add it to every call, including `contentUri` (§ 4)
7	**Checking only for HTTP 429**	Throttled calls treated as hard failures	`AF429` can arrive as HTTP 403 — check the body (§ 9.3)
8	**Collecting `Audit.AzureActiveDirectory` alongside Route A**	Every sign-in twice, in two shapes `dedup` cannot match	Leave it out while Route A runs (§ 3.1)
9	**Auditing off**	Every call fails or returns nothing	Check in Exchange Online PowerShell, not Security & Compliance (§ 7.1)
10	**Microsoft Graph permissions instead of Office 365 Management APIs**	`AF10001`	The permission sits under **Office 365 Management APIs** (§ 7.2)
11	**Restarting a subscription to "reset" the collector**	Permanent gap	Content published while stopped cannot be retrieved (§ 14.3)
12	**Staleness measured on `timestamp`**	False outage alerts every quiet hour	Measure on `m365.collected_at` (§ 14.1)
13	**Staleness check with no `events == 0` branch**	A collector that never ran reports `OK`	Test the count explicitly (§ 14.1)
14	**`CreationTime` parsed as local time**	Every record shifted by the UTC offset	It is UTC with no suffix — append `Z` (§ 5)

#	Gotcha	Symptom	Fix
15	**`Drop record` in Processing**	Raw discarded and metrics never built	`No storage assignment` in the Bucket assignment stage (§ 11.2)
16	**`=` or `IN` on `dt.security_context`**	An IAM policy that matches nothing	Read FAQ-24 before writing the policy (§ 11.3)
17	**Running the `m365` Collector receiver on its defaults**	Sign-ins twice; a gap after every restart	Set `logs.azureAD: false` and a `storage` extension (§ 8.1)

Gotchas 1, 3, 5 and 13 share FAQ-23's shape: each produces a confident, plausible answer — metrics flowing, `200 OK`, a first page of results, a health check that says `OK` — rather than an error.

---

17. Recommended Approach

Use what exists first. Install the Hub extension if you want Microsoft 365 service health and adoption metrics, and send Entra ID sign-ins and directory audit through Azure log forwarding. Neither needs a collector.

Then build one collector for the workload audit, in Python. Subscribe to `Audit.Exchange`, `Audit.SharePoint` and `Audit.General` — plus `DLP.All` if you need it, and `Audit.AzureActiveDirectory` only if Entra ID cannot take the other route. Checkpoint per content type on `contentCreated`, follow every `NextPageUri`, and send `PublisherIdentifier` on every call. Use a Workflow instead only after measuring that your tenant drains inside one run.

Consider the `m365` Collector receiver only for a small commercial-cloud tenant. Turn off `azureAD`, give it a `storage` extension, and first confirm that your listings fit on one page — v1.2.0 does not follow `NextPageUri` (§ 8.1). Re-check its gaps on each release.

Get the timestamp rule right before anything else. Send the event time as `m365.creation_time` on every record, stamp `timestamp` from it only when it is under 23 hours old, and treat an ingest `200` as a failure. Then prove it by stopping the collector for more than a day.

Count each record once. Duplicates are part of Microsoft's contract. Remove them within a run, and `dedup audit.id` in every report.

Alert on the collector, on the right clock. Measure freshness on collection time, test it by stopping the collector, and route gap reports to someone who will record them. Microsoft keeps content for seven days; after that, an unnoticed outage is permanent.

What "done" looks like: sign-ins and workload audit in one dedicated bucket with signed-off retention and readers; a collector that has survived a deliberate day-long stop without losing a record; and a staleness alert that has been tested by breaking it on purpose.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.