

IAM-04: Policy Authoring and Management

Series: IAM — IAM Administration | **Notebook:** 4 of 12 | **Created:** January 2026 |
Last Updated: 08/03/2026

Mastering Dynatrace Policy Syntax

Policies are the heart of Dynatrace Gen3 IAM. They define what actions users can perform. This notebook provides a comprehensive guide to policy authoring, from basic syntax to advanced patterns.

Table of Contents

1. [Policy Fundamentals](#)
 2. [Policy Statement Syntax](#)
 3. [Services and Actions Reference](#)
 4. [Condition Expressions](#)
 5. [Common Policy Patterns](#)
 6. [Default vs Custom Policies](#)
 7. [Policy Testing and Validation](#)
 8. [GitOps for Policies](#)
 9. [Policy Assessment Queries](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Gen3 IAM enabled
Permissions	<code>account-iam-admin</code> to create/modify policies
Prior Knowledge	IAM-01 through IAM-03

Sprint 1.337 (April 2026): Permission Updates

New `frontend.name` permission field for Grail RUM data. Sprint 1.337 introduced `frontend.name` as a permission-relevant field on `metric` and `dt.entity.*` Smartscape data. You can now scope IAM policies to specific frontend applications at the **record-level** and **field-level**, instead of relying on bucket-only isolation:

```
ALLOW storage:metric:read
  WHERE storage:metric.frontend.name == "checkout-web"
  OR    storage:metric.frontend.name == "account-web";

ALLOW storage:user.sessions:read
  WHERE storage:user.sessions.frontend.name STARTS WITH "public-";
```

Use this when multiple teams own different frontends but share a single RUM tenant — fine-grained ABAC without splitting buckets.

`dt.security_context` is now a primary field at source. OneAgent now enriches all telemetry with `dt.security_context` (and `dt.cost.*`) as **top-level primary fields** at ingestion on Latest Dynatrace tenants. The `MATCH()` ABAC patterns covered later in this notebook (and across IAM-04, IAM-05, IAM-07) work directly against the primary-field value with no need for OpenPipeline parse processors on OneAgent-instrumented data. The structured-context spread completed in PR #50 already aligns with this — the sprint-1.337 change makes the field guarantee tenant-wide rather than pipeline-dependent.

1. Policy Fundamentals

Policies define **what actions** users can perform. They're separate from boundaries (which control **what data** users can see).

One policy language for all authorization. Dynatrace's security policies are the single configuration surface for every authorization requirement — environment access, Grail data access, and AppEngine capabilities alike. What used to be split between classic *roles* (RBAC — "viewer", "admin") and attribute-based permissions is now unified: classic roles are expressed as `environment:roles:*` permissions inside the same `ALLOW` grammar as `storage:*` and `app-engine:*`. There is no separate RBAC subsystem to manage — one policy, one language, optionally scoped by `WHERE` conditions or boundaries.

Policy Scope Levels

Scope	Applies To	Managed By
Account	All environments in account	Account Admin
Environment	Single environment	Environment Admin

Policy Assignment

Policies are assigned to **groups**, not individual users:

```
Group: dt-checkout-editors
├─ Policy: environment-editor (default)
├─ Policy: log-management (custom)
└─ Boundary: checkout-services-only
```


Policy Types

Type	Description	Examples
Default Policies	Pre-built by Dynatrace	<code>environment-viewer</code> , <code>environment-editor</code>

Type	Description	Examples
Custom Policies	Created by you	log-read-only , dashboard-manager

2. Policy Statement Syntax

Every policy contains one or more **statements** that define allowed actions.



Basic Statement Structure

```
ALLOW <service>:<resource>:<action>;
```

Examples:

Statement	Meaning
ALLOW storage:logs:read	Read all logs
ALLOW storage:spans:read	Read all spans
ALLOW settings:objects:write	Modify settings
ALLOW document:documents:read	Read dashboards/notebooks

Wildcards Are NOT Supported

Live-verified (Account Management API, 07/2026): wildcards in permissions are **rejected** with Invalid policy statement. ALLOW or DENY must be followed by a permission in the form `service-name:resource:action`. This applies to every position: `storage:*:read`, `storage:logs:*`, and `*:*:read` all fail.

Enumerate permissions explicitly — comma-separated lists keep statements compact:

Instead of (invalid)	Write (valid)
<code>ALLOW</code> <code>storage:*:read</code>	<code>ALLOW storage:logs:read, storage:spans:read,</code> <code>storage:metrics:read, storage:events:read,</code> <code>storage:bizevents:read, storage:entities:read,</code> <code>storage:system:read, storage:buckets:read;</code>
<code>ALLOW</code> <code>storage:logs:*</code>	<code>ALLOW storage:logs:read, storage:logs:write;</code>
<code>ALLOW</code> <code>settings:objects:*</code>	<code>ALLOW settings:objects:read, settings:objects:write;</code>

Also live-verified — storage WRITE permissions accept no WHERE conditions.

`ALLOW storage:logs:write WHERE storage:bucket-name = "x"` is rejected (Invalid condition name for permission `storage:logs:write`). Storage writes are all-or-nothing: only `storage:logs:write`, `storage:metrics:write`, `storage:events:write` exist (`storage:spans:write` and `storage:bizevents:write` do **not**), and none can be scoped by `WHERE`. Scope *reads* with conditions; grant *writes* unscoped and only to trusted ingest paths — or control ingest via OpenPipeline instead. (Settings writes are different: `settings:objects:write` **does** accept `settings:schemaId` and `settings:dt.security_context` conditions.)

Multiple Statements

Combine statements to build complete policies:

```
ALLOW storage:logs:read;
ALLOW storage:spans:read;
ALLOW storage:metrics:read;
ALLOW document:documents:read;
```

Note: Statements are additive. Multiple ALLOW statements combine to grant broader access.

3. Services and Actions Reference

Core Services

Service	Purpose	Common Resources
storage	Grail data access	logs , spans , metrics , events , bizevents
settings	Configuration objects	objects , schemas
document	Notebooks, dashboards	documents , directShares
environment	Environment management	extensions , activeGates , oneAgents
state	Entity management	entities , problems
automation	Workflows, AutomationEngine	workflows , calendars

Common Actions

Action	Description
read	View/query data
write	Create/modify data
delete	Remove data
share	Share with others

Action	Description
execute	Run (for workflows)

Storage Service Details

Resource	Description	Actions
logs	Log records	read, write
spans	Distributed traces	read, write
metrics	Metric data	read, write
events	Platform events	read, write
bizevents	Business events	read, write
entities	Smartscape entities (Classic surface)	read
system	System/audit data (query events, internal)	read
buckets	Storage buckets	read, write, delete

Specialized record tables follow the same `storage:<table>:read` shape — e.g.

`storage:security.events:read` , `storage:user.sessions:read` ,
`storage:user.replays:read` , `storage:application.snapshots:read` ,
`storage:smartscape:read` .

Storage condition attributes: scope `storage:*` reads with `storage:bucket-name` (e.g. `startsWith "default_"`), `storage:table-name` (e.g. `= "logs"`), and `storage:dt.security_context` (the Gen3 record-level scoping spine — see IAM-05).

Example — default-bucket monitoring read:

```
ALLOW storage:buckets:read WHERE storage:bucket-name startsWith "default_";
ALLOW storage:logs:read, storage:events:read, storage:metrics:read,
      storage:entities:read, storage:bizevents:read, storage:spans:read;
```

For system/audit data (the *Storage All System Data Read* managed policy), add

`storage:system:read` :

```
ALLOW storage:buckets:read;
ALLOW storage:system:read;
ALLOW storage:logs:read, storage:events:read, storage:metrics:read,
      storage:entities:read, storage:bizevents:read, storage:spans:read;
```

Settings Service Details

Resource	Description	Actions
objects	Configuration objects	read, write, delete
schemas	Schema definitions	read

Document Service Details

Resource	Description	Actions
documents	Dashboards, notebooks	read, write, delete, share
directShares	Direct share links	write

Environment Roles (`environment:roles:*`)

Classic environment roles are expressed as IAM permissions. This is how RBAC folds into the one policy language (see §1). All are scoped by the `environment:management-zone` attribute **except** `agent-install` and `configure-request-capture-data` (management-zone not applicable).

Permission	Classic role meaning
<code>environment:roles:viewer</code>	Access environment (read)
<code>environment:roles:manage-settings</code>	Change monitoring settings
<code>environment:roles:agent-install</code>	Download / install OneAgent
<code>environment:roles:logviewer</code>	View logs
<code>environment:roles:view-sensitive-request-data</code>	View sensitive request data
<code>environment:roles:configure-request-capture-data</code>	Configure capture of sensitive data

Permission	Classic role meaning
<code>environment:roles:replay-sessions-with-masking</code>	Replay session data (masked)
<code>environment:roles:replay-sessions-without-masking</code>	Replay session data (unmasked)
<code>environment:roles:view-security-problems</code>	View security problems
<code>environment:roles:manage-security-problems</code>	Manage security problems

Privacy / compliance note. The session-replay pair and `view-sensitive-request-data` are the load-bearing least-privilege controls here — `replay-sessions-without-masking` and `view-sensitive-request-data` expose unmasked end-user data and should be granted narrowly (specific groups, specific management zones), never bundled into a broad editor policy.

Extensions Service Details

Extension authorization splits into *definitions* (the extension package) and *configurations* (a deployed instance):

Permission	Description	Condition attributes
<code>extensions:definitions:read / :write</code>	The extension package itself	<code>extensions:extension-name</code>
<code>extensions:configurations:read / :write</code>	A deployed configuration	<code>extensions:host</code> , <code>extensions:host-group</code> , <code>extensions:ag-group</code> , <code>extensions:management-zone</code>
<code>extensions:configuration.actions:write</code>	Run configuration actions	(as above)
<code>extensions:discovery.jmx:read / extensions:discovery.pmi:read</code>	Discovery surfaces	—

Extension-ownership delegation — let a team manage *their* extension on *their* hosts without granting tenant-wide extension rights:

```
ALLOW extensions:definitions:read, extensions:definitions:write
WHERE extensions:extension-name = "com.acme.payments";
ALLOW extensions:configurations:read, extensions:configurations:write
WHERE extensions:host-group startsWith "payments-";
```

AppEngine Service Details (`app-engine:*`)

Permission	Description
<code>app-engine:apps:run</code>	See / open an app (controls navigation visibility)
<code>app-engine:functions:run</code>	Execute AppEngine functions
<code>app-engine:apps:install</code>	Install apps
<code>app-engine:apps:delete</code>	Remove apps

App-level scoping uses the `shared:app-id` attribute (operators `=`, `!=`, `IN`, `NOT IN`, `startsWith`, `NOT startsWith`). Example — let developers install only their own apps, and grant Workflows access:

```
ALLOW app-engine:apps:install, app-engine:apps:delete WHERE shared:app-id
startsWith "my.";
ALLOW app-engine:apps:run WHERE shared:app-id = "dynatrace.automations";
```

4. Condition Expressions

Conditions add fine-grained control to policy statements.

Boundary standardization (canonical pattern). The examples in this section illustrate condition operators with literal values for clarity. In production, the literal value should be a `${bindParam:NAME}` placeholder bound at policy-binding time, with the boundary field standardized on `dt.security_context` (Gen3) or management zone (Classic). See [IAM REFERENCE.md § Policy Parameterization and](#)

Boundary Standardization for the canonical pattern, and IAM-10: Templated Policy-Group Assignments for binding mechanics.

Condition Syntax

```
ALLOW <service>:<resource>:<action> WHERE <condition>;
```

Condition Operators

Operator	Description	Example
=	Equals	storage:dt.security_context = "team-a"
!=	Not equals	storage:dt.security_context != "restricted"
IN	In list	storage:dt.security_context IN ("team-a", "team-b")
startsWith	Prefix match	settings:schemaId startsWith "builtin:alerting"
contains	Contains substring	settings:schemaId contains "custom"
MATCH	Wildcard pattern match	storage:dt.security_context MATCH('*/app:easytrade')

MATCH() VS startsWith : Use MATCH() for flexible wildcard patterns anywhere in the value (e.g. middle segment). Use startsWith when matching a fixed leading prefix.

Common Condition Fields

Domain	Field	Description
Storage	storage:dt.security_context	Security context on data
Storage	storage:bucket-name	Bucket name (for team isolation)
Settings	settings:schemaId	Settings schema

Domain	Field	Description
Settings	<code>settings:dt.security_context</code>	Security context on settings

Security Context vs Bucket Conditions

Default to `dt.security_context` for general team-scoped access. Bucket conditions work well when the bucket already exists for a specific reason — compliance separation, retention isolation, or hard cost attribution. For routine team-scoped access without one of those scenarios, the `dt.security_context` condition alone is usually simpler. See **IAM REFERENCE.md § Bucket-Match Overlay** for the scenario gating.

Condition Type	Use Case	Example
Security Context	Flexible, changeable access	<code>storage:dt.security_context = "team-a"</code>
Bucket	Hard data isolation	<code>storage:bucket-name = "team-a-logs"</code>
Both	Defense in depth	Both conditions in policy + boundary

Condition Examples

Scope to security context (exact):

```
ALLOW storage:logs:read WHERE storage:dt.security_context = "checkout"
```

Scope to specific bucket (team isolation):

```
ALLOW storage:logs:read WHERE storage:bucket-name = "checkout_logs"
```

Scope to settings schema:

```
ALLOW settings:objects:write WHERE settings:schemaId startsWith  
"builtin:alerting"
```

Multiple conditions (AND):

```
ALLOW storage:logs:read WHERE storage:dt.security_context = "team-a" AND
storage:bucket-name = "team-a-logs"
```

Multiple values (IN):

```
ALLOW storage:logs:read WHERE storage:dt.security_context IN ("team-a", "team-
b", "shared")
```

Multiple buckets (team access):

```
ALLOW storage:logs:read WHERE storage:bucket-name IN ("checkout_logs",
"shared_logs")
ALLOW storage:spans:read WHERE storage:bucket-name IN ("checkout_spans",
"shared_spans")
```

Parameterizing the boundary (recommended for production):

The examples above use literal values to illustrate operators. In production, replace each literal with a `${bindParam:NAME}` placeholder and supply the value at binding time. One policy template, N bindings — instead of N near-identical policies.

Recommended (parameterized):

```
ALLOW storage:logs:read WHERE storage:dt.security_context =
"${bindParam:team}";
```

Resolved when bound with `{"parameters": {"team": "checkout"}}`:

```
ALLOW storage:logs:read WHERE storage:dt.security_context = "checkout";
```

The same approach applies to bucket-name overlays, MATCH() patterns, and IN lists. See **IAM-10: Templated Policy-Group Assignments** for binding mechanics and **IAM REFERENCE.md § Policy Parameterization and Boundary Standardization** for the full canonical pattern (including the Gen3 = `dt.security_context` standardization rule and the change-management caveat).

Structured Security Context Design

For organisations with multiple teams and cross-cutting access requirements, a flat single-value context (e.g. `"checkout"`) scales poorly. Instead, encode multiple dimensions into a structured key-value string:

```
comp:<component>/bu:<business-unit>/app:<application>;
```

Examples:

Entity	Security Context
Easytrade application	comp:app/bu:digital/app:easytrade
Easytrade database	comp:db/bu:digital/app:easytrade
Easytrade load balancer	comp:lb/bu:digital/app:easytrade

 **Dimension order matters for Smartscape/Classic entities.** Place the dimension you need transversal access on **first** — typically `comp` for infrastructure teams. Use `MATCH('comp:db*')` to anchor at the start (equivalent to a prefix match). **Do not use `startsWith` on `storage:smartscape:read` or `storage:entities:read`** — there is a known bug; `MATCH` with a trailing `*` is the supported alternative.

Using `MATCH()` for flexible policy conditions:

App team — access all components within their application:


```
// 3rd Gen signals
ALLOW storage:metrics:read WHERE storage:dt.security_context
MATCH('*/app:easytrade');
ALLOW storage:logs:read     WHERE storage:dt.security_context
MATCH('*/app:easytrade');
ALLOW storage:spans:read    WHERE storage:dt.security_context
MATCH('*/app:easytrade');
```

Transversal database team — access all database components across all applications:

```
// 3rd Gen signals
ALLOW storage:metrics:read WHERE storage:dt.security_context
MATCH('comp:db*');
ALLOW storage:logs:read WHERE storage:dt.security_context
MATCH('comp:db*');
ALLOW storage:spans:read WHERE storage:dt.security_context
MATCH('comp:db*');
// Smartscape/Classic entities – use MATCH (NOT startsWith – known bug)
ALLOW storage:entities:read WHERE storage:dt.security_context
MATCH('comp:db*');
ALLOW storage:smartscape:read WHERE storage:dt.security_context
MATCH('comp:db*');
```

Note: `MATCH()` is the recommended operator for both Grail signals (`storage:logs:read` etc.) and Smartscape/Classic entity access (`storage:smartscape:read` , `storage:entities:read`). Anchor at the start with a trailing `*` to mimic prefix matching (e.g. `MATCH('comp:db*')`). `startsWith` has a known bug on `storage:smartscape:read` and is no longer recommended. See [IAM-05](#) for the corresponding boundary patterns.

5. Common Policy Patterns



Pattern 1: Read-Only Access

```
// Read all data, no write access (wildcards are not supported – enumerate)
ALLOW storage:logs:read, storage:spans:read, storage:metrics:read,
      storage:events:read, storage:bizevents:read, storage:entities:read,
      storage:system:read, storage:buckets:read;
ALLOW document:documents:read;
ALLOW state:app-states:read, state:user-app-states:read;
ALLOW settings:objects:read, settings:schemas:read;
```

Pattern 2: Team-Scoped Access (Security Context)

Recommended (parameterized): One policy bound per team via binding parameters. Boundary field standardized on `dt.security_context` (Gen3 universal scoping field). See **IAM REFERENCE.md § Policy Parameterization and Boundary Standardization**.

```
// Parameterized – one policy template, bound per team via binding parameters.
// Reads are condition-scoped; storage writes cannot take WHERE conditions
// (live-verified) – grant them separately and unscoped only where justified.
ALLOW storage:logs:read, storage:spans:read, storage:metrics:read,
      storage:events:read, storage:bizevents:read
      WHERE storage:dt.security_context = "${bindParam:team}";
ALLOW settings:objects:read, settings:objects:write
      WHERE settings:dt.security_context = "${bindParam:team}";
```

Resolved when bound with `{"parameters": {"team": "checkout-team"}} :`

```
ALLOW storage:logs:read, storage:spans:read, storage:metrics:read,
      storage:events:read, storage:bizevents:read
      WHERE storage:dt.security_context = "checkout-team";
ALLOW settings:objects:read, settings:objects:write
      WHERE settings:dt.security_context = "checkout-team";
```

Bind the same template to each team with a different `team` value rather than duplicating the policy. See **IAM-10** for binding mechanics.

Pattern 3: Bucket-Based Team Isolation (Recommended for Strict Isolation)

Use this pattern when teams must have physically separated data with no cross-team visibility.

When this pattern fits: bucket boundaries are well-suited to scenarios where the bucket already exists for a specific reason — compliance separation (PCI/HIPAA), retention isolation, or hard cost attribution. For general team-scoped access without one of those reasons, **Pattern 2 (`dt.security_context`) is usually the simpler choice**. When the bucket-based scenario does apply, parameterize per-team bucket names via binding parameters (shared bucket statements stay as literals — "shared" is a singleton, not a per-team scope). See **IAM REFERENCE.md § Bucket-Match Overlay** for scenario gating, and **§ Policy Parameterization and Boundary Standardization** for the full pattern.

```
// Parameterized – per-team data permissions (one policy template, bound per team)
ALLOW storage:logs:read    WHERE storage:bucket-name = "${bindParam:log-bucket}";
ALLOW storage:spans:read   WHERE storage:bucket-name = "${bindParam:span-bucket}";
ALLOW storage:metrics:read WHERE storage:bucket-name = "${bindParam:metric-bucket}";

// Shared buckets – literal (Class C: shared is a singleton, not per-team)
ALLOW storage:logs:read   WHERE storage:bucket-name = "shared_logs";
ALLOW storage:spans:read  WHERE storage:bucket-name = "shared_spans";
```

Writes cannot be bucket-scoped (live-verified). `storage:logs:write` WHERE `storage:bucket-name = ...` is rejected by the API, and `storage:spans:write` does not exist at all — spans are ingest-only. If a team needs log/metric/event write access, grant it unscoped in a separate, deliberately-assigned policy, and control WHERE data lands via OpenPipeline routing — not via IAM conditions.

Resolved when bound with `{"parameters": {"log-bucket": "checkout_logs", "span-bucket": "checkout_spans", "metric-bucket": "checkout_metrics"}}` :

```
ALLOW storage:logs:read WHERE storage:bucket-name = "checkout_logs";
ALLOW storage:spans:read WHERE storage:bucket-name = "checkout_spans";
ALLOW storage:metrics:read WHERE storage:bucket-name = "checkout_metrics";
ALLOW storage:logs:read WHERE storage:bucket-name = "shared_logs";
ALLOW storage:spans:read WHERE storage:bucket-name = "shared_spans";
```

Bucket + Security Context (Defense in Depth):

Recommended (parameterized): Both fields parameterized; the bucket-match clause is the Gen3 data overlay on the `dt.security_context` spine.

```
// Parameterized defense-in-depth
ALLOW storage:logs:read WHERE storage:bucket-name = "${bindParam:log-bucket}"
AND storage:dt.security_context = "${bindParam:team}";
```

Resolved when bound with `{"parameters": {"log-bucket": "checkout_logs", "team": "checkout"}} :`

```
ALLOW storage:logs:read WHERE storage:bucket-name = "checkout_logs"
AND storage:dt.security_context = "checkout";
```

Pattern 4: Log Management

```
// Read all logs, manage log pipelines
ALLOW storage:logs:read;
ALLOW settings:objects:read, settings:objects:write
WHERE settings:schemaId startsWith "builtin:logmonitoring";
```

Pattern 5: Dashboard Creator

```
// Create and share dashboards (documents:share does not exist – sharing is
direct-shares)
ALLOW document:documents:read, document:documents:write,
document:documents:delete;
ALLOW document:direct-shares:read, document:direct-shares:write;
ALLOW storage:logs:read, storage:spans:read, storage:metrics:read,
storage:events:read, storage:bizevents:read;
```

Pattern 6: Alerting Manager

```
// Manage alerting configuration
ALLOW settings:objects:read, settings:objects:write
  WHERE settings:schemaId startsWith "builtin:alerting";
ALLOW settings:objects:read, settings:objects:write
  WHERE settings:schemaId startsWith "builtin:problem";
ALLOW automation:workflows:read, automation:workflows:write,
automation:workflows:run;
```

Pattern 7: SRE On-Call Access

```
// Broad read access + problem/event visibility
// (state:problems does not exist – Davis problems are events in Grail)
ALLOW storage:logs:read, storage:spans:read, storage:metrics:read,
  storage:events:read, storage:bizevents:read, storage:entities:read;
ALLOW document:documents:read;
ALLOW settings:objects:read;
```

Pattern 8: Security Auditor

```
// Read everything, write nothing (enumerated – wildcards are not supported)
ALLOW storage:logs:read, storage:spans:read, storage:metrics:read,
  storage:events:read, storage:bizevents:read, storage:entities:read,
  storage:system:read, storage:buckets:read;
ALLOW settings:objects:read, settings:schemas:read;
ALLOW document:documents:read, document:direct-shares:read,
document:environment-shares:read;
ALLOW state:app-states:read, state:user-app-states:read;
```

Pattern 9: Compliance-Scoped Access (Bucket-Based)

```
// Access only to PCI-compliant data bucket
ALLOW storage:logs:read WHERE storage:bucket-name = "pci_logs";
ALLOW storage:spans:read WHERE storage:bucket-name = "pci_spans";
// No access to non-PCI buckets
```

Pattern 10: Tiered Environment Access (Management Zones)

Grant full access in lower environments and read-only in production — using `environment:roles:*` permissions scoped by the `environment:management-zone` attribute (the classic-role surface, complementary to the Grail patterns above):

```
// Full access in dev and hardening
ALLOW environment:roles:viewer, environment:roles:manage-settings
WHERE environment:management-zone IN ("dev", "hardening");

// Read-only in production
ALLOW environment:roles:viewer
WHERE environment:management-zone IN ("prod");
```

This is the canonical app-developer shape: iterate freely in non-prod, observe-only in prod. Add `environment:roles:logviewer` to the prod statement if the team needs production log access without settings rights.

Choosing Between Patterns

Requirement	Recommended Pattern
Teams can share data flexibly	Pattern 2 (Security Context)
Teams must never see each other's data	Pattern 3 (Buckets)
Compliance/regulatory data separation	Pattern 9 (Bucket-Based)
Temporary project access	Pattern 2 (Security Context)
Cost allocation by team	Pattern 3 (Buckets)
Full in non-prod, read-only in prod	Pattern 10 (Management Zones)

6. Default vs Custom Policies

Built-in Default Policies

Dynatrace provides default policies for common use cases:

Policy	Description	Use Case
<code>environment-viewer</code>	Read-only access	Stakeholders, auditors
<code>environment-editor</code>	Read + write (most things)	Team members
<code>environment-admin</code>	Full environment control	Administrators
<code>account-viewer</code>	View account settings	Cross-team visibility
<code>account-iam-admin</code>	Manage IAM	IAM administrators

Reach for Managed Policies First

Default to the policies Dynatrace ships before authoring custom ones. Most standard scenarios — monitoring read access, app users, automation access — are already covered by managed policies; hand-rolling a custom policy for these is extra surface to maintain and audit. Author custom policies only when no managed policy fits (team-scoped data, specific single-capability grants, compliance constraints).

Commonly used managed policies for Grail and AppEngine onboarding:

Managed policy	Grants
Storage Default Monitoring Read	Read on default buckets; auto-adjusts as new tables are added
Storage <code><table></code> Read	Per-table read (e.g. <i>Storage logs Read</i>)
Storage All System Data Read	System/audit + query-event data
Storage All Grail Data Read	Unrestricted Grail read
AppEngine User	<code>apps:run</code> + <code>functions:run</code> (+ document/state access)
AppEngine Admin	Full app lifecycle (install/delete) + settings
AppEngine Developer	App install/delete scoped to own apps
AutomationEngine Access	Workflows app + automation run access

Environment-admin groups typically receive *AppEngine Admin*, *AutomationEngine Access*, and *Storage All Grail Data Read* automatically. Verify the exact managed-

policy set in your tenant's Account Management — the catalog evolves.

When to Use Default Policies

Scenario	Recommendation
Simple permission levels	Use defaults
Quick onboarding	Use defaults
Standard roles	Use defaults

When to Create Custom Policies

Scenario	Recommendation
Team-scoped data access	Custom policy
Specific capability (logs only)	Custom policy
Compliance requirements	Custom policy
Least privilege needs	Custom policy

Combining Default and Custom

Groups can have both:

```
Group: dt-checkout-editors
├─ Default Policy: environment-viewer (base access)
└─ Custom Policy: checkout-write-access (team-specific writes)
```

This provides base read access plus targeted write permissions.

7. Policy Testing and Validation

Testing Methodology

1. Create policy in test environment (or account)

2. **Assign to test group** with a test user
3. **Verify intended access** works
4. **Verify unintended access** is blocked
5. **Document results** and adjust

Common Validation Tests

Test	Expected Result
Query logs in scope	Success
Query logs out of scope	Empty or denied
Modify settings in scope	Success
Modify settings out of scope	Denied
Create dashboard	Success (if allowed)

Debugging Policy Issues

If access isn't working as expected:

1. **Check group membership** - Is user in the correct group?
2. **Check policy assignment** - Is policy assigned to group?
3. **Check boundary** - Does boundary allow access to the entity?
4. **Check conditions** - Do condition values match?
5. **Check security context** - Is data tagged correctly?

See **IAM-09: Troubleshooting Access Issues** for detailed debugging.

8. GitOps for Policies

Version control your policies for audit trail and peer review.

Policy as Code Structure

```
iam-config/
├─ policies/
│   ├─ team-checkout-access.yaml
│   ├─ team-payments-access.yaml
│   ├─ sre-oncall.yaml
│   └─ security-auditor.yaml
├─ groups/
│   └─ group-assignments.yaml
└─ README.md
```

Policy YAML Format (Monaco)

Recommended (parameterized): One policy template, bound to each team's group with a different `team` value. Source-control the policy AND the bindings — that's how you safely shepherd parameter-shape changes across N bindings (see **IAM REFERENCE.md § Change-Management Caveat**).

Parameterized policy:

```
# policies/team-data-access.yaml
configs:
  - id: team-data-access-policy
    type:
      settings:
        schema: builtin:iam.policy
    config:
      name: team-data-access
      description: "Parameterized: data access scoped to ${bindParam:team}"
      statements:
        - effect: ALLOW
          service: storage
          resource: "*"
          action: read
          condition: "storage:dt.security_context = '${bindParam:team}'"
        - effect: ALLOW
          service: storage
          resource: "*"
          action: write
          condition: "storage:dt.security_context = '${bindParam:team}'"
```


Bindings (one per team, same policy, different parameter values):

```
# bindings/team-bindings.yaml
configs:
  - id: checkout-team-binding
    type:
      settings:
        schema: builtin:iam.binding
      config:
        policyId: "{{ .team-data-access-policy }}"
        groupId: "{{ .dt-checkout-team }}"
        parameters:
          team: "checkout"
  - id: payments-team-binding
    type:
      settings:
        schema: builtin:iam.binding
      config:
        policyId: "{{ .team-data-access-policy }}"
        groupId: "{{ .dt-payments-team }}"
        parameters:
          team: "payments"
```

Adding a new team is a binding (values-file edit), not a policy review. The policy YAML stays unchanged. See **IAM-10: Templated Policy-Group Assignments** for the full Monaco + IAM API integration pattern.

GitOps Workflow

1. **Branch:** Create feature branch for policy change
2. **Edit:** Modify policy YAML
3. **Review:** Submit PR, get peer review
4. **Test:** Deploy to test environment
5. **Approve:** Merge after validation
6. **Deploy:** CI/CD applies to production

Benefits

Benefit	Description
Audit Trail	Git history shows who changed what, when

Benefit	Description
Peer Review	Changes reviewed before apply
Rollback	Easy revert to previous state
Documentation	Policies are self-documenting
Consistency	Same process for all changes

See **AUTOM-03: Monaco** for full implementation details.

9. Policy Assessment Queries

Use these queries to audit policy-related activity.

```
// Track policy changes in audit log
fetch logs, from: now() - 30d
| filter matchesPhrase(log.source, "audit")
| filter matchesPhrase(content, "policy")
| fields timestamp, content
| sort timestamp desc
| limit 50
```

```
// Find access denied events
fetch logs, from: now() - 7d
| filter matchesPhrase(log.source, "audit")
| filter matchesPhrase(content, "denied") or matchesPhrase(content,
"unauthorized")
| fields timestamp, content
| sort timestamp desc
| limit 100
```

```
// Summarize permission-related audit events
fetch logs, from: now() - 7d
| filter matchesPhrase(log.source, "audit")
| filter matchesPhrase(content, "permission") or matchesPhrase(content,
"policy")
| summarize eventCount = count(), by:{content}
| sort eventCount desc
| limit 20
```

Next Steps

With policies defined, complete your access control setup:

Recommended Path

1. **IAM-05: Boundary Design Patterns** - Control what entities users can see
2. **IAM-06: User Lifecycle and Provisioning** - Automate user management
3. **IAM-07: Audit Logging and Compliance** - Monitor policy effectiveness
4. **IAM-10: Templated Policy-Group Assignments** - Reusable parameterized policies for team/environment scoping at scale
5. **IAM-11: Policy Persona Workshop** - Design persona-based policies from stakeholder requirements using a structured 6-goal methodology

Policy Checklist

Before moving on, ensure you have:

- ☐ Understood policy statement syntax
- ☐ Reviewed services and actions reference
- ☐ Designed policies for your team structure
- ☐ Decided on default vs custom policy usage
- ☐ Established testing methodology
- ☐ Considered GitOps for policy management

Summary

In this notebook, you learned:

- Policy fundamentals and scope levels
- Statement syntax: `ALLOW <service>:<resource>:<action>`
- Services and actions reference
- Condition expressions for fine-grained control
- Seven common policy patterns

- When to use default vs custom policies
 - Policy testing and validation approaches
 - GitOps for policy version control
-

References

- [Policy Statements Reference](#)
- [Permission Management](#)
- [Account Policies](#)

References

- [IAM policy statements reference \(DT docs\)](#) — authoritative permission + attribute catalog
 - [Tailored access management, Part 1 — one configuration for all authorization requirements \(Dynatrace News\)](#)
 - [Tailored access management, Part 2 — onboard users to Grail and AppEngine \(Dynatrace News\)](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.