

K8S-01: Kubernetes Monitoring Fundamentals

Series: K8S — Kubernetes Monitoring | **Notebook:** 1 of 14 | **Created:** January 2026
| **Last Updated:** 10/02/2026

Introduction to Kubernetes Observability with Dynatrace

Kubernetes introduces unique observability challenges: ephemeral workloads, dynamic scaling, complex networking, and multi-layer abstractions. Dynatrace provides comprehensive Kubernetes monitoring through the DynaKube operator, which deploys and manages monitoring components automatically.

Table of Contents

- [1. Kubernetes Observability Challenges](#)
 - [2. Dynatrace Monitoring Architecture](#)
 - [3. Entity Model for Kubernetes](#)
 - [4. Data Sources and Signals](#)
 - [5. Key Metrics and Dimensions](#)
 - [6. Your First Kubernetes Queries](#)
 - [7. Next Steps](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Kubernetes monitoring enabled
Kubernetes Cluster	Any distribution (EKS, AKS, GKE, OpenShift, etc.)
Permissions	To run the queries: <code>storage:metrics:read</code> , <code>storage:events:read</code> , <code>storage:smartscape:read</code> (Smartscape nodes) — plus the bucket permissions for the data you query
Knowledge	Basic Kubernetes concepts (pods, deployments, services)

1. Kubernetes Observability Challenges

Kubernetes environments present unique monitoring requirements:

Challenge	Description	Dynatrace Solution
Ephemeral Workloads	Pods come and go constantly	Entity relationships preserved across restarts
Dynamic Scaling	Replicas change based on load	Automatic discovery of new instances
Multi-Layer Stack	Infra → K8s → App complexity	Unified view from cluster to code
Distributed Services	Microservices across namespaces	End-to-end distributed tracing
Resource Constraints	CPU/memory limits and requests	Resource utilization vs. limits monitoring
Network Complexity	Service mesh, ingress, CNI	Network flow and latency analysis

Traditional vs. Cloud-Native Monitoring

Aspect	Traditional	Kubernetes
Identity	IP address, hostname	Labels, selectors, namespaces
Lifecycle	Long-lived servers	Short-lived pods
Configuration	Static files	Dynamic ConfigMaps, Secrets
Networking	Fixed topology	Service discovery, DNS
Scaling	Manual or scheduled	HPA, VPA, KEDA

2. Dynatrace Monitoring Architecture

Dynatrace monitors Kubernetes through multiple components:

DynaKube Operator Components


Component	Purpose	Deployment Mode
OneAgent	Full-stack monitoring (processes, code)	DaemonSet or application-only
ActiveGate	Routing; <code>kubernetes-monitoring</code> capability queries the K8s API and cAdvisor	StatefulSet in cluster
Prometheus Integration	Custom metrics ingestion	Optional

Deployment Modes

Mode	Use Case	OneAgent	Code Modules
<code>cloudNativeFullStack</code>	Full visibility, K8s-native	Privileged DaemonSet	Injected via webhook
<code>classicFullStack</code> (legacy)	Traditional deployment — avoid for new deployments	Privileged DaemonSet	Loaded from host

Mode	Use Case	OneAgent	Code Modules
applicationMonitoring	App-only, no infra	None	Injected via webhook
hostMonitoring	Infra-only	DaemonSet	None

Data Flow



3. Entity Model for Kubernetes

Dynatrace creates entities for each Kubernetes resource and maintains relationships between them.

Kubernetes Node Types (Smartscape)


Queries in this series use **Smartscape** node types (`smartscapeNodes "K8S_..."`). The classic `dt.entity.*` model they replace is shown for reference:

Smartscape node type	Classic entity (dt.entity.*)	Represents
K8S_CLUSTER	kubernetes_cluster	Cluster
K8S_NODE	kubernetes_node	Worker node
K8S_NAMESPACE	cloud_application_namespace	Namespace
K8S_DEPLOYMENT , K8S_STATEFULSET , K8S_DAEMONSET , K8S_REPLICASET , K8S_JOB , K8S_CRONJOB	cloud_application (one classic type for all)	Workloads
K8S_POD	cloud_application_instance	Pod
K8S_SERVICE	kubernetes_service	Kubernetes Service object
SERVICE	service	Detected application service

Smartscape also models objects the classic model never had — K8S_INGRESS , K8S_CONFIGMAP , K8S_PERSISTENTVOLUMECLAIM , K8S_HORIZONTALPODAUTOSCALER , K8S_DYNAKUBE and more.

Dictionary: `fetch dt.semantic_dictionary.models | filter data_object == "smartscape.nodes"` — 23 K8S_* node types and their classic_models , read 10/02/2026.

Entity Relationships



Entity Naming

Resource	Dynatrace Entity Name Pattern
Cluster	Cluster name from kubeconfig
Namespace	namespace - name
Deployment	deployment - name in namespace
Pod	pod - name (ephemeral, tied to PGI)
Service	Auto-detected from traffic patterns

4. Data Sources and Signals

Dynatrace collects multiple signal types from Kubernetes:

Metrics and Events

Collector	What it gathers
ActiveGate with the <code>kubernetes-monitoring</code> capability	Cluster topology and state from the Kubernetes API; node- and container-level metrics via cAdvisor; Kubernetes events
OneAgent (full-stack / host monitoring)	Host and process metrics, code-level traces
Prometheus scraping (optional)	Application metrics from annotated pods

Logs

Log Type	Collected by	Use Case
Container logs (stdout/stderr)	Log module — integrated with OneAgent, or the standalone Kubernetes Log module	Application debugging
Kubernetes events	ActiveGate, from the Kubernetes API — stored as events (<code>event.provider == "KUBERNETES_EVENT"</code>), not logs	Scheduling, scaling, errors
Node / system logs	OneAgent log module	Infrastructure issues

Sources: [Kubernetes platform monitoring \(DT docs\)](#) — “Uses the Kubernetes API and cAdvisor to get node- and container-level metrics and Kubernetes events”; [Kubernetes log monitoring \(DT docs\)](#).

Traces

Trace Source	Coverage
OneAgent auto-instrumentation	Supported languages/frameworks
OpenTelemetry	Custom instrumentation
Service mesh	Istio, Linkerd sidecars

5. Key Metrics and Dimensions

Grail Metric Keys (what DQL queries)

Metric	Grain	Description
<code>dt.kubernetes.container.cpu_usage</code>	Container	CPU used (millicores)
<code>dt.kubernetes.container.memory_working_set</code>	Container	Working-set memory (bytes)
<code>dt.kubernetes.container.cpu_throttled</code>	Container	CPU throttled
<code>dt.kubernetes.container.requests_cpu / .limits_cpu</code>	Container	CPU requests / limits
<code>dt.kubernetes.container.requests_memory / .limits_memory</code>	Container	Memory requests / limits
<code>dt.kubernetes.container.restarts / .oom_kills</code>	Container	Restarts and OOM kills (counts)
<code>dt.kubernetes.node.cpu_allocatable / .memory_allocatable / .pods_allocatable</code>	Node	Allocatable capacity
<code>dt.kubernetes.node.conditions</code>	Node	Node conditions
<code>dt.kubernetes.pods</code> , <code>dt.kubernetes.workloads</code> , <code>dt.kubernetes.nodes</code>	Cluster objects	Object counts
<code>dt.kubernetes.workload.pods_desired</code> , <code>dt.kubernetes.workload.conditions</code>	Workload	Desired pods, conditions
<code>dt.kubernetes.pod.network_received_data / _transmitted_data</code>	Pod	Network bytes

Key list read with `metrics | filter startsWith(metric.key, "dt.kubernetes")` on the validation tenant, 10/02/2026 (24 `dt.kubernetes.*` keys; OneAgent container keys are under `dt.containers.*`).

Classic Metric Keys (Metrics API, Data Explorer)

Metric	Description	Unit
<code>builtin:containers.cpu.usagePercent</code>	CPU usage vs. limit	Percent
<code>builtin:containers.memory.usagePercent</code>	Memory usage vs. limit	Percent
<code>builtin:kubernetes.workload.requests_cpu / limits_cpu</code>	CPU requests / limits	Millicores
<code>builtin:kubernetes.workload.requests_memory / limits_memory</code>	Memory requests / limits	Bytes

These are `builtin:` keys — the classic Metrics namespace, not Grail. They are correct for the classic Metrics API, Data Explorer and classic dashboards, but **DQL reads Grail only**: `metrics | filter startsWith(metric.key, "builtin:")` returns **zero rows over 7 days** on a live tenant (08/11/2026), while `startsWith(metric.key, "dt.kubernetes")` returns 27 keys in the same command. A `timeseries avg(builtin:kubernetes.workload.requests_cpu)` therefore charts nothing and reports no error.

For DQL, use the Grail equivalents — and note the **grain differs, not just the prefix**: Grail emits requests and limits at **container** grain

(`dt.kubernetes.container.requests_cpu`, `.requests_memory`, `.limits_cpu`, `.limits_memory`), summed by you across `k8s.workload.name` to reach a workload figure. There is no `dt.kubernetes.workload.requests_*`. Full Grail key list and the derivation pattern: **K8S-08 §2**.

Request and limit metrics changed what they count in ActiveGate 1.343 / 1.345 — they carry an *effective* pod value: from 1.343 init containers are accounted for, and from 1.345 pod-level `spec.resources` (Kubernetes 1.34+) and pod overhead too, which can move a limit *down*. Reserved figures can step at either upgrade with no workload change; usage metrics are unaffected. This matters whenever you compare a request/limit trend across an upgrade boundary — see K8S-08 §2 before drawing capacity conclusions from one.

6. Your First Kubernetes Queries

Let's explore common queries for Kubernetes monitoring.

```
// Verify Kubernetes clusters are reporting (smartscape topology)
smartscapeNodes "K8S_CLUSTER"
| fields entity.name = name, tags
| sort entity.name asc

// Legacy alternative (deprecated for new content):
// fetch dt.entity.kubernetes_cluster
// | fields entity.name, tags
// | sort entity.name asc
```

```
// Count Kubernetes nodes
smartscapeNodes "K8S_NODE"
| summarize nodeCount = count()

// Legacy alternative:
// fetch dt.entity.kubernetes_node
// | summarize nodeCount = count()
```

```
// List Kubernetes namespaces
smartscapeNodes "K8S_NAMESPACE"
| fields entity.name = name, tags
| sort entity.name asc
| limit 50

// Legacy alternative:
// fetch dt.entity.cloud_application_namespace
// | fields entity.name, tags
// | sort entity.name asc
// | limit 50
```

```
// Container CPU usage – highest consumers
// Sort on a scalar. `sort` on the timeseries array itself compares arrays
// element by element,
// so the order has nothing to do with the average – reduce with arrayAvg()
// first.
timeseries cpu = avg(dt.kubernetes.container.cpu_usage), from:-1h,
  by:{k8s.cluster.name, k8s.namespace.name, k8s.pod.name, k8s.container.name}
| fieldsAdd avgCpuMillicores = round(arrayAvg(cpu), decimals: 0)
| fields k8s.cluster.name, k8s.namespace.name, k8s.pod.name,
k8s.container.name, avgCpuMillicores
| sort avgCpuMillicores desc
| limit 20
```

```
// Container memory usage approaching limits
// "Approaching the limit" needs the limit: compare working set with
// limits_memory per container.
// Containers with no memory limit drop out (limitMiB is null) – they cannot
// be OOM-killed for
// exceeding a limit, only by node pressure.
timeseries {
  used = avg(dt.kubernetes.container.memory_working_set),
  lim = avg(dt.kubernetes.container.limits_memory)
}, from:-1h, by:{k8s.cluster.name, k8s.namespace.name, k8s.pod.name,
k8s.container.name}
| fieldsAdd usedMiB = round(arrayAvg(used) / 1048576, decimals: 0),
  limitMiB = round(arrayAvg(lim) / 1048576, decimals: 0)
| filter limitMiB > 0
| fieldsAdd pctOfLimit = round(100 * usedMiB / limitMiB, decimals: 1)
| fields k8s.cluster.name, k8s.namespace.name, k8s.pod.name,
k8s.container.name, usedMiB, limitMiB, pctOfLimit
| sort pctOfLimit desc
| limit 20
```

```

// Kubernetes events - recent warnings and errors
// Data object corrected 08/12/2026. Kubernetes events are NOT logs. This cell
scraped
// `fetch logs` for `log.source` containing "kubernetes" or for content
substrings like "BackOff" -
// no log.source matches, and kubelet event text is not in the log stream, so
it returned nothing
// while the cluster emitted 231,296 Kubernetes events in the same window.
// They arrive as `fetch events | filter event.provider ==
"KUBERNETES_EVENT`, STRUCTURED:
//   dt.kubernetes.event.reason           Unhealthy · BackOff · Killing ·
FailedScheduling ·
//                                           FailedMount · BackoffLimitExceeded
· EvictionThresholdMet ...
//   dt.kubernetes.event.message           the human-readable text
//   status                                "WARN" for Kubernetes Warning
events, "INFO" for Normal
//   dt.kubernetes.event.involved_object.kind / .name
//   dt.kubernetes.event.count / .first_seen / .last_seen
//   plus k8s.cluster.name · k8s.namespace.name · k8s.pod.name ·
k8s.workload.name · k8s.node.name
// NOTE: event.type is CUSTOM_INFO on every one of these - it is NOT
"Warning". The Warning/Normal
// split is in status ("WARN" / "INFO"). dt.kubernetes.event.important was
"true" on all 204,078
// events over 30 days (10/02/2026), so it separates nothing. Enumerate
reasons with:
//   fetch events, from:-24h | filter event.provider == "KUBERNETES_EVENT"
//   | summarize n = count(), by:{dt.kubernetes.event.reason} | sort n desc
fetch events, from:-1h
| filter event.provider == "KUBERNETES_EVENT"
| filter status == "WARN"
| fields timestamp, k8s.cluster.name, k8s.namespace.name,
dt.kubernetes.event.reason, dt.kubernetes.event.message
| sort timestamp desc
| limit 25

```

```
// Pod restarts - find crashlooping workloads
// Data object corrected 08/12/2026. Kubernetes events are NOT logs. This cell
scraped
// `fetch logs` for `log.source` containing "kubernetes" or for content
substrings like "BackOff" -
// no log.source matches, and kubelet event text is not in the log stream, so
it returned nothing
// while the cluster emitted 231,296 Kubernetes events in the same window.
// They arrive as `fetch events | filter event.provider ==
"KUBERNETES_EVENT`, STRUCTURED:
//   dt.kubernetes.event.reason           Unhealthy · BackOff · Killing ·
FailedScheduling ·
//                                           FailedMount · BackoffLimitExceeded
· EvictionThresholdMet ...
//   dt.kubernetes.event.message           the human-readable text
//   status                                "WARN" for Kubernetes Warning
events, "INFO" for Normal
//   dt.kubernetes.event.involved_object.kind / .name
//   dt.kubernetes.event.count / .first_seen / .last_seen
//   plus k8s.cluster.name · k8s.namespace.name · k8s.pod.name ·
k8s.workload.name · k8s.node.name
// NOTE: event.type is CUSTOM_INFO on every one of these – it is NOT
"Warning". The Warning/Normal
// split is in status ("WARN" / "INFO"). dt.kubernetes.event.important was
"true" on all 204,078
// events over 30 days (10/02/2026), so it separates nothing. Enumerate
reasons with:
//   fetch events, from:-24h | filter event.provider == "KUBERNETES_EVENT"
//   | summarize n = count(), by:{dt.kubernetes.event.reason} | sort n desc
fetch events, from:-24h
| filter event.provider == "KUBERNETES_EVENT"
| filter in(dt.kubernetes.event.reason, {"BackOff", "BackoffLimitExceeded",
"Killing"})
| summarize events = count(), by:{k8s.namespace.name, k8s.pod.name,
dt.kubernetes.event.reason}
| sort events desc
| limit 25
```

7. Next Steps

Now that you understand Kubernetes monitoring fundamentals, proceed to:

Next Notebook	Topic
K8S-02: DynaKube Operator Deployment	Install and configure the operator
K8S-03: GitOps for DynaKube	Manage DynaKube with ArgoCD/Flux
K8S-04: Cluster Health Monitoring	Deep-dive into cluster metrics

Summary

In this notebook, you learned:

- Kubernetes observability challenges and how Dynatrace addresses them
- Dynatrace monitoring architecture (OneAgent, ActiveGate, DynaKube)
- Entity model for Kubernetes resources
- Data sources: metrics, logs, and traces
- Key metrics for container and cluster monitoring
- Basic DQL queries for Kubernetes data using `smartscapeNodes` (modern) with `dt.entity.*` legacy alternatives shown as comments

References

- [Setup on Kubernetes \(DT docs\)](#) — top-level entry point for all K8s monitoring docs
 - [How it works \(DT docs\)](#) — architecture, components, and data flow
 - [Quickstart \(DT docs\)](#) — minimum viable deployment
 - [Reference \(DT docs\)](#) — DynaKube parameters, feature flags, network, security, storage, workload mutation
 - [Dynatrace Operator \(Dynatrace GitHub\)](#) — source, releases, and Helm chart
 - [Effective pod resources \(DT docs\)](#) — how request/limit metrics represent the enforced pod value
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.