

K8S-07: Kubernetes Events and Log Ingestion

Series: K8S — Kubernetes Monitoring | **Notebook:** 7 of 14 | **Created:** January 2026
| **Last Updated:** 10/05/2026

Capturing and Analyzing Kubernetes Events and Logs

Kubernetes events and container logs provide crucial insights for debugging and operational awareness. This notebook covers event monitoring, log ingestion configuration, and analysis patterns in Dynatrace.

Table of Contents

1. [Kubernetes Events Overview](#)
 2. [Event Ingestion Configuration](#)
 3. [Container Log Collection](#)
 4. [OpenPipeline for K8s Logs](#)
 5. [Event Analysis Patterns](#)
 6. [Log Analysis Patterns](#)
 7. [Alerting on Events and Logs](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with log ingestion enabled
DynaKube	ActiveGate with <code>kubernetes-monitoring</code>
Permissions	<code>storage:logs:read</code> , <code>storage:events:read</code> ; Settings write access to change Kubernetes monitoring settings
Knowledge	K8S-01 Fundamentals

1. Kubernetes Events Overview

Event Types


Kubernetes events are first-class objects that record what happened in the cluster.

Type	Description	Examples
Normal	Routine operations	Scheduled, Pulled, Created
Warning	Potential issues	FailedScheduling, BackOff, Unhealthy

Event Sources

Source	Events Generated
Scheduler	Scheduling decisions, failures
Kubelet	Pod lifecycle, probe results
Controller Manager	ReplicaSet scaling, Deployment rollouts
Custom Controllers	CRD reconciliation events

Event Lifecycle



Important: Kubernetes events are short-lived. Dynatrace captures them for long-term storage and analysis. On the validation tenant (10/02/2026) they were stored in the `default_davis_k8s_ops_events` bucket, retention 35 days — check `fetch dt.system.buckets` for your own.

2. Event Ingestion Configuration

ActiveGate Kubernetes Monitoring

Events are read from the Kubernetes API by the ActiveGate with the `kubernetes-monitoring` capability:

```
spec:
  activeGate:
    capabilities:
      - kubernetes-monitoring
      - routing
```

Event Filtering

Which events are ingested is a **Kubernetes monitoring setting** on the cluster's connection (Kubernetes app → cluster → Settings), not a DynaKube field:

Setting	What it does
Monitor events	<i>"All events are monitored unless event filters are specified. All ingested events are subject to licensing by default."</i>
Filter events	<i>"Include only events specified by Events Field Selectors"</i>
Events field selectors	Kubernetes field-selector expressions, for example <code>type=Warning</code> or <code>involvedObject.namespace=checkout</code>
Include important events	<i>"Automatically include all events that are relevant for Davis"</i>

Recommended Configuration

Use Case	Configuration
Full visibility	Monitor events on, no filters
Warning focus	Filter events on, field selector <code>type=Warning</code> , Include important events on
Cost control	As above, plus field selectors limited to the namespaces you need

Sources: [Kubernetes monitoring settings schema \(DT docs\)](#) — *"Automatically include all events that are relevant for Davis"*.

```

// Recent Kubernetes events
// Data object corrected 08/12/2026. Kubernetes events are NOT logs. This cell
scraped
// `fetch logs` for `log.source` containing "kubernetes" or for content
substrings like "BackOff" –
// no log.source matches, and kubelet event text is not in the log stream, so
it returned nothing
// while the cluster emitted 231,296 Kubernetes events in the same window.
// They arrive as `fetch events | filter event.provider ==
"KUBERNETES_EVENT`, STRUCTURED:
//   dt.kubernetes.event.reason           Unhealthy · BackOff · Killing ·
FailedScheduling ·
//                                           FailedMount · BackoffLimitExceeded
· EvictionThresholdMet ...
//   dt.kubernetes.event.message           the human-readable text
//   status                                "WARN" for Kubernetes Warning
events, "INFO" for Normal
//   dt.kubernetes.event.involved_object.kind / .name
//   dt.kubernetes.event.count / .first_seen / .last_seen
//   plus k8s.cluster.name · k8s.namespace.name · k8s.pod.name ·
k8s.workload.name · k8s.node.name
// NOTE: event.type is CUSTOM_INFO on every one of these – it is NOT
"Warning". The Warning/Normal
// split is in status ("WARN" / "INFO"). dt.kubernetes.event.important was
"true" on all 204,078
// events over 30 days (10/02/2026), so it separates nothing. Enumerate
reasons with:
//   fetch events, from:-24h | filter event.provider == "KUBERNETES_EVENT"
//   | summarize n = count(), by:{dt.kubernetes.event.reason} | sort n desc
fetch events, from:-1h
| filter event.provider == "KUBERNETES_EVENT"
| fields timestamp, k8s.cluster.name, k8s.namespace.name,
dt.kubernetes.event.reason, dt.kubernetes.event.message
| sort timestamp desc
| limit 50

```

```

// Warning-class events only
// Data object corrected 08/12/2026. Kubernetes events are NOT logs. This cell
scraped
// `fetch logs` for `log.source` containing "kubernetes" or for content
substrings like "BackOff" –
// no log.source matches, and kubelet event text is not in the log stream, so
it returned nothing
// while the cluster emitted 231,296 Kubernetes events in the same window.
// They arrive as `fetch events | filter event.provider ==
"KUBERNETES_EVENT"`, STRUCTURED:
//   dt.kubernetes.event.reason           Unhealthy · BackOff · Killing ·
FailedScheduling ·
//                                           FailedMount · BackoffLimitExceeded
· EvictionThresholdMet ...
//   dt.kubernetes.event.message           the human-readable text
//   status                                "WARN" for Kubernetes Warning
events, "INFO" for Normal
//   dt.kubernetes.event.involved_object.kind / .name
//   dt.kubernetes.event.count / .first_seen / .last_seen
//   plus k8s.cluster.name · k8s.namespace.name · k8s.pod.name ·
k8s.workload.name · k8s.node.name
// NOTE: event.type is CUSTOM_INFO on every one of these – it is NOT
"Warning". The Warning/Normal
// split is in status ("WARN" / "INFO"). dt.kubernetes.event.important was
"true" on all 204,078
// events over 30 days (10/02/2026), so it separates nothing. Enumerate
reasons with:
//   fetch events, from:-24h | filter event.provider == "KUBERNETES_EVENT"
//   | summarize n = count(), by:{dt.kubernetes.event.reason} | sort n desc
fetch events, from:-1h
| filter event.provider == "KUBERNETES_EVENT"
| filter status == "WARN"
| summarize events = count(), by:{dt.kubernetes.event.reason,
k8s.namespace.name}
| sort events desc
| limit 25

```

3. Container Log Collection

Log Collection Methods

Method	Source	Configuration
OneAgent Log module (full-stack)	Container stdout/stderr	Included in the host OneAgent; add <code>spec.logMonitoring: {}</code>
Kubernetes Log module (no host OneAgent)	Container stdout/stderr	DaemonSet deployed by the Operator from <code>spec.logMonitoring</code>
Fluent Bit / other forwarders	Any file or stream	Log ingest API or OTLP endpoint

What the Log modules collect

"It only captures logs that are written to the container's stdout/stderr streams." Logs an application writes to files inside the container need a forwarder or a custom log source.

Log Attributes

Attribute	Source	Example
<code>k8s.namespace.name</code>	Container metadata	<code>checkout</code>
<code>k8s.pod.name</code>	Container metadata	<code>checkout-api-abc123</code>
<code>k8s.container.name</code>	Container metadata	<code>api</code>
<code>k8s.workload.name</code>	Container metadata	<code>checkout-api</code>
<code>k8s.cluster.name</code>	Container metadata	<code>prod-eu-1</code>
<code>loglevel</code>	Parsed from content	<code>ERROR</code> , <code>WARN</code> , <code>INFO</code>

DynaKube Log Configuration

```
spec:
  logMonitoring: {}
  templates:
    logMonitoring:      # only for the standalone Kubernetes Log module
    imageRef:
      repository: public.ecr.aws/dynatrace/dynatrace-logmodule
      tag: &lt;tag&gt;
```

Sources: [Kubernetes log monitoring \(DT docs\)](#) — *"It only captures logs that are written to the container's stdout/stderr streams."*

```
// Container logs by namespace
fetch logs, from:-1h
| filter isNotNull(k8s.namespace.name)
| summarize logCount = count(), by:{k8s.namespace.name}
| sort logCount desc
| limit 15
```

```
// Error logs with Kubernetes context
fetch logs, from:-1h
| filter loglevel == "ERROR" or loglevel == "SEVERE"
| filter isNotNull(k8s.namespace.name)
| fields timestamp, k8s.namespace.name, k8s.pod.name, content
| sort timestamp desc
| limit 30
```

4. OpenPipeline for K8s Logs

Log Processing Pipeline

OpenPipeline processes logs between ingest and storage. A record arrives at an ingest source, **dynamic routing** sends it to a pipeline whose matcher it satisfies, the pipeline's **processing** stage transforms it, and the **storage** stage assigns its bucket. Pipelines are configured in the OpenPipeline settings (UI, Settings API or Monaco), not as YAML in the cluster.

Common Processing Rules

Processor	Use Case	Example
DQL	Extract fields	<code>parse content,</code> <code>"JSON:parsed"</code>
Fields rename / remove	Tidy fields	Drop a high-cardinality attribute
Masking (DQL <code>replacePattern</code>)	Remove sensitive data	Card numbers, tokens
Drop record	Drop logs	Debug logs from noisy namespaces
Storage assignment	Choose the bucket	Retention by namespace or team

Example: route Kubernetes logs and parse JSON

- Routing matcher: `isNull(k8s.namespace.name)`
- Processing — DQL processor, matcher `startsWith(content, "{") :`

```
parse content, "JSON:parsed"  
| fieldsFlatten parsed
```

Example: drop debug logs

- Processing — Drop record processor, matcher: `loglevel == "DEBUG"`

Building, testing and ordering pipelines is covered in the OPLOGS series.

5. Event Analysis Patterns

Key Event Reasons to Monitor

Reason	Meaning	Action
FailedScheduling	Pod can't be scheduled	Check resource availability

Reason	Meaning	Action
FailedMount	Volume mount failed	Check PV/PVC config
BackOff	Container restart backoff	Check logs, fix crash
Unhealthy	Probe failed	Check probe config, app health
Evicted	Pod evicted from node	Check node pressure
OOM kill	Container exceeded its memory limit	Not an event reason — read the <code>dt.kubernetes.container.oom_kills</code> metric (K8S-04 §5)

```

// Failed scheduling events
// Data object corrected 08/12/2026. Kubernetes events are NOT logs. This cell
scraped
// `fetch logs` for `log.source` containing "kubernetes" or for content
substrings like "BackOff" –
// no log.source matches, and kubelet event text is not in the log stream, so
it returned nothing
// while the cluster emitted 231,296 Kubernetes events in the same window.
// They arrive as `fetch events | filter event.provider ==
"KUBERNETES_EVENT"`, STRUCTURED:
//   dt.kubernetes.event.reason           Unhealthy · BackOff · Killing ·
FailedScheduling ·
//                                           FailedMount · BackoffLimitExceeded
· EvictionThresholdMet ...
//   dt.kubernetes.event.message           the human-readable text
//   status                                "WARN" for Kubernetes Warning
events, "INFO" for Normal
//   dt.kubernetes.event.involved_object.kind / .name
//   dt.kubernetes.event.count / .first_seen / .last_seen
//   plus k8s.cluster.name · k8s.namespace.name · k8s.pod.name ·
k8s.workload.name · k8s.node.name
// NOTE: event.type is CUSTOM_INFO on every one of these – it is NOT
"Warning". The Warning/Normal
// split is in status ("WARN" / "INFO"). dt.kubernetes.event.important was
"true" on all 204,078
// events over 30 days (10/02/2026), so it separates nothing. Enumerate
reasons with:
//   fetch events, from:-24h | filter event.provider == "KUBERNETES_EVENT"
//   | summarize n = count(), by:{dt.kubernetes.event.reason} | sort n desc
fetch events, from:-7d
| filter event.provider == "KUBERNETES_EVENT"
| filter dt.kubernetes.event.reason == "FailedScheduling"
| fields timestamp, k8s.namespace.name,
dt.kubernetes.event.involved_object.name, dt.kubernetes.event.message
| sort timestamp desc
| limit 25

```

```

// Restart back-off events (BackOff, BackoffLimitExceeded)
// CrashLoopBackOff is not a reason of its own: it appears in the message text
of BackOff events.
// For restart counts, use the dt.kubernetes.container.restarts metric (K8S-05
$2).
// Data object corrected 08/12/2026. Kubernetes events are NOT logs. This cell
scraped
// `fetch logs` for `log.source` containing "kubernetes" or for content
substrings like "BackOff" –
// no log.source matches, and kubelet event text is not in the log stream, so
it returned nothing
// while the cluster emitted 231,296 Kubernetes events in the same window.
// They arrive as `fetch events | filter event.provider ==
"KUBERNETES_EVENT"`, STRUCTURED:
//   dt.kubernetes.event.reason           Unhealthy · BackOff · Killing ·
FailedScheduling ·
//                                           FailedMount · BackoffLimitExceeded
· EvictionThresholdMet ...
//   dt.kubernetes.event.message           the human-readable text
//   status                                "WARN" for Kubernetes Warning
events, "INFO" for Normal
//   dt.kubernetes.event.involved_object.kind / .name
//   dt.kubernetes.event.count / .first_seen / .last_seen
//   plus k8s.cluster.name · k8s.namespace.name · k8s.pod.name ·
k8s.workload.name · k8s.node.name
// NOTE: event.type is CUSTOM_INFO on every one of these – it is NOT
"Warning". The Warning/Normal
// split is in status ("WARN" / "INFO"). dt.kubernetes.event.important was
"true" on all 204,078
// events over 30 days (10/02/2026), so it separates nothing. Enumerate
reasons with:
//   fetch events, from:-24h | filter event.provider == "KUBERNETES_EVENT"
//   | summarize n = count(), by:{dt.kubernetes.event.reason} | sort n desc
fetch events, from:-24h
| filter event.provider == "KUBERNETES_EVENT"
| filter in(dt.kubernetes.event.reason, {"BackOff", "BackoffLimitExceeded"})
| summarize events = count(), by:{k8s.namespace.name,
dt.kubernetes.event.involved_object.name}
| sort events desc
| limit 25

```

```

// Volume mount and storage failures
// Data object corrected 08/12/2026. Kubernetes events are NOT logs. This cell
scraped
// `fetch logs` for `log.source` containing "kubernetes" or for content
substrings like "BackOff" –
// no log.source matches, and kubelet event text is not in the log stream, so
it returned nothing
// while the cluster emitted 231,296 Kubernetes events in the same window.
// They arrive as `fetch events | filter event.provider ==
"KUBERNETES_EVENT`, STRUCTURED:
//   dt.kubernetes.event.reason           Unhealthy · BackOff · Killing ·
FailedScheduling ·
//                                           FailedMount · BackoffLimitExceeded
· EvictionThresholdMet ...
//   dt.kubernetes.event.message           the human-readable text
//   status                                "WARN" for Kubernetes Warning
events, "INFO" for Normal
//   dt.kubernetes.event.involved_object.kind / .name
//   dt.kubernetes.event.count / .first_seen / .last_seen
//   plus k8s.cluster.name · k8s.namespace.name · k8s.pod.name ·
k8s.workload.name · k8s.node.name
// NOTE: event.type is CUSTOM_INFO on every one of these – it is NOT
"Warning". The Warning/Normal
// split is in status ("WARN" / "INFO"). dt.kubernetes.event.important was
"true" on all 204,078
// events over 30 days (10/02/2026), so it separates nothing. Enumerate
reasons with:
//   fetch events, from:-24h | filter event.provider == "KUBERNETES_EVENT"
//   | summarize n = count(), by:{dt.kubernetes.event.reason} | sort n desc
fetch events, from:-7d
| filter event.provider == "KUBERNETES_EVENT"
| filter in(dt.kubernetes.event.reason, {"FailedMount", "FailedAttachVolume",
"FreeDiskSpaceFailed", "EvictionThresholdMet"})
| fields timestamp, k8s.namespace.name, dt.kubernetes.event.reason,
dt.kubernetes.event.message
| sort timestamp desc
| limit 25

```

```

// Event frequency by reason (last 24h)
// Data object corrected 08/12/2026. Kubernetes events are NOT logs. This cell
scraped
// `fetch logs` for `log.source` containing "kubernetes" or for content
substrings like "BackOff" –
// no log.source matches, and kubelet event text is not in the log stream, so
it returned nothing
// while the cluster emitted 231,296 Kubernetes events in the same window.
// They arrive as `fetch events | filter event.provider ==
"KUBERNETES_EVENT`, STRUCTURED:
//   dt.kubernetes.event.reason           Unhealthy · BackOff · Killing ·
FailedScheduling ·
//                                           FailedMount · BackoffLimitExceeded
· EvictionThresholdMet ...
//   dt.kubernetes.event.message           the human-readable text
//   status                                "WARN" for Kubernetes Warning
events, "INFO" for Normal
//   dt.kubernetes.event.involved_object.kind / .name
//   dt.kubernetes.event.count / .first_seen / .last_seen
//   plus k8s.cluster.name · k8s.namespace.name · k8s.pod.name ·
k8s.workload.name · k8s.node.name
// NOTE: event.type is CUSTOM_INFO on every one of these – it is NOT
"Warning". The Warning/Normal
// split is in status ("WARN" / "INFO"). dt.kubernetes.event.important was
"true" on all 204,078
// events over 30 days (10/02/2026), so it separates nothing. Enumerate
reasons with:
//   fetch events, from:-24h | filter event.provider == "KUBERNETES_EVENT"
//   | summarize n = count(), by:{dt.kubernetes.event.reason} | sort n desc
fetch events, from:-24h
| filter event.provider == "KUBERNETES_EVENT"
| summarize events = count(), by:{dt.kubernetes.event.reason, status}
| sort events desc
| limit 25

```

6. Log Analysis Patterns

Error Log Investigation

```
// Pattern: Find errors with full context
fetch logs, from:-1h
// status == "ERROR" also counts SEVERE, CRITICAL and FATAL logs; loglevel ==
"ERROR" misses them
| filter status == "ERROR"
| filter k8s.namespace.name == "checkout"
| fields timestamp, k8s.pod.name, content
| sort timestamp desc
| limit 50
```

Log Volume Analysis

```
// Pattern: Identify noisy pods
fetch logs, from:-1h
| summarize count = count(), by:{k8s.pod.name}
| sort count desc
| limit 10
```

Exception Tracking

```
// Pattern: Find stack traces
fetch logs, from:-1h
| filter matchesPhrase(content, "Exception") or matchesPhrase(content,
"Traceback")
| fields timestamp, k8s.namespace.name, content
| sort timestamp desc
| limit 20
```

```
// Log volume by pod (find noisy pods)
fetch logs, from: now() - 1h
| filter isNotNull(k8s.pod.name)
| summarize logCount = count(), by:{k8s.pod.name}
| sort logCount desc
| limit 15
```

```
// Exception and error messages
fetch logs, from:-1h
| filter matchesPhrase(content, "Exception") or matchesPhrase(content,
"error") or matchesPhrase(content, "failed")
| filter isNotNull(k8s.namespace.name)
| fields timestamp, k8s.namespace.name, k8s.pod.name, content
| sort timestamp desc
| limit 30
```

```
// Log level distribution by namespace
fetch logs, from: now() - 1h
| filter isNotNull(k8s.namespace.name) and isNotNull(loglevel)
| summarize count = count(), by:{k8s.namespace.name, loglevel}
| sort count desc
| limit 30
```

7. Alerting on Events and Logs

Event-Based Alerts

Alert	Condition	Severity
Failed Scheduling	FailedScheduling events > 5 in 10 min	Warning
Crash Loop	BackOff events (reason BackOff)	Warning
OOM Kills	<code>dt.kubernetes.container.oom_kills</code> > 0 (a metric, not an event)	Critical
Volume Failures	FailedMount events	Critical

Log-Based Alerts

Alert	Condition	Severity
Error Spike	Error log count > baseline	Warning
Critical Errors	Specific error patterns	Critical
No Logs	Log volume drops to 0	Warning

Custom alerts on logs and events

Build the alert as a **custom alert on a DQL query** — a `makeTimeseries` count of matching log records or Kubernetes events — with a static threshold or an adaptive baseline. Where you create it depends on your tenant version: *"Starting with Dynatrace version 1.344, custom alerts have moved to Settings. Because Anomaly Detection is deprecated, we highly recommend that you use Settings to access your existing configurations and create new ones."* SaaS 1.344 rolls out to tenants in stages; on earlier versions the **Anomaly Detection** app is where custom alerts are created. ALERT-02 covers choosing the detector, and ALERT-03 routing the problem it raises.

Sources: [Anomaly Detection \(DT docs\)](#) — *"Starting with Dynatrace version 1.344, custom alerts have moved to Settings."*

Next Steps

With event and log monitoring configured, proceed to:

Next Notebook	Topic
K8S-08: DQL for Kubernetes	Advanced query patterns
K8S-09: Troubleshooting	Debugging K8s monitoring

Summary

In this notebook, you learned:

- Kubernetes event types and sources
- Event ingestion configuration via DynaKube
- Container log collection methods
- OpenPipeline for log processing
- Event analysis patterns for common issues
- Log analysis patterns for debugging

- [Alerting strategies for events and logs](#)
-

References

- [Kubernetes log monitoring \(DT docs\)](#)
 - [Set up Dynatrace on Kubernetes \(DT docs\)](#)
 - [Logs \(DT docs\)](#)
 - [Log processing with OpenPipeline \(DT docs\)](#)
 - [OpenPipeline \(DT docs\)](#)
 - [Kubernetes app — events view \(DT docs\)](#)
 - [Dynatrace Operator releases \(Dynatrace GitHub\)](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.