M2S-07: Step 7 — Enable: User Enablement and Communication

> **Series:** M2S — Managed to SaaS Migration | **Notebook:** 7 of 9 | **Phase:** Run | **Step:** Enable | **Created:** March 2026 | **Last Updated:** 10/24/2026

A successful migration is measured not by the technical cutover but by whether every team in the organization can use the new platform effectively. Step 7 focuses on communication, training, documentation, and establishing the support structures that ensure adoption. Without deliberate enablement, teams will struggle with new URLs, unfamiliar interfaces, and unanswered questions — undermining the value of the migration.

> **M2S Migration Journey — 3 Phases / 9 Steps**
>
> **Plan:** 1. Discover | 2. Strategize | 3. Design
>
> **Upgrade:** 4. Prepare | 5. Execute | 6. Integrate
>
> **Run: 7. Enable** | 8. Expand | 9. Optimize

Table of Contents

1. Communication Plan
2. Persona-Based Training
3. Documentation Updates
4. Self-Service Resources
5. Support Channels
6. Measuring Enablement Success
7. Step Completion Checklist

Prerequisites

Requirement	Details
Steps 1–6 Complete	Discovery, strategy, design, preparation, execution, and integration phases finished
SaaS Tenant Stable	All OneAgents reporting, integrations reconnected
IAM Configured	User accounts, groups, and policies set up in SaaS
Stakeholder List	Complete inventory of Dynatrace users by role and team
Communication Channels	Access to organizational email, Slack/Teams, intranet

1. Communication Plan

The migration affects every Dynatrace user in the organization. A structured communication plan ensures no team is surprised, confused, or left behind. Tailor the message to the audience — executives need a different message than SRE engineers.

Communication Timeline

When	Audience	Message	Channel
Pre-migration (2–4 weeks before)	All stakeholders	Migration timeline, what changes, what stays the same, how to prepare	Email + town hall
During migration	Platform team	Wave status updates, issues found, resolution progress	Slack/Teams war room
Day of cutover	All users	New URLs, login process, where to get help	Email + Slack/Teams announcement
Post-migration (+1 week)	All users	Feature highlights, training schedule, FAQ link	Email + intranet
Post-migration (+2 weeks)	All users	Adoption metrics, success stories, remaining training sessions	Email

When	Audience	Message	Channel
Post-migration (+30 days)	Management	Migration summary, adoption dashboard, next steps (Step 9)	Executive briefing

Audience-Specific Messaging

Each audience cares about different things. Craft messages that answer their specific questions:

Audience	Key Questions	Message Focus
Executives	What's the ROI? Is it done?	Business value, cost savings, reduced operational burden
Platform Admins	How do I manage it? What's different?	New admin workflows, IAM policies, Settings 2.0, Grail buckets
SRE/Operations	Where are my alerts? How do I troubleshoot?	New URLs, Notebooks replacing USQL, Workflow-based alerting
Developers	Does my pipeline still work? Where's my data?	API endpoint changes, tracing, log analysis, deployment markers
Security Team	Is it compliant? Where are the audit logs?	Application security, audit logging, compliance dashboards

Communication Templates

Pre-Migration Email (Example)

```
Subject: Dynatrace Platform Migration — Moving to SaaS [Timeline]

Team,

We are migrating our Dynatrace monitoring platform from Managed (self-hosted)
to SaaS (cloud-hosted). This migration brings improved performance, automatic
updates, and access to new capabilities including Grail data lakehouse,
Notebooks, and the AutomationEngine.

What changes:
- New URL: https://{tenant}.apps.dynatrace.com (replaces your current
bookmark)
- Login: SSO via your corporate identity provider
- Query language: DQL replaces USQL for advanced queries

What stays the same:
- Your dashboards, management zones, and alerting rules are migrated
- OneAgent monitoring continues without interruption
- Historical data stays in Managed — it does not move to SaaS, so keep the
  Managed environment readable until it is decommissioned

Timeline: [Insert dates]
Training: [Insert schedule link]
Questions: [Insert support channel]
```

Post-Migration Email (Example)

```
Subject: Dynatrace Migration Complete — Your New Platform is Live

Team,

The Dynatrace migration to SaaS is complete. Your new environment is live at:

  https://{tenant}.apps.dynatrace.com

Action required:
1. Update your bookmarks to the new URL
2. Log in using SSO (same corporate credentials)
3. Verify your dashboards and alerts are working

Resources:
- Quick Start Guide: [link]
- FAQ: [link]
- Support channel: #dynatrace-migration-support
- Training schedule: [link]

If anything looks wrong, report it in #dynatrace-migration-support
immediately.
```

Verify User Access to SaaS

Before sending the post-migration announcement, confirm that users can actually log in and see data:

```
// Verify real user traffic is arriving on the SaaS tenant after cutover.
//
// Corrected 08/27/2026. This cell previously queried
`dt.entity.user_session_synthetic`,
// which DOES NOT EXIST — it raises ENTITY_DATA_OBJECT_UNDEFINED, returns
ok:true with
// zero rows, and so always printed "No user sessions — verify IAM
configuration" on
// every environment. Controls at the time: synthetic_test 117,
multiprotocol_monitor 4,
// user.sessions 3,262 — so the zero was the identifier, not an empty tenant.
//
// It also conflated two questions. User sessions are RUM data; they cannot
tell you
// whether IAM groups are configured. Group membership is an Account
Management concern
// — see IAM-04 and IAM-05 for that check.
//
// Corrected 10/05/2026. Synthetic monitors create user sessions too, so a
bare count()
// can report "real users arriving" on a tenant with no real-user RUM at all —
on the
// validation tenant all 4,702 sessions in 24 h were synthetic. Split on
dt.rum.user_type
// (values real_user / robot / synthetic in the semantic dictionary).
fetch user.sessions, from:-24h
| summarize {realUserSessions = countIf(dt.rum.user_type == "real_user"),
             syntheticSessions = countIf(dt.rum.user_type == "synthetic"),
             robotSessions = countIf(dt.rum.user_type == "robot")}
| fieldsAdd status = if(realUserSessions &gt; 0,
                        then: "Real user sessions arriving on SaaS",
                        else: if(syntheticSessions &gt; 0,
                                 then: "Only synthetic sessions — real-user
RUM not arriving yet",
                                 else: "No user sessions — check RUM
enablement and OneAgent/RUM JS injection"))
```

2. Persona-Based Training

One-size-fits-all training does not work. Each persona interacts with Dynatrace differently, so training must be tailored to their workflows. Deliver training within the first two weeks of migration to prevent users from forming bad habits or giving up on the platform.

Training Matrix

Persona	Training Focus	Key Topics	Duration	Format
Platform Admins	SaaS administration	IAM policies, Settings 2.0, Grail bucket management, OpenPipeline configuration, token management	2 hours	Hands-on workshop
SRE/Operations	Monitoring and alerting	Notebooks, DQL queries, Workflow-based alerting, problem analysis, SLO management	2 hours	Hands-on workshop
Developers	Observability integration	Distributed tracing in Notebooks, log analysis with DQL, deployment markers, release tracking	1.5 hours	Hands-on workshop
Management	Executive dashboards	Dashboard navigation, SLO reporting, business analytics, sharing and scheduling	1 hour	Demo + walkthrough
Security Team	Security features	Application Security module, vulnerability analysis, audit log queries, compliance dashboards	1.5 hours	Hands-on workshop

What Changed: Managed vs. SaaS Key Differences

Training must address the specific changes users will encounter:

Feature	Managed Experience	SaaS Experience
URL	`https://{managed-server}/e/{env-id}/`	`https://{tenant}.apps.dynatrace.com`

Feature	Managed Experience	SaaS Experience
Login	Local users or LDAP	SSO via identity provider
Navigation	Classic UI	Unified platform UI with app launcher
Query language	USQL (limited)	DQL (full Grail access)
Data analysis	Custom charts, data explorer	Notebooks (interactive, shareable)
Alerting	Notification rules + alerting profiles	Workflows (AutomationEngine)
Configuration	Classic settings pages and Configuration API	Settings 2.0 (schema-based) and the Settings app
Data storage	Elasticsearch + Cassandra	Grail data lakehouse
Updates	Manual cluster updates (downtime)	Automatic, zero-downtime updates

DQL Transition Guide

The biggest workflow change for power users is the transition from USQL to DQL. Include these common translations in training:

USQL Pattern	DQL Equivalent
`SELECT * FROM usersession`	`fetch user.sessions, from:-24h` — the time range is mandatory in DQL
`WHERE country = 'US'`	`| filter country == "..."` (classic RUM on Grail) or `| filter geo.country.iso_code == "US"` (New RUM)
`GROUP BY browserFamily`	`| summarize c = count(), by:{browserFamily}` (classic) or `by:{browser.name}` (New RUM)
`ORDER BY duration DESC`	`| sort duration desc`
`LIMIT 10`	`| limit 10`

Which column applies depends on whether the tenant runs classic RUM on Grail or New RUM — the field names differ. **WEBRUM-09** has the full grammar and field mapping for both, including the fields with no New RUM equivalent.

> **Tip:** DQL is far more powerful than USQL. Frame the transition as an upgrade, not a disruption. Teams that learn DQL gain access to all Grail data sources (logs, spans, events, metrics, entities) through a single query language.

Hands-On Training: First DQL Queries

Use these queries in training sessions to give users immediate hands-on experience with the SaaS platform:

```
// Training query 1: Explore recent log data (replaces USQL session queries)
fetch logs, from:-1h
| summarize logCount = count(), by:{loglevel}
| sort logCount desc
```

```
// Training query 2: View active detected problems
fetch dt.davis.problems, from:-24h
| filter event.status == "ACTIVE"
| fieldsKeep display_id, event.name, event.category, timestamp
| sort timestamp desc
| limit 10
```

```
// Training query 3: Host CPU overview (replaces Managed data explorer)
timeseries avgCpu = avg(dt.host.cpu.usage), from:-1h, by:{dt.entity.host}
| fieldsAdd avgCpuValue = arrayAvg(avgCpu)
| sort avgCpuValue desc
| limit 10
```

```
// Training query 4: Service response time analysis (distributed tracing)
// Dividing a duration by 1ms turns it into a plain number of milliseconds
fetch spans, from:-1h
| filter span.kind == "server"
| summarize {avgDurationMs = avg(duration) / 1ms, p95DurationMs =
percentile(duration, 95) / 1ms, requestCount = count()}, by:
{dt.entity.service}
| sort p95DurationMs desc
| limit 10
```

3. Documentation Updates

Every internal document that references Dynatrace must be updated. Stale documentation with Managed URLs, old login procedures, or outdated screenshots will confuse users and generate unnecessary support requests.

Documentation Inventory

Document	Updates Required
Runbooks	New SaaS URLs, updated troubleshooting procedures, revised escalation paths
Architecture diagrams	New network topology showing SaaS endpoints, ActiveGate placement, data flow
Onboarding guides	SaaS-specific login instructions, navigation walkthrough, first-day tasks
Disaster recovery plans	Updated recovery procedures (SaaS manages infrastructure; focus shifts to data and configuration backup)
API documentation	New base URLs (`{tenant}.live.dynatrace.com`), token management procedures
Migration records	Archive SaaS Upgrade Assistant deploy result CSVs for audit trail
Security documentation	Updated SSO configuration, token policies, audit log access
ITSM integration guides	Updated webhook URLs, ServiceNow/Jira connector configuration

URL Replacement Checklist

Search all documentation for these patterns and replace:

Find	Replace With
`https://{managed-server}/e/{env-id}/`	`https://{tenant}.apps.dynatrace.com/`

Find	Replace With
`https://{managed-server}/e/{env-id}/api/v2/`	`https://{tenant}.live.dynatrace.com/api/v2/`
`https://{managed-server}/e/{env-id}/api/config/v1/`	`https://{tenant}.live.dynatrace.com/api/config/v1/`
`https://{managed-server}/api/cluster/v2/`	N/A — rewrite to use Account Management API
References to "Managed cluster"	"SaaS environment"
References to "cluster node"	Remove (SaaS manages infrastructure)

> **Important:** Do not just replace URLs. Review the surrounding context. Some procedures (e.g., cluster restarts, Cassandra maintenance) no longer apply in SaaS and should be removed entirely.
>
> **The `/api/config/v1/` row is a stopgap, not a destination.** The rewrite above is correct for Managed → SaaS. But most of `/api/config/v1` stops answering when the tenant later upgrades to the latest Dynatrace, so documentation repointed at it will need a second pass. If you are editing these pages anyway, flag every `/api/config/v1/` reference for a Gen3 follow-up — or resolve it now against AUTOM-02's schema catalog and rewrite straight to the Settings 2.0 or platform equivalent. Same for `/api/v2/` procedures that read entities, metrics, logs, events or tags; the ingest endpoints are unaffected.

4. Self-Service Resources

Not every user will attend training. Self-service resources let users learn at their own pace and serve as a reference after training is complete. Create and share the following:

Quick Start Guide

A one-page guide covering the first five minutes on the new platform:

1. **Log in** — Navigate to `https://{tenant}.apps.dynatrace.com` and authenticate via SSO
2. **Navigate** — Use the app launcher (left sidebar) to find Dashboards, Notebooks, Problems, Services
3. **Run your first query** — Open Notebooks > New Notebook > Add DQL cell > paste: `fetch logs, from:-1h | summarize count = count(), by:{loglevel} | sort count desc`
4. **Find your dashboard** — Go to Dashboards and search by name
5. **Check alerts** — Go to Problems to see active Dynatrace-detected issues

DQL Cheat Sheet

Common queries that replace everyday Managed workflows:

Task	DQL Query
Find error logs	`fetch logs, from:-1h | filter status == "ERROR" | limit 50`
List active problems	`fetch dt.davis.problems, from:-24h | filter event.status == "ACTIVE"`
Check host CPU	`timeseries avg(dt.host.cpu.usage), from:-1h, by:{dt.entity.host}`
View service requests	`fetch spans, from:-1h | filter span.kind == "server" | summarize count(), by:{dt.entity.service}`
Count log volume	`fetch logs, from:-24h | summarize count = count(), by:{log.source} | sort count desc`
Find slow transactions	`fetch spans, from:-1h | filter span.kind == "server" | filter duration > 5s | limit 20`

Frequently Asked Questions

Publish answers to the questions you know users will ask:

Question	Answer
Where do I log in?	`https://{tenant}.apps.dynatrace.com` — use your corporate SSO credentials

Question	Answer
What happened to my dashboards?	Migrated via SaaS Upgrade Assistant — find them in Dashboards (use search)
How do I query data now?	DQL in Notebooks replaces USQL. See the DQL Cheat Sheet above
Where are my management zones?	Migrated by the Upgrade Assistant where your archive carried them — find them by searching *Management zones* in Settings. For data access on the new platform, see the IAM and ORGNZ series
Why do problem counts look different?	Dynatrace Intelligence is relearning baselines on the new SaaS environment — in community practice roughly one to four weeks depending on the signal. FAQ-25 § 5 has a query that measures each host's history depth
Can I still use the old Managed URL?	Until decommission, yes — it still holds the history for every migrated host, but no new data for them. After decommission, no
Where did USQL go?	USQL is not available in SaaS. Use DQL in Notebooks — it is more powerful and covers all data sources
How do I create a new API token?	Go to Account Management > Identity & Access Management > OAuth clients, or use Settings > Access Tokens for environment tokens
Why does the UI look different?	SaaS uses the unified platform UI. The app launcher on the left provides access to all features
Who do I contact for help?	Use the #dynatrace-migration-support channel (Slack/Teams) or attend weekly office hours

Video Walkthroughs

Record short (5–10 minute) screen recordings covering:

Topic	Audience	Content
SaaS login and navigation	All users	SSO login, app launcher, finding dashboards

Topic	Audience	Content
Creating your first Notebook	SRE/Developers	New notebook, DQL cells, visualization, sharing
Workflow-based alerting	SRE/Operations	Creating a workflow trigger, conditional routing, Slack notification
IAM and token management	Platform Admins	Groups, policies, environment tokens, OAuth clients
Dashboard management	Management	Finding dashboards, sharing, scheduling PDF exports

5. Support Channels

Establish clear support channels before the migration goes live. Users need to know exactly where to go when they have questions. Ambiguity leads to frustration and shadow IT workarounds.

Support Channel Matrix

Channel	Purpose	Owner	Response Time
Slack/Teams channel (`#dynatrace-migration-support`)	Day-to-day questions, quick troubleshooting	Platform team	< 2 hours during business hours
Office hours (weekly 30-min drop-in)	Live Q&A, screen sharing, hands-on help	Platform team	Scheduled
Jira/ServiceNow	Formal requests: access issues, configuration changes, feature requests	SRE team	Per SLA
Dynatrace University	Self-paced learning, certifications	Individual	Self-service
Dynatrace Community	Peer support, best practices, shared queries	Community	Best-effort

Channel	Purpose	Owner	Response Time
Dynatrace Support	Platform issues, bugs, outage investigation	Dynatrace	Per contract SLA

Escalation Path

Define a clear escalation path so support requests do not get stuck:

```
User question
  → Slack/Teams channel (platform team triages)
    → If access issue: IAM admin resolves
    → If configuration issue: Platform team resolves
    → If platform bug: Dynatrace Support ticket
    → If feature request: Backlog (Jira/ServiceNow)
```

Transition Support Timeline

Support intensity should decrease over time as users become self-sufficient:

Period	Support Level	Activities
Week 1–2	High (hyper-care)	Daily office hours, immediate Slack response, proactive check-ins
Week 3–4	Medium	Bi-weekly office hours, same-day Slack response
Month 2–3	Standard	Weekly office hours, standard SLA response
Month 4+	Business as usual	Monthly office hours, standard support channels

6. Measuring Enablement Success

Enablement is not complete when training is delivered — it is complete when users are actively and effectively using the platform. Measure adoption with concrete metrics and set targets that demonstrate real usage.

Enablement Metrics

Metric	Target	How to Measure
Active users (weekly)	≥ Managed baseline within 30 days	Platform audit events (`dt.system.events`)
Notebooks created	10+ in the first month	Notebook listing / document count
DQL queries executed	Increasing week-over-week trend	Query audit (if enabled)
Support tickets	Decreasing trend after week 2	ITSM system
Training attendance	80%+ of target personas	Training records
Dashboard views	≥ Managed baseline	Dashboard analytics
Workflow automations created	At least 5 within first month	Workflow listing
Managed environment logins	Declining to zero	Managed access logs

Tracking Active Users

Monitor platform adoption using audit events. On SaaS these are platform audit events in `dt.system.events` (`event.kind == "AUDIT_EVENT"`), not records in the generic `events` table. Grail-based audit logging is on by default. Reading it needs `storage:system:read` for `AUDIT_EVENT` records and `storage:buckets:read` on the `dt_system_events` bucket:

> **Sources:** Audit logs (DT docs) — *"Grail-based audit logging is active by default in all Dynatrace environments."*

```
// Track daily active users over the past 7 days (sign-ins recorded in the platform audit log)
fetch dt.system.events, from:-7d
| filter event.kind == "AUDIT_EVENT" and event.provider == "API_GATEWAY" and event.type == "LOGIN"
| summarize activeUsers = countDistinct(user.id), by:{day = bin(timestamp, 24h)}
| sort day desc
```

```
// Track audit events by provider — shows which platform surfaces are being used
fetch dt.system.events, from:-7d
| filter event.kind == "AUDIT_EVENT"
| summarize eventCount = count(), by:{event.provider}
| sort eventCount desc
| limit 15
```

Monitoring Data Flow Health

Verify that the platform is healthy and data is flowing — users cannot adopt a platform that is not working correctly:

```
// Verify log ingestion is healthy — compare volume to expected baseline
fetch logs, from:-24h
| makeTimeseries logVolume = count(default: 0), interval: 1h
```

```
// Count monitored hosts — should match or exceed Managed baseline
// from:-24h: without it, only entities seen in the default 2-hour window are counted — Step 1 counts over 24 h too.
fetch dt.entity.host, from:-24h
| summarize hostCount = count()

// Alternative: Smartscape on Grail (entity.name → name)
// smartscapeNodes HOST
// | summarize hostCount = count()
```

```
// Count monitored services — compare to Managed inventory
// from:-24h: without it, only entities seen in the default 2-hour window are
counted — Step 1 counts over 24 h too.
fetch dt.entity.service, from:-24h
| summarize serviceCount = count()

// Alternative: Smartscape on Grail (entity.name → name)
// smartscapeNodes SERVICE
// | summarize serviceCount = count()
```

Adoption Dashboard

Build a migration adoption dashboard that tracks these metrics over time. Share it with management as evidence that the migration is succeeding. Key tiles:

Tile	Visualization	Data Source
Active users per day	Line chart	Audit log events
Monitored hosts	Single value	Entity count
Monitored services	Single value	Entity count
Log ingestion volume	Line chart	Log count timeseries
Problems detected	Line chart	detected problems timeseries
Managed logins remaining	Line chart (target: zero)	Managed access logs

7. Step Completion Checklist

Do not proceed to Step 8 (Expand) until all items are confirmed.

Checkpoint	Status
Pre-migration communication sent to all stakeholders	[ ]
Post-migration announcement sent with new URLs and login instructions	[ ]
Platform admin training delivered	[ ]
SRE/operations training delivered	[ ]

Checkpoint	Status
Developer training delivered	[ ]
Management dashboard walkthrough completed	[ ]
Security team training delivered	[ ]
All runbooks updated with SaaS URLs and procedures	[ ]
Architecture diagrams updated to reflect SaaS topology	[ ]
Onboarding guides updated for SaaS	[ ]
API documentation updated with new base URLs	[ ]
Disaster recovery plans revised for SaaS model	[ ]
Quick Start Guide published and shared	[ ]
DQL Cheat Sheet published and shared	[ ]
FAQ published and shared	[ ]
Support channels established and announced	[ ]
Escalation path documented and communicated	[ ]
Enablement success metrics baselined	[ ]
Active user count meets or exceeds Managed baseline	[ ]

Next Step

> **M2S-08: Step 8 — Expand: Adopt New SaaS Capabilities** — Adopt the SaaS-exclusive capabilities that were never available in Managed: Grail, Notebooks, OpenPipeline, workflow automation, Dynatrace Assist, and platform apps.

Additional Resources

- Dynatrace University — Self-paced learning and certifications

- Dynatrace Community — Peer support and shared best practices
- DQL Documentation — Full DQL reference
- Notebooks Documentation — Interactive data exploration
- Workflows Documentation — AutomationEngine for alerting and remediation
- IAM Documentation — User and access management in SaaS

Summary

In Step 7, you:

- Executed a structured communication plan targeting all stakeholders with audience-specific messaging
- Delivered persona-based training for platform admins, SRE/operations, developers, management, and security
- Updated all internal documentation (runbooks, architecture diagrams, onboarding guides, API docs) to reflect SaaS
- Published self-service resources including a Quick Start Guide, DQL Cheat Sheet, and FAQ
- Established support channels with clear escalation paths and a hyper-care schedule
- Defined and baselined enablement success metrics to track adoption

> **Key Takeaway:** The migration is only as successful as the enablement effort behind it. Technology changes are straightforward; people changes require communication, training, and ongoing support. Invest in enablement now to avoid months of low adoption, frustrated users, and shadow IT workarounds.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.