

MZ2POL-00: SDK Management Zone Analysis Tool

Series: MZ2POL — Management Zone to Policy Migration | **Notebook:** 1 of 10 |
Created: December 2025 | **Last Updated:** 07/21/2026

Purpose: Query Management Zone configurations via the Dynatrace SDK, analyze entity assignments, and assess security context coverage to support migration planning.

What This Notebook Does

1. **Export all MZ configurations** - Get all MZ rules in a flat table format via Settings API
 2. **Analyze MZ rule patterns** - Understand what types of rules are in use
 3. **Check entity security context coverage** - Find entities missing `dt.security_context`
 4. **Identify most-used MZs** - Prioritize migration based on usage
 5. **Find tag patterns** - Discover common tag structures for segment creation
 6. **Kubernetes namespace analysis** - K8s namespaces for segment variables
 7. **Host group analysis** - Host groups for vertical MZ migration
 8. **Migration readiness summary** - Combined view of readiness indicators
-

Table of Contents

1. [Export All Management Zone Configurations](#)
2. [MZ Rule Pattern Analysis](#)
3. [Entity Security Context Coverage](#)
4. [Entity Counts by Management Zone](#)

5. [Tag Pattern Analysis](#)
 6. [Kubernetes Namespace Analysis](#)
 7. [Host Group Analysis](#)
 8. [Migration Readiness Summary](#)
 9. [Related Notebooks](#)
 10. [Documentation](#)
-

1. Export All Management Zone Configurations

This TypeScript function queries all MZs via the Settings API and flattens their rules into an analyzable table.

```

import { settingsObjectsClient } from "@dynatrace-sdk/client-classic-environment-v2";

// Type definitions
interface MzAttributeRuleCondition {
  key: string;
  operator: string;
  tag?: string;
  stringValue?: string;
  caseSensitive?: boolean;
}

interface MzAttributeRule {
  entityType: string;
  serviceToHostPropagation?: boolean;
  serviceToPGPropagation?: boolean;
  conditions?: MzAttributeRuleCondition[];
}

interface MzRule {
  enabled: boolean;
  type: string;
  entitySelector?: string;
  attributeRule: MzAttributeRule[];
}

interface MzDefinition {
  name: string;
  description?: string;
  rules: MzRule[];
}

// Fetch all MZs with pagination
async function getAllMZs(nextPageKey?: string, mzList: MzDefinition[] = []):
Promise<MzDefinition[]> {
  const config: any = nextPageKey
    ? { nextPageKey }
    : { schemaIds: "builtin:management-zones", pageSize: 100, fields:
"objectId,value" };

  const resp = await settingsObjectsClient.getSettingsObjects(config);
  mzList.push(...resp.items.map(elem => elem.value as MzDefinition));

  return resp.nextPageKey ? getAllMZs(resp.nextPageKey, mzList) : mzList;
}

// Flatten rules into records

```

```

function flattenRules(mzList: MzDefinition[]): any[] {
  const records: any[] = [];

  mzList.forEach(mz => {
    mz.rules?.forEach(rule => {
      if (rule.attributeRule) {
        const attr = rule.attributeRule;
        const base = {
          mzName: mz.name,
          description: mz.description || "",
          ruleType: rule.type,
          ruleEnabled: rule.enabled,
          entityType: attr.entityType || "",
          propagateToHost: attr.serviceToHostPropagation || false,
          propagateToPG: attr.serviceToPGPropagation || false
        };

        if (attr.conditions?.length) {
          attr.conditions.forEach(c => {
            records.push({
              ...base,
              conditionKey: c.key,
              operator: c.operator,
              value: c.stringValue || c.tag || "",
              caseSensitive: c.caseSensitive || false
            });
          });
        } else {
          records.push({ ...base, conditionKey: "", operator: "", value: "",
            caseSensitive: false });
        }
      }
    })
  })

  if (rule.type === "SELECTOR") {
    records.push({
      mzName: mz.name,
      description: mz.description || "",
      ruleType: "SELECTOR",
      ruleEnabled: rule.enabled,
      entityType: "",
      propagateToHost: false,
      propagateToPG: false,
      conditionKey: "entitySelector",
      operator: "SELECTOR",
      value: rule.entitySelector || "",
      caseSensitive: false
    });
  }
}

```

```
    });  
  });  
  
  return records;  
}  
  
export default async function() {  
  const mzList = await getAllMZs();  
  return flattenRules(mzList);  
}
```

2. MZ Rule Pattern Analysis

Analyze what types of MZ rules are in use to understand migration complexity.

```

import { settingsObjectsClient } from "@dynatrace-sdk/client-classic-environment-v2";

interface MzRule {
  enabled: boolean;
  type: string;
  entitySelector?: string;
  attributeRule?: { entityType: string; conditions?: { key: string; operator: string }[] }[];
}

interface MzDefinition {
  name: string;
  rules: MzRule[];
}

async function getAllMZs(nextPageKey?: string, mzList: MzDefinition[] = []): Promise<MzDefinition[]> {
  const config: any = nextPageKey
    ? { nextPageKey }
    : { schemaIds: "builtin:management-zones", pageSize: 100, fields: "objectId,value" };
  const resp = await settingsObjectsClient.getSettingsObjects(config);
  mzList.push(...resp.items.map(elem => elem.value as MzDefinition));
  return resp.nextPageKey ? getAllMZs(resp.nextPageKey, mzList) : mzList;
}

export default async function() {
  const mzList = await getAllMZs();

  // Count patterns
  const stats = {
    totalMZs: mzList.length,
    totalRules: 0,
    byRuleType: {} as Record<string, number>,
    byEntityType: {} as Record<string, number>,
    byConditionKey: {} as Record<string, number>,
    byOperator: {} as Record<string, number>,
    selectorRules: 0,
    meRules: 0,
    tagBasedRules: 0,
    nameBasedRules: 0
  };

  mzList.forEach(mz => {
    mz.rules?.forEach(rule => {
      stats.totalRules++;
    });
  });
}

```

```

    stats.byRuleType[rule.type] = (stats.byRuleType[rule.type] || 0) + 1;

    if (rule.type === "SELECTOR") {
        stats.selectorRules++;
    } else if (rule.type === "ME" && rule.attributeRule) {
        stats.meRules++;
        const attr = rule.attributeRule;
        stats.byEntityType[attr.entityType] =
(stats.byEntityType[attr.entityType] || 0) + 1;

        attr.conditions?.forEach(c => {
            stats.byConditionKey[c.key] = (stats.byConditionKey[c.key] || 0) +
1;
            stats.byOperator[c.operator] = (stats.byOperator[c.operator] || 0)
+ 1;

            if (c.key.includes("TAGS")) stats.tagBasedRules++;
            if (c.key.includes("NAME")) stats.nameBasedRules++;
        });
    }
    });
});

// Convert to table format
const results = [
    { category: "Summary", metric: "Total Management Zones", count:
stats.totalMZs },
    { category: "Summary", metric: "Total Rules", count: stats.totalRules },
    { category: "Summary", metric: "Avg Rules per MZ", count:
Math.round(stats.totalRules / stats.totalMZs * 10) / 10 },
    { category: "Rule Type", metric: "SELECTOR rules", count:
stats.selectorRules },
    { category: "Rule Type", metric: "ME (attribute) rules", count:
stats.meRules },
    { category: "Pattern", metric: "Tag-based rules", count:
stats.tagBasedRules },
    { category: "Pattern", metric: "Name-based rules", count:
stats.nameBasedRules }
];

// Add top entity types
Object.entries(stats.byEntityType)
    .sort((a, b) => b[1] - a[1])
    .slice(0, 10)
    .forEach(([type, count]) => {
        results.push({ category: "Entity Type", metric: type, count });
    });

```

```

// Add top condition keys
Object.entries(stats.byConditionKey)
  .sort((a, b) => b[1] - a[1])
  .slice(0, 10)
  .forEach(([key, count]) => {
    results.push({ category: "Condition Key", metric: key, count });
  });

return results;
}

```

3. Entity Security Context Coverage

Check how many entities have `dt.security_context` set vs those relying only on `managementZones` .

Service Coverage

```

fetch dt.entity.service
| summarize
    total = count(),
    withSecurityContext = countIf(isNotNull(dt.security_context)),
    withMZs = countIf(isNotNull(managementZones)),
    noAccess = countIf(isNull(dt.security_context) AND
    isNull(managementZones))
| fieldsAdd
    securityContextPercent = round(100.0 * withSecurityContext / total,
    decimals: 2),
    mzPercent = round(100.0 * withMZs / total, decimals: 2)

```


Host Coverage

```
fetch dt.entity.host
| summarize
    total = count(),
    withSecurityContext = countIf(isNotNull(dt.security_context)),
    withMZs = countIf(isNotNull(managementZones))
| fieldsAdd
    securityContextPercent = round(100.0 * withSecurityContext / total,
    decimals: 2),
    mzPercent = round(100.0 * withMZs / total, decimals: 2)
```

4. Entity Counts by Management Zone

See how many entities are assigned to each MZ to prioritize migration.

Services per Management Zone

```
fetch dt.entity.service
| expand mz = managementZones
| filter isNotNull(mz)
| summarize entityCount = count(), by: {managementZone = mz}
| sort entityCount desc
| limit 30
```

5. Tag Pattern Analysis

Discover common tag patterns that could be used for segment creation.

Most Common Host Tags

```
fetch dt.entity.host
| expand tag = tags
| filter isNotNull(tag)
| summarize count = count(), by: {tag}
| sort count desc
| limit 50
```

Tag Key Frequency

Extract key from `key:value` format to see which tag keys are most common.

```
fetch dt.entity.host
| expand tag = tags
| filter isNotNull(tag)
| filter contains(tag, ":")
| parse tag, "LD:tagKey ':' LD:tagValue"
| summarize count = count(), uniqueValues = countDistinct(tagValue), by:
{tagKey}
| sort count desc
| limit 30
```

6. Kubernetes Namespace Analysis

K8s namespaces are a common source for segment variables.

```
fetch dt.entity.cloud_application_namespace
| fields namespace = entity.name, id
| sort namespace asc
| limit 30
```

7. Host Group Analysis

Host groups are key for vertical MZ migration (environment-based access control).

```
fetch dt.entity.host_group  
| fields hostGroup = entity.name, id
```

8. Migration Readiness Summary

Combined view of migration readiness indicators.

```

import { settingsObjectsClient } from "@dynatrace-sdk/client-classic-environment-v2";
import { queryExecutionClient } from "@dynatrace-sdk/client-query";

async function getMZCount(): Promise<number> {
  const resp = await settingsObjectsClient.getSettingsObjects({
    schemaIds: "builtin:management-zones",
    pageSize: 1,
    fields: "totalCount"
  });
  return resp.totalCount || 0;
}

async function runDQL(query: string): Promise<any[]> {
  const resp = await queryExecutionClient.queryExecute({
    body: {
      query,
      requestTimeoutMilliseconds: 60000,
      maxResultRecords: 10
    }
  });
  return resp.result?.records || [];
}

export default async function() {
  const mzCount = await getMZCount();

  // Get service coverage
  const serviceCoverage = await runDQL(`
    fetch dt.entity.service
    | summarize
      total = count(),
      withSecContext = countIf(isNotNull(dt.security_context)),
      withMZ = countIf(isNotNull(managementZones))
  `);

  // Get host coverage
  const hostCoverage = await runDQL(`
    fetch dt.entity.host
    | summarize
      total = count(),
      withSecContext = countIf(isNotNull(dt.security_context))
  `);

  const svc = serviceCoverage[0] || {};
  const host = hostCoverage[0] || {};

```

```

return [
  { category: "Management Zones", metric: "Total MZs to migrate", value:
mzCount, status: mzCount > 50 ? "HIGH" : "OK" },
  { category: "Services", metric: "Total services", value: svc.total || 0,
status: "INFO" },
  { category: "Services", metric: "With security_context", value:
svc.withSecContext || 0, status: (svc.withSecContext / svc.total) > 0.8 ?
"GOOD" : "NEEDS_WORK" },
  { category: "Services", metric: "With MZ assignment", value: svc.withMZ
|| 0, status: "INFO" },
  { category: "Hosts", metric: "Total hosts", value: host.total || 0,
status: "INFO" },
  { category: "Hosts", metric: "With security_context", value:
host.withSecContext || 0, status: (host.withSecContext / host.total) > 0.8
? "GOOD" : "NEEDS_WORK" }
];
}

```

Next Steps

This tool is an inventory aid, not a converter. There is **no automatic Management Zone to Segment conversion** in Dynatrace — no button, no API, no wizard. Unlike Metric Events, which have a supported automatic upgrade path to Anomaly Detectors, every Management Zone must be re-expressed by hand as a segment or an IAM policy. The classifications below tell you *how much work each zone is* and *which construct it maps to*; they do not perform the migration. See MZ2POL-05 for the mapping scenarios and MZ2POL-04 for the policy/boundary path.

Based on the analysis above:

1. **High MZ count?** -> Prioritize by usage (query `dt.sfm.server.management_zones.queries_counter`)
2. **Low security_context coverage?** -> Plan host tag deployment or OpenPipeline enrichment
3. **Many tag-based MZ rules?** -> Map to segment filters using `matchesValue(tags, "key:value")`

4. **Many SELECTOR rules?** -> Expect the highest per-zone effort; entity includes support only `=` and `in()` , so many selectors need the fallback tag (MZ2POL-05 §4.6)
5. **Host groups in use?** -> Use `dt.host_group.id` in boundaries for vertical MZ migration

Migration Priority:

Priority	Criteria	Action
1	MZs with >100 entities	Migrate first - highest impact
2	Tag-based MZ rules	Lowest-effort mapping (still manual) — MZ2POL-05 §4.1
3	Host group-based MZs	Use boundary on <code>dt.host_group.id</code>
4	Entity selector MZs	Hand-map to segment filters; check the operator restrictions in MZ2POL-05
5	Complex/nested MZs	Manual analysis required

Related Notebooks

- **MZ2POL-01:** Introduction - Why Migrate
- **MZ2POL-03:** Assessment and Planning
- **MZ2POL-05:** Migrating Management Zone Filtering to Segments
- **MZ2POL-06:** Migration Execution
- **MZ2POL-07:** Validation and Troubleshooting

Documentation

- [Management Zones Documentation](#)
- [Upgrade from RBAC to IAM Policies](#)

- [Grant Access to Entities with Security Context](#)

MZ2POL-00: SDK Management Zone Analysis Tool - Part of the MZ to Policies/Boundaries/Segments Series

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.