

MZ2POL-05: Migrating Management Zone Filtering to Segments

Series: MZ2POL — Management Zone to Policy Migration | **Notebook:** 6 of 10 |

Created: December 2025 | **Last Updated:** 10/05/2026

Overview

Management Zones do three jobs: **access control**, **filtering**, and **alerting** (see MZ2POL-09 for the third). Most migration guidance — including the rest of this series — covers the **access-control** job, which becomes IAM policies, boundaries, and

`dt.security_context`. This notebook covers the second job: **filtering and scoping what people see**, which becomes **Segments**.

If your Management Zones are mostly used to scope dashboards, filter views, and give teams a "my stuff only" lens — rather than to enforce who may read what or who gets paged — this notebook is the one you need, and you can read it without working through the access-control notebooks first.

The work is **manual**. There is no button that turns a Management Zone into a segment. What this notebook gives you is the mapping: for each way an MZ is commonly defined, the corresponding way to get the same scope in Grail — plus the four constraints that make some MZs impossible to reproduce exactly.

Table of Contents

1. [Which Job Is Your MZ Doing?](#)
2. [There Is No Automatic Conversion](#)
3. [Segment Mechanics: Where They Are Documented](#)
4. [Migration Scenarios](#)
5. [Conversion Blockers](#)

6. [Mapping MZ Rules to Segment Filters](#)
7. [Testing Your Mapping with DQL](#)
8. [Coexistence: Classic Apps and Parallel Running](#)
9. [Visibility and Permissions](#)
10. [Migration Checklist](#)

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Grail — segments are a Grail-only construct
Existing Management Zones	An inventory of the zones you intend to migrate (see MZ2POL-00 and MZ2POL-03)
Permissions	<code>storage:filter-segments:read</code> and <code>storage:filter-segments:write</code> to create segments; <code>storage:entities:read</code> and <code>storage:logs:read</code> to run the validation queries in §7
Enrichment in place	Host tags, host groups, or Kubernetes metadata must already reach Grail — see §4 for which scenario applies
Background reading	ORGNZ-08 and ORGNZ-10 for segment mechanics; this notebook covers the migration, not the feature

1. Which Job Is Your MZ Doing?

Before mapping anything, classify each Management Zone by what it is actually being used for. The replacement construct differs, and picking the wrong one produces either a security gap or a lot of unnecessary IAM work.

What the MZ is used for	Replacement	Covered in
Restricting who may read which data	IAM policy + boundary on <code>dt.security_context</code>	MZ2POL-04
Scoping what a user sees in apps and dashboards	Segment	This notebook
Deciding who gets paged — alerting profile scope	Routing dimensions on a problem-triggered workflow	MZ2POL-09
More than one of the above	Each job migrates independently — they are separate layers	The matching notebook(s) above

The three are not alternatives and they do not substitute for each other:

Segments are not access control. A segment changes what is *shown*, not what is *permitted*. Anyone with `storage:filter-segments:read` can apply any segment, and a segment never widens what a user is allowed to query — the IAM policy still decides that. If the MZ was load-bearing for security, a segment alone does **not** replace it. Segments also do not replace alerting-profile scoping — see MZ2POL-09.

A useful heuristic

If removing the Management Zone would let someone see data they are **not allowed** to see, it is doing the access-control job — go to MZ2POL-04. If removing it would only make their dashboards noisy, it is doing the filtering job, and this notebook covers it. If removing it would change who gets paged, it is doing the alerting job — go to MZ2POL-09.

Most MZ estates contain far more of the filtering kind than teams expect. Run the inventory in MZ2POL-00 and MZ2POL-03 first; it classifies the estate for you.

2. There Is No Automatic Conversion

There is no automatic Management Zone to Segment conversion. No button, no API endpoint, no wizard. Unlike Metric Events, which have a supported automatic

upgrade path to Anomaly Detectors, Management Zones must be re-expressed as segments by hand.

This matters for planning. A 200-zone estate is 200 hand-authored segment definitions unless you consolidate first — and consolidation is usually the right move, because MZ estates typically accumulate zones that no longer have owners or users.

What tooling does exist

Tool	What it does	What it does not do
MZ2POL-00 SDK analysis tool	Exports <code>builtin:management-zones</code> , classifies each rule (tag-based / name-based / SELECTOR / host-group / complex-nested), scores migration readiness	Emits no segments — it is an inventory and triage aid
Filter Segments API (<code>POST /platform/storage/filter-segments/v1/filter-segments</code> , <code>storage:filter-segments:write</code> ; SDK: <code>client-filter-segment-management</code> (Dynatrace Developer))	Creates segments programmatically once you have written the definition	Does not derive the definition from an MZ. Segments are not a Settings 2.0 object — there is no <code>builtin:filter-segments</code> schema (the Settings API returns 404 for it)
Terraform <code>dynatrace_segment</code> / Monaco <code>segments</code>	Manages segment definitions as code	Same — you supply the DQL

The practical sequence is therefore: **inventory (MZ2POL-00)** → **consolidate** → **hand-author the definitions using the scenarios below** → **deploy as code (ORGNZ-10 §10, AUTOM-04)**.

Consolidate before you convert: one segment per *dimension*, not per *value*

This is the single most important planning decision, and it changes the arithmetic of the whole migration.

A Management Zone is a **static definition of one value** — `Easytrade` , `Hipstershop` , `Easytravel` are three separate zones, and onboarding a fourth app means authoring a fourth zone. A segment does not work that way: a segment carries **variables**, and a variable's values come from a live DQL query. One segment named `app` , whose variable enumerates every distinct application value, covers every app that exists now and every app added later — with no further configuration.

So the migration is **not** N zones → N segments:

	Management Zones	Segments
Unit of definition	One per value	One per dimension
Typical count	Dozens to hundreds	Usually three to eight
New app onboarded	Author a new zone	Appears automatically, if enrichment is present

An estate with dozens of zones across three dimensions — app, stage, platform — converges on **three segments**, not dozens. If your target count is climbing past roughly eight, that is the signal that you are converting values instead of dimensions, and it is worth stopping to revisit the dimension design.

Segment sprawl is the common failure mode. Teams migrate zone-for-zone, end up with dozens of overlapping segments, and reproduce exactly the maintenance burden they were trying to escape — with the added downside that users now face a long, confusing segment picker.

This rule is canonical in **ORGNZ-10** (§ *Best Practice: One Segment Per Dimension*) and **ORGNZ-99** (rule 46); this notebook applies it to the migration.

Other reductions worth making

- Zones with no query activity in the last 90 days: delete rather than convert.
- Zones that duplicate a dimension already carried by a Primary Grail field (host group, K8s namespace, cloud tags) — those come nearly free via §4.2 and §4.4.

Sources: [Best practice examples: from Management Zones to Segments \(DT docs\)](#), [Segments \(DT](#)

3. Segment Mechanics: Where They Are Documented

This notebook covers the **migration**, not segment mechanics in general. For how segments actually work, the canonical treatment in this corpus is the ORGNZ series:

Topic	Read
Segments vs. buckets, include structure, design patterns, limits	ORGNZ-08 — Grail Segments
Filter syntax and operators, include rules, variables, host-group segments, visibility, segments-as-code, Query API, Davis problem includes	ORGNZ-10 — Advanced Segment Definitions
Managing segments with Terraform and Monaco	AUTOM-04, AUTOM-97

What you need to carry into the migration from those notebooks is a short list of constraints, because they determine which Management Zones can be reproduced faithfully and which cannot:

Constraint	Value	Why it matters when converting
Includes per segment	20	A complex MZ with many rule types may not fit in one segment
Include blocks per data type	1	You cannot add two <code>logs</code> includes — combine with <code>OR</code> inside one
Expressions per filter condition	10	Caps how many conditions a single MZ rule can become
Segments applied per query	10	Limits how far you can decompose one MZ into many small segments
Segments per environment	10,000	Rarely binding, but relevant for very large estates

Constraint	Value	Why it matters when converting
Values in a variable dropdown	10,000	Relevant for tenant-wide variables such as host group
Selected variable values per segment	100	Classic entities; constrains variable-driven consolidation
Filter operators	= and in() only	There is no separate contains operator — see below
Wildcards	* inside the value	Gives starts-with / contains / ends-with matching
Classic entity properties	Only entity.name accepts wildcards	Every other entity property is exact-equals only

Include blocks are ORed. Data matching *any* include is in scope. Within a single include, conditions are ANDed. An MZ whose rules were implicitly ANDed across entity types will not behave the same way if you map each rule to its own include.

Start with `Data (all types)`

Segments offer a `Data (all types)` include block alongside the type-specific ones, and it is the **recommended** starting point. It applies your filter conditions across signals *and* Smartscape entities in one rule, which means a single include often replaces what would otherwise be several — and it keeps you well clear of the 20-include ceiling.

Reach for a type-specific include only when a dimension genuinely differs by data type, or when you need one of the cases `Data (all types)` does not cover: the `events` include for Davis problems (§5.4), or a **classic entity** include.

Classic entity includes are a backward-compatibility path, not the default.

Dynatrace states that entity relationships in segments "are only supported for backward compatibility with classic entities." They remain useful for higher-cardinality classic entities — Kubernetes workloads filtered by their related clusters, for example — but a migration should not reach for them first. See §8 for what this changed about the Services app.

Sources: [Segment limits \(DT docs\)](#), [Include data in segments \(DT docs\)](#) — verbatim: entity relationships in segments *"are only supported for backward compatibility with classic entities"*, [Filter Smartscape nodes with segments \(DT docs\)](#) — the `Data (all types)` include block is the recommended form.

4. Migration Scenarios

Dynatrace documents a set of fully supported upgrade scenarios, organized by **how the Management Zone was defined** rather than by what it was called. Find the row matching your MZ's rule type and follow that scenario.

#	If the MZ is defined by...	Segment approach	Extra tagging needed?
4.1	Auto-tagging on process / host / service name or metadata	Set the tag at the source as a host tag	Yes — move tagging to OneAgent
4.2	Host group membership	Filter on <code>dt.host_group.id</code>	No — enriched automatically
4.3	Systematic host-group naming (<code><CMDBID>-<app>-<component>-<stage></code>)	Parse the name with DPL in a variable definition	No
4.4	Kubernetes cluster or namespace	Filter on the K8s enrichment fields	No — enriched automatically
4.5	Host tags / properties set at OneAgent install	Filter on the tag, allowing for the <code>[Environment]</code> prefix	No

#	If the MZ is defined by...	Segment approach	Extra tagging needed?
4.6	Anything else (rules that do not fit above)	Introduce a Segment tag whose value is the MZ name	Yes — the guide gives this for zones built on names or metadata without auto-tags; this notebook uses it as the general fallback
4.7	Extension data	Filter on the extension enrichment attributes	No
4.8	Cloud-native / application-only injection	Set <code>DT_TAGS</code> and <code>OTEL_RESOURCE_ATTRIBUTES</code> on the process	Yes — environment variables

The guide describes these as its *"fully supported scenarios"* — its stated purpose is *"sharing best practices for fully supported scenarios."* Treat them as the supported set today, not an exhaustive account of every MZ shape.

Sources: [Best practice examples: from Management Zones to Segments \(DT docs\)](#).

4.1 Auto-tagging on process, host, or service names

The MZ pattern: an auto-tagging rule derives a tag from a process name, host name, service name, or their metadata, and the Management Zone selects on that tag.

The move: stop deriving the tag from observed names and **set it as a host tag at the source**, during OneAgent deployment. Tags applied this way propagate into Grail as field values on the signals themselves, which is what makes them usable in segment filter conditions across data types.

This is the single most common scenario and also the one that requires the most upfront work, because it changes where tagging happens — from a rule evaluated in Dynatrace to a value supplied at install time. See FAQ-02 for tagging-source strategy and ORGNZ-10 §3 for the enrichment prerequisite and a coverage-audit query.

Watch the derived-data caveat. Custom tags do **not** all propagate to derived data such as service metrics and span-based calculations. Only `dt.security_context`,

`dt.cost.costcenter` , and `dt.cost.product` carry through. See §5.2.

4.2 Host group membership

The MZ pattern: the zone selects hosts by host group.

The move: no tagging work at all. Every signal from a OneAgent that belongs to a host group is automatically enriched with `dt.host_group.id` . Filter on it directly.

This is the cheapest scenario to migrate and a good first candidate for a pilot conversion. FAQ-01 covers host-group naming strategy, which determines how much you can extract in 4.3.

4.3 Extraction from systematic host-group naming

The MZ pattern: host groups follow a structured naming convention — for example `<CMDBID>-<AppName>-<component>-<deploymentEnvironment>` — and separate MZs slice on different parts of the name.

The move: parse the components out with DPL in a **segment variable definition**, then filter on the extracted value. One segment with a variable replaces the whole family of MZs.

```
// Segment variable definition – extract dimensions from systematic host-group
naming.
// Pattern: <CMDBID>-<AppName>-<component>-<deploymentEnvironment>
<CMDBID>-<AppName>-<component>-<deploymentEnvironment>
// Live-verified 07/21/2026. smartsceneNodes scans 0 bytes – this costs
nothing to run.
smartsceneNodes "HOST"
| filter isNotNull(dt.host_group.id)
| parse dt.host_group.id, "LD:CMDBID '-' LD:app '-' LD:component '-' LD:stage"
| filter isNotNull(CMDBID)
| fields CMDBID
| dedup CMDBID
| sort CMDBID asc
```

A segment variable definition is ordinary DQL. It is not restricted to the classic entity model — `smartsceneNodes` works and is the preferred form. The published upgrade-guide examples use `fetch dt.entity.host_group` , which still works but queries the deprecated classic entity model; the `smartsceneNodes` equivalents

above and below were verified against a live tenant on 07/21/2026 and scan **zero bytes**.

The real constraints on a variable query are different ones: the result must be **columnar** — the first column becomes the primary variable and later columns become secondary variables — high-volume data types such as `logs` should be avoided, and **neither variable names nor values may contain `*`** (though a wildcard may *follow* a variable reference in a condition, as `$stage*`).

There is no `HOST_GROUP` Smartscape node type. Host groups are a *field* (`dt.host_group.id`) on `HOST` nodes, not nodes of their own — which is why the query above starts from `HOST` and dedups, rather than fetching host groups directly.

If your host groups are *not* systematically named, this scenario is unavailable and you fall back to 4.6. That is a strong argument for fixing host-group naming before migrating rather than after — see FAQ-01.

4.4 Kubernetes cluster and namespace

The MZ pattern: the zone scopes to a Kubernetes cluster or namespace.

The move: Kubernetes signals are automatically enriched with cluster and namespace information — no additional tagging. Build the variable from the namespace entity and filter on the enrichment fields.

That is enough for simple setups. For production the guide recommends a richer baseline: *"For most production environments, use Namespace annotations and labels as the baseline."* Map the namespace labels or annotations your teams already maintain through Kubernetes telemetry enrichment, and filter segments on those ([upgrade guide \(DT docs\)](#)).

```
// Segment variable definition – Kubernetes namespaces.
// Live-verified 07/21/2026: K8S_NAMESPACE is a real Smartscape node type; 0
bytes scanned.
smartscapeNodes "K8S_NAMESPACE"
| fields namespace = name
| dedup namespace
| sort namespace asc
```

For the filter conditions themselves, the Primary Grail Fields are `k8s.cluster.name` (cluster-scoped) and `k8s.namespace.name` (namespace-scoped). ORGNZ-10 §5 walks the full four-step build for both the host-group and Kubernetes variants.

Limitation: Primary Grail Tags do not yet enrich Kubernetes **metrics or events** — they work for logs, spans, and topology. A segment that must scope K8s metrics needs a different include shape.

4.5 Host tags and properties

The MZ pattern: the zone selects on tags set during OneAgent installation.

The move: filter on the tag directly — but account for the prefix. Host tags added via OneAgent appear **with an [Environment] prefix**, so a tag whose key is `App` must be matched as `[Environment]App`.

```
// Segment variable definition — distinct values of a primary tag.
// On Smartscape nodes `tags` is a RECORD: access keys with unquoted bracket notation.
// Live-verified 07/21/2026; 0 bytes scanned.
smartscapeNodes "HOST"
| fieldsAdd app = tags[primary_tags.application]
| filter isNotNull(app)
| fields app
| dedup app
| sort app asc
```

Two different tag shapes — this is the trap. On **Smartscape** nodes, `tags` is a **record**: keys are raw (`primary_tags.application`) and you read them with unquoted bracket notation, `tags[primary_tags.application]`. On the **classic** entity model (`fetch dt.entity.host | expand tag = tags`), `tags` is an **array of "key:value" strings**, and host tags added via OneAgent surface there with an `[Environment]` prefix — so a tag whose key is `App` reads as `"[Environment]App:<value>"`.

Mixing the two produces a filter that validates but silently matches nothing. Running `matchesValue(tags, "[Environment]App:x")` against `smartscapeNodes "HOST"` returns the verifier warning *"this parameter should be a string, an array or a smartscape id, but was a record"* followed by *"the query will always return an empty result as the condition can't be true"* — confirmed 07/21/2026. Prefer the Smartscape form and the record accessor.

`dt.security_context` is also worth knowing: on Smartscape nodes it is an **array**, so enumerating its distinct values needs `expand` (see §4.6).

4.6 Everything else — the `segment` tag fallback

The MZ pattern: any zone whose definition does not reduce to one of the patterns above — hand-maintained entity lists, complex nested rules, SELECTOR-based rules.

The move: introduce a **new tag with a key such as `segment`**, whose value is the name of the **Management Zone**. Apply it to the entities the zone covered, then define a segment that filters on that tag. The guide gives this move for its non-auto-tagging scenario — zones defined by process name, host name, service name, or metadata; using it for every zone that fits no other scenario is this notebook's extension of it.

This is deliberately low-cleverness. It reproduces the zone's membership without trying to re-derive its logic, which is exactly what you want for zones whose rules have accumulated exceptions nobody can explain. The tag becomes the explicit, greppable statement of membership that the MZ rule was only implying.

Two consequences worth accepting up front:

- Membership is now **maintained where the tag is set**, not by a rule that re-evaluates. Entities added later need the tag.
- In community practice it works well as a migration destination and ages less well as a long-term model, because membership is hand-maintained rather than derived from the estate. Where a real dimension exists (host group, namespace, environment), prefer 4.2–4.5 and treat 4.6 as the fallback it is.

4.7 Extensions

The MZ pattern: the zone scopes extension-collected data.

The move: filter on the extension enrichment attributes — **security context**, **product**, and **cost center**. These are the same three keys that survive into derived data (§5.2), which is not a coincidence.

4.8 Cloud-native and application-only injection

The MZ pattern: the zone covers workloads instrumented by application-only injection, where there is no host OneAgent to carry host tags.

The move: set the enrichment on the process itself, via environment variables —

`dt.security_context` , `DT_TAGS` , and `OTEL_RESOURCE_ATTRIBUTES` .

They are not interchangeable. `DT_TAGS` and `OTEL_RESOURCE_ATTRIBUTES` look similar but differ semantically — most visibly, `OTEL_RESOURCE_ATTRIBUTES` uses a **colon** where `DT_TAGS` uses a **space** as separator. Setting one and assuming the other is covered is a reliable way to produce a segment that half-works.

Sources: [Best practice examples: from Management Zones to Segments \(DT docs\)](#) — for non-auto-tagging zones, a new tag *"with the value set to the name of the management zone"*, [Segment data by Kubernetes clusters \(DT docs\)](#).

4.9 Shortcut: reuse `dt.security_context` as a segment dimension

If you are migrating access control at the same time (MZ2POL-04), `dt.security_context` is already being populated across your signals. That makes it available as a ready-made segment dimension at no extra enrichment cost — the values are already there.

The caveat is conceptual, and worth stating plainly: **filtering on `dt.security_context` does not make the segment an access-control mechanism.** The field happens to be convenient and universally populated; the segment is still just a view filter, and the IAM policy is still what decides what the user may read (§1). Use it because the enrichment already exists, not because it sounds secure.

On Smartscape nodes `dt.security_context` is an **array**, so enumerating its distinct values requires `expand` :

```
// Segment variable definition – enumerate distinct security-context values.
// dt.security_context is an ARRAY on Smartscape nodes, hence the expand.
// Live-verified 07/21/2026; 0 bytes scanned.
smartscapeNodes "PROCESS"
| fields dt.security_context
| expand dt.security_context
| filter isNotNull(dt.security_context)
| dedup dt.security_context
| sort dt.security_context asc
```

5. Conversion Blockers

Four constraints make some Management Zones impossible to reproduce exactly. Find these before you start authoring, not after — each one changes the design rather than just the syntax.

5.1 Segments cannot express exclusions

Management Zone rules support exclusions. Segment includes do not. There is no way to say "everything except team-X" in a segment.

Any MZ built as *broad rule minus exceptions* has to be inverted into an explicit statement of what to include. For a zone defined as "all production hosts except the three legacy boxes," the segment must enumerate or characterize the included set positively.

This is the single most common reason a conversion is not one-to-one, and it is why the fallback tag in §4.6 exists: applying a `Segment` tag to exactly the right entities sidesteps the need to express the exclusion at all.

5.2 Custom tags do not fully propagate to derived data

Derived data — service metrics, span-based calculations — does **not** inherit arbitrary custom tags. Only three keys carry through:

- `dt.security_context`
- `dt.cost.costcenter`
- `dt.cost.product`

A segment that filters correctly on logs and spans may therefore return unfiltered or empty results on service metrics. If the Management Zone was used to scope metric-based dashboards, verify that surface specifically — it is the one most likely to silently misbehave after conversion.

5.3 Entity properties other than `entity.name` are equals-only

Segment includes offer only two operators: `=` and `in()`. There is no distinct `contains` operator — substring and suffix matching are expressed by putting a **wildcard** `*` **inside the value**.

The restriction that bites during migration is *where* wildcards are allowed on classic entities:

Entity property	Matching available
<code>entity.name</code>	Starts-with only — <code>"payment*"</code> . No contains (<code>"*payment*"</code>) or ends-with (<code>"*-prod"</code>)
Every other property	Exact equals only — no wildcards
Any property not suggested in the filter field	Not allowed

These restrictions are specific to classic-entity includes. The *Segment limits* page states they do not apply to Smartscape on Grail, where the general `*` wildcard rules hold — verify the Smartscape include in your own tenant before relying on it.

Tag conditions are a separate case: a tag filter tests **membership in the tag set** rather than a substring of a string field, so verify tag-based entity includes against your own tenant before assuming wildcard behavior.

So a *service name starts with payment* rule converts (`entity.name = "payment*"`), but a *service name contains payment* rule does **not**, on a classic-entity include — nor does `host property X contains Y` . For a contains rule, or for the equals-only properties, the options are, in order of preference: rely on an exact value or an `in()` set, use a Smartscape include, move the condition onto a signal include on a field that carries the same value, or fall back to the `Segment` tag (§4.6).

Wildcards may also follow a variable name in a starts-with condition — `foo = $bar*` — which is what makes the variable-driven consolidation in §2 practical. Variable *values* themselves may never contain `*` .

Put wildcards next to a separator. Dynatrace's guidance for wildcard matching is to place the `*` immediately before or after a word separator — `-` , `_` , `.` , or `/` — because `MATCH("db-tech-*")` evaluates more efficiently than `MATCH("db-tech*")` . This is a performance rule, not a correctness one, but it is a good reason to design host-group and security-context values with explicit separators between dimensions rather than running them together. The same rule applies to IAM boundary conditions — see MZ2POL-04.

5.4 Davis problems need an events include

Entity includes alone do **not** filter the problem feed. To scope problems, the segment needs an `events` include with `event.kind = "DAVIS_PROBLEM"` plus the scoping

condition. Entity includes then additionally scope the affected-hosts and affected-services lists shown inside a problem.

If the Management Zone was used to give a team a filtered problem view — a very common use — this is the include that does the work. ORGNZ-10 §12 has the full include shape and a worked example.

Sources: [Segment limits \(DT docs\)](#) — classic entities: "starts-with, contains, and ends-with are not supported", "entity.name supports starts-with", "Other properties only support equals", "Only selected properties (suggested in the filter field) are allowed", "these limits don't apply to Smartscape on Grail",
Best practice examples: [from Management Zones to Segments \(DT docs\)](#), [Problems app \(DT docs\)](#).

6. Mapping MZ Rules to Segment Filters

Rule-level mapping, once you have chosen a scenario from §4. The **Include type** column matters: segment filter conditions are evaluated per data object, so an entity condition and a signal condition are not interchangeable.

MZ rule	Segment filter	Include type	Notes
Host group equals	<code>dt.host_group.id = "<group>"</code>	Signal + entity	Cheapest mapping (§4.2)
Host group starts with	<code>dt.host_group.id = "<prefix>*"</code>	Signal	Wildcard inside the value
Host tag equals	<code>tags = "[Environment]<key>:<value>"</code>	Entity	Mind the <code>[Environment]</code> prefix (§4.5)
Kubernetes namespace	<code>k8s.namespace.name = "<namespace>"</code>	Signal	Auto-enriched (§4.4)
Kubernetes cluster	<code>k8s.cluster.name = "<cluster>"</code>	Signal	Auto-enriched
Service name equals	<code>entity.name = "<name>"</code>	Entity	Use <code>in()</code> for a set

MZ rule	Segment filter	Include type	Notes
Service name starts with	<code>entity.name = " <text>*"</code>	Entity	Works — <code>entity.name</code> supports starts-with (§5.3)
Service name contains	—	—	Not expressible on classic entities — starts-with only (§5.3); use a Smartscape or signal include, or the fallback tag
Other entity property contains	—	—	Not expressible — equals-only (§5.3)
Cloud provider / arbitrary tag	<code>tags = "<key>: <value>"</code>	Entity	Requires the tag to be set at source (§4.1)
Anything else	<code>tags = "Segment:<MZ name>"</code>	Entity	The fallback (§4.6)
Broad rule minus exceptions	—	—	Not expressible — invert to positive (§5.1)
Problem feed scoping	<code>event.kind = "DAVIS_PROBLEM" and <condition></code>	Events	Required for problem filtering (§5.4)

Worked conversions

MZ: hosts with tag `env = production`

Segment: entity include on `dt.entity.host` with `tags = "`

`[Environment]env:production"`, plus signal includes on logs and spans with the corresponding field condition if the segment must scope those too.

MZ: all hosts in host groups beginning `PROD-`

Segment: signal include with `dt.host_group.id = "PROD-*"` — wildcard inside the value, one include, no tagging work.

MZ: production Kubernetes namespace, and its problems

Segment: three includes — `k8s.namespace.name = "production"` on logs, the same on

spans, and an `events` include with `event.kind = "DAVIS_PROBLEM"` and `k8s.namespace.name = "production"`.

MZ: "Payments platform" — a hand-maintained list of 40 services

Segment: tag those 40 services `Segment:Payments platform` and filter on it. Do not attempt to re-derive the list from naming.

7. Testing Your Mapping with DQL

Validate the scope **before** creating the segment. Every query below is ad-hoc validation DQL — it uses the modern `smartscapeNodes` form, unlike the segment-variable definitions in §4, which must use the classic entity model.

The question each answers: *does this filter select the same entities the Management Zone did?*

```
// Validate a host-group mapping (scenario 4.2).
// Compare this count against the host count in the Management Zone.
smartscapeNodes "HOST"
| filter dt.host_group.id == "PROD-PAYMENTS"
| fields name, dt.host_group.id
| sort name asc
| limit 50
```

```
// Validate a prefix-wildcard host-group mapping.
// Entity includes support prefix wildcards only -- verify the prefix is
selective enough.
smartscapeNodes "HOST"
| filter startsWith(dt.host_group.id, "PROD-")
| summarize hosts = count(), by:{dt.host_group.id}
| sort hosts desc
| limit 50
```

```
// Validate a tag-based mapping (scenarios 4.1, 4.5, 4.6).
// `tags` is a RECORD on Smartscape nodes -- use unquoted bracket access, NOT
matchesValue().
smartscapeNodes "HOST"
| fieldsAdd env = tags[primary_tags.environment]
| filter env == "production"
| fields name, env, dt.host_group.id
| sort name asc
| limit 50
```

```
// Coverage audit -- which hosts would the mapping MISS?
// Any host without the tag falls outside the segment. Review this list before
cutover:
// these are the hosts that silently disappear from the filtered view.
smartscapeNodes "HOST"
| fieldsAdd env = tags[primary_tags.environment]
| filter isNull(env)
| fields name, dt.host_group.id
| sort name asc
| limit 100
```

```
// Validate a Kubernetes namespace mapping against signal data (scenario 4.4).
// Entity-level checks do not prove the signal include works -- test the
signal directly.
fetch logs, from:-1h
| filter k8s.namespace.name == "production"
| summarize records = count(), by:{k8s.cluster.name, k8s.namespace.name}
| sort records desc
| limit 20
```

```
// Validate the Davis problem include shape (blocker 5.4).
// Entity includes alone do NOT filter the problem feed -- this is the
condition that does.
// On problem records k8s.namespace.name is an ARRAY: use in("<value>",
field), never ==.
fetch events, from:-24h
| filter event.kind == "DAVIS_PROBLEM"
      and in("production", k8s.namespace.name)
| fields timestamp, event.name, event.status
| sort timestamp desc
| limit 20
```

On problem records, entity-derived fields are arrays. A problem aggregates several events, so `k8s.namespace.name` on a `DAVIS_PROBLEM` record holds a list. `k8s.namespace.name == "production"` against it is always false, and Grail attaches no warning — on the validation tenant (10/05/2026) the `==` form returned 0 of 15,738 problem records for a namespace where `in("<namespace>", k8s.namespace.name)` returned 6. Use `in("<value>", field)` in DQL. The segment include itself (`k8s.namespace.name = "production"` in §6) is segment syntax, which accepts array fields, and does not need this change.

Validate the signal surface, not just the entity surface. A filter that selects the right hosts proves nothing about whether logs, spans, and metrics carry the same field. Derived metrics in particular may not (§5.2). Run at least one signal-level check per data type the segment is meant to scope.

8. Coexistence: Classic Apps and Parallel Running

Segments and Management Zones coexist, which makes a staged migration practical.

Surface	Scoping mechanism during migration
Grail-based apps (Logs, Distributed Traces, Problems, Dashboards, Notebooks, SLOs)	Segments
Services app — Explorer tab	Segments. The tab is built on time-series queries and Smartscape 2.0 entities , and ships ready-made segments that join service entities to their metric data — so you can filter by AWS region, Kubernetes namespace, or host group without authoring a join
Services app — Explorer (Deprecated) tab	The classic-entity surface; a classic entity include is what scopes it
Classic apps	Management Zones — segments do not apply
Alerting profiles	Management Zones, for as long as you keep the profiles. Segment support for alerting profiles is not pending — the exit is a problem-triggered workflow, not a segment

Guidance you may encounter that is now out of date. Workshop material from earlier in the segment rollout instructed teams to extend every segment with a hand-built classic-entity include — typically an `application:$app` tag filter — on the grounds that the Services app could not otherwise be filtered. That is no longer the general case: the current Explorer tab uses Smartscape 2.0 entities, and a `Data (all types)` include covers it. Build the classic entity include when you specifically need the deprecated tab or a high-cardinality classic entity relationship, not as a routine step for every segment.

The practical consequence: **do not delete the Management Zones when the segments go live.** Keep both running until the classic surfaces that depend on MZs are themselves retired or migrated. MZ2POL-06 covers the parallel-running period, cutover sequencing, and rollback; MZ2POL-07 covers post-cutover validation.

Order of operations

1. Build and validate segments alongside the existing MZs (this notebook).
2. Move Grail-app users onto segments; leave MZs untouched.
3. Confirm no classic surface still depends on the MZ.
4. Retire the MZ.

Reversing steps 3 and 4 is the mistake that produces a broken alerting profile discovered at 3 a.m.

9. Visibility and Permissions

A Management Zone was itself a permission object. A segment is not, and the mental model has to change accordingly.

Visibility settings

There are exactly two:

Setting	Effect
Unlisted (default)	The segment appears only in the owner's segment list

Setting	Effect
Anyone in the environment	The segment appears in everyone's segment list

There is no share-with-specific-group mechanism. Visibility controls only whether a segment is *listed* in someone's picker. Unlisted segments still become reachable when referenced from a shared notebook or dashboard — anyone with access to that app gets read access to the segment.

Visibility is not access control. Anyone holding `storage:filter-segments:read` can access any segment regardless of its visibility setting. Applying a segment never grants data the user's IAM policy does not already permit.

The common team pattern that follows from this: application teams keep their segments **Unlisted** and surface them by reference from team-owned dashboards, rather than adding noise to every user's picker.

Permission scopes

Scope	Action
<code>storage:filter-segments:read</code>	View and use segments
<code>storage:filter-segments:write</code>	Create and edit segments
<code>storage:filter-segments:share</code>	Share with others
<code>storage:filter-segments:delete</code>	Delete segments
<code>storage:filter-segments:admin</code>	Manage segment permissions

ORGNZ-10 §6 maps these to the default Dynatrace user policies and covers the governance model; ORGNZ-10 §10 covers restricting write scope when segments are managed as code.

Sources: [Visibility of segments \(DT docs\)](#) — verbatim: any segment can be accessed with `storage:filter-segments:read` regardless of configured visibility, [Dynatrace default policies reference \(DT docs\)](#).

10. Migration Checklist

Inventory and triage

- [] Ran the MZ2POL-00 analysis tool; have a classified inventory of every zone
- [] Split the estate: access-control zones (→ MZ2POL-04) vs. filtering zones (→ this notebook) vs. both
- [] Identified zones with no recent usage — delete rather than convert
- [] Counted **dimensions, not zones** — the target is one segment per dimension (usually 3-8 total), not one per Management Zone

Per zone

- [] Matched the zone to a §4 scenario, or explicitly assigned it the §4.6 fallback
- [] Checked it against all four §5 blockers — especially exclusions
- [] Confirmed the required enrichment exists (host tag, host group, K8s field, `Segment` tag)
- [] Verified the scope with DQL (§7) at both entity **and** signal level
- [] Confirmed the segment fits within the §3 limits
- [] Added the `events` include if the zone was used for problem filtering
- [] Used `Data (all types)` unless a type-specific or classic-entity include is genuinely required

Rollout

- [] Segment created, named per convention, visibility set deliberately
- [] Dashboards and notebooks repointed from MZ filtering to the segment
- [] Users know the segment exists and how to select it
- [] MZ left in place for classic apps and alerting profiles
- [] Segment promoted to code (Settings API / Terraform / Monaco) if it is long-lived
- [] MZ retired only after confirming no classic surface depends on it

Summary

1. **Classify first** — Management Zones do three jobs: access control, filtering, and alerting. Only the filtering job becomes a segment; access control becomes IAM policy and boundary (MZ2POL-04), and alerting becomes routing dimensions on a problem-triggered workflow (MZ2POL-09).
2. **There is no automatic conversion** — every segment is hand-authored. Consolidate the estate before converting it.

3. **Match the zone to a scenario** — how the MZ was *defined* determines the approach, from free (host group, Kubernetes) to tagging work (name-based auto-tagging) to the `Segment` tag fallback.
4. **Check the blockers** — exclusions cannot be expressed, custom tags do not reach derived data, entity includes lack `contains` , and problems need an `events` include.
5. **Validate both surfaces** — entity-level and signal-level, before cutover.
6. **Coexist deliberately** — classic apps and alerting profiles still need the Management Zone. Retire it last. The Services app's current Explorer tab, though, is Smartscape-backed and needs no classic-entity include.

Next Steps

Continue to **MZ2POL-06: Migration Execution** for the phased rollout, parallel-running period, cutover, and rollback procedures. For segment mechanics in depth, read **ORGNZ-08** and **ORGNZ-10**; for managing segments as code, **AUTOM-04**. For the alerting job, see **MZ2POL-09**.

Additional Resources

- [Best practice examples: from Management Zones to Segments \(DT docs\)](#)
 - [Segments \(DT docs\)](#)
 - [Segment limits \(DT docs\)](#)
 - [Visibility of segments \(DT docs\)](#)
 - [Include data in segments \(DT docs\)](#)
 - [Segment data by Kubernetes clusters \(DT docs\)](#)
 - [Filter Smartscape nodes with segments \(DT docs\)](#)
 - [Services app \(DT docs\)](#)
 - [Dynatrace default policies reference \(DT docs\)](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.