MZ2POL-99: Best Practice Summary

> **Series:** MZ2POL — Management Zone to Policy Migration | **Notebook:** 99 | **Created:** March 2026 | **Last Updated:** 10/13/2026

Overview

This notebook consolidates the actionable practices from the whole MZ2POL series — notebooks 00–13, the MZ2POL-14 workshop and the MZ2POL-95 LAB — into one checklist. Each row gives the practice, the concrete setting or value, a priority, and the notebook that explains and sources it.

It was rewritten on 10/13/2026. The previous version recommended practices the series has since corrected: splitting zone and Grail conditions into separate boundaries as a rule, describing boundary lines as always OR-combined, measuring security-context coverage on entities, setting security context with auto-tagging rules, and a fixed 2–4-week parallel run. Those rows are gone; where the correction matters, the row that replaces it says so.

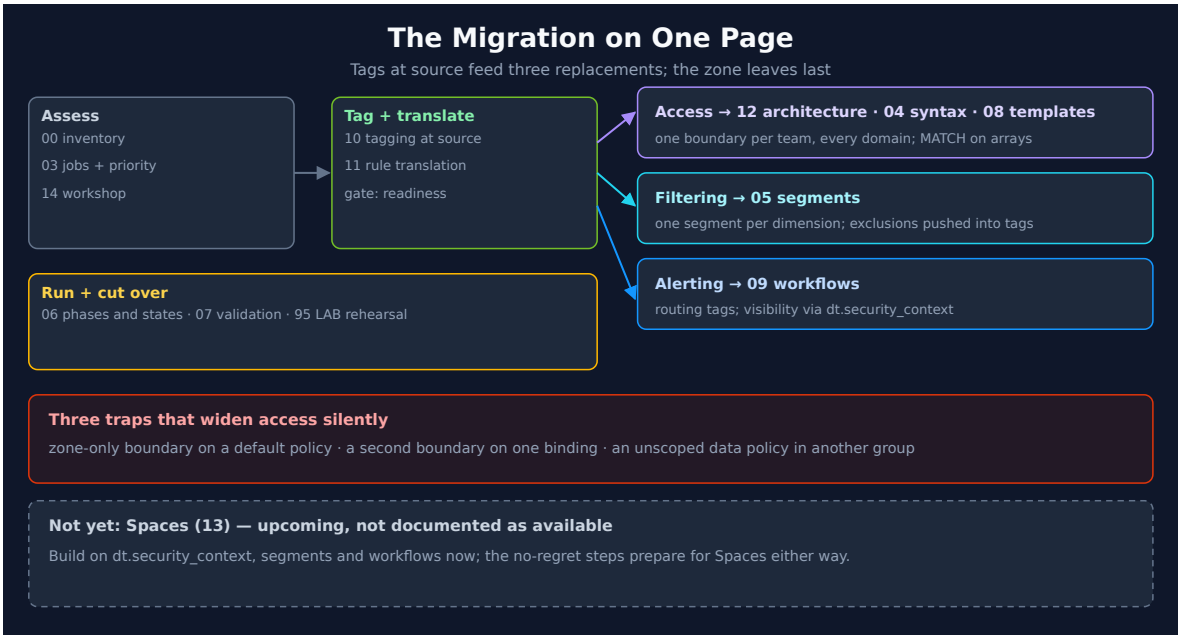


Table of Contents

1. Model and Assessment
2. Tagging at Source
3. Rule Translation
4. Access: Policies and Boundaries
5. Settings Access
6. Filtering: Segments
7. Buckets
8. Alerting
9. Execution and Cutover
10. Validation and Ongoing Checks
11. Spaces

Prerequisites

Requirement	Details
Use	A checklist alongside the series — every row names the notebook that explains and sources it
Status	Reflects the series as of 10/13/2026, including the corrections in MZ2POL-04, -06 and -12

1. Model and Assessment

Practice	Setting / value	Priority	Notebook
Classify every zone by job before anything else	Access (removing it exposes data), filtering (removing it adds noise), alerting (it decides who is paged); a zone can do several	Critical	03, 05 §1, 14
Delete unused zones instead of converting them	No queries, role assignments, profiles or dashboards in 90 days → archive the export, delete	Critical	05 §2, 14

Practice	Setting / value	Priority	Notebook
Know what does not carry forward	Auto-tagging rules have no effect on Grail data; zones are not supported by any platform feature or Latest app	Critical	10 §2
Prioritize by risk	P1 security/compliance access · P2 team access and alerting scope · P3 filtering only · P4 unused	Critical	03
Export rules with the SDK tool	`builtin:management-zones` flattened per rule	Recommended	00
Inventory dependencies per zone	Role assignments, alerting profiles, dashboards, SLOs, API calls, automation	Critical	03, 06
Run the disposition workshop with the owners	Six sessions, six tables, stop list	Recommended	14

2. Tagging at Source

Practice	Setting / value	Priority	Notebook
Put the context on the data before building anything	Every replacement reads a field on the record; enrichment is forward-only	Critical	10 §1, §8
Use built-in fields first	`dt.host_group.id`, `k8s.namespace.name`, `k8s.cluster.name`, cloud account fields — no configuration	Critical	10 §5
Promote standard tags	Cloud tags, K8s labels, plain host tags live only on Smartscape nodes — promote to `primary_tags.*` or a special field	Critical	10 §3
Write access dimensions twice, in one step	`primary_tags.team` (filter, route) **and** `dt.security_context` (IAM) from the same source	Critical	10 §4

Practice	Setting / value	Priority	Notebook
Climb the enrichment ladder from the top	Built-in → source (`--set-host-tag`, `DT_TAGS`, K8s annotations, OTel attributes) → central Ingest enrichment configuration (OneAgent 1.343+) → OpenPipeline primary Grail tag rule	Critical	10 §5
Guard OpenPipeline tag rules	Condition `isNull(primary_tags.<key>)` — rules overwrite existing values by default	Critical	10 §5
Set routing-critical values at source or centrally	Docs disagree on whether OpenPipeline tag rules run before routing	Recommended	10 §5
Constrain values	Allowlists for stage, region, classification; Grail compares exactly	Critical	10 §4
Stay under the tag cap	Fewer than 20 primary tags per host or process — extras are dropped silently	Critical	10 §8
Move host properties to host tags	From OneAgent 1.345 host `primary_tags` come from host tags	Recommended	11 §3
Pass the readiness gate	Parity by value, per-zone acceptance, coverage by host group (incl. no-host data), drift, derived-data reach	Critical	10 §9–§10

3. Rule Translation

Practice	Setting / value	Priority	Notebook
Translate rules, not zones	A zone is ORed rules of ANDed conditions; each rule may land elsewhere	Critical	11 §1
Take the first yes on the decision path	Built-in field → set at source → central rule → OpenPipeline tag rule → explicit membership tag	Critical	11 §2
Push exclusions into enrichment	`NOT` in the central-rule condition; filter on a positive value	Critical	11 §5

Practice	Setting / value	Priority	Notebook
Compute what no surface can evaluate	Regex, numeric comparisons, IP ranges → a value set at the source	Critical	11 §4
Replace propagation flags deliberately	Set the value where it is true (host or process); propagation is built in and downward	Recommended	10 §7
Treat selector zones by purpose	Entity lists → Smartscape queries; data scope → a tag	Recommended	11 §7
Keep the per-rule worksheet	Zone, rule, step, target, owner, segment, boundary, verified	Critical	11 §10

4. Access: Policies and Boundaries

Practice	Setting / value	Priority	Notebook
Scope Grail data by record fields	Grail does not evaluate zones; `dt.security_context` is the only field every table accepts	Critical	12 §1
Default policies for "what", boundaries for "where"	Standard / Pro / Admin User + Read Logs, Read Spans, Read Metrics, Read Events, Settings Reader/Writer	Critical	04 §3, 12 §6
One boundary covers every domain its policies grant	`storage:` line, `settings:` line, and `environment:management-zone` while classic pages are used. *Corrects the earlier "never mix, two bindings" rule* — mixing is documented; a missing domain is unscoped	**Critical**	04 §6, 12 §4
Never attach a zone-only boundary to a default policy	Default policies carry classic and Grail statements; the Grail ones come out unscoped	**Critical**	12 §4.2

Practice	Setting / value	Priority	Notebook
Don't add a second boundary for shared data	Boundaries are computed separately; add the shared value to the same condition	Critical	12 §4.1, §6
Use `MATCH` on `storage:dt.security_context`	Events hold it as an array; `=`, `IN`, `startsWith` never match arrays	Critical	12 §6, 08 §2
Repeated lines on one condition multiply (OR)	*Corrects "every line is OR-combined"* — lines on different conditions each scope their own permissions	Recommended	04 §4–§5
Expect no negation on record fields	Record conditions take `=`, `IN`, `startsWith`, `MATCH`; negation exists only on `storage:bucket-name`	Recommended	04 §4, 11 §4
Keep team members out of unscoped data groups	One unscoped table permission anywhere wins	Critical	12 §3
No record-level conditional `DENY`	It runs as an unconditional DENY on the whole table	Critical	04 §2, 12 §3
Entities scope only by security context	`storage:entities:read` accepts `dt.security_context` and `entity.type`	Recommended	12 §4.1
Groups by team × role from the IdP	Not one group per zone	Critical	12 §6
Template when teams are many	`${bindParam:…}` via REST API; bound policies editable only while parameters are unchanged	Recommended	08
Inventory classic roles that see all zones	`agent-install`, `configure-request-capture-data`	Recommended	12 §8

5. Settings Access

Practice	Setting / value	Priority	Notebook
Replace zone-scoped settings conditions	`environment:management-zone` and MZ-derived security contexts stop matching settings as the Smartscape transition arrives (rolling out)	Critical	12 §5
Use edge-set security context for settings	`settings:dt.security_context` matches only where the context was set at the source	Critical	12 §5
Address host groups by name	`settings:scope = "dt.host_group.id:<name>"` or `settings:entity.hostGroup`	Recommended	12 §5
Fall back to schema scoping	`settings:schemaId` where no edge-set context exists (e.g. process groups)	Recommended	12 §5

6. Filtering: Segments

Practice	Setting / value	Priority	Notebook
Segments are views, never access	Anyone with `storage:filter-segments:read` can apply any segment	Critical	05 §1, §9
One segment per dimension	Variables from live queries; usually three to eight segments	Critical	05 §2
Start with `Data (all types)`	Classic-entity includes only for backward compatibility	Recommended	05 §3
Use only `=` and `in()`	Wildcard `*` inside the value; no exclusions	Critical	05 §3
Add an events include for problems	`event.kind = "DAVIS_PROBLEM"` AND the condition	Critical	05 §5.4
Filter derived data on primary fields or tags	Standard tags do not reach service metrics	Critical	05 §5.2

Practice	Setting / value	Priority	Notebook
Validate at entity and signal level	Before creating the segment	Critical	05 §7

7. Buckets

Practice	Setting / value	Priority	Notebook
Grant bucket and table permissions together	`storage:buckets:read` plus the table permission; record conditions on top	Critical	12 §1
Use buckets for isolation or retention, not for views	Views are segments; bucket design is in ORGNZ-02/03	Recommended	02 §3.5, ORGNZ-03
Route to buckets with values present before routing	Source or central enrichment	Recommended	10 §5

8. Alerting

Practice	Setting / value	Priority	Notebook
Replace zone-scoped profiles with problem-triggered workflows	Trigger filters on affected-entity tags plus an optional DQL matcher; segments do not scope alerting	Critical	09 §2
Tag entities before building workflows	Routing tags at full coverage	Critical	09 §4
Map delay to Minimum duration	Fixed values; decide rounding per profile	Recommended	09 §6.1
Keep visibility on `dt.security_context`	With `MATCH` — problem records hold it as an array	Critical	09 §7
Prove parity, disable, soak, then delete zones	Disable and delete in separate change windows	Critical	09 §8

Practice	Setting / value	Priority	Notebook
Test zone deletion on a profile in non-prod	Fail-open vs fail-closed is undocumented	Critical	09 §8

9. Execution and Cutover

Practice	Setting / value	Priority	Notebook
Follow the phase order with gates	Assess → tag + translate (gate) → access / filtering / alerting → parallel run → cutover → cleanup	Critical	06 §1
Move each team through four states	Today → dual boundary → verified → zone line removed	Critical	06 §6, 12 §9
Check "today" first	A zone-only boundary on a default policy may already grant all Grail data	Critical	12 §9
Exit states on evidence, not on a calendar	Tag gate, effective permissions, alert parity. *Replaces the fixed "2–4 weeks"*	Critical	06 §6
Pilot every boundary change on one group	Boundary edits apply immediately to every binding	Critical	12 §9
Rehearse on the LAB estate	Steps 1–8	Recommended	95
Keep a rollback line per batch	Re-add the zone line; re-enable the profile	Critical	06 §9
Delete zones last, in their own window	After profiles are disabled and dashboards repointed	Critical	06 §8, 09 §8

10. Validation and Ongoing Checks

Practice	Setting / value	Priority	Notebook
Measure security context on **records**, not entities	Per table: logs, spans, events. *Replaces entity coverage* — entity contexts can be copied from zone names and read 100% while data reads far less	**Critical**	06 §3, 12 §10
Check the type per table	Any `array` row needs `MATCH`	Critical	12 §10
Compare zone and tag per zone on data	The acceptance test: match / in zone, tag missing / tagged, not in zone	Critical	10 §9, 06 §10
Resolve effective permissions per persona	Account API, one pilot user each; no `storage:*:read` without a condition	Critical	12 §10
Test as a pilot user, never an admin	Admin grants hide boundary problems	Critical	07 §5
Re-run drift and coverage on a schedule	New producers arrive unenriched	Recommended	10 §9

11. Spaces

Practice	Setting / value	Priority	Notebook
Do not design access around Spaces yet	`dt.space` is an `experimental` dictionary entry; no docs page, IAM condition or release entry as of 10/13/2026	Critical	13 §1–§2
Do not keep zones waiting for Spaces	Zones are not evaluated for Grail data	Recommended	13 §4
Prepare with no-regret steps	Team values at source, one boundary per team, segments per dimension, ownership tags, inventory of delegated responsibility	Recommended	13 §4

Practice	Setting / value	Priority	Notebook
Re-check each platform release	Dictionary query with a control row	Optional	13 §5

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.