

ONBRD-00: Architect's Sequence & Dependency Runbook

Series: ONBRD — Dynatrace Onboarding | **Reference:** 00 — Architect's Sequence & Dependency Runbook | **Created:** May 2026 | **Last Updated:** 09/18/2026

Overview

Sequence-and-dependency runbook for an architect or experienced tenant lead onboarding a new Dynatrace tenant in 2026. Assumes you have onboarded a tenant before — this is order, dependencies, decisions, not a GUI tour.

Per-domain depth lives in the canonical series (linked at each step). For recommended defaults, the decision matrix, validation DQL, anti-patterns, and the cross-series next-steps map, see **ONBRD-99: Best Practice Summary**.

Table of Contents


1. [The Sequence](#)
 2. [Phase 1 — Foundation](#)
 3. [Phase 2 — Organize](#)
 4. [Phase 3 — Operate](#)
 5. [Phase Checkpoints](#)
 6. [Common Variants & Skip Rules](#)
 7. [After ONBRD](#)
-

Prerequisites

Requirement	Details
Tenant URL + admin credentials	from procurement / account team
Account-level decisions in place	tenant region, SaaS vs Managed, account hierarchy. If not yet decided, see NR2DT-00: Tenant Prerequisites — Minimum Setup Before Migration — the same checklist generalizes beyond New Relic migrations.
Reading time	30–45 minutes for the runbook; the actual onboarding is 1–4 weeks depending on scope and skip rules applied.

1. The Sequence

ONBRD is 10 steps grouped into 3 phases. Some are sequential, some parallel-able after their predecessor. Two phase gates separate the phases.



Phase	Steps	Phase Output
Phase 1 — Foundation	01 → 02 → {03, 04, 05}	Hosts reporting; tenant lead can query data

Phase	Steps	Phase Output
Phase 2 — Organize	06 → 07	Entities tagged and scoped; <code>dt.security_context</code> populated
Phase 3 — Operate	{08, 09, 10}	Validation DQL set + at least one workflow + at least one dashboard

2. Phase 1 — Foundation

Step 01 — First Steps

- **Decision:** Confirm access; record tenant URL; plan token strategy for downstream automation.
- **Depends on:** procurement complete, tenant provisioned.
- **2026 specifics:** Plan for Platform Tokens (`dt0s16` , `Authorization: Bearer`) as the default for new automation per sprint-1.337. Settings v2 / Configuration as Code (Terraform / Monaco) over Configuration API for new work. **Extensions 2.0** for any custom integrations (Extensions Framework 1.0 reached end of support on 2025-03-31).
- **Output:** confirmed tenant access; first Platform Token decision recorded.
- **Deep dive:** **ONBRD-99: Best Practice Summary** § 1 Recommended Defaults for the 2026 token / API / extension stack.

Step 02 — IAM and Authentication

- **Decision:** SSO method (SAML / OIDC); initial group structure; first parameterized policies.
- **Depends on:** Step 01 (access confirmed).
- **2026 specifics:** Issue the first Platform Token (`dt0s16` , `Authorization: Bearer`) — Classic API Tokens (`dt0c01`) only for legacy installer-download paths. Use parameterized policies bound to groups via binding parameters; the `dt.security_context` field is the standardized boundary for Gen3 IAM scoping (configured fully in Step 06).
- **Output:** SSO live; admin group + first parameterized policy in place; first Platform Token issued.

- **Deep dive:** IAM series for full IAM administration (15 notebooks); **FAQ-02: Tagging — Sources, Standards, and Strategy** for `dt.security_context` strategy.

Step 03 — ActiveGate (parallel-able with 04 + 05 after 02)

- **Decision:** Need ActiveGate? Required for >500 hosts, hybrid / on-prem, or cloud-API polling at scale. Defer if pure SaaS-only with no on-prem footprint.
- **Depends on:** Step 02 (Platform Token to register AG).
- **2026 specifics:** Sizing baseline 10–20 GB; 2–3 AGs per zone for load distribution + failover. OneAgent attribute enrichment (OneAgent 1.333+) emits primary fields on every signal at ingest — verify AG version supports it.
- **Output:** at least one AG registered; cluster connectivity verified; AG visible in Deployment Status.
- **Deep dive:** CLOUD series for AG-routed cloud integrations; AUTOM series for AG config-as-code.

Step 04 — Cloud & SaaS Integrations (parallel-able with 03 + 05 after 02)

- **Decision:** Which clouds (AWS / Azure / GCP) and integration mechanism per cloud.
- **Depends on:** Step 02 (Platform Token); optionally Step 03 (AG for legacy paths).
- **2026 specifics:** Clouds app for AWS (GA) and Azure (direct connection since SaaS 1.337 — no ActiveGate for metric polling); the GCP Clouds-app connection is in **Preview** — verify it has reached your tenant; the classic ActiveGate-based GCP integration remains the working path until then. AWS Lambda primary-tag propagation via `DT_TAGS` env var (sprint-1.337). IAM Role pattern for production multi-account.
- **Output:** cloud entities visible in Smartscape; cloud tags arriving as `aws.tag.*` / `azure.tag.*` / `gcp.label.*`.
- **Deep dive:** CLOUD series for per-provider depth; **FAQ-02: Tagging — Sources, Standards, and Strategy** for cross-cloud tag normalization.

Step 05 — OneAgent Deployment (parallel-able with 03 + 04 after 02)

- **Decision:** Deployment method (direct, package manager, Operator for K8s); initial scope (pilot or full); host-tag taxonomy.
- **Depends on:** Step 02 (Platform Token for installer download).

- **2026 specifics: Set primary fields/tags at install** via the single `--set-host-tag` form the June-2026 tags hub documents for both: `oneagentctl --set-host-tag="primary_tags.<key>=<value>"` (prefix written explicitly — it is never added automatically) and `oneagentctl --set-host-tag="dt.security_context=<value>"` — not retroactively. Primary tags emit on every signal (metrics / spans / logs / events) at ingest. Kubernetes: Dynatrace Operator + **Cloud Native FullStack** mode (Classic FullStack deprecated for new deployments). DynaKube **v1beta5** baseline; **v1beta6** as canary-validated upgrade. From **OneAgent 1.337**, Windows OneAgent requires **Npcap** for Network Agent metrics (WinPcap is no longer supported).
- **Output:** OneAgent installed on at least one host group; primary tags propagating to every signal at ingest.
- **Deep dive:** K8S series for DynaKube + Cloud Native FullStack; **FAQ-01: Why you need a good Host Group naming strategy** for host-group naming; **FAQ-02: Tagging — Sources, Standards, and Strategy** for primary-fields strategy; **FAQ-03: OneAgent vs OpenTelemetry — A Decision Framework** for OneAgent vs OTel decision.

3. Phase 2 — Organize

Step 06 — Organizing Your Environment

- **Decision:** Tag taxonomy (env / team / app / compliance / cost-center dimensions); host-group naming convention; bucket strategy; segment definitions.
- **Depends on:** Step 05 (OneAgent hosts reporting); Step 04 (cloud tags arriving) if cross-cloud normalization is needed.
- **2026 specifics:** Primary fields/tags via `oneagentctl --set-host-tag="primary_tags.<key>=<value>"` and `--set-host-tag="dt.security_context=<value>"` (June-2026 tags-hub form; the `oneagentctl` reference still documents the older `--set-host-property` form for `dt.*` keys). Use **Segments + `dt.security_context`** for data scoping — *not* legacy Management Zones. **MZ-on-calculated-metrics** is on the May-2026 deprecation list. Bucket strategy: `dt.security_context` is the default for general data access; buckets only for compliance / retention / hard cost / hostile multi-tenancy.
- **Output:** entities grouped meaningfully; `dt.security_context` populated on signals; segments defined for primary scopes; bucket strategy documented.
- **Deep dive:** ORGNZ series — buckets, segments, `dt.security_context` (full depth, 11 notebooks); IAM series — parameterized policies bound to `dt.security_context`; **FAQ-01: Why you need a good Host Group naming strategy** — host-group naming; **FAQ-02: Tagging — Sources, Standards, and Strategy** — tagging strategy.

Step 07 — Understanding Your Data

- **Decision:** Which signals matter for the workloads in scope (spans / logs / metrics / events / business events). Confirm primary fields propagate.
- **Depends on:** Step 06 (entities grouped); Step 05 (OneAgent producing signals).
- **2026 specifics:** OneAgent attribute enrichment (OneAgent 1.333+) emits primary fields on every signal at ingest — spot-check `dt.security_context` on logs / spans / metrics with the validation DQL in **ONBRD-99: Best Practice Summary** § 3.
- **Output:** confirmed which signal types are flowing; one validation DQL per signal type saved.
- **Deep dive:** SPANS series, OPLOGS series, OPIPE series, OPMIG series — per-signal-type depth.

4. Phase 3 — Operate

Step 08 — Validation Queries (parallel-able with 09 + 10 after 07)

- **Decision:** Which canonical queries become the deployment's validation set (host count, service inventory, log volume per source, primary-field propagation).
- **Depends on:** Step 07 (signals understood); Step 06 (entities grouped).
- **Output:** validation DQL set saved as a notebook or dashboard (the architect's smoke-test for the tenant).
- **Deep dive:** **ONBRD-99: Best Practice Summary** § 3 Validation Queries for the canonical set; SPANS, DBMON, WEBRUM, BIZEV for per-domain DQL patterns.

Step 09 — Workflows & Alerts (parallel-able with 08 + 10 after 07)

- **Decision:** Which Davis-detected problems become workflows; which custom anomaly detectors are needed; notification targets; on-call routing.
- **Depends on:** Step 07 (data flowing); Step 06 (entities grouped for scoping).
- **2026 specifics:** Workflows over Alerting Profiles. Davis Anomaly Detectors over legacy metric-events (where the signal is supported).
- **Output:** at least one workflow with notification target; first alert fires.
- **Deep dive:** WFLOW series — workflows + notifications; AIOPS series — Davis problems + anomaly detection + Davis CoPilot.

Step 10 — Dashboards (parallel-able with 08 + 09 after 07)

- **Decision:** Tier strategy (Executive / Operations / Engineering); per-tier dashboard structure; sharing model.
- **Depends on:** Step 07 (data understood); Step 06 (entities grouped for filters).
- **Output:** at least one dashboard per tier in scope; shared with the right groups.
- **Deep dive:** DASH series — dashboard strategy + executive reporting + tile-type decision.

5. Phase Checkpoints

Verify before moving from one phase to the next. The DQL queries in **ONBRD-99: Best Practice Summary** § 3 Validation Queries make these gates concrete.

G1 — Foundation → Organize

The tenant has at least one host reporting and the tenant lead role can query that data.

- *Verify:* `smartscapeNodes "HOST"` returns host count > 0; tenant lead role can run that query.
- *If not:* stay in Phase 1 — most likely OneAgent installer didn't run, or the Platform Token used by the installer lacks `InstallerDownload` scope.

G2 — Organize → Operate

Entities are grouped meaningfully and `dt.security_context` is populated on signals.

- *Verify:* `fetch logs, from:-1h | summarize count(), by:{dt.security_context}` returns at least one non-null value.
- *If not:* OneAgent install or OpenPipeline enrichment didn't propagate the field — fix before Phase 3 (filter scoping in dashboards / alerts will not work without it).

After Phase 3 — exit ONBRD

ONBRD ends. The next checkpoint is *"do you need a deeper series for the domains in scope?"* — see **ONBRD-99: Best Practice Summary** § 5 Where to Go Deeper for the 32-series cross-reference map.

6. Common Variants & Skip Rules

ONBRD — Variants & Skip Rules

The default path runs all 10 steps. Common situations modify it as shown.

Default — Standard New Tenant <i>No special starting context.</i> Modified path: ▶ All 10 steps in order. ▶ Phase 1: 01 → 02 → {03/04/05} ▶ Phase 2: 06 → 07 ▶ Phase 3: {08, 09, 10} parallel No skip rules apply. Use the canonical sequence and gates G1, G2 as written.	Existing SSO from another DT tenant <i>Org has SAML/OIDC live elsewhere.</i> Modified path: ▶ Step 02: skip IDP-config; reuse the existing IDP. ▶ Still configure groups + parameterized policies in this tenant. All other steps unchanged. Saves 1–3 days vs full IDP setup.	Pure SaaS-only, no on-prem <i>All workloads in cloud SaaS.</i> Modified path: ▶ Step 03 (ActiveGate) optional — defer until cloud-API polling at scale demands it. ▶ All other steps unchanged. Reassess Step 03 if the host count grows past ~500 or extensions appear.
Cloud-only / Serverless <i>No VMs / containers under your control.</i> Modified path: ▶ Step 05 (OneAgent) mostly N/A. ▶ Focus on Step 04 (Clouds app) and OpenTelemetry for app code. ▶ See FAQ-03 for OneAgent vs OTel. Phase 2 still needed — primary tags via OpenPipeline enrichment, not at install.	Kubernetes-first <i>Workloads run primarily on K8s.</i> Modified path: ▶ Step 05 collapses into K8S series. ▶ Operator + DynaKube replaces per-host install. ▶ Cloud Native FullStack mode default. DynaKube v1beta5 baseline; v1beta6 as canary-validated upgrade.	Migrating from another tool <i>Replacing NR / Splunk / Sumo / Managed.</i> Modified path: ▶ Run ONBRD foundation in parallel with the migration series: NR2DT / S2D / SL2DT / M2S / S2S. ▶ Migration series owns cutover. ONBRD provides the destination tenant shape; migration provides the workload.

Variant	When	What to do
Existing SSO already configured	New tenant for an org with SAML/OIDC live on another Dynatrace tenant	Skip the IDP-config part of Step 02; reuse the existing IDP. Still configure groups + parameterized policies in the new tenant.
Pure SaaS-only, no on-prem	All workloads in cloud; no private/on-prem footprint	Step 03 (ActiveGate) is optional — defer until cloud-API polling demands it.
Cloud-only, no hosts	Serverless / managed services only; no VMs or containers under your control	Step 05 (OneAgent) is mostly inapplicable — focus on Step 04 (cloud integrations) and OpenTelemetry for application code. See FAQ-03: OneAgent vs OpenTelemetry — A Decision Framework .
Kubernetes-first	Workloads run primarily on K8s	Step 05 collapses into K8S series — Operator + DynaKube replaces per-host install.
Migrating from another tool	Replacing New Relic / Splunk / Sumo Logic / managed Dynatrace	Run ONBRD foundation in parallel with the migration series (NR2DT, S2D, SL2DT, M2S, S2S).

Variant	When	What to do
Multi-tenant org	Several Dynatrace tenants for different BUs / environments	Step 02 expands — see IAM-08: Multi-Environment IAM for cross-tenant patterns.

7. After ONBRD

ONBRD ends after Phase 3. The next set of decisions is per-domain — observability, automation, AI/Davis, dashboards-at-scale, etc.

Reference points to keep open:

- **ONBRD-99: Best Practice Summary** — the architect's reference card: recommended defaults, decision matrix, validation queries, anti-patterns, and the **32-series cross-reference map** for what to do after ONBRD.
 - **-START-HERE- › Foundation Module** — cross-series Foundation Module reading order (ONBRD + ORGNZ + IAM in parallel) with priorities and skip rules.
 - **-START-HERE- › Doorway 1 — Net New** — doorway selection for greenfield vs migration paths (Sub-Paths A–E).
 - **ADOPT series** — overall maturity and adoption roadmap once ONBRD is complete.
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against the current [Dynatrace documentation](#).