

ONBRD-10: Building Dashboards

Series: ONBRD — Dynatrace Onboarding | **Notebook:** 10 of 10 | **Created:** December 2025 | **Last Updated:** 10/02/2026

Visualizing Your Data

Dashboards provide at-a-glance visibility into your environment's health and performance. This notebook covers dashboard creation, common visualization patterns, and sharing with your team.

Table of Contents

1. [Dashboards vs Notebooks](#)
 2. [Creating Your First Dashboard](#)
 3. [Common Tile Types](#)
 4. [Dashboard Patterns](#)
 5. [Useful Queries for Dashboards](#)
 6. [Sharing and Permissions](#)
 7. [Next Steps](#)
-

Prerequisites

- Viewer access or higher
- Familiarity with DQL (ONBRD-08)
- Data in your environment to visualize

1. Dashboards vs Notebooks

Dynatrace offers two visualization tools:

Feature	Dashboard	Notebook
Purpose	Monitoring at a glance	Interactive analysis
Layout	Fixed grid tiles	Sequential document
Auto-refresh	Yes	Manual
Interactivity	Click-through links	Query editing
Sharing	Wall displays, reports	Investigation documentation
Best for	NOC screens, status pages	Troubleshooting, exploration

When to Use Each

Scenario	Use
Team status display	Dashboard
Incident investigation	Notebook
Executive summary	Dashboard
Root cause analysis	Notebook
SLA reporting	Dashboard
Ad-hoc queries	Notebook

2. Creating Your First Dashboard

Location: Observe and explore → Dashboards → Create dashboard

Dashboard Creation Steps

1. Click "Create dashboard"
2. Name your dashboard (e.g., "Production Overview")

3. Add tiles by clicking "+" or dragging
4. Configure each tile
5. Arrange and resize as needed
6. Save

Dashboard Settings

Setting	Description
Name	Dashboard title
Time frame	Default time range
Segment	Default filter scope
Refresh rate	Auto-refresh interval
Owner	Dashboard owner
Sharing	Who can view/edit

3. Common Tile Types

Data Tiles

Tile Type	Use Case
Single value	Key metrics (uptime, error count)
Graph	Time series data
Table	Lists and details
Top list	Ranked entities
Pie chart	Distribution
Honeycomb	Entity health grid

Content Tiles

Tile Type	Use Case
Markdown	Documentation, links
Problems	Active problem list
Host health	Infrastructure status
Service health	Service status
SLO	Service level objectives

Visualization Types

Visualization Selection Guide

Choose the right tile type for your data

What do you want to show?

Single number

Trend over time

Compare values

Distribution

Top N items

Detailed list

Single Value

Line/Area Chart

Bar Chart

Pie Chart

Top List

Table

Entity status

Geographic

Active problems

SLO status


Honeycomb

Map

Problems Tile

SLO Tile

4. Dashboard Patterns



Pattern: Executive Summary

Key elements:

- **KPI row:** Uptime, Errors, Avg Response, Active Problems
- **Trend chart:** Response time over 24h
- **Problem list:** Current issues
- **Service health:** Status indicators for key services

Pattern: Service Dashboard

Key elements:

- **KPI row:** Requests, Error Rate, P95 Response, Apdex
- **Timeline:** Request rate and errors over time
- **Top endpoints:** Most called API endpoints
- **Dependencies:** Downstream service health

Pattern: Infrastructure Dashboard

Key elements:

- **KPI row:** Host count, Avg CPU, Avg Memory, Disk Alerts

- **Honeycomb views:** Visual host health for CPU and memory
- **Top consumers:** Hosts using most resources

5. Useful Queries for Dashboards

These queries work well as dashboard tiles.

Service Health Queries

```
// Total request count (Single Value tile)
fetch spans, from: now() - 1h
| filter span.kind == "server"
| summarize request_count = count()
```

```
// Error rate percentage (Single Value tile)
fetch spans, from: now() - 1h
| filter span.kind == "server"
| summarize
    total = count(),
    errors = countIf(span.status_code == "error")
| fieldsAdd error_rate = round(100.0 * errors / total, decimals: 2)
```

```
// Top services by request count (Top List tile)
fetch spans, from: now() - 1h
| filter span.kind == "server"
| summarize requests = count(), by: {service.name}
| sort requests desc
| limit 10
```

```
// Service error rates (Table tile)
fetch spans, from: now() - 1h
| filter span.kind == "server"
| summarize
    requests = count(),
    errors = countIf(span.status_code == "error"),
    by: {service.name}
| fieldsAdd error_rate = round(100.0 * errors / requests, decimals: 2)
| sort error_rate desc
| limit 15
```

Infrastructure Queries

```
// Host count (Single Value tile)
fetch dt.entity.host
| filter state == "RUNNING"
| summarize host_count = count()

// Smartscape note (dt.entity.* is deprecated but still functional): this
query uses the
// classic-only field state, which has NO Smartscape node equivalent
// (Smartscape expresses liveness via node lifetime, not a state field). Keep
the classic
// query above for state detail. For monitoring mode, do not use
monitoringMode – it is
// empty on Kubernetes and Fargate hosts; read billing events instead (ONBRD-
05 § 6).
```

```
// Hosts by OS type (Pie Chart tile)
fetch dt.entity.host
| summarize count = count(), by: {osType}
| sort count desc

// Smartscape equivalent (dt.entity.* is deprecated but still functional):
//   smartscapeNodes "HOST"
//   | summarize count = count(), by: {os.type}
//   | sort count desc
// Caveat: Smartscape reflects CURRENT live topology and can report fewer
entities
// than the classic entity store; for a pre-migration discovery inventory keep
the
// classic query above.
// Field maps: osType -> os.type (LINUX -> OS_TYPE_LINUX).
```

```
// Host inventory (Table tile)
fetch dt.entity.host
| fields name = entity.name, os = osType, state, cpuCores
| sort name
| limit 25

// Smartscape note (dt.entity.* is deprecated but still functional): this
query uses the
// classic-only field state, which has NO Smartscape node equivalent
// (Smartscape expresses liveness via node lifetime, not a state field). Keep
the classic
// query above for state detail. For monitoring mode, do not use
monitoringMode – it is
// empty on Kubernetes and Fargate hosts; read billing events instead (ONBRD-
05 § 6).
// Other fields do map: osType -> os.type (LINUX -> OS_TYPE_LINUX);
cpuCores -> cores; entity.name -> name.
```

Log Queries

```
// Error log count (Single Value tile)
fetch logs, from: now() - 1h
| filter loglevel == "ERROR"
| summarize error_count = count()
```

```
// Log volume by severity (Pie Chart tile)
fetch logs, from: now() - 1h
| summarize count = count(), by: {loglevel}
| sort count desc
```

```
// Recent errors (Table tile)
fetch logs, from: now() - 1h
| filter loglevel == "ERROR"
| fields timestamp, log.source, content
| sort timestamp desc
| limit 10
```

Problem Queries

```
// Active problem count (Single Value tile)
fetch dt.davis.problems, from: now() - 30d
| filter event.status == "ACTIVE"
| summarize active_problems = count()
```

```
// Problems by status (Pie Chart tile)
fetch dt.davis.problems, from: now() - 7d
| summarize problem_count = count(), by: {event.status}
```

```
// Recent problem list (Table tile)
fetch dt.davis.problems, from: now() - 24h
| fields timestamp, display_id, event.name, event.status
| sort timestamp desc
| limit 10
```

6. Sharing and Permissions

Sharing Options

Sharing Level	Who Can Access
Private	Only you
Shared with specific users	Named users
Shared with groups	User groups
Public	Anyone with environment access

Segment Filtering

Dashboards can be filtered using segments:

- Set a default segment in dashboard settings
- Viewers can switch segments if they have access
- Data respects viewer's permissions

Dashboard Presets

Create presets for common views:

- 1. Set segment filter
- 2. Set time range
- 3. Save as preset with descriptive name

Exporting Dashboards

Format	Use Case
PDF	Reporting, documentation
JSON	Backup, migration
Link	Sharing (respects permissions)

7. Next Steps

Congratulations! You've completed the onboarding series.

What You've Learned

Notebook	Key Topics
ONBRD-01	Tenant access, navigation
ONBRD-02	IAM, SAML, Platform Tokens, parameterized policies
ONBRD-03	ActiveGate deployment
ONBRD-04	Cloud & SaaS integrations
ONBRD-05	OneAgent deployment, primary tags at source
ONBRD-06	Tags, segments, naming conventions, <code>dt.security_context</code>
ONBRD-07	Data model, topology
ONBRD-08	DQL fundamentals
ONBRD-09	Workflows alerting

Notebook	Key Topics
ONBRD-10	Dashboard creation

Where to Go Next — by Topic

Goal	Topic Series
Dashboard strategy + executive reporting	DASH series (8 notebooks)
Davis AI / anomaly detection / RCA	AIOPS series (8 notebooks)
Workflow automation + AI tasks	WFLOW series (12 notebooks)
Deepen DQL — spans / logs / OpenPipeline	SPANS, OPLOGS, OPMIG, OPIPE
Synthetic monitoring	SYNTH series
Web RUM / Mobile RUM	WEBRUM, MOBL
Business events & funnel analysis	BIZEV series
Database monitoring	DBMON series
Cloud integration deep dives	CLOUD series
Kubernetes monitoring	K8S series
OpenTelemetry integration	OTEL series
IAM administration depth	IAM series
Configuration automation / GitOps	AUTOM series
Access control modernization	MZ2POL series (Management Zones → Policies)
Migration playbooks	NRLC + NR2DT (New Relic), SL2DT (Sumo), S2D (Splunk), M2S (Managed → SaaS)
Platform maturity / adoption	ADOPT series

Goal	Topic Series
Official training	Dynatrace University
Community	community.dynatrace.com

Dashboard Checklist

- ☐ First dashboard created (using **Dashboards (new)**, not legacy Classic Dashboards)
 - ☐ Key metrics visualized
 - ☐ Problem overview included
 - ☐ SLO / SLI tile evaluated for service-health views
 - ☐ Dashboard shared with team
 - ☐ Segment filter set
-

Summary

In this notebook, you learned:

- Difference between Dashboards (new) and Classic Dashboards
 - Difference between dashboards and notebooks
 - How to create and configure dashboards
 - Common tile types and when to use them
 - Dashboard patterns for different use cases
 - DQL queries that work well in dashboards
 - That `dt.davis.problems` queries use `event.status` (`ACTIVE` / `CLOSED`), not `status` (`OPEN`)
 - How to share dashboards with your team
-

References

- [Dashboards](#)
- [Dashboards and notebooks \(DT docs\)](#)

- [Notebooks](#)
 - [Dynatrace Query Language \(DT docs\)](#)
 - [Dashboards and notebooks \(DT docs\)](#)
 - [Dynatrace Community](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.