

# ONBRD-10: Building Dashboards

---

**Series:** ONBRD — Dynatrace Onboarding | **Notebook:** 10 of 10 | **Created:** December 2025 | **Last Updated:** 07/24/2026

## Visualizing Your Data

---

Dashboards provide at-a-glance visibility into your environment's health and performance. This notebook covers dashboard creation, common visualization patterns, and sharing with your team.

---

## Table of Contents

---

1. [Dashboards vs Notebooks](#)
  2. [Creating Your First Dashboard](#)
  3. [Common Tile Types](#)
  4. [Dashboard Patterns](#)
  5. [Useful Queries for Dashboards](#)
  6. [Sharing and Permissions](#)
- 

## Prerequisites

---

- Viewer access or higher
- Familiarity with DQL (ONBRD-08)
- Data in your environment to visualize

# 1. Dashboards vs Notebooks

---

Dynatrace offers two visualization tools:

| Feature       | Dashboard                 | Notebook                     |
|---------------|---------------------------|------------------------------|
| Purpose       | Monitoring at a glance    | Interactive analysis         |
| Layout        | Fixed grid tiles          | Sequential document          |
| Auto-refresh  | Yes                       | Manual                       |
| Interactivity | Click-through links       | Query editing                |
| Sharing       | Wall displays, reports    | Investigation documentation  |
| Best for      | NOC screens, status pages | Troubleshooting, exploration |

## When to Use Each

| Scenario               | Use       |
|------------------------|-----------|
| Team status display    | Dashboard |
| Incident investigation | Notebook  |
| Executive summary      | Dashboard |
| Root cause analysis    | Notebook  |
| SLA reporting          | Dashboard |
| Ad-hoc queries         | Notebook  |

# 2. Creating Your First Dashboard

---

**Location:** Observe and explore → Dashboards → Create dashboard

## Dashboard Creation Steps

1. Click "Create dashboard"
2. Name your dashboard (e.g., "Production Overview")

3. Add tiles by clicking "+" or dragging
4. Configure each tile
5. Arrange and resize as needed
6. Save

## Dashboard Settings

| Setting      | Description           |
|--------------|-----------------------|
| Name         | Dashboard title       |
| Time frame   | Default time range    |
| Segment      | Default filter scope  |
| Refresh rate | Auto-refresh interval |
| Owner        | Dashboard owner       |
| Sharing      | Who can view/edit     |

## 3. Common Tile Types

---

### Data Tiles

| Tile Type    | Use Case                          |
|--------------|-----------------------------------|
| Single value | Key metrics (uptime, error count) |
| Graph        | Time series data                  |
| Table        | Lists and details                 |
| Top list     | Ranked entities                   |
| Pie chart    | Distribution                      |
| Honeycomb    | Entity health grid                |

# Content Tiles

| Tile Type      | Use Case                 |
|----------------|--------------------------|
| Markdown       | Documentation, links     |
| Problems       | Active problem list      |
| Host health    | Infrastructure status    |
| Service health | Service status           |
| SLO            | Service level objectives |

# Visualization Types

Visualization Selection Guide

Choose the right tile type for your data

What do you want to show?

Single number

→

Single Value

Trend over time

→

Line/Area Chart

Compare values

→

Bar Chart

Distribution

→

Pie Chart

Top N items

→

Top List

Detailed list

→

Table

Entity status

Honeycomb

Geographic

Map


Active problems

Problems Tile

SLO status

SLO Tile

## 4. Dashboard Patterns



### Pattern: Executive Summary

Key elements:

- **KPI row:** Uptime, Errors, Avg Response, Active Problems
- **Trend chart:** Response time over 24h
- **Problem list:** Current issues
- **Service health:** Status indicators for key services

### Pattern: Service Dashboard

Key elements:

- **KPI row:** Requests, Error Rate, P95 Response, Apdex
- **Timeline:** Request rate and errors over time
- **Top endpoints:** Most called API endpoints
- **Dependencies:** Downstream service health

### Pattern: Infrastructure Dashboard

Key elements:

- **KPI row:** Host count, Avg CPU, Avg Memory, Disk Alerts

- **Honeycomb views:** Visual host health for CPU and memory
- **Top consumers:** Hosts using most resources

## 5. Useful Queries for Dashboards

---

These queries work well as dashboard tiles.

### Service Health Queries

```
// Total request count (Single Value tile)
fetch spans, from: now() - 1h
| filter span.kind == "server"
| summarize request_count = count()
```

```
// Error rate percentage (Single Value tile)
fetch spans, from: now() - 1h
| filter span.kind == "server"
| summarize
    total = count(),
    errors = countIf(span.status_code == "error")
| fieldsAdd error_rate = round(100.0 * errors / total, decimals: 2)
```

```
// Top services by request count (Top List tile)
fetch spans, from: now() - 1h
| filter span.kind == "server"
| summarize requests = count(), by: {service.name}
| sort requests desc
| limit 10
```

```
// Service error rates (Table tile)
fetch spans, from: now() - 1h
| filter span.kind == "server"
| summarize
    requests = count(),
    errors = countIf(span.status_code == "error"),
    by: {service.name}
| fieldsAdd error_rate = round(100.0 * errors / requests, decimals: 2)
| sort error_rate desc
| limit 15
```

## Infrastructure Queries

```
// Host count (Single Value tile)
fetch dt.entity.host
| filter state == "RUNNING"
| summarize host_count = count()

// Smartscape note (dt.entity.* is deprecated but still functional): this
query uses the
// classic-only fields state / monitoringMode, which have NO Smartscape node
equivalent
// (Smartscape expresses liveness via node lifetime, not a state field). Keep
the classic
// query above for state / monitoring-mode detail.
```

```
// Hosts by OS type (Pie Chart tile)
fetch dt.entity.host
| summarize count = count(), by: {osType}
| sort count desc

// Smartscape equivalent (dt.entity.* is deprecated but still functional):
//   smartscapeNodes "HOST"
//   | summarize count = count(), by: {os.type}
//   | sort count desc
// Caveat: Smartscape reflects CURRENT live topology and can report fewer
entities
// than the classic entity store; for a pre-migration discovery inventory keep
the
// classic query above.
// Field maps: osType -> os.type (LINUX -> OS_TYPE_LINUX).
```

```
// Host inventory (Table tile)
fetch dt.entity.host
| fields name = entity.name, os = osType, state, cpuCores
| sort name
| limit 25

// Smartscape note (dt.entity.* is deprecated but still functional): this
query uses the
// classic-only fields state / monitoringMode, which have NO Smartscape node
equivalent
// (Smartscape expresses liveness via node lifetime, not a state field). Keep
the classic
// query above for state / monitoring-mode detail.
// Other fields do map: osType -> os.type (LINUX -> OS_TYPE_LINUX);
cpuCores -> cores; entity.name -> name.
```

## Log Queries

```
// Error log count (Single Value tile)
fetch logs, from: now() - 1h
| filter loglevel == "error"
| summarize error_count = count()
```

```
// Log volume by severity (Pie Chart tile)
fetch logs, from: now() - 1h
| summarize count = count(), by: {loglevel}
| sort count desc
```

```
// Recent errors (Table tile)
fetch logs, from: now() - 1h
| filter loglevel == "error"
| fields timestamp, log.source, content
| sort timestamp desc
| limit 10
```

## Problem Queries

```
// Active problem count (Single Value tile)
fetch dt.davis.problems, from: now() - 30d
| filter event.status == "ACTIVE"
| summarize active_problems = count()
```



```
// Problems by status (Pie Chart tile)
fetch dt.davis.problems, from: now() - 7d
| summarize problem_count = count(), by: {event.status}
```

```
// Recent problem list (Table tile)
fetch dt.davis.problems, from: now() - 24h
| fields timestamp, display_id, title, event.status
| sort timestamp desc
| limit 10
```

## 6. Sharing and Permissions

---

### Sharing Options

| Sharing Level              | Who Can Access                 |
|----------------------------|--------------------------------|
| Private                    | Only you                       |
| Shared with specific users | Named users                    |
| Shared with groups         | User groups                    |
| Public                     | Anyone with environment access |

### Segment Filtering

Dashboards can be filtered using segments:

- Set a default segment in dashboard settings
- Viewers can switch segments if they have access
- Data respects viewer's permissions

### Dashboard Presets

Create presets for common views:

1. Set segment filter
2. Set time range
3. Save as preset with descriptive name

## Exporting Dashboards

| Format | Use Case                       |
|--------|--------------------------------|
| PDF    | Reporting, documentation       |
| JSON   | Backup, migration              |
| Link   | Sharing (respects permissions) |

## 7. Next Steps

---

Congratulations! You've completed the onboarding series.

### What You've Learned

| Notebook | Key Topics                                                           |
|----------|----------------------------------------------------------------------|
| ONBRD-01 | Tenant access, navigation                                            |
| ONBRD-02 | IAM, SAML, Platform Tokens, parameterized policies                   |
| ONBRD-03 | ActiveGate deployment                                                |
| ONBRD-04 | Cloud & SaaS integrations                                            |
| ONBRD-05 | OneAgent deployment, primary tags at source                          |
| ONBRD-06 | Tags, segments, naming conventions, <code>dt.security_context</code> |
| ONBRD-07 | Data model, topology                                                 |
| ONBRD-08 | DQL fundamentals                                                     |
| ONBRD-09 | Workflows alerting                                                   |
| ONBRD-10 | Dashboard creation                                                   |

## Where to Go Next — by Topic

| Goal                                     | Topic Series                                                               |
|------------------------------------------|----------------------------------------------------------------------------|
| Dashboard strategy + executive reporting | DASH series (8 notebooks)                                                  |
| Davis AI / anomaly detection / RCA       | AIOPS series (8 notebooks)                                                 |
| Workflow automation + AI tasks           | WFLOW series (10 notebooks)                                                |
| Deepen DQL — spans / logs / OpenPipeline | SPANS, OPLOGS, OPMIG, OPIPE                                                |
| Synthetic monitoring                     | SYNTH series                                                               |
| Web RUM / Mobile RUM                     | WEBRUM, MOBL                                                               |
| Business events & funnel analysis        | BIZEV series                                                               |
| Database monitoring                      | DBMON series                                                               |
| Cloud integration deep dives             | CLOUD series                                                               |
| Kubernetes monitoring                    | K8S series                                                                 |
| OpenTelemetry integration                | OTEL series                                                                |
| IAM administration depth                 | IAM series                                                                 |
| Configuration automation / GitOps        | AUTOM series                                                               |
| Access control modernization             | MZ2POL series (Management Zones → Policies)                                |
| Migration playbooks                      | NRLC + NR2DT (New Relic), SL2DT (Sumo), S2D (Splunk), M2S (Managed → SaaS) |
| Platform maturity / adoption             | ADOPT series                                                               |
| Official training                        | Dynatrace University                                                       |
| Community                                | <a href="https://community.dynatrace.com">community.dynatrace.com</a>      |

## Dashboard Checklist

- ☐ First dashboard created (using **Dashboards (new)**, not legacy Classic Dashboards)
  - ☐ Key metrics visualized
  - ☐ Problem overview included
  - ☐ SLO / SLI tile evaluated for service-health views
  - ☐ Dashboard shared with team
  - ☐ Segment filter set
- 

## Summary

---

In this notebook, you learned:

- Difference between Dashboards (new) and Classic Dashboards
  - Difference between dashboards and notebooks
  - How to create and configure dashboards
  - Common tile types and when to use them
  - Dashboard patterns for different use cases
  - DQL queries that work well in dashboards
  - That `dt.davis.problems` queries use `event.status` ( `ACTIVE` / `CLOSED` ), not `status` ( `OPEN` )
  - How to share dashboards with your team
- 

## References

---

- [Dashboards](#)
- [Dashboards \(new\)](#)
- [Notebooks](#)
- [DQL Examples](#)
- [Dashboard Sharing](#)
- [Dynatrace Community](#)

---

*This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.*