

ONBRD-99: Best Practice Summary

Series: ONBRD — Dynatrace Onboarding | **Reference:** 99 — Best Practice Summary | **Created:** March 2026 | **Last Updated:** 09/18/2026

Overview

Reference card for an architect or tenant lead onboarding a new 2026 Dynatrace tenant: recommended defaults, the key decision matrix, validation queries, anti-patterns, and the cross-series next-steps map.

For the **sequence + dependencies + per-step decisions**, see **ONBRD-00: Architect's Sequence & Dependency Runbook**.

Table of Contents

- 1. [Recommended Defaults \(2026\)](#)
- 2. [Decision Matrix](#)
- 3. [Validation Queries](#)
- 4. [Anti-Patterns to Avoid](#)
- 5. [Where to Go Deeper — Topic Series Map](#)

Prerequisites

Requirement	Details
Audience	Architect or tenant lead onboarding a new 2026 Dynatrace tenant

Requirement	Details
Used alongside	ONBRD-00: Architect's Sequence & Dependency Runbook for sequence + dependencies

1. Recommended Defaults (2026)

A new 2026 tenant should default to these choices unless there is a specific reason not to.

Area	Default	Why
API token	Platform Token (<code>dt0s16</code>) with <code>Authorization: Bearer</code>	Sprint-1.337 default; aligns with Gen3 IAM model. Classic API Tokens (<code>dt0c01</code> , <code>Authorization: Api-Token</code>) only for legacy paths.
Configuration	Settings v2 / Configuration as Code (Terraform <code>dynatrace_settings</code> , Monaco v2)	Sprint-1.337 announced Configuration API endpoints have Settings v2 equivalents. Plan automation around Settings v2.
Extensions	Extensions 2.0 (managed via the Dynatrace API Application → Extensions surface)	Current extensions framework; EF1.0 end of support 2025-03-31.
Tagging at source	Primary fields/tags at OneAgent install via <code>oneagentctl --set-host-tag="primary_tags.<key>=<value>"</code> and <code>--set-host-tag="dt.security_context=<value>"</code> (single tags-hub form, June 2026; prefix written explicitly)	OneAgent attribute enrichment (OneAgent 1.333+) emits these on every signal at ingest.
Boundary field	<code>dt.security_context</code> for data + IAM scoping	Gen3 standard; segments + this field replace legacy Management Zones.

Area	Default	Why
IAM policies	Parameterized policies bound to groups via binding parameters	Avoids N-copies-of-similar-policy maintenance burden.
K8s deployment	Dynatrace Operator + Cloud Native FullStack	Classic FullStack deprecated for new deployments.
K8s CRD baseline	DynaKube v1beta5 baseline; v1beta6 canary	v1beta5 is field-tested baseline; v1beta6 for newer features (DB extensions preview) once canary-validated.
Alerting	Workflows + Davis Anomaly Detectors	Alerting Profiles + metric-events being phased out.
Logs	OpenPipeline	Future extension versions install ingest assets requiring updated configuration API (Nov 2025+).
Windows OneAgent	Npcap for network insight (OneAgent 1.337+)	Required for Network Agent metrics; WinPcap is no longer supported.

2. Decision Matrix

Decisions an architect makes during onboarding, with default and when to deviate.

Decision	Default	When to deviate
ActiveGate yes/no	Yes if >500 hosts, hybrid/on-prem, or cloud-API polling at scale	Pure SaaS-only with no on-prem footprint can defer until cloud-API polling demands it
ActiveGate sizing	10–20 GB; 2–3 AGs per zone	Larger only when load testing shows it is needed

Decision	Default	When to deviate
Per-cloud integration	Clouds app for AWS (GA) and Azure (SaaS 1.337+, no AG for metric polling); GCP Clouds-app connection in Preview — verify it has reached your tenant; the classic AG-based GCP integration remains the working path until then	If existing CloudWatch / Azure Monitor pipelines exist, evaluate migration effort vs leaving in place
OneAgent rollout sequence	Pilot 5–10 hosts → expand by host group → full	Skip pilot if you have onboarded the same workload class on another tenant
Tag taxonomy	env / team / app / dt.security_context / dt.cost.costcenter	Add compliance dimensions (*-pci , *-pii) only if a hard audit boundary exists
Host group naming	<env>-<app> or <env>-<workload-type> per FAQ-01	Customer-specific patterns acceptable; avoid all-hosts / default / hardware-trait names
Bucket strategy	dt.security_context for general data access; buckets only for compliance / retention / hard cost / hostile multi-tenancy	See FAQ-02 for the buckets-vs- dt.security_context decision
OTel coexistence	OneAgent for instrumentation; OTel for app-code spans / business events / serverless	Per FAQ-03 OneAgent vs OpenTelemetry decision framework
Workflow trigger	Davis problems (default); custom DQL-event triggers (advanced)	Custom DQL triggers when business-event-driven alerting is required
Dashboard tier	Operations + Executive (mandatory); Engineering (recommended)	Skip Executive if there is no exec stakeholder for this tenant
Notification channel for new tenant	Slack / email / PagerDuty silent channel for first 1–2 weeks	Switch to live channels only after dual-alert window validates volume parity

3. Validation Queries

A small DQL set that verifies each phase complete. Save these as a notebook in the new tenant — they are the architect's smoke-test for the deployment.

Phase 1 — Foundation complete

Hosts reporting (Gate G1):

```
// Hosts reporting (Gate G1 – Foundation -&gt; Organize)
smartscapeNodes "HOST"
| summarize host_count = count()
```

```
// Services discovered
smartscapeNodes "SERVICE"
| summarize service_count = count()
```

Phase 2 — Organize complete

Primary fields propagating; `dt.security_context` populated (Gate G2):

```
// dt.security_context populated on logs (Gate G2 – Organize -&gt; Operate)
fetch logs, from:-1h
| summarize { c = count() }, by:{dt.security_context}
| sort c desc
```

```
// Primary tags propagating
fetch logs, from:-1h
| filter isNotNull(primary_tags.environment)
| summarize { c = count() }, by:{primary_tags.environment}
```

Phase 3 — Operate complete

Workflows running; Davis surfacing problems with the expected scopes:

```
// Davis problems detected (last 7 days)
fetch dt.davis.problems, from:-7d
| filter event.status == "ACTIVE" or event.status == "CLOSED"
| summarize { c = count() }, by:{event.status}
```

4. Anti-Patterns to Avoid

Anti-pattern	Why it bites later
Single host group (<code>all-hosts</code> / <code>default</code>)	Per-host-group thresholds, alerting profiles, and IAM policies become impossible to scope
Setting tags retroactively (after install)	Primary-tag-at-install propagation gets skipped; signals before the retroactive setting do not carry the tag
Many copies of similar IAM policies (one per group with hardcoded scope)	Maintenance grows with org churn; use parameterized policies bound via binding parameters
Legacy Management Zones for new tenant	MZ-on-calculated-metrics is on the May-2026 deprecation list; segments + <code>dt.security_context</code> are the modern path
Auto-tagging rules as the primary tagging strategy	Computed at view time, not at ingest; doesn't propagate to logs / business events uniformly
<code>fetch dt.entity.*</code> queries in new content	Modern path is <code>smartscapeNodes "<TYPE>"</code> ; <code>dt.entity.*</code> works but is deprecated
Buckets used as the primary access-control mechanism	Buckets are for compliance / retention / hard cost / hostile multi-tenancy; default for general data access is <code>dt.security_context</code>
Classic FullStack on Kubernetes	Deprecated for new deployments; use Cloud Native FullStack
Configuration API for new automation	Settings v2 equivalents exist as of sprint-1.337; plan new pipelines on Settings v2 / Terraform / Monaco v2
Naming host groups by hardware traits (<code>4cpu-hosts</code> , <code>region-only</code>)	Encodes properties that change; rename burden as workloads scale

5. Where to Go Deeper — Topic Series Map

ONBRD covers the foundation. These topic series cover each domain in depth:

Domain	Topic Series	Notebooks
IAM administration	IAM	15
Data organization (buckets, segments, security context)	ORGNZ	11
Frequently asked questions	FAQ	growing collection
OpenPipeline log processing	OPLOGS	9
Classic Logs → OpenPipeline migration	OPMIG	10
OpenPipeline beyond logs (spans, metrics, events, bizevents)	OPIPE	7
Distributed tracing & spans	SPANS	9
Kubernetes monitoring & DynaKube	K8S	15
Cloud integration deep dives (AWS, Azure, GCP)	CLOUD	9
OpenTelemetry integration	OTEL	9
Workflows & alert notifications	WFLOW	12
Dynatrace Intelligence (Causal/Predictive/Generative AI, Davis)	AIOPS	8
Dashboard strategy & executive reporting	DASH	8
Configuration automation & GitOps	AUTOM	14
Management Zone → Policy migration	MZ2POL	11
Web Real User Monitoring	WEBRUM	10
Native mobile monitoring	MOBL	13
Synthetic monitoring	SYNTH	7
Business events & funnel analytics	BIZEV	8
Database monitoring	DBMON	7
Platform maturity & adoption roadmap	ADOPT	7
Managed → SaaS migration	M2S	11

Domain	Topic Series	Notebooks
SaaS → SaaS migration	S2S	11
New Relic → Dynatrace (procedural runbook)	NR2DT	11
New Relic → Dynatrace (component deep dives)	NRLC	9
Sumo Logic → Dynatrace	SL2DT	11
Splunk → Dynatrace	S2D	10

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against the current [Dynatrace documentation](#).