OPLOGS-02: Migration to OpenPipeline

> **Series:** OPLOGS — OpenPipeline Logs | **Notebook:** 2 of 8 | **Created:** December 2025 | **Last Updated:** 11/06/2026

Planning and Executing Your Log Migration

This notebook guides you through assessing your current log environment and planning migration to OpenPipeline v2.0.

Table of Contents

1. Migration Assessment
2. Migration Paths
3. Planning Your OpenPipeline Configuration
4. Migration Validation
5. Common Migration Patterns
6. Migration Checklist
7. 📝 Summary
8. ➡️ Next Steps
9. 📚 References

Prerequisites

- ✅ Access to a Dynatrace environment with log data

- ✅ Completed OPLOGS-01 fundamentals
- ✅ Knowledge of your current log sources and volumes

1. Migration Assessment

Before migrating, you need to understand your current log landscape.

Key Questions to Answer:

1. **What log sources do I have?** (OneAgent, API, OTLP)
2. **What is my log volume?** (records/hour, GB/day)
3. **Which logs need parsing?** (unstructured content)
4. **What sensitive data exists?** (PII, credentials)
5. **What can be dropped?** (debug, health checks)

```
// Assessment Query 1: Log sources inventory
fetch logs, from: now() - 24h
| summarize {
    log_count = count(),
    unique_hosts = countDistinct(dt.entity.host)
  }, by: {dt.openpipeline.source, log.source}
| sort log_count desc
| limit 20
```

```
// Assessment Query 2: Log volume by hour (capacity planning)
fetch logs, from: now() - 24h
| makeTimeseries {log_count = count()}, interval: 1h
```

```
// Assessment Query 3: Log level distribution
// status groups levels (SEVERE/CRITICAL/FATAL… → ERROR; DEBUG/TRACE/NOTICE → INFO),
// so DEBUG/TRACE are split out of info to keep each record in exactly one bucket
fetch logs, from: now() - 24h
| summarize {
    total = count(),
    errors = countIf(status == "ERROR"),
    warnings = countIf(status == "WARN"),
    info = countIf(status == "INFO" AND NOT in(loglevel, {"DEBUG", "TRACE"})),
    debug = countIf(in(loglevel, {"DEBUG", "TRACE"})),
    none = countIf(status == "NONE")
  }
| fieldsAdd debug_pct = round((toDouble(debug) / toDouble(total)) * 100, decimals: 1)
| fieldsAdd none_pct = round((toDouble(none) / toDouble(total)) * 100, decimals: 1)
```

```
// Assessment Query 4: Identify logs that need parsing (loglevel = NONE)
fetch logs, from: now() - 1h
| filter loglevel == "NONE" OR status == "NONE"
| fieldsAdd content_preview = substring(content, from: 0, to: 80)
| summarize {count = count()}, by: {content_preview, dt.openpipeline.source}
| sort count desc
| limit 15
```

```
// Assessment Query 5: Potential drop candidates (health checks, heartbeats)
fetch logs, from: now() - 24h
| summarize {
    total = count(),
    health_checks = countIf(matchesPhrase(content, "health") OR matchesPhrase(content, "healthcheck")),
    heartbeats = countIf(matchesPhrase(content, "heartbeat")),
    debug_logs = countIf(loglevel == "DEBUG"),
    // one countIf, so a record matching two conditions is counted once
    droppable = countIf(matchesPhrase(content, "health") OR matchesPhrase(content, "healthcheck") OR matchesPhrase(content, "heartbeat") OR loglevel == "DEBUG")
  }
| fieldsAdd savings_pct = round((toDouble(droppable) / toDouble(total)) * 100, decimals: 1)
```

2. Migration Paths

The one thing every path has in common

Log records keep arriving through the same ingest sources (OneAgent, `/api/v2/logs/ingest`, OTLP) whichever path you take. What changes is **where they are processed**. Until a dynamic route sends a record to one of your pipelines, it falls through to the default route — and on a tenant where the classic pipeline exists, that means the classic processing rules:

> *"Data that doesn't match is processed by the default route and continues to be processed by the classic pipeline until you turn off the rules."*

So the migration is: build pipelines, add routes for each data set, compare the output with the classic rules, then turn the classic rules off. Records the classic pipeline processed carry the pipeline id `logs:default` — the query below measures how much is still on it. `dt.openpipeline.source` and `dt.openpipeline.pipelines` are populated for classic records too, so an `isNotNull()` check on them proves nothing.

Path A: OneAgent

1. Add dynamic routes for the OneAgent data sets (for example by `k8s.namespace.name` or `log.source`) to your pipelines
2. Move each data set's classic processing rules into its pipeline
3. Watch the `logs:default` share fall to zero for those sources

Path B: Log Ingest API

1. Keep using the `/api/v2/logs/ingest` endpoint — no client change
2. Route the API data sets (`dt.openpipeline.source == "/api/v2/logs/ingest"`) to your pipelines
3. Move their processing rules

Path C: OTLP

1. Point OTLP exporters at `/api/v2/otlp/v1/logs`

2. Route by `dt.openpipeline.source == "/api/v2/otlp/v1/logs"` (or resource attributes) to your pipelines
3. Configure OTLP-specific processing there

OPMIG-09 covers the full classic-pipeline cut-over, including the business-events scope.

> **Sources:** Upgrade from classic pipeline to OpenPipeline (DT docs) — the quote above, and the `in(dt.openpipeline.pipelines, "logs:default")` filter for records the classic pipeline processed.

```
// How much log volume is still processed by the classic pipeline, by ingest source.
// Classic records carry the pipeline id "logs:default"; dt.openpipeline.* is populated
// for them too, so isNotNull(dt.openpipeline.pipelines) cannot tell the paths apart.
fetch logs, from: now() - 1h
| summarize {total = count(), classic = countIf(in(dt.openpipeline.pipelines, "logs:default"))}, by: {dt.openpipeline.source}
| fieldsAdd classic_pct = round(100.0 * toDouble(classic) / toDouble(total), decimals: 1)
| sort total desc
```

3. Planning Your OpenPipeline Configuration

Pipeline Strategy

Use Case	Pipeline Configuration
Unmatched data	Falls to the default route — the classic pipeline (`logs:default`) where it exists; watch its share
Custom Parsing	Create pipeline with DPL parse rules
PII Masking	Masking processors first in the Processing stage (there is no separate masking stage)
Cost Reduction	Add drop rules for noise

Use Case	Pipeline Configuration
Custom Routing	Route to specific buckets

Bucket Strategy

Log Type	Bucket	Retention
Production Errors	`prod_errors`	90 days
Application Logs	`default_logs`	35 days
Debug/Verbose	Drop or 7 days	Minimal
Audit/Compliance	`audit_logs`	365+ days

```
// Analyze current bucket usage
fetch logs, from: now() - 24h
| summarize {
    log_count = count(),
    error_count = countIf(loglevel == "ERROR" OR status == "ERROR")
  }, by: {dt.system.bucket}
| fieldsAdd error_rate = round((toDouble(error_count) / toDouble(log_count)) *
100, decimals: 2)
| sort log_count desc
```

4. Migration Validation

After configuring OpenPipeline, validate that:

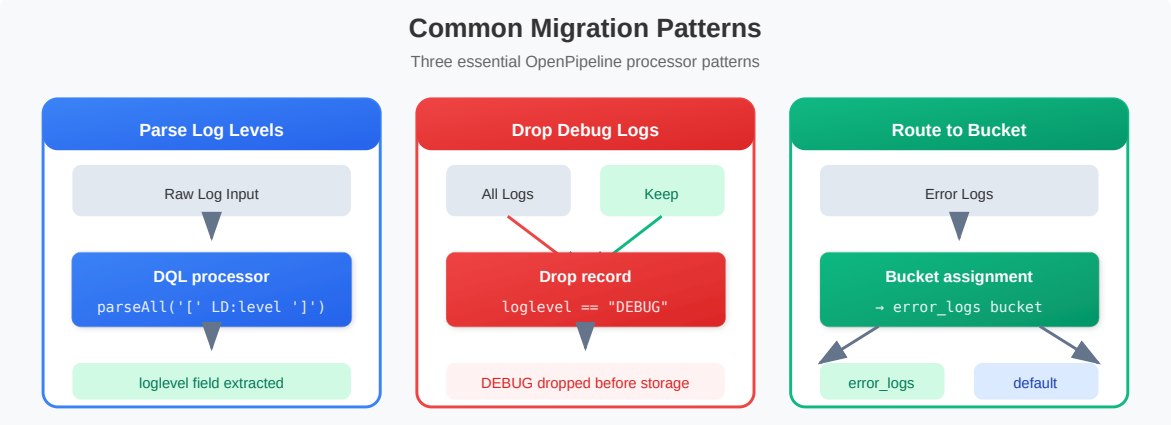
1. ✅ All expected log sources are present
2. ✅ Log counts are consistent
3. ✅ Parsing rules extract expected fields
4. ✅ Masking rules are applied correctly
5. ✅ Routing sends logs to correct buckets

```
// Validation Query 1: Confirm all sources are flowing, and none is left on
the classic pipeline
fetch logs, from: now() - 1h
| summarize {
    total = count(),
    classic = countIf(in(dt.openpipeline.pipelines, "logs:default")),
    has_bucket = countIf(isNotNull(dt.system.bucket))
  }
| fieldsAdd openpipeline_pct = round((toDouble(total - classic) /
toDouble(total)) * 100, decimals: 1)
```

```
// Validation Query 2: Check log volume consistency (compare hours)
fetch logs, from: now() - 4h
| makeTimeseries {log_count = count()}, by: {dt.openpipeline.source},
interval: 1h
```

```
// Validation Query 3: Verify entity context is preserved
fetch logs, from: now() - 1h
| summarize {
    total = count(),
    with_host = countIf(isNotNull(dt.entity.host)),
    with_process = countIf(isNotNull(dt.entity.process_group)),
    with_k8s_cluster = countIf(isNotNull(dt.entity.kubernetes_cluster)),
    with_k8s_namespace = countIf(isNotNull(k8s.namespace.name))
  }
```

5. Common Migration Patterns



Pattern 1: Parse Log Levels from Content

When logs have `loglevel = NONE`, configure a DQL processor to extract the level from content patterns like `[INFO]` or `[ERROR]`.

Processor Configuration:

- **Type:** DQL
- **Matcher:** `loglevel == "NONE"`
- **Statement:**

```
// parseAll collects every bracketed token: "[2026-10-12 13:20:22][worker:4]
[INFO]: ..." as well as "2026-10-02 12:00:01 [ERROR] ..."
fieldsAdd parsed_levels = parseAll(content, "'[' LD:l ']'")
| fieldsAdd parsed_level =
arrayRemoveNulls(iCollectArray(if(in(upper(parsed_levels[]), {"EMERGENCY",
"ALERT", "CRITICAL", "SEVERE", "ERROR", "FATAL", "WARN", "NOTICE", "INFO",
"DEBUG", "TRACE"}), upper(parsed_levels[])))[0]
| fieldsAdd loglevel = coalesce(parsed_level, loglevel)
| fieldsRemove parsed_levels, parsed_level
```

> `parseAll` collects every bracketed token; the first one that is a supported level wins, so a leading timestamp or thread name (`[2026-…][worker:4][INFO]`) does not hide the level. Anything else leaves the record unchanged. Do not use a single `parse content, "DATA? '[' LD:level ']'"` for this: `DATA?` consumes ahead and the capture lands on a later bracket (`[main] [ERROR] x` → `x`), so on the validation tenant it recovered **0 of 2,171** bracketed NONE-level lines in 30 minutes (10/12/2026). Measure the parse rate at query time (the next cell) before deploying: ingest-time parsing is forward-only.
>
> `parseAll` , `arrayRemoveNulls` , `coalesce` , `in` and `upper` are listed among the OpenPipeline DQL processor functions, and `iCollectArray` among its iterative operators. The statement above was executed at query time (10/12/2026); verify it saves in your pipeline before relying on it.
>
> **Sources:** DQL functions in OpenPipeline (DT docs) — *"These functions apply exclusively to the DQL processor."*, DQL operators in OpenPipeline DQL processor (DT docs) — *"Collects the results of an iterative expression into an array."*

Pattern 2: Drop Debug Logs

Reduce storage costs by dropping DEBUG-level logs before storage.

Processor Configuration:
- **Type:** Drop record
- **Matcher:** `loglevel == "DEBUG"`

Pattern 3: Route Errors to Dedicated Bucket

Send error logs to a dedicated bucket with longer retention for compliance.

Bucket Routing Configuration:
- **Matcher:** `status == "ERROR"` — `status` groups SEVERE, CRITICAL, FATAL… with ERROR; `loglevel == "ERROR"` would miss them
- **Bucket:** `error_logs`
- **Retention:** 90 days

```
// Simulate parsing log levels from content
fetch logs, from: now() - 1h
| filter loglevel == "NONE" OR status == "NONE"
| fieldsAdd parsed_levels = parseAll(content, "'[' LD:l ']'")
| fieldsAdd parsed_level =
arrayRemoveNulls(iCollectArray(if(in(upper(parsed_levels[]), {"EMERGENCY",
"ALERT", "CRITICAL", "SEVERE", "ERROR", "FATAL", "WARN", "NOTICE", "INFO",
"DEBUG", "TRACE"}), upper(parsed_levels[])))[0]
| filter isNotNull(parsed_level)
| fields timestamp, content, parsed_level
| limit 10
```

```
// Identify logs that would be dropped (DEBUG + health checks)
fetch logs, from: now() - 24h
| filter loglevel == "DEBUG"
     OR matchesPhrase(content, "healthcheck")
     OR matchesPhrase(content, "health check")
| summarize {would_drop = count()}, by: {dt.openpipeline.source}
| sort would_drop desc
```

6. Migration Checklist

Pre-Migration

- [ ] Document current log sources and volumes
- [ ] Identify logs requiring custom parsing
- [ ] List sensitive data patterns to mask
- [ ] Define drop rules for noise reduction
- [ ] Plan bucket strategy and retention

During Migration

- [ ] Configure OpenPipeline settings
- [ ] Create custom pipelines as needed
- [ ] Set up masking rules
- [ ] Configure bucket routing
- [ ] Test with sample data

Post-Migration

- [ ] Validate all log sources flowing
- [ ] Verify volume consistency
- [ ] Confirm parsing rules working
- [ ] Test masking is applied
- [ ] Update dashboards and alerts
- [ ] Update saved queries to use new fields

📝 Summary

In this notebook, you learned:

✅ **Assessment queries** to understand your current log environment
✅ **Migration paths** for OneAgent, API, and OTLP sources
✅ **Pipeline planning** strategies for processing and routing

✅ **Validation queries** to confirm successful migration
✅ **Common patterns** for parsing, dropping, and routing

➡️ Next Steps

Continue to **OPLOGS-03: OpenPipeline Processing** to configure parsing, enrichment, extraction and bucket routing.

📚 References

- OpenPipeline (DT docs)
- Log Monitoring API v2 - POST ingest logs (DT docs)
- Ingest OTLP logs (DT docs)

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.