

📊 Analytics & Dashboards

> **Series:** OPL0GS | **Notebook:** 7 of 8 | **Created:** December 2025

Aggregation, Time Series, and Visualization Queries

This notebook covers aggregation functions, time series analysis, statistical patterns, and dashboard-ready queries for log analytics.

Table of Contents

1. Aggregation Functions
2. Time Series Analysis
3. Statistical Analysis
4. Dashboard-Ready Queries
5. Trend Analysis
6. Log Pattern Analysis
7. Operational Dashboard Queries

Prerequisites

- Access to a Dynatrace environment with log data
- Completed OPL0GS-01 through OPL0GS-06
- Familiarity with dashboard concepts

1. Aggregation Functions

! [Aggregation Functions]

(


```

ob3I9Im1pZGRsZSI+UGVyY2VudGlsZTwvdGV4dD4KCiAgPHRleHQgeD0iNjIwIiB5PSIxMjUiIGZv
bnQtZmFtaWx5PSJtb25vc3BhY2UiIGZvbnQtc2l6ZT0iMTEiIGZvbnQtd2VpZ2h0PSJib2xkIiBma
WxsPSIjZDk3NzA2Ij5wZXJjZW50aWx1KGyNTApPC90ZXh0PgogIDx0ZXh0IHg9IjYyMCiGeT0iMT
QwIiBmb250LWZhbWlseT0iQXJpYwsiHnhbnMtc2VyaWYiIGZvbnQtc2l6ZT0iMTAiIGZpbGw9IiM
2NDc00GIiPk1lZGlhbiAocDUwKTtwvdGV4dD4KCiAgPHRleHQgeD0iNjIwIiB5PSIxNjUiIGZvbnQt
ZmFtaWx5PSJtb25vc3BhY2UiIGZvbnQtc2l6ZT0iMTEiIGZvbnQtd2VpZ2h0PSJib2xkIiBmaWxsP
SIjZDk3NzA2Ij5wZXJjZW50aWx1KGyNTUpPC90ZXh0PgogIDx0ZXh0IHg9IjYyMCiGeT0iMTgwII
Bmb250LWZhbWlseT0iQXJpYwsiHnhbnMtc2VyaWYiIGZvbnQtc2l6ZT0iMTAiIGZpbGw9IiM2NDc
00GIiPk1dGggcGVyY2VudGlsZTwvdGV4dD4KCiAgPHRleHQgeD0iNjIwIiB5PSiyMDUiIGZvbnQt
ZmFtaWx5PSJtb25vc3BhY2UiIGZvbnQtc2l6ZT0iMTEiIGZvbnQtd2VpZ2h0PSJib2xkIiBmaWxsP
SIjZDk3NzA2Ij5wZXJjZW50aWx1KGyNTkpPC90ZXh0PgogIDx0ZXh0IHg9IjYyMCiGeT0iMjIwIi
Bmb250LWZhbWlseT0iQXJpYwsiHnhbnMtc2VyaWYiIGZvbnQtc2l6ZT0iMTAiIGZpbGw9IiM2NDc
00GIiPk5dGggcGVyY2VudGlsZTwvdGV4dD4KCiAgPHRleHQgeD0iNjIwIiB5PSiyNDUiIGZvbnQt
ZmFtaWx5PSJtb25vc3BhY2UiIGZvbnQtc2l6ZT0iMTEiIGZvbnQtd2VpZ2h0PSJib2xkIiBmaWxsP
SIjZDk3NzA2Ij5zdgRkZXYoZmlbGQpPC90ZXh0PgogIDx0ZXh0IHg9IjYyMCiGeT0iMjYwIiBmb2
50LWZhbWlseT0iQXJpYwsiHnhbnMtc2VyaWYiIGZvbnQtc2l6ZT0iMTAiIGZpbGw9IiM2NDc00GI
iPln0YW5kYXjkIGRldmlhdGvbjwvdGV4dD4KCiAgPHJly3QgeD0iNjIwIiB5PSiyNzUiHdpZHRo
PSIxMzAiIGhlaWdodD0iMTgiIHJ4PSI0IiBmaWxsPSIjZmVmM2M3Ii8+CiAgPHRleHQgeD0iNjg1I
iB5PSiyODgiIGZvbnQtZmFtaWx5PSJBcmhbCwgc2Fucy1zZXJpZiIgZm9udC1zaXplPSIxMCiGZm
lsbD0iIzkyNDawZSIgdGV4dC1hbmnob3I9Im1pZGRsZSI+UmV0dXJuczogZG91YmxlPC90ZXh0Pgo
8L3N2Zz4K)

```

Common Aggregation Functions

Function	Description	Example
`count()`	Count records	`summarize {total = count()}`
`countIf(condition)`	Conditional count	`countIf(loglevel == "ERROR")`
`countDistinct(field)`	Unique values	`countDistinct(k8s.pod.name)`
`sum(field)`	Sum values	`sum(bytes)`
`avg(field)`	Average	`avg(response_time)`
`min(field)`	Minimum	`min(timestamp)`
`max(field)`	Maximum	`max(timestamp)`
`takeFirst(field)`	First value	`takeFirst(content)`
`takeLast(field)`	Last value	`takeLast(content)`
`collectDistinct(field)`	Array of unique values	`collectDistinct(loglevel)`

🚫 Important: Named Aliases Required

```

```dql
// ✅ CORRECT – Named aliases
| summarize {count = count(), errors = countIf(...)}, by: {...}
| sort count desc

// ❌ WRONG – Anonymous aggregation can't be used in sort

```

```

| summarize count(), by: {...}
| sort count() desc // ERROR!
```

```python
// Basic aggregations
fetch logs, from: now() - 1h
| summarize {
 total_logs = count(),
 unique_hosts = countDistinct(dt.entity.host),
 unique_namespaces = countDistinct(k8s.namespace.name),
 unique_pods = countDistinct(k8s.pod.name)
}
```

```python
// Aggregations with grouping
fetch logs, from: now() - 1h
| summarize {
 total = count(),
 errors = countIf(loglevel == "ERROR"),
 warnings = countIf(loglevel == "WARN"),
 info = countIf(loglevel == "INFO")
}, by: {k8s.namespace.name}
| sort total desc
| limit 15
```

```python
// Error rate calculation
fetch logs, from: now() - 1h
| summarize {
 total = count(),
 errors = countIf(loglevel == "ERROR")
}, by: {k8s.namespace.name}
| filter total > 100 // Minimum sample size
| fieldsAdd error_rate_pct = round((errors * 100.0) / total, decimals: 2)
| sort error_rate_pct desc
| limit 15
```

## 2. Time Series Analysis
```

The `makeTimeseries` command creates time-bucketed data for trend visualization.

```

```python
// Log volume over time (5-minute buckets)
```

```

fetch logs, from: now() - 6h
| makeTimeseries {log_count = count()}, interval: 5m
```

```python
// Error trend over time
fetch logs, from: now() - 6h
| makeTimeseries {
 total = count(),
 errors = countIf(loglevel == "ERROR"),
 warnings = countIf(loglevel == "WARN")
}, interval: 5m
```

```python
// Time series by dimension (namespace)
fetch logs, from: now() - 6h
| filter isNotNull(k8s.namespace.name)
| makeTimeseries {
 log_count = count()
}, by: {k8s.namespace.name}, interval: 10m
```

```python
// Error rate time series by host
fetch logs, from: now() - 6h
| filter isNotNull(dt.entity.host)
| makeTimeseries {
 total = count(),
 errors = countIf(loglevel == "ERROR")
}, by: {dt.entity.host}, interval: 15m
```

## 3. Statistical Analysis

```python
// Log volume statistics by source
fetch logs, from: now() - 24h
| summarize {
 total = count(),
 first_seen = min(timestamp),
 last_seen = max(timestamp)
}, by: {dt.openpipeline.source}
| fieldsAdd duration_hours = (last_seen - first_seen) / 3600000000000
| fieldsAdd logs_per_hour = round(total / duration_hours, decimals: 0)
| sort total desc
```

```

```

```python
// Percentile analysis (if numeric field available)
fetch logs, from: now() - 1h
| fieldsAdd content_length = stringLength(content)
| summarize {
 avg_length = avg(content_length),
 min_length = min(content_length),
 max_length = max(content_length),
 total_logs = count()
}, by: {k8s.namespace.name}
| sort avg_length desc
| limit 15
```

```python
// Hourly distribution analysis
fetch logs, from: now() - 24h
| fieldsAdd hour_bucket = bin(timestamp, 1h)
| summarize {log_count = count()}, by: {hour_bucket}
| sort hour_bucket asc
```

```python
// Daily distribution analysis
fetch logs, from: now() - 7d
| fieldsAdd day_bucket = bin(timestamp, 1d)
| summarize {
 log_count = count(),
 error_count = countIf(loglevel == "ERROR")
}, by: {day_bucket}
| sort day_bucket asc
```

```

4. Dashboard-Ready Queries

These queries are optimized for dashboard tiles and visualizations.

```

```python
// Single Value: Total Logs (Last Hour)
fetch logs, from: now() - 1h
| summarize {total_logs = count()}
```

```python
// Single Value: Error Count (Last Hour)
fetch logs, from: now() - 1h
| filter loglevel == "ERROR"
| summarize {error_count = count()}
```

```

```
```
```python
// Pie Chart: Log Level Distribution
fetch logs, from: now() - 1h
| summarize {count = count()}, by: {loglevel}
| sort count desc
```

```python
// Bar Chart: Top Error Sources
fetch logs, from: now() - 1h
| filter loglevel == "ERROR"
| summarize {error_count = count()}, by: {k8s.namespace.name}
| sort error_count desc
| limit 10
```

```python
// Line Chart: Log Trend (6 Hours)
fetch logs, from: now() - 6h
| makeTimeseries {
    errors = countIf(loglevel == "ERROR"),
    warnings = countIf(loglevel == "WARN"),
    info = countIf(loglevel == "INFO")
}, interval: 5m
```

```python
// Table: Top Error Messages
fetch logs, from: now() - 1h
| filter loglevel == "ERROR"
| fieldsAdd error_preview = substring(content, from: 0, to: 100)
| summarize {
    occurrences = count(),
    first_seen = min(timestamp),
    last_seen = max(timestamp)
}, by: {error_preview, k8s.namespace.name}
| sort occurrences desc
| limit 20
```

```python
// Heat Map: Errors by Hour and Namespace
fetch logs, from: now() - 24h
| filter loglevel == "ERROR"
| fieldsAdd hour_bucket = bin(timestamp, 1h)
| summarize {error_count = count()}, by: {hour_bucket, k8s.namespace.name}
```

```

| sort hour_bucket asc
```

5. Trend Analysis

```python
// Compare current hour to previous hour
fetch logs, from: now() - 2h
| fieldsAdd period = if(timestamp > now() - 1h, "current", else: "previous")
| summarize {count = count()}, by: {period, loglevel}
| sort period asc, count desc
```

```python
// New error patterns (appeared in last hour)
fetch logs, from: now() - 1h
| filter loglevel == "ERROR"
| fieldsAdd error_sig = substring(content, from: 0, to: 80)
| summarize {
    count = count(),
    first_seen = min(timestamp)
}, by: {error_sig, k8s.namespace.name}
| filter first_seen > now() - 30m // First seen in last 30 minutes
| sort first_seen desc
| limit 15
```

```python
// Volume anomaly detection (compare to baseline)
fetch logs, from: now() - 6h
| fieldsAdd time_bucket = bin(timestamp, 15m)
| summarize {log_count = count()}, by: {time_bucket, k8s.namespace.name}
| sort time_bucket desc
| limit 100
```

6. Log Pattern Analysis

```python
// Top log patterns (by content prefix)
fetch logs, from: now() - 1h
| fieldsAdd pattern = substring(content, from: 0, to: 60)
| summarize {count = count()}, by: {pattern}
| sort count desc
| limit 25
```

```python

```

```

// Unique log levels and statuses
fetch logs, from: now() - 1h
| summarize {
    loglevels = collectDistinct(loglevel),
    statuses = collectDistinct(status),
    sources = collectDistinct(dt.openpipeline.source)
}
```
```
```python
// Log diversity score (unique patterns per namespace)
fetch logs, from: now() - 1h
| fieldsAdd pattern = substring(content, from: 0, to: 50)
| summarize {
 total_logs = count(),
 unique_patterns = countDistinct(pattern)
}, by: {k8s.namespace.name}
| fieldsAdd diversity_ratio = round((unique_patterns * 100.0) / total_logs,
decimals: 2)
| sort diversity_ratio desc
| limit 15
```
```
7. Operational Dashboard Queries

```python
// Operations Summary (last 1h)
fetch logs, from: now() - 1h
| summarize {
    total_logs = count(),
    error_count = countIf(loglevel == "ERROR"),
    warn_count = countIf(loglevel == "WARN"),
    unique_hosts = countDistinct(dt.entity.host),
    unique_pods = countDistinct(k8s.pod.name)
}
| fieldsAdd error_rate = round((error_count * 100.0) / total_logs, decimals:
2)
```
```
```python
// Health scorecard by namespace
fetch logs, from: now() - 1h
| filter isNotNull(k8s.namespace.name)
| summarize {
 total = count(),
 errors = countIf(loglevel == "ERROR"),
 warnings = countIf(loglevel == "WARN")
}, by: {k8s.namespace.name}

```

```
| fieldsAdd error_rate = round((errors * 100.0) / total, decimals: 2)
| fieldsAdd health_status = if(error_rate > 10, "CRITICAL",
| else: if(error_rate > 5, "WARNING",
| else: "HEALTHY"))
| sort error_rate desc
````

```python
// Ingestion pipeline summary
fetch logs, from: now() - 1h
| summarize {
 log_count = count(),
 unique_pipelines = countDistinct(dt.openpipeline.pipelines)
}, by: {dt.openpipeline.source, dt.system.bucket}
| sort log_count desc
````

---
```

📈 Summary

In this notebook, you learned:

- ✓ **Aggregation functions** – count, countIf, countDistinct, sum, avg
- ✓ **Time series** – makeTimeseries with intervals
- ✓ **Statistical analysis** – hourly/daily patterns, percentiles
- ✓ **Dashboard queries** – Single values, charts, tables
- ✓ **Trend analysis** – Period comparison, anomaly detection
- ✓ **Pattern analysis** – Log clustering and diversity

➡️ Next Steps

Continue to **OPLOGS-08: Security & Data Protection** for data protection patterns.

📖 References

- [DQL Aggregation Functions]
(<https://docs.dynatrace.com/docs/platform/grail/dynatrace-query-language/functions/aggregation-functions>)
- [DQL makeTimeseries]
(<https://docs.dynatrace.com/docs/platform/grail/dynatrace-query-language/commands/makeTimeseries>)

- [Dynatrace Dashboards] (<https://docs.dynatrace.com/docs/observe-and-explore/dashboards-and-notebooks/dashboards-new>)