OPMIG-01: OpenPipeline Migration Guide: Part 1

> **Series:** OPMIG — OpenPipeline Migration | **Notebook:** 1 of 10 | **Created:** December 2025 | **Last Updated:** 10/02/2026

Introduction & Why Migrate from Classic to OpenPipeline v2.0

> 🆕 **New Addition (March 2026)**
>
> **Companion Series**
> - **OPLOGS** — Already migrated? Deep-dive into log processing with **OPLOGS-01: OpenPipeline Fundamentals**
> - **OPIPE** — Need spans, metrics, or event processing? See **OPIPE-01: OpenPipeline as a Multi-Scope Platform**

Table of Contents

1. What is OpenPipeline?
2. Classic vs OpenPipeline: Key Differences
3. Benefits of Migrating
4. OpenPipeline Configuration Scopes
5. API Migration: Classic → OpenPipeline v2.0
6. Real-World Migration Scenarios
7. Understanding Your Current State
8. Migration Readiness Assessment

9. Migration Readiness Checklist
10. OpenPipeline Limits & Constraints

Learning Objectives

By the end of this notebook, you will:

- ✅ Understand what OpenPipeline is and why it replaces Classic log ingestion
- ✅ Learn the key architectural differences between Classic and OpenPipeline v2.0
- ✅ **Understand API endpoint compatibility and migration requirements**
- ✅ **Know OpenPipeline limits and constraints**
- ✅ Identify the benefits of migrating (cost, security, performance)
- ✅ **Review real-world migration scenarios**
- ✅ Assess your readiness to migrate

Prerequisites

Requirement	Details
Dynatrace Environment	Dynatrace SaaS with Grail — OpenPipeline runs in the SaaS environment; Managed is not covered by this series
Classic Log Ingestion	Existing logs flowing via `/api/v2/logs/ingest`
API Access	`logs.read` and `logs.ingest` token scopes
Knowledge	Basic Dynatrace familiarity; no OpenPipeline experience required

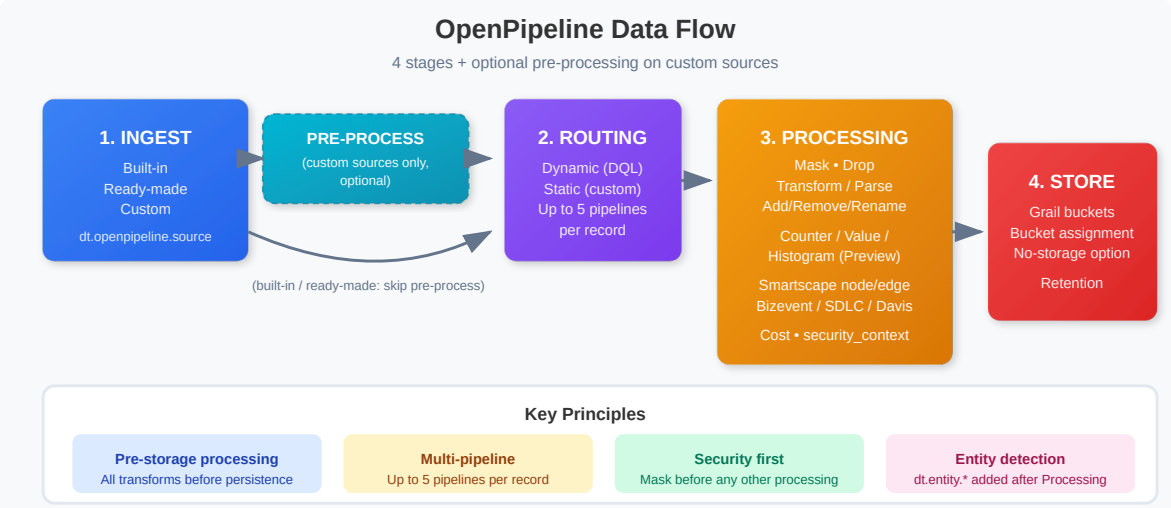
---

What is OpenPipeline?

OpenPipeline is Dynatrace's unified data handling solution that seamlessly ingests and processes data from different sources, at any scale, and in any format.

> 💡 **Key Insight:** OpenPipeline is NOT just for logs! It's a comprehensive data processing framework that handles **logs, spans, metrics, events, business events, security events, and more**.

OpenPipeline Architecture Overview



Core Capabilities

Capability	Description
Unified Ingestion	Single solution for logs, spans, metrics, events, and business events
Dynamic Routing	Route data to specific pipelines based on matching conditions
Real-time Processing	Transform, enrich, and mask data at ingestion time
Metric Extraction	Create metrics from any data source for long-term analytics
Event Generation	Generate custom events and business events from incoming data
Bucket Management	Control retention and cost with targeted storage routing

Capability	Description
Security & Compliance	Mask sensitive data before storage

Classic vs OpenPipeline: Key Differences

Understanding the fundamental differences helps you plan your migration effectively.

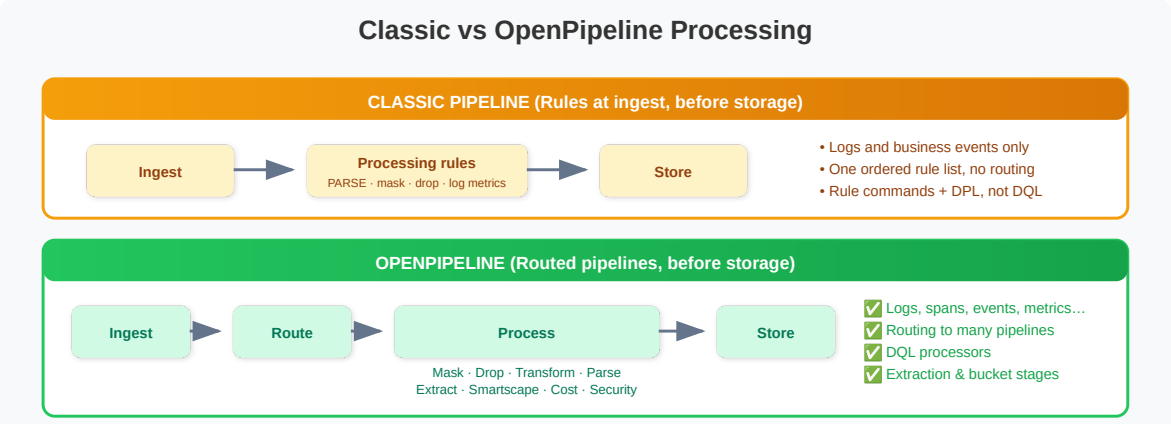
Feature Comparison

Feature	Classic pipeline	OpenPipeline
Data Types	Logs and business events	Logs, spans, metrics, events, business events, and more
Processing Location	At ingest, before storage	At ingest, before storage
Rule language	Processing-rule commands (`PARSE` , `FIELDS_ADD` , `FILTER_OUT` , …) with DPL	DQL processors with DPL, plus no-code processors
Routing	One ordered list of rules, each with a matcher	Dynamic routing to many pipelines, pipeline groups
Extraction	Log metrics	Metric, Davis event, business event, SDLC event and Smartscape stages
Data Masking	At ingest (mask, drop or rename attributes in a rule)	At ingest (Processing-stage processors)
Bucket Assignment	—	Bucket assignment stage, or No storage assignment
Ownership	Shared settings	Owner-based access per custom pipeline
Configuration	Settings → Log Monitoring → processing rules	Settings → Process and contextualize → OpenPipeline

Both pipelines process records **before** they are written to Grail. The move to OpenPipeline is not about *when* processing happens; it is about scope (every signal type, not just logs), language (DQL instead of rule commands), structure (routed pipelines with fixed stages instead of one rule list) and what you can extract.

> **Sources:** Log processing with classic pipeline (DT docs) — *"Log processing occurs as log data arrives in the Dynatrace SaaS environment and before it is written to disk (stored)."*; Customize incoming log data with log processing rules (DT docs) — rule examples for *Drop a log event*, *Mask attributes* and log metrics; Upgrade from classic pipeline to OpenPipeline (DT docs) — classic pipeline per configuration scope, *Logs or Business events*.

Processing Stage Comparison



> ⚠️ **Important:** With OpenPipeline, as with the classic pipeline, data processing happens **before** storage — dropping unwanted data and masking sensitive information happen before anything is written to Grail. What OpenPipeline adds is the ability to do it per pipeline, for every signal type, and to send records to buckets with different retention (or to no storage at all).

Benefits of Migrating

1. 💰 Cost Optimization

Drop unwanted data before storage:
- Filter out debug logs, health checks, and noise
- How much you save depends on how much of your volume is noise — the assessment queries in this notebook measure that for your own data before you commit to a number
- Route high-volume, low-value data to shorter retention buckets

2. 🔐 Security & Compliance

Mask sensitive data at ingestion:
- PII (Personal Identifiable Information) never touches storage
- Credit card numbers, SSNs, emails masked before persistence
- Meet GDPR, HIPAA, PCI-DSS requirements

3. 📊 Enhanced Analytics

Extract metrics with dimensions:
- Create custom metrics from any log pattern
- Build long-term trend dashboards
- Enable business KPI tracking from technical data

4. ⚡ Improved Query Performance

Parse once, query fast:
- Structured fields extracted at ingestion
- No runtime parsing overhead
- Faster dashboards and alerts

5. 🎯 Unified Data Processing

Single solution for all data types:
- Consistent processing for logs, spans, metrics, events
- Centralized configuration management
- Simplified operations and governance

6. 🔄 Real-time Enrichment

Add context at ingestion:
- Add environment tags (prod, staging, dev)
- Enrich with business context
- Standardize field names across sources

OpenPipeline Configuration Scopes

OpenPipeline supports multiple **configuration scopes** - each handling a different data type:

Configuration Scope	Data Type	Use Case
Logs	Log records	Application logs, system logs, audit logs
Spans	Distributed traces	Span processing, trace enrichment
Metrics	Time-series data	Metric ingestion and transformation
Events	Platform events	Generic events, detected events, SDLC events
Business Events	Business analytics	User journeys, transactions, conversions
Security Events	Security data	Vulnerability findings, compliance events
System Events	Infrastructure	System-level events and alerts

Ingest Sources per Scope

Each scope supports different ingest sources:

Logs:
- OneAgent log ingestion
- Generic log API (`/api/v2/logs/ingest`)
- OTLP logs
- Fluent integrations

Spans:
- OneAgent distributed tracing

- OTLP spans
- OpenTelemetry collectors

Metrics:
- OneAgent metrics
- OTLP metrics
- Metric ingestion API

> 💡 **Tip:** Access OpenPipeline configuration at: **Settings → Process and contextualize → OpenPipeline**

API Migration: Classic → OpenPipeline v2.0

API Endpoint Compatibility

Good news: **The API endpoint remains the same!**

Endpoint	Classic	OpenPipeline	Notes
Logs	`/api/v2/logs/ingest`	`/api/v2/logs/ingest`	✅ No change required
OTLP Logs	Not available	`/api/v2/otlp/v1/logs`	✅ New endpoint
Spans	`/api/v2/otlp/v1/traces`	`/api/v2/otlp/v1/traces`	✅ No change required

> 💡 **Migration Tip:** You don't need to update your API calls! OpenPipeline automatically receives data sent to `/api/v2/logs/ingest`. The difference is in **how** the data is processed after ingestion.

Ingestion Methods Comparison

Method	Classic	OpenPipeline	`dt.openpipeline.source` Value
OneAgent	✅ Supported	✅ Supported	`oneagent`
Generic Log API	✅ Supported	✅ Supported	`/api/v2/logs/ingest`
OTLP Protocol	⚠️ Limited	✅ Full support	`/api/v2/otlp/v1/logs`
Fluent Bit	✅ Via API	✅ Via API	the API path it sends to (usually `/api/v2/logs/ingest`)
Fluentd	✅ Via API	✅ Via API	the API path it sends to (usually `/api/v2/logs/ingest`)
Logstash	✅ Via API	✅ Via API	the API path it sends to (usually `/api/v2/logs/ingest`)
Vector	✅ Via API	✅ Via API	the API path it sends to (usually `/api/v2/logs/ingest`)

For a built-in API source, `dt.openpipeline.source` holds the endpoint **path**, not a short name — a validation tenant showed `oneagent`, `/api/v2/otlp/v1/logs` and `/api/v2/logs/ingest` over 24 hours (09/28/2026). Run the *Analyze log volume by OpenPipeline source* query below to see your own values before writing a routing condition on this field.

> **Sources:** Data flow (DT docs) — ingest sources *"are defined by a name and a path (dt.openpipeline.source)"*.

Required Token Permissions

Scope	Classic	OpenPipeline	Notes
`logs.ingest`	✅ Required	✅ Required	Same permission
`metrics.ingest`	N/A	⚠️ Optional	Only if using metric extraction

Scope	Classic	OpenPipeline	Notes
`events.ingest`	N/A	⚠️ Optional	Only if using event extraction

Migration Strategy for API Clients

Scenario	Action Required
Using OneAgent	✅ No code changes
Using `/api/v2/logs/ingest`	✅ No code changes
Using custom log shippers	⚠️ Optional
Want to leverage new features	⚠️ Configure pipelines

Real-World Migration Scenarios

> These are **illustrative** scenarios, not customer case studies. The savings percentages assume the volume mix described in each challenge; measure your own mix with the assessment queries below before you plan to a number.

Scenario 1: E-Commerce Platform

Challenge:
- 50M logs/day (70% debug logs)
- Payment logs contain credit card numbers
- Need to track order metrics from logs

OpenPipeline Solution:
1. **Drop debug logs** → Reduce volume by 70%
2. **Mask credit cards** → PCI-DSS compliance
3. **Extract payment metrics** → `order.amount`, `order.count`
4. **Route to tiered buckets** → High-value logs = 90 days, others = 7 days

Result: 70% cost savings, PCI compliance, new business metrics

Scenario 2: Healthcare SaaS

Challenge:
- HIPAA compliance required
- Logs contain patient IDs, MRNs, SSNs
- Need audit trail for 7 years

OpenPipeline Solution:
1. **Mask all PHI fields** → Patient IDs, SSNs, MRNs
2. **Route audit logs** → Dedicated bucket with 2555-day retention
3. **Extract security events** → Failed auth attempts, data access
4. **Drop health checks** → Reduce noise

Result: HIPAA compliance, 7-year audit retention, 40% cost savings

Scenario 3: FinTech Startup (ELK Migration)

Challenge:
- Migrating from ELK stack to Dynatrace
- Custom log formats (not JSON)
- Need APM + logs correlation

OpenPipeline Solution:
1. **Parse custom log formats** → DPL patterns for structured extraction
2. **Extract request IDs** → Correlate with traces
3. **Create SLI metrics** → Error rate, latency percentiles
4. **Unified observability** → Logs + APM in one platform

Result: ELK replacement, APM correlation, unified observability

Scenario 4: Global Retailer (Multi-Region)

Challenge:
- Multi-region deployment (US, EU, APAC)
- GDPR compliance for EU customers
- 100+ microservices

OpenPipeline Solution:
1. **Environment-based routing** → prod/staging/dev to different buckets
2. **Mask PII for EU** → Email, IP addresses for EU region logs
3. **Service-level pipelines** → Dedicated processing per critical service
4. **Metric extraction** → Service-level SLIs

Result: Multi-region compliance, per-service observability, 50% cost reduction

---

Understanding Your Current State

Before migrating, you need to understand your current log ingestion landscape. The following queries help you assess what you're working with.

```
// Identify your current log sources and volume
// This shows which sources are sending the most logs
fetch logs, from: now() - 7d
| summarize {log_count = count()}, by: {log.source}
| sort log_count desc
| limit 25
```

```
// Check which logs are already processed by OpenPipeline vs Classic
// This helps identify migration progress
fetch logs, from: now() - 24h
| fieldsAdd pipeline_type = if(isNotNull(dt.openpipeline.pipelines),
"OpenPipeline",
                            else: if(isNotNull(dt.openpipeline.source),
"OpenPipeline",
                            else: "Classic"))
| summarize {log_count = count()}, by: {pipeline_type}
| sort log_count desc
```

```
// Analyze log volume by OpenPipeline source
// Identify the ingestion methods being used
fetch logs, from: now() - 24h
| summarize {log_count = count()}, by: {dt.openpipeline.source}
| sort log_count desc
```

```
// Check current bucket distribution
// Understand where your logs are being stored
fetch logs, from: now() - 24h
| summarize {log_count = count()}, by: {dt.system.bucket}
| sort log_count desc
```

```
// Identify which pipelines are processing your logs
// Shows custom pipelines already configured
fetch logs, from: now() - 24h
| filter isNotNull(dt.openpipeline.pipelines)
| summarize {log_count = count()}, by: {dt.openpipeline.pipelines}
| sort log_count desc
```

Migration Readiness Assessment

Use these queries to assess your migration readiness and identify areas that need attention.

```
// Check parsing coverage - how many logs have structured data?
// Low coverage indicates need for parsing pipelines
fetch logs, from: now() - 24h
| fieldsAdd has_structured_data = isNotNull(loglevel) OR isNotNull(status)
| summarize {
    total = count(),
    structured = countIf(has_structured_data),
    unstructured = countIf(NOT has_structured_data)
  }
| fieldsAdd coverage_pct = round((toDouble(structured) / toDouble(total)) *
100, decimals: 1)
```

```
// Identify logs that could be dropped to save costs
// Debug logs and health checks are common candidates
fetch logs, from: now() - 24h
| summarize {
    total = count(),
    debug_logs = countIf(loglevel == "DEBUG" OR status == "DEBUG"),
    info_logs = countIf(loglevel == "INFO" OR status == "INFO"),
    health_checks = countIf(contains(toString(content), "health") OR
contains(toString(content), "heartbeat")),
    metrics_endpoints = countIf(contains(toString(content), "/metrics") OR
contains(toString(content), "/prometheus"))
  }
| fieldsAdd droppable = debug_logs + health_checks + metrics_endpoints
| fieldsAdd potential_savings_pct = round((toDouble(droppable) /
toDouble(total)) * 100, decimals: 1)
```

```
// Find logs with potential PII that needs masking
// Look for common patterns that might contain sensitive data
fetch logs, from: now() - 1h
| filter contains(toString(content), "email")
    OR contains(toString(content), "password")
    OR contains(toString(content), "ssn")
    OR contains(toString(content), "credit")
    OR contains(toString(content), "@")
| summarize {potentially_sensitive = count()}, by: {log.source}
| sort potentially_sensitive desc
| limit 20
```

```
// Analyze log level distribution
// Helps identify noise reduction opportunities
fetch logs, from: now() - 24h
| summarize {log_count = count()}, by: {loglevel}
| sort log_count desc
```

```
// Check log volume trends over time
// Understand your ingestion patterns
fetch logs, from: now() - 7d
| makeTimeseries {log_count = count()}, interval: 1h
```

Migration Readiness Checklist

Based on your assessment queries, complete this checklist:

Discovery Phase

- [ ] Identified all log sources (`log.source` values)
- [ ] Documented current log volume by source
- [ ] Identified which logs are already on OpenPipeline
- [ ] Analyzed current bucket usage

Planning Phase

- [ ] Identified logs that can be dropped (debug, health checks)

- [ ] Identified logs requiring parsing (unstructured content)
- [ ] Identified logs with sensitive data requiring masking
- [ ] Planned bucket strategy (retention periods, cost tiers)

Configuration Phase

- [ ] Created custom pipelines for each use case
- [ ] Configured dynamic routing rules
- [ ] Set up parsing processors (DQL/DPL)
- [ ] Configured masking for sensitive data
- [ ] Set up metric extraction where needed
- [ ] Configured bucket routing

Validation Phase

- [ ] Tested pipelines with sample data
- [ ] Verified parsing produces expected fields
- [ ] Confirmed masking works correctly
- [ ] Validated metrics are being extracted
- [ ] Checked data appears in correct buckets

OpenPipeline Limits & Constraints

Before migrating, understand these key limits:

Data Size Limits

Limit	Value	What Happens if Exceeded
Max record size (after processing)	16 MB	Record is **dropped**
Working memory per record	Limited — no figure published †	Record dropped once processing memory is exhausted

Limit	Value	What Happens if Exceeded
Log attribute size	32 KB (4,096 characters per attribute in an event template)	Attribute is **truncated**
Max field name length	255 characters †	Field creation fails
Max string field length	32 KB †	Content truncated

Processing Limits

Limit	Value	Impact
Max extractions per record	5 pipelines	Beyond 5, data extraction stops; the record is still processed and stored
Max processors per pipeline	1,000 processors (100 in a base pipeline)	Cannot add more processors
Max DQL processor script length	8,192 characters	Split into multiple processors
Max DQL commands per processor	10 commands †	Split into multiple processors
Max parse operations per processor	100 patterns †	Create multiple parse processors
Processing timeout	30 seconds †	Record dropped if processing exceeds

Timestamp Constraints

Data Type	Earliest accepted timestamp	Timestamp more than 10 min in the future
Logs	Ingest time minus 24 hours — **older records are dropped** before processing. From **SaaS 1.348** (pre-release; staged tenant rollout planned from 09/22/2026) logs are accepted up to 72 hours in the past — verify the new window has reached your tenant before relying on it	**Adjusted** to ingest time + 10 min

Data Type	Earliest accepted timestamp	Timestamp more than 10 min in the future
Spans	60 minutes past (end time) †	Not adjusted (the adjustment doesn't apply to spans)
Events	Ingest time minus 24 hours — older records are dropped	**Adjusted** to ingest time + 10 min

> ⚠️ **Important:** Always send data with recent timestamps. Historical data imports require special considerations.

Routing & Pipeline Limits

Limit	Value	Notes
Max custom pipelines	100	Per configuration scope (includes built-in and custom)
Max dynamic routes	100 routes	Per configuration scope (per `/reference/limits`)
Max conditions per route	10 conditions †	Use AND/OR to combine
Processor matching condition length	4,096 characters	Settings API configurations; 1,500 for legacy configurations

† Not stated on the *OpenPipeline limits* page (read 09/28/2026) — treat as community-reported and verify in your tenant before planning to it.

Field Restrictions

Read-Only Fields (Cannot be modified in pipelines):
- `dt.ingest.*` - Ingestion metadata
- `dt.openpipeline.*` - Pipeline processing metadata
- `dt.retain.*` - Retention information
- `dt.system.*` - System metadata (including bucket)

Entity Fields (Added **after** Processing stage — examples from the documented list):
- `dt.entity.service`
- `dt.entity.kubernetes_cluster`, `dt.entity.kubernetes_node`, `dt.entity.kubernetes_service`
- `dt.entity.cloud_application`, `dt.entity.cloud_application_instance`
- `dt.source_entity`

OPMIG-02 carries the full list. `dt.entity.host` and `dt.entity.process_group` are **not** on it.

> 💡 **Design Tip:** The listed entity fields are NOT available during routing or processing. They're added automatically by Dynatrace after the Processing stage, so they can be used in the later stages — such as Bucket assignment and the extraction stages — but not in routing conditions or Processing-stage processors.
>
> **Sources:**
>
> - OpenPipeline limits (DT docs) — *"The following fields are added after the Processing stage when Dynatrace runs its entity detection."*; *"If the timestamp is more than 10 minutes in the future, it's adjusted to the ingest server time plus 10 minutes."*; *"The maximum size of a record after processing is 16 MB."*; *"You can extract data on a single record in a maximum of five different pipelines"*
> - What's new in SaaS 1.348 (DT docs) — *"The log ingestion pipeline now accepts log records with timestamps up to 72 hours in the past, extended from the previous 24-hour limit."*

Next Steps

Now that you understand OpenPipeline and have assessed your current state, continue with the migration series:

Notebook	Focus Area
OPMIG-02	OpenPipeline Architecture & Key Concepts
OPMIG-03	Migration Assessment & Planning
OPMIG-04	Pipeline Configuration Fundamentals

Notebook	Focus Area
OPMIG-05	Routing & Bucket Management
OPMIG-06	Processing, Parsing & Transformation
OPMIG-07	Metric & Event Extraction
OPMIG-08	Security, Masking & Compliance
OPMIG-09	Troubleshooting & Validation

References

- OpenPipeline Documentation
- OpenPipeline Limits
- Processing Examples
- Log Processing Tutorial
- Ingest API Reference

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.