

# OPMIG-02: OpenPipeline Migration

## Guide: Part 2

---

**Series:** OPMIG — OpenPipeline Migration | **Notebook:** 2 of 10 | **Created:** December 2025 | **Last Updated:** 06/23/2026

## Architecture & Key Concepts

---

## Table of Contents

---

1. [Data Flow Architecture](#)
  2. [Processing Stages in Detail](#)
  3. [Pipeline Types](#)
  4. [Processor Types](#)
  5. [Dynatrace Pattern Language \(DPL\)](#)
  6. [Complete OpenPipeline Limits Reference](#)
  7. [Entity Field Availability Timeline](#)
  8. [Key Fields & Metadata](#)
  9. [Exploring Your Pipeline Configuration](#)
  10. [Understanding Processing Order](#)
  11. [Summary: Key Architecture Concepts](#)
- 

## Learning Objectives

---

By the end of this notebook, you will:

-  Understand the complete OpenPipeline data flow architecture

-  Learn detailed processing stage execution order
-  Master DPL (Dynatrace Pattern Language) fundamentals
-  **Know ALL OpenPipeline limits in comprehensive detail**
-  Understand processor types and capabilities
-  **Learn entity field availability timeline**
-  Explore your pipeline configuration with DQL

---

## Prerequisites

---

| Requirement           | Details                                                     |
|-----------------------|-------------------------------------------------------------|
| Dynatrace Environment | SaaS or Managed with Grail enabled                          |
| API Access            | <code>logs.read</code> token scope for DQL validation       |
| Knowledge             | OPMIG-01 or familiarity with Classic log ingestion concepts |


---

## Data Flow Architecture

---

Understanding how data flows through OpenPipeline is essential for designing effective pipelines.

# Complete Data Flow



**Doc alignment (May 2026):** OpenPipeline's official `/concepts/data-flow` documentation describes a **four-stage flow** (Ingest → Routing → Processing → Storage) with optional pre-processing on custom sources. Earlier versions of this notebook described a five-stage flow with a separate "Extraction" stage; metric, event, Smartscape, and cost-allocation extraction processors actually live *within* the Processing stage.

## Key Principles

- Pre-storage Processing:** All transformations happen BEFORE data is written to Grail.
- Order Matters:** Processors execute in the order defined within each pipeline.
- Multi-pipeline Support:** A single record can be processed by up to 5 pipelines.
- Entity Detection:** Entity fields ( `dt.entity.*` ) are added AFTER the Processing stage.
- Immutable Storage:** Once stored, data cannot be modified (mask early!).

# Processing Stages in Detail

---

## Stage 1: Ingest

Records enter via an ingest source. The source is recorded as

```
dt.openpipeline.source .
```

| Source Type       | Owner                    | Pre-processing | Routing           |
|-------------------|--------------------------|----------------|-------------------|
| <b>Built-in</b>   | OpenPipeline (view-only) | No             | Dynamic only      |
| <b>Ready-made</b> | Extensions               | No             | Dynamic only      |
| <b>Custom</b>     | User-defined             | Yes (optional) | Dynamic or static |

## Stage 2: Pre-processing (*optional, custom sources only*)

A DQL transform applied to records from a custom source *before* routing. Use this to normalize raw data from third-party shippers into a common shape that all downstream pipelines can consume.

## Stage 3: Routing

The routing stage determines which pipeline(s) process each incoming record.

| Aspect               | Description                                                                             |
|----------------------|-----------------------------------------------------------------------------------------|
| <b>Purpose</b>       | Match incoming data to appropriate pipelines                                            |
| <b>Timing</b>        | Before any pipeline processing                                                          |
| <b>Configuration</b> | Dynamic routing rules with matching conditions, or static assignment for custom sources |
| <b>Fallback</b>      | Unmatched data goes to the default pipeline                                             |

### Matching Condition Examples:

```
k8s.namespace.name == "production"
log.source == "nginx"
contains(content, "payment")
dt.openpipeline.source == "oneagent"
```

## Stage 4: Processing

All in-pipeline work happens here. Processors execute in the order defined in the pipeline. Functional groups within Processing:

| Group                                  | Processors                                                                                             | Purpose                                         |
|----------------------------------------|--------------------------------------------------------------------------------------------------------|-------------------------------------------------|
| <b>Masking</b>                         | DQL with <code>replacePattern</code>                                                                   | Redact PII before anything else                 |
| <b>Filtering</b>                       | Drop record                                                                                            | Remove unwanted records early                   |
| <b>Field &amp; record manipulation</b> | Add/Remove/Rename fields, DQL, Parse, Technology                                                       | Transform and enrich                            |
| <b>Metric extraction</b>               | Counter, Value, Histogram (Preview); sampling-aware variants for spans                                 | Generate timeseries from records                |
| <b>Smartscape topology</b>             | Smartscape node, Smartscape edge                                                                       | Populate the Smartscape topology graph          |
| <b>Event extraction</b>                | Business event, SDLC event, Davis event                                                                | Convert records into events                     |
| <b>Cost &amp; security</b>             | DPS Cost Allocation — Cost Center, DPS Cost Allocation — Product, Set <code>dt.security_context</code> | Attribute cost / apply record-level access      |
| <b>Storage assignment</b>              | Bucket assignment, No storage assignment                                                               | Place the record in a bucket, or skip retention |

Recommended order *within* Processing: mask → drop → transform/parse → extract → assign cost/security → bucket assignment.

## Stage 5: Storage

The storage stage routes data to Grail buckets. Use the **No storage assignment** processor to skip retention entirely (useful when a record is only needed to extract a metric or event).

| Aspect         | Description                                                  |
|----------------|--------------------------------------------------------------|
| Bucket Routing | Direct data to specific buckets                              |
| Retention      | Different retention periods per bucket                       |
| Cost Control   | Route high-volume data to shorter retention or to no-storage |

---

## Pipeline Types

---

OpenPipeline supports three types of pipelines:

### 1. Default Pipeline

- **Purpose:** Handles data that doesn't match any custom pipeline
- **Configuration:** Minimal processing, default bucket
- **Modification:** Can add processors but cannot delete

### 2. Custom Pipelines

- **Purpose:** Targeted processing for specific data sources or patterns
- **Configuration:** Full control over all processing stages
- **Routing:** Requires matching conditions in dynamic routing

### 3. Built-in Pipelines

- **Purpose:** Pre-configured for common technologies
- **Examples:** Kubernetes logs, AWS logs, Azure logs
- **Modification:** Can extend but not fundamentally change

 **Best Practice:** Create focused pipelines for specific use cases rather than one large pipeline with complex conditionals.

---

# Processor Types

## DQL Processor

The DQL processor uses DQL commands to transform data. Available commands:

| Command                   | Purpose                   | Example                                               |
|---------------------------|---------------------------|-------------------------------------------------------|
| <code>fieldsAdd</code>    | Add new fields            | <code>fieldsAdd environment = "prod"</code>           |
| <code>fieldsRemove</code> | Remove fields             | <code>fieldsRemove sensitive_field</code>             |
| <code>fieldsRename</code> | Rename fields             | <code>fieldsRename old_name = new_name</code>         |
| <code>parse</code>        | Extract with DPL patterns | <code>parse content, "LD:prefix<br/>INT:count"</code> |

### fieldsAdd Examples:

```
// Static value
| fieldsAdd environment = "production"

// Conditional value
| fieldsAdd severity = if(loglevel == "ERROR", "critical", else: "normal")

// Computed value
| fieldsAdd message_length = stringLength(content)

// From existing field
| fieldsAdd short_host = substring(host.name, 0, 10)
```

## Drop Processor

Removes records matching specified conditions:

```
Matching Condition: loglevel == "DEBUG"
Result: All DEBUG logs are dropped before storage
```

## Technology Processors

Built-in parsers for common log formats:

| Technology | Parsed Fields                                |
|------------|----------------------------------------------|
| Apache     | client_ip, method, path, status, bytes       |
| Nginx      | Similar to Apache with nginx-specific fields |
| JSON       | Flattens JSON structure to fields            |
| Syslog     | facility, severity, hostname, message        |

## Dynatrace Pattern Language (DPL)

DPL is a powerful pattern matching language used in the `parse` command.

### Core Matchers

| Matcher                | Description              | Example Match                               |
|------------------------|--------------------------|---------------------------------------------|
| <code>INT</code>       | Integer number           | <code>42</code> , <code>-17</code>          |
| <code>LONG</code>      | Long integer             | <code>1234567890123</code>                  |
| <code>DOUBLE</code>    | Decimal number           | <code>3.14</code> , <code>-0.5</code>       |
| <code>IPADDR</code>    | IPv4 or IPv6 address     | <code>192.168.1.1</code> , <code>::1</code> |
| <code>IPV4ADDR</code>  | IPv4 address only        | <code>10.0.0.1</code>                       |
| <code>IPV6ADDR</code>  | IPv6 address only        | <code>2001:db8::1</code>                    |
| <code>TIMESTAMP</code> | Timestamp with format    | <code>2024-01-15T10:30:00Z</code>           |
| <code>LD</code>        | Line data (to delimiter) | Any text until delimiter                    |
| <code>DATA</code>      | Any data (greedy)        | Consumes remaining text                     |
| <code>SPACE</code>     | Whitespace               | Spaces, tabs                                |
| <code>NSPACE</code>    | Non-whitespace           | Any non-space chars                         |
| <code>WORD</code>      | Word characters          | <code>hello</code> , <code>user123</code>   |
| <code>JSON</code>      | JSON structure           | <code>{"key": "value"}</code>               |



| Matcher | Description | Example Match   |
|---------|-------------|-----------------|
| EOL     | End of line | Line terminator |

## Pattern Syntax

| Element         | Syntax                         | Description                |
|-----------------|--------------------------------|----------------------------|
| Export to field | <code>MATCHER:fieldname</code> | Extract and name the field |
| Match only      | <code>MATCHER</code>           | Match but don't extract    |
| Optional        | <code>MATCHER?</code>          | Matcher is optional        |
| Literal         | <code>'exact text'</code>      | Match literal string       |
| Alternatives    | <code>('opt1'\ 'opt2')</code>  | Match either option        |
| Quantifier      | <code>MATCHER{2,5}</code>      | Match 2-5 times            |

## DPL Examples

```
// Extract IP and port from log
| parse content, "IPADDR:client_ip ':' INT:port"

// Extract user ID with flexible prefix
| parse content, "('user='|'userId='|'user_id=')LD:user_id"

// Parse Apache-style log
| parse content, "IPADDR:client_ip SPACE '-' SPACE LD:user SPACE '['
LD:timestamp ']"

// Extract JSON payload
| parse content, "LD JSON:payload"

// Parse with optional port
| parse content, "IPADDR:ip (':' INT:port)?"

// Extract error code
| parse content, "'error_code=' INT:error_code"
```

## Timestamp Format Patterns

| Symbol | Meaning          | Example |
|--------|------------------|---------|
| yyyy   | Year (4 digits)  | 2024    |
| MM     | Month (01-12)    | 01      |
| dd     | Day (01-31)      | 15      |
| HH     | Hour 24h (00-23) | 14      |
| mm     | Minute (00-59)   | 30      |
| ss     | Second (00-59)   | 45      |
| SSS    | Milliseconds     | 123     |

## Complete OpenPipeline Limits Reference

This comprehensive reference contains ALL OpenPipeline limits you need to know for planning your migration.

### Data Size & Volume Limits

| Limit                               | Value          | Behavior When Exceeded                     |
|-------------------------------------|----------------|--------------------------------------------|
| Max record size (after processing)  | 16 MB          | Record is <b>dropped</b> , logged in audit |
| Max record size (before processing) | 10 MB          | Ingestion rejected with 413 error          |
| Working memory per record           | 16 MB          | Processing fails, record dropped           |
| Log attribute size                  | 32 KB          | Attribute value <b>truncated</b>           |
| Max field name length               | 255 characters | Field creation fails                       |
| Max string field length             | 32 KB          | Content truncated                          |

| Limit                    | Value         | Behavior When Exceeded  |
|--------------------------|---------------|-------------------------|
| Max array size           | 1000 elements | Array truncated         |
| Max nesting depth (JSON) | 10 levels     | Deeper levels flattened |

## Processing Limits

| Limit                              | Value          | Impact                                                                                                                      |
|------------------------------------|----------------|-----------------------------------------------------------------------------------------------------------------------------|
| Max pipelines per record           | 5              | A record can be processed by up to 5 different pipelines. After 5, data extraction stops but the record is still persisted. |
| Max processors per pipeline        | 1,000          | Cannot add more processors to pipeline                                                                                      |
| Max DQL commands per processor     | 10 commands    | Split complex logic into multiple processors                                                                                |
| Max parse operations per processor | 100 patterns   | Create additional parse processors                                                                                          |
| Max fields per record              | 500 fields     | Additional fields ignored                                                                                                   |
| Processing timeout per record      | 30 seconds     | Record dropped if exceeded                                                                                                  |
| Max processor name length          | 100 characters | Validation error                                                                                                            |

## Timestamp Constraints

| Data Type       | Accepted Range                     | Records Outside Range |
|-----------------|------------------------------------|-----------------------|
| Logs            | 24 hours past to 10 minutes future | <b>Dropped</b>        |
| Spans           | 60 minutes past (end time)         | <b>Dropped</b>        |
| Events          | 24 hours past to 10 minutes future | <b>Dropped</b>        |
| Business Events | 24 hours past to 10 minutes future | <b>Dropped</b>        |

| Data Type | Accepted Range                 | Records Outside Range |
|-----------|--------------------------------|-----------------------|
| Metrics   | 1 hour past to 1 minute future | Dropped               |

**⚠ Critical:** Historical data imports require workarounds. Contact Dynatrace support for backfilling options.

## Pipeline Group Limits (*per group*)

| Limit                      | Value | Notes                                 |
|----------------------------|-------|---------------------------------------|
| Pipeline slots per group   | 10    | Per <a href="#">/reference/limits</a> |
| Member pipelines per group | 1,000 | Per <a href="#">/reference/limits</a> |

## Pipeline & Routing Limits

| Limit                    | Value          | Scope                                       |
|--------------------------|----------------|---------------------------------------------|
| Max custom pipelines     | 100 pipelines  | Per configuration scope (logs, spans, etc.) |
| Max dynamic routes       | 3,000 routes   | Per configuration scope                     |
| Max conditions per route | 10 conditions  | Combine with AND/OR operators               |
| Max pipeline name length | 100 characters | Validation error                            |
| Max route name length    | 100 characters | Validation error                            |

## Extraction Limits

| Limit                               | Value | Notes                            |
|-------------------------------------|-------|----------------------------------|
| Max metric extractions per pipeline | 10    | Value + counter metrics combined |
| Max event extractions per pipeline  | 5     | All event types combined         |

| Limit                                 | Value          | Notes                |
|---------------------------------------|----------------|----------------------|
| Max bizevent extractions per pipeline | 3              | Business events only |
| Max dimensions per metric             | 10 dimensions  | Keep cardinality low |
| Metric key length                     | 250 characters | Validation error     |
| Event type name length                | 100 characters | Validation error     |

## DQL & DPL Limits

| Limit                               | Value                                                                                    | Context                                                  |
|-------------------------------------|------------------------------------------------------------------------------------------|----------------------------------------------------------|
| DQL processor script length         | 8,192 characters                                                                         | Per <a href="#">/reference/limits</a> (May 2026 doc)     |
| Processor matching condition length | 4,096 characters (Settings API); 1,500 for legacy Configurations API / Classic pipelines | Per <a href="#">/reference/limits</a> (Jun 17, 2026 doc) |
| Max DPL pattern length              | 4 KB                                                                                     | Per parse pattern                                        |
| Max captured groups per parse       | 100 groups                                                                               | Use multiple parse operations                            |
| Max alternatives in pattern         | 50 alternatives                                                                          | <code>(opt1\ opt2\ ...\ opt50)</code>                    |

## Bucket & Storage Limits

| Limit                | Value                | Notes                     |
|----------------------|----------------------|---------------------------|
| Max custom buckets   | 50 buckets           | Per environment           |
| Min retention period | 1 day                | Per bucket                |
| Max retention period | 2555 days (~7 years) | Per bucket                |
| Bucket name length   | 100 characters       | Alphanumeric + underscore |

## Rate Limits

| Endpoint       | Limit             | Per       |
|----------------|-------------------|-----------|
| Log Ingest API | 500 requests/min  | Per token |
| OTLP Endpoint  | 1000 requests/min | Per token |
| Config API     | 100 requests/min  | Per token |

## Field & Matching Restrictions

**Read-Only Fields** (Cannot be modified):

```
dt.ingest.*      - Ingestion metadata
dt.openpipeline.* - Pipeline processing metadata
dt.retain.*      - Retention information
dt.system.*      - System metadata (bucket, etc.)
timestamp        - Record timestamp (can parse new, not modify original)
```

**Reserved Field Prefixes** (Cannot create):

```
dt.*      - Reserved for Dynatrace
dynatrace.* - Reserved for Dynatrace
```

**Entity Fields** (Available ONLY after Processing stage):

```
dt.entity.service
dt.entity.host
dt.entity.process_group
dt.entity.process_group_instance
dt.entity.kubernetes_cluster
dt.entity.cloud_application
dt.entity.cloud_application_namespace
```

💡 **Design Pattern:** Entity fields are added by Dynatrace AFTER the Processing stage. You cannot:

- Use them in routing conditions
- Use them in matching conditions during processing
- Modify them with fieldsAdd/fieldsRename

You CAN use them in:

- Extraction stage (metrics, events)
- Storage stage (bucket routing based on entity)


## Workarounds for Common Limit Issues

| Problem                          | Workaround                                                |
|----------------------------------|-----------------------------------------------------------|
| <b>&gt;50 processors needed</b>  | Split into multiple pipelines, use multi-pipeline routing |
| <b>&gt;10 DQL commands</b>       | Break into multiple processors (order matters!)           |
| <b>&gt;100 parse patterns</b>    | Use multiple parse processors in sequence                 |
| <b>&gt;16MB after processing</b> | Drop unnecessary fields, reduce field sizes               |
| <b>&gt;5 pipelines needed</b>    | Consolidate processing logic, use conditional processors  |
| <b>&gt;10 metric dimensions</b>  | Reduce cardinality, use separate metrics                  |

## Entity Field Availability Timeline

Understanding WHEN entity fields become available is critical for pipeline design.

## Processing Stage Timeline



## Practical Examples

### ✗ WRONG - Cannot route based on entity

Dynamic Route:

Condition: dt.entity.service == "SERVICE-123" ✗ FAILS

Why: Entity fields not available during routing

### ✗ WRONG - Cannot use entity in processing

DQL Processor:

fieldsAdd service\_name = dt.entity.service ✗ FAILS

Why: Entity fields not available during processing



## ✅ CORRECT - Use entity in metric extraction

```
Metric Extraction:  
  Key: log.request.duration  
  Value: duration_ms  
  Dimensions: dt.entity.service, dt.entity.host  ✅ WORKS
```

Why: Entity fields available in extraction stage

## ✅ CORRECT - Use entity for bucket routing

```
Bucket Routing (in Storage stage):  
  Condition: dt.entity.service == "CRITICAL-SERVICE"  
  Bucket: critical_logs  ✅ WORKS
```

Why: Entity fields available in storage stage

## Design Patterns

### Pattern 1: Route by source, extract metrics by entity

```
Routing: log.source == "payment-service" (uses source field)  
Processing: Parse payment data  
Extraction: Create metric with dt.entity.service dimension  ✅
```

### Pattern 2: Enrich with custom field, then use entity

```
Processing: fieldsAdd app_tier = "frontend" (custom field)  
Extraction: Dimensions: app_tier, dt.entity.service  ✅
```

### Pattern 3: Cannot mix entity with routing

```
❌ Route based on entity → Use workaround:  
  - Route based on source/content patterns instead  
  - Use service name string (not entity ID) if available
```

# Key Fields & Metadata

## OpenPipeline-Specific Fields

| Field                     | Description                       | Example Value              |
|---------------------------|-----------------------------------|----------------------------|
| dt.openpipeline.source    | Data source identifier            | oneagent , generic , otlp  |
| dt.openpipeline.pipelines | Pipeline(s) that processed record | ["custom-pipeline-1"]      |
| dt.system.bucket          | Grail storage bucket              | default_logs , custom_logs |

## Log Fields

| Field        | Description                             |
|--------------|-----------------------------------------|
| timestamp    | Log timestamp (primary time field)      |
| content      | Log message content                     |
| loglevel     | Log level (ERROR, WARN, INFO, DEBUG)    |
| status       | Status string (alternative to loglevel) |
| log.source   | Log source identifier                   |
| log.iostream | Stream type (stdout, stderr)            |

## Entity Context Fields

| Field                            | Description                |
|----------------------------------|----------------------------|
| dt.entity.host                   | Host entity ID             |
| dt.entity.process_group          | Process group entity ID    |
| dt.entity.process_group_instance | Process instance entity ID |
| dt.entity.service                | Service entity ID          |
| host.name                        | Host name (string)         |

| Field                                | Description             |
|--------------------------------------|-------------------------|
| <code>process.executable.name</code> | Process executable name |

**⚠ Important:** Entity fields ( `dt.entity.*` ) are added AFTER the Processing stage. You cannot use them in routing or processing conditions.

## Exploring Your Pipeline Configuration

Use these queries to understand how your environment is configured.

```
// View data sources currently sending to OpenPipeline
// Shows the distribution of data by ingestion source
fetch logs, from: now() - 24h
| summarize {record_count = count()}, by: {dt.openpipeline.source}
| sort record_count desc
```

```
// Analyze which pipelines are processing your logs
// Helps verify routing is working correctly
fetch logs, from: now() - 24h
| filter isNotNull(dt.openpipeline.pipelines)
| summarize {record_count = count()}, by: {dt.openpipeline.pipelines}
| sort record_count desc
```

```
// Check bucket distribution for stored logs
// Verify data is routing to expected buckets
fetch logs, from: now() - 24h
| summarize {record_count = count()}, by: {dt.system.bucket}
| sort record_count desc
```

```
// Analyze pipeline processing by source and pipeline
// Shows the relationship between sources and pipelines
fetch logs, from: now() - 24h
| filter isNotNull(dt.openpipeline.pipelines)
| summarize {record_count = count()}, by: {dt.openpipeline.source,
dt.openpipeline.pipelines}
| sort record_count desc
| limit 25
```


```
// Check for span processing through OpenPipeline
// Spans also flow through OpenPipeline
fetch spans, from: now() - 24h
| summarize {span_count = count()}, by: {dt.openpipeline.source}
| sort span_count desc
```

```
// View pipeline processing over time
// Helps identify volume patterns and processing trends
fetch logs, from: now() - 24h
| filter isNotNull(dt.openpipeline.pipelines)
| makeTimeseries {record_count = count()}, by: {dt.openpipeline.pipelines},
interval: 1h
```

```
// Identify logs going to default pipeline (may need custom routing)
// High volume in default may indicate missing routing rules
fetch logs, from: now() - 24h
| filter isNull(dt.openpipeline.pipelines) OR dt.openpipeline.pipelines == ""
| summarize {unrouted_count = count()}, by: {log.source}
| sort unrouted_count desc
| limit 20
```

## Understanding Processing Order

The order of processors within the Processing stage is critical:



💡 **Tip:** Masking is applied FIRST so sensitive data is protected even if subsequent processors fail. The numbered groups above are processor *categories within the Processing stage* — they are not separate pipeline stages.

## Summary: Key Architecture Concepts

| Concept                 | Key Points                                                                                                                                       |
|-------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Data Flow</b>        | Ingest → (Pre-process) → Route → Process → Store (4 stages + optional pre-processing per <a href="/concepts/data-flow">/concepts/data-flow</a> ) |
| <b>Pre-storage</b>      | All processing happens before data is persisted                                                                                                  |
| <b>Source Types</b>     | Built-in / Ready-made / Custom — only custom sources support pre-processing and static routing                                                   |
| <b>Routing</b>          | Dynamic (DQL matcher) or static (custom sources only)                                                                                            |
| <b>Processing Order</b> | Mask → Drop → Transform → Extract → Cost/Security → Storage assignment                                                                           |
| <b>Entity Detection</b> | Happens AFTER processing, BEFORE most extraction processors run                                                                                  |
| <b>Multi-pipeline</b>   | One record can be processed by up to 5 pipelines ( <a href="/reference/limits">/reference/limits</a> )                                           |
| <b>DPL</b>              | Powerful pattern language for parsing                                                                                                            |
| <b>Buckets</b>          | Control retention and cost at the storage assignment step; or skip with No storage assignment                                                    |

## Next Steps

Now that you understand OpenPipeline architecture, continue with:

| Notebook | Focus Area                           |
|----------|--------------------------------------|
| OPMIG-03 | Migration Assessment & Planning      |
| OPMIG-04 | Pipeline Configuration Fundamentals  |
| OPMIG-05 | Routing & Bucket Management          |
| OPMIG-06 | Processing, Parsing & Transformation |

---

## References

- [OpenPipeline Data Flow](#)
- [OpenPipeline Processing](#)
- [Dynatrace Pattern Language](#)
- [OpenPipeline Limits](#)
- [DPL Architect Tool](#)

---

*Last Updated: May 6, 2026*

---

*This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.*