

ORGNZ-99: Best Practice Summary & DQL Reference

Series: ORGNZ — Organize Data: Buckets, Segments, Security | **Notebook:** 99 | **Created:** March 2026 | **Last Updated:** 09/24/2026

Overview

This notebook consolidates every actionable best practice from the ORGNZ series (notebooks 01-10) into a single definitive reference. Each section provides the exact setting, value, or action to take — and a **Validation Queries** subsection with executable DQL to verify your configuration against a live Grail environment.

Table of Contents

1. [Bucket Design](#)
2. [Retention Strategy](#)
3. [Data Routing](#)
4. [IAM and Access Control](#)
5. [Security Context](#)
6. [Segment Design](#)
7. [Metadata Enrichment](#)
8. [Query Performance](#)
9. [Governance and Operations](#)

Each section ends with a **Validation Queries** block — executable DQL to verify your configuration against a live Grail environment.

1. Bucket Design

#	Best Practice	Recommended Setting/Value	Priority	Category
1	Keep per-bucket ingest under 1 TB/day	Target: <1 TB/day per bucket; split at >2 TB/day (Dynatrace guidance)	Critical	Performance
2	One data type per bucket	Each bucket stores exactly one table type (logs, metrics, spans, events, bizevents, security.events)	Critical	Architecture
3	Use consistent naming convention	Format: <code><provider>_<type>_<retention></code> or <code><team>_<type>_<retention></code> (e.g., <code>aws_logs_35d</code> , <code>team_platform_logs</code>)	Critical	Naming
4	Bucket names: lowercase, letter-first, underscores/hyphens only	Regex: <code>^[a-z][a-z0-9_-]*\$</code> ; never use <code>default_*</code> prefix	Critical	Naming
5	Plan bucket names before creation	Bucket names are immutable after creation; display names can be changed	Critical	Planning
6	Stay within the custom-bucket environment default — 250 from SaaS 1.346 (80 on earlier versions)	Budget buckets across teams; request increase only if justified	Recommended	Planning
7	Create dedicated buckets per cost center	One bucket per business unit/LOB for direct GiB/day cost attribution	Recommended	Cost
8	Create separate buckets for compliance data	Dedicated bucket (e.g., <code>compliance_audit_logs</code>) with extended retention for regulatory requirements	Recommended	Compliance
9	Create short-retention buckets for debug logs	Bucket: <code>debug_logs_7d</code> with 7-day retention for high-volume verbose logs	Recommended	Cost
10	Split high-volume buckets to maintain queryability	If bucket exceeds 1 TB/day, split by team, environment, or application	Recommended	Performance

Validation Queries — Bucket Design

List all buckets with data type, retention, and estimated storage:

```
// Inventory of all Grail buckets – verify data types, retention, and naming convention
fetch dt.system.buckets
| fields name, display_name, dt.system.table, retention_days, estimated_uncompressed_bytes
| sort dt.system.table asc, name asc
```

List all built-in buckets (use to compare your environment to the expected defaults):

```
// Discover all built-in Grail buckets – compare against expected defaults
fetch dt.system.buckets
| filter startsWith(name, "default_") or startsWith(name, "dt_")
| fields name, display_name, dt.system.table, retention_days
| sort dt.system.table asc, name asc
```

2. Retention Strategy

#	Best Practice	Recommended Setting/Value	Priority	Category
11	Set debug/development log retention to 3-7 days	Bucket retention: 7 days maximum for DEBUG-level logs	Critical	Cost
12	Set standard operational log retention to 35 days	Bucket retention: 35 days for application and infrastructure logs	Recommended	Operations
13	Set performance analysis retention to 60-90 days	Bucket retention: 90 days for logs/events used in quarterly trend analysis (metrics retention is fixed at ~15 months in <code>default_metrics</code>)	Recommended	Operations
14	Set audit/compliance log retention to 365-3657 days	Bucket retention: match regulatory requirement (SOX: 7 years = 2,555 days; HIPAA: 6 years = 2,190 days)	Critical	Compliance
15	Set span retention to 10-35 days	Bucket retention: 10 days for default APM; 35 days if long-running incident correlation needed	Recommended	Operations
16	Set security event retention to 90-365 days	Bucket retention: 365 days for security investigation timelines	Recommended	Security
17	Never lower retention without understanding data loss	Reducing retention immediately deletes data older than the new setting; this is irreversible	Critical	Operations

Validation Queries — Retention Strategy

Compare retention settings across all bucket types:

```
// Compare bucket retention configurations – verify each tier matches its compliance requirement
fetch dt.system.buckets
| fields name, dt.system.table, retention_days
| sort retention_days desc
```

3. Data Routing

#	Best Practice	Recommended Setting/Value	Priority	Category
18	Route data to custom buckets via OpenPipeline	Configure Storage-stage Bucket assignment processors (first match wins) with DQL matchers on <code>loglevel</code> , <code>dt.host_group.id</code> , <code>log.source</code> , <code>dt.cost.*</code>	Critical	Architecture
19	Decide where unmatched records go	There is no <code>default:</code> route setting; data with no bucket assignment goes to the scope's default bucket (<code>default_logs</code> for logs) — check unmatched records in the pipeline preview	Critical	Reliability
20	Route DEBUG logs to short-retention buckets	Bucket assignment matcher <code>loglevel == "DEBUG"</code> → bucket <code>debug_logs_7d</code>	Recommended	Cost
21	Route audit logs to compliance buckets	Bucket assignment matcher <code>matchesPhrase(log.source, "audit")</code> → bucket <code>audit_logs_365d</code>	Recommended	Compliance
22	Data cannot be moved between buckets after ingestion	Route new data via OpenPipeline; old data stays in the original bucket until retention expires	Critical	Architecture

Validation Queries — Data Routing

Verify routing rules are directing data to the expected buckets:

```
// Audit log volume by bucket – verify routing rules are directing data correctly
fetch logs, from:-1h
| summarize recordCount = count(), by:{dt.system.bucket}
| sort recordCount desc
```

```
// Track log ingest trends by bucket over the last day – identify routing anomalies
fetch logs, from:-24h
| summarize recordCount = count(), by:{time_bucket = bin(timestamp, 1h), dt.system.bucket}
| sort time_bucket desc
```

4. IAM and Access Control

#	Best Practice	Recommended Setting/Value	Priority	Category
23	Always include <code>storage:buckets:read</code> in every data access policy	Without bucket read, users cannot query any data; always pair with table permissions	Critical	Security
24	Assign policies to groups, never to individual users	Create IAM groups per team/role; assign all policies at the group level	Critical	Governance
25	Use the two-step policy pattern	Step 1: <code>ALLOW storage:buckets:read WHERE ...</code> ; Step 2: <code>ALLOW storage:logs:read</code> (or other table types)	Critical	Security
26	Use <code>STARTSWITH</code> for team bucket prefixes	Policy: <code>ALLOW storage:buckets:read WHERE storage:bucket-name STARTSWITH "team_platform_"</code> — automatically includes future buckets with matching prefix	Recommended	Scalability
27	Use <code>MATCH</code> operator for array fields	When <code>dt.security_context</code> holds an array, use <code>MATCH</code> (not <code>=</code> , <code>STARTSWITH</code> , or <code>IN</code> — those always return false for arrays)	Critical	Security
28	Start with least privilege	Grant minimum required access; expand only when justified	Critical	Security
29	Use policy boundaries for reusable conditions	Define regional/environmental restrictions once as a boundary; apply to multiple policies	Recommended	Governance
30	Limit policy boundaries to 10 restrictions per boundary	Maximum 10 conditions per boundary; create additional boundaries if more are needed	Recommended	Architecture
31	Use fieldsets for sensitive fields	Grant <code>storage:fieldsets:read</code> on specific fieldsets (e.g. <code>builtin-sensitive-spans</code> , or a custom fieldset for log fields) only where needed — users without it don't see those fields. There is no per-field <code>DENY</code> permission	Recommended	Security
32	Test policies with sample users before deployment	Create a test user in the target group and verify query results before rolling out to the team	Critical	Operations

#	Best Practice	Recommended Setting/Value	Priority	Category
33	Document every IAM policy	Include name, description, tags, target group, and purpose in a policy registry	Recommended	Governance

Validation Queries — IAM and Access Control

Verify bucket access and data distribution across teams:

```
// Verify bucket access – distribution of accessible log data by bucket
fetch logs, from:-1h
| summarize count = count(), by:{dt.system.bucket}
| sort count desc
```

```
// Check namespace-based access – verify namespace-scoped policies return expected data
fetch logs, from:-1h
| filter isNotNull(k8s.namespace.name)
| summarize count = count(), by:{k8s.namespace.name}
| sort count desc
```

```
// Check host group-based access – verify host-group-scoped policies return expected data
fetch logs, from:-1h
| filter isNotNull(dt.host_group.id)
| summarize count = count(), by:{dt.host_group.id}
| sort count desc
```

5. Security Context

#	Best Practice	Recommended Setting/Value	Priority	Category
34	Use <code>dt.security_context</code> for record-level access control	Set at the source first (host tag, <code>DT_TAGS</code> , K8s namespace labels, OTel resource attribute); OpenPipeline Set security context processor for computed values	Critical	Security
35	Use hierarchical encoding in security context values	Format: <org>/<department>/<team> (e.g., <code>acme/engineering/platform</code>); enables department-level access with <code>MATCH ('acme/engineering/*')</code>	Critical	Scalability
36	Use type prefixes for security context values	Prefix with <code>team:</code> , <code>lob:</code> , <code>env:</code> , <code>cloud:</code> (e.g., <code>team:checkout</code> , <code>lob:finance</code> , <code>env:production</code>)	Recommended	Naming

#	Best Practice	Recommended Setting/Value	Priority	Category
37	Plan security context hierarchy before implementation	Context values are assigned to records at ingest time; changing schema requires re-ingesting data	Critical	Planning
38	Use security context instead of additional buckets when at the bucket-count limit	Security context enables multi-team isolation within shared buckets without consuming bucket quota	Recommended	Scalability
39	Set security context on entities via Grail Security Context settings	Settings → Topology model → Grail Security Context; this scopes the entities only, not their logs/spans/metrics — use a host tag or primary-field enrichment for those	Recommended	Security
40	Always use MATCH for array-valued security context in IAM policies	<code>ALLOW storage:logs:read WHERE storage:dt.security_context MATCH ("team-a*")</code>	Critical	Security

Validation Queries — Security Context

Verify `dt.security_context` assignment, distribution, and hierarchical patterns:

```
// View logs with security context – confirm dt.security_context is being set
fetch logs, from:-1h
| filter isNotNull(dt.security_context)
| fields timestamp, content, dt.security_context
| limit 20
```

```
// Summarize log volume by security context – verify naming conventions and distribution
fetch logs, from:-1h
| filter isNotNull(dt.security_context)
| summarize recordCount = count(), by:{dt.security_context}
| sort recordCount desc
| limit 20
```

```
// Filter logs by specific security context – test access control boundary for a team
fetch logs, from:-1h
| filter dt.security_context == "team:platform"
| limit 50
```

```
// Check hierarchical security context – verify parent-level wildcard access works
fetch logs, from:-1h
| filter isNotNull(dt.security_context)
| filter startsWith(dt.security_context, "finance/")
| summarize count = count(), by:{dt.security_context}
| sort count desc
```

```
// Audit security context coverage — percentage of logs with dt.security_context assigned
fetch logs, from:-1h
| summarize
    total = count(),
    withContext = countIf(isNotNull(dt.security_context)),
    withoutContext = countIf(isNull(dt.security_context))
| fieldsAdd coverage = round(toDouble(withContext) / toDouble(total) * 100, decimals: 2)
```

6. Segment Design

#	Best Practice	Recommended Setting/Value	Priority	Category
41	Use segments for dynamic data views, not access control	Segments filter at query time; they do not enforce security boundaries — IAM policies do	Critical	Architecture
42	Stay within the 20-include limit per segment	Maximum 20 include blocks per segment; plan rule complexity accordingly	Critical	Architecture
43	One include block per data type	Cannot have two <code>logs</code> includes; combine conditions with <code>OR</code> within a single include	Critical	Architecture
44	Use Primary Grail Fields in segment filters	Filter on <code>dt.host_group.id</code> , <code>k8s.namespace.name</code> , <code>k8s.cluster.name</code> , <code>aws.account.id</code> , <code>azure.subscription</code> , <code>gcp.project.id</code> — these are indexed and propagate across all signal types	Critical	Performance
45	Use prefix-based naming for cross-signal filtering	Name host groups, namespaces, etc. with prefixes (e.g., <code>prod-web-tier</code> not <code>web-tier-prod</code>) so <code>starts-with</code> works on both entity and signal includes	Recommended	Naming
46	Create one segment per dimension	Separate segments for platform, application, environment, and region — users combine them as needed	Recommended	Architecture
47	Add event includes for detected problem filtering	Define events include with <code>event.kind = "DAVIS_PROBLEM"</code> — entity includes alone do not filter problems	Critical	Troubleshooting
48	Set official segments to Public visibility	Platform team-owned segments: set to Public; application team segments: share with their group	Recommended	Governance

#	Best Practice	Recommended Setting/Value	Priority	Category
49	Use consistent segment naming prefixes	Prefixes: team- , env- , app- , region- (e.g., team-platform-infra , env-production)	Recommended	Naming
50	Test segment filter conditions with DQL before creating the segment	Run the equivalent filter clause as a DQL query; verify record counts match expectations	Critical	Operations
51	Limit segment variable dropdowns to 10,000 values	Maximum 10,000 values per variable; 100 selected values per segment	Recommended	Architecture
52	Use entity relationship traversal for comprehensive segments	Top-down: host → process group → service; bottom-up: service → process group → host	Recommended	Architecture
53	Segments reduce scanned bytes	Applying a segment injects filters at query time, reducing data scanned and improving performance on high-volume buckets	Recommended	Performance

Validation Queries — Segment Design

Simulate segment filter conditions before committing to a segment definition:

```
// Simulate host-group-based segment filter on logs
fetch logs, from:-1h
| filter isNotNull(dt.host_group.id)
| summarize count = count(), by:{dt.host_group.id}
| sort count desc
| limit 10
```

```
// Simulate tag-based segment filter on services
fetch dt.entity.service
| expand tags
| filter contains(tags, "team:platform")
| fields entity.name, tags
| limit 20

// Smartscape equivalent (dt.entity.* is deprecated but still functional):
//   smartscapeNodes "SERVICE"
//   | expand tags
//   | filter contains(tags, "team:platform")
//   | fields name, tags
//   | limit 20
// Caveat: Smartscape reflects CURRENT live topology and can report fewer entities than
// the classic entity store; for a pre-migration inventory keep the classic query above.
// Note: entity tags are not a flat "tags" field on Smartscape (resolve via getNodeField).
```

```
// Simulate combined segment filter – verify host group + namespace combination returns
expected logs
fetch logs, from:-1h
| filter isNotNull(dt.host_group.id) and isNotNull(k8s.namespace.name)
| summarize count = count(), by:{dt.host_group.id, k8s.namespace.name}
| sort count desc
| limit 20
```

7. Metadata Enrichment

#	Best Practice	Recommended Setting/Value	Priority	Category
54	Enrich all signals with Primary Grail Fields before building segments	Verify <code>dt.host_group.id</code> , <code>k8s.namespace.name</code> , <code>k8s.cluster.name</code> are present on logs, spans, and metrics	Critical	Architecture
55	Set <code>dt.security_context</code> via a host tag for dedicated infrastructure	Host tag <code>dt.security_context=<team></code> (installation, <code>oneagentctl --set-host-tag</code> , or Deployment Status)	Recommended	Security
56	Set <code>dt.cost.costcenter</code> and <code>dt.cost.product</code> on hosts	Host tags: <code>dt.cost.costcenter=<center></code> , <code>dt.cost.product=<product></code> — these propagate to service metrics	Recommended	Cost

#	Best Practice	Recommended Setting/Value	Priority	Category
57	Use <code>DT_TAGS</code> environment variable for shared infrastructure	Set <code>DT_TAGS=primary_tags.team=<team></code> per process when multiple apps share the same host	Recommended	Architecture
58	Restart applications after enrichment changes for spans	Logs: no restart needed (OneAgent auto-enriches); Spans: application restart required; Metrics: OneAgent restart applies automatically	Critical	Operations
59	Verify enrichment coverage before building segments	Run DQL audit query to check field coverage percentage across signal types; fields with 0% coverage need enrichment	Critical	Operations
60	Use Primary Grail Tags (<code>primary_tags.*</code>) when standard fields do not match your organization	Set via host tags (<code>hostautotag.conf</code> / <code>oneagentctl --set-host-tag</code>) or the <code>DT_TAGS</code> environment variable; propagates across logs, spans, and topology. Host properties no longer work for this — see row 60a	Recommended	Architecture
60a	Do not set host-level <code>primary_tags.*</code> via <code>hostcustomproperties.conf</code>	Breaking, OneAgent 1.345: "the OS Agent derives <code>primary_tags</code> from host tags (<code>hostautotag.conf</code>) rather than from host properties (<code>hostcustomproperties.conf</code>) at the host level." Existing primary tags configured through host properties silently stop being promoted once a host reaches 1.345 — audit before the fleet upgrades, and migrate them to host tags	Critical	Migration

Validation Queries — Metadata Enrichment

Audit primary Grail field coverage before building segment definitions:

```
// Audit primary Graill field coverage on logs – low-coverage fields need enrichment before
segment use
fetch logs, from:-1h
| summarize
    total = count(),
    with_namespace = countIf(isNotNull(k8s.namespace.name)),
    with_cluster = countIf(isNotNull(k8s.cluster.name)),
    with_host_group = countIf(isNotNull(dt.host_group.id)),
    with_security_ctx = countIf(isNotNull(dt.security_context))
| fieldsAdd
    ns_pct = round(toDouble(with_namespace) / toDouble(total) * 100, decimals: 1),
    cluster_pct = round(toDouble(with_cluster) / toDouble(total) * 100, decimals: 1),
    hg_pct = round(toDouble(with_host_group) / toDouble(total) * 100, decimals: 1),
    sec_pct = round(toDouble(with_security_ctx) / toDouble(total) * 100, decimals: 1)
```

```
// Audit primary Graill field coverage on spans – spans require app restart to pick up new
enrichment
fetch spans, from:-1h
| summarize
    total = count(),
    with_namespace = countIf(isNotNull(k8s.namespace.name)),
    with_cluster = countIf(isNotNull(k8s.cluster.name)),
    with_host_group = countIf(isNotNull(dt.host_group.id)),
    with_service = countIf(isNotNull(dt.entity.service))
| fieldsAdd
    ns_pct = round(toDouble(with_namespace) / toDouble(total) * 100, decimals: 1),
    cluster_pct = round(toDouble(with_cluster) / toDouble(total) * 100, decimals: 1),
    hg_pct = round(toDouble(with_host_group) / toDouble(total) * 100, decimals: 1),
    svc_pct = round(toDouble(with_service) / toDouble(total) * 100, decimals: 1)
```

```
// Variable DQL: list host groups – use as segment variable for host-group-scoped filters
fetch dt.entity.host_group
| fields host_group = entity.name, id
| sort host_group asc

// Note: dt.entity.* is deprecated. There is no Smartscape "HOST_GROUP" node type; host-
group
// membership is exposed as the dt.host_group.id field on HOST nodes. Keep the classic
query above.
```

```
// Variable DQL: list Kubernetes namespaces – use as segment variable for namespace-scoped filters
fetch dt.entity.cloud_application_namespace
| fields namespace = entity.name
| fieldsAdd tag = concat("Namespace:", namespace)
| sort namespace asc

// Smartscape equivalent (dt.entity.* is deprecated but still functional):
//   smartscapeNodes "K8S_NAMESPACE"
//   | fields namespace = name
//   | fieldsAdd tag = concat("Namespace:", namespace)
//   | sort namespace asc
// Caveat: Smartscape reflects CURRENT live topology and can report fewer entities than
// the classic entity store; for a pre-migration inventory keep the classic query above.
// Note: entity tags are not a flat "tags" field on Smartscape (resolve via getNodeField).
```

8. Query Performance

#	Best Practice	Recommended Setting/Value	Priority	Category
61	Understand the default 500 GB scan limit	<code>fetch</code> stops at 500 GB scanned by default (<code>scanLimitGBytes</code>); at 1 TB/day ingest an unmodified query reaches back ~12 hours — raise <code>scanLimitGBytes</code> or narrow the query	Critical	Performance
62	Target specific buckets in queries	Use <code>fetch logs, bucket:{"team_logs", "app_logs_*"}</code> to avoid scanning all buckets	Recommended	Performance
63	Use exact match (<code>=</code>) over pattern match in segment filters	<code>=</code> is faster than <code>starts-with</code> ; use exact values when known	Recommended	Performance
64	Minimize segment filter expressions	Fewer conditions = better query performance; maximum 10 expressions per filter	Recommended	Performance
65	Use OpenPipeline pre-processing for complex segment logic	If segment conditions would be complex, pre-enrich a simple field (e.g., <code>dt.security_context</code>) and filter on that	Recommended	Performance
66	Maximum 1,000 records returned per query	Use aggregations (<code>summarize</code> , <code>makeTimeseries</code>) for larger datasets	Recommended	Performance

#	Best Practice	Recommended Setting/Value	Priority	Category
67	Maximum 10 segments per query	Do not stack more than 10 segments; design broader segments instead of many narrow ones	Recommended	Architecture

Validation Queries — Query Performance

Use bucket targeting to control scan scope and stay within the 500 GB scan limit:

```
// Query from multiple buckets simultaneously – reduce scan scope by naming targets explicitly
fetch logs, from:-1h, bucket:{"default_logs", "audit_logs", "security_logs"}
| limit 100
```

```
// Filter by bucket within query – useful when bucket list is dynamic
fetch logs, from:-1h
| filter dt.system.bucket == "default_logs"
| filter loglevel == "ERROR"
| limit 50
```

9. Governance and Operations

#	Best Practice	Recommended Setting/Value	Priority	Category
68	Schedule regular access reviews	Quarterly review of all IAM policies and group memberships; remove unnecessary permissions	Critical	Governance
69	Document all bucket purposes	Maintain a registry mapping each bucket to its owner, data type, retention, cost center, and routing rule	Recommended	Governance
70	Document all security context values	Maintain a registry of all <code>dt.security_context</code> values, their meaning, and which IAM policies reference them	Recommended	Governance
71	Platform team owns infrastructure segments	Platform/SRE team creates and maintains public environment and infrastructure segments	Recommended	Governance
72	Application teams own app-specific segments	Each team creates segments for their applications; share with their IAM group	Recommended	Governance

#	Best Practice	Recommended Setting/Value	Priority	Category
73	Review segments quarterly	Remove unused segments; verify filter conditions still match data; check for empty results	Recommended	Operations
74	Use combined approach for defense in depth	Buckets for retention/cost + Security context for access control + Segments for dynamic views — all three together	Critical	Architecture
75	Implement in phases	Phase 1: Plan → Phase 2: Buckets + routing → Phase 3: IAM policies → Phase 4: Segments → Phase 5: Governance	Recommended	Operations
76	Selective record deletion is not possible	Cannot delete individual records; wait for bucket retention to expire, or truncate the entire bucket	Critical	Operations
77	Bucket deletion destroys all data permanently	Deleting a bucket removes ALL records immediately and irreversibly	Critical	Operations

Validation Queries — Governance and Operations

Routine audit queries for operational health checks:

```
// Audit bucket data distribution – verify all expected buckets are receiving data
fetch logs, from:-1h
| summarize
    recordCount = count(),
    by:{dt.system.bucket}
| sort recordCount desc
```

```
// Audit bucket retention settings – verify retention tiers match compliance requirements
fetch dt.system.buckets
| fields name, dt.system.table, retention_days
| sort dt.system.table asc, retention_days desc
```

Audit bucket management events — who created, updated, truncated, or deleted a bucket and when:

```
// Audit all bucket management actions – create, update, truncate, delete
// Bucket changes are rare: widen to from:-365d for a full year of history
fetch dt.system.events, from:-90d
| filter event.kind == "AUDIT_EVENT" and event.category == "BUCKET_MANAGEMENT"
| fields timestamp, event.type, resource, user.id, event.outcome
| sort timestamp desc
```

This notebook was AI-generated from Dynatrace documentation and enterprise best practices. It is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.