

OTEL-02: OpenTelemetry Collector Architecture

Series: OTEL — OpenTelemetry Integration | **Notebook:** 2 of 8 | **Created:** January 2026 | **Last Updated:** 10/02/2026

Understanding the OTel Collector Pipeline

The OpenTelemetry Collector is the backbone of OTel deployments—a vendor-agnostic service for receiving, processing, and exporting telemetry data. This notebook covers its architecture, components, and configuration.

Table of Contents

- 1. [Pipeline Architecture](#)
- 2. [Receivers](#)
- 3. [Processors](#)
- 4. [Exporters](#)
- 5. [Extensions](#)
- 6. [Connectors](#)
- 7. [Configuration Basics](#)

Prerequisites

Requirement	Details
Knowledge	OTEL-01: OpenTelemetry Fundamentals

Requirement	Details
Tools	YAML familiarity for configuration

1. Collector Overview

What is the Collector?

The OpenTelemetry Collector is a standalone service that:

Function	Description
Receives	Ingests telemetry from multiple sources
Processes	Transforms, filters, enriches data
Exports	Sends to one or more backends

Collector Distributions

Distribution	Use Case	Components
Core	Minimal, essential only	Basic receivers/exporters
Contrib	Extended functionality	Many community components
Custom	Your specific needs	Build with ocb (OTel Collector Builder)
Vendor	Pre-configured for vendor	Dynatrace, AWS, etc.


Why Use a Collector?

Benefit	Description
Decoupling	Apps don't need backend-specific SDKs
Processing	Centralized filtering, sampling, enrichment
Fan-out	Send to multiple backends
Buffering	Retry and batching for reliability

Benefit	Description
Security	Centralize credentials and auth

2. Pipeline Architecture

Data Flow



Pipeline Types

Pipeline	Data Type	Example
traces	Distributed traces	Request spans
metrics	Quantitative measurements	CPU usage
logs	Log records	Application logs

Each signal type has its own pipeline(s) with separate receivers, processors, and exporters.

3. Receivers

Receivers ingest telemetry data from external sources.

Common Receivers

Receiver	Protocol	Signals	Use Case
<code>otlp</code>	OTLP (gRPC/HTTP)	Traces, Metrics, Logs	OTel SDKs
<code>jaeger</code>	Jaeger (gRPC/HTTP)	Traces	Jaeger clients
<code>zipkin</code>	Zipkin HTTP	Traces	Zipkin clients
<code>prometheus</code>	Prometheus scrape	Metrics	Prometheus targets
<code>file_log</code> (<code>filelog</code> before contrib v0.149.0)	File tailing	Logs	Log files
<code>host_metrics</code> (<code>hostmetrics</code> before contrib v0.151.0)	Host stats	Metrics	CPU, memory, disk

Component renames are version-dependent. Collector-contrib renamed `filelog` → `file_log` in **v0.149.0** and `hostmetrics` → `host_metrics`, `spanmetrics` → `span_metrics` in **v0.151.0**; each component's `metadata.yaml` keeps the old name as `deprecated_type`, so both spellings load on those versions and later. Older Collectors — including Dynatrace Collector builds based on an earlier upstream — know **only** the old names. If one config must run on a mixed fleet, use the old names until every Collector is on v0.151.0+.

Sources: [filelogreceiver metadata.yaml \(OpenTelemetry Collector Contrib GitHub\)](#) — `type: file_log`, `deprecated_type: filelog`; same pattern in `hostmetricsreceiver` and `spanmetricsconnector`. Release versions read from the tagged files, 10/02/2026.

OTLP Receiver Configuration

```
receivers:
  otlp:
    protocols:
      grpc:
        endpoint: 0.0.0.0:4317
      http:
        endpoint: 0.0.0.0:4318
```

Prometheus Receiver

```
receivers:
  prometheus:
    config:
      scrape_configs:
        - job_name: 'my-app'
          scrape_interval: 15s
          static_configs:
            - targets: ['app:8080']
```

4. Processors

Processors transform, filter, or enrich telemetry data.

Essential Processors

Processor	Purpose	When to Use
batch	Batch data for efficiency	Always
memory_limiter	Prevent OOM	Always
resourcedetection	Add resource attributes	K8s, cloud
attributes	Add/modify attributes	Enrichment
filter	Drop unwanted data	Cost reduction
transform	OTTL-based transformations	Complex processing

Batch Processor (Critical)

```
processors:
  batch:
    timeout: 10s
    send_batch_size: 1000
    send_batch_max_size: 1500
```

Memory Limiter (Critical)

```
processors:
  memory_limiter:
    check_interval: 1s
    limit_mib: 400
    spike_limit_mib: 100
```

Attributes Processor

```
processors:
  attributes:
    actions:
      - key: environment
        value: production
        action: upsert
      - key: api.key
        action: delete
```

Filter Processor

```
processors:
  filter:
    traces:
      span:
        - 'attributes["url.path"] == "/health"' # url.path, not the full URL
```

Masking Sensitive Data (redaction / transform)

Telemetry can carry PII — tokens in headers, emails in attributes, user IDs in log bodies. Strip it in **the collector, before egress**, so it never reaches the backend. Two processors

do this:

Processor	Use
redaction	Straightforward allow/block of attribute values; redacts matching values completely
transform	OTTL-based; full or partial redaction via pattern replacement, per signal

`redaction` — block values matching a pattern (e.g. a leaked Dynatrace token) and keep a summary of what was scrubbed:

```
processors:
  redaction:
    allow_all_keys: true
    blocked_values:
      - "dt0[a-z]0[1-9]\\.[A-Za-z0-9]{24}\\.[A-Za-z0-9]{64}" # Dynatrace token
    shape
    summary: info
```

`transform` — partial masking of specific attributes across traces, metrics, and logs (separate `trace_statements` / `metric_statements` / `log_statements` blocks):

```
processors:
  transform:
    trace_statements:
      - context: span
        statements:
          - replace_all_patterns(attributes, "value", "[\\w.]+@[\\w.]+",
            "***@***") # mask emails
```

Two rules that bite:

- **Order matters** — place redaction/transform **before** any routing or exporting so downstream steps only ever see sanitized data.
- **Span names are not attributes** in OTLP — masking attribute values does not touch the span name; handle it separately (e.g. a name-rewriting statement / `set_semconv_span_name`).

This is the collector-side complement to OpenPipeline ingest-time masking (OPLOGS / OPIPE): use collector redaction when data must be scrubbed **before it leaves your network**, OpenPipeline masking when it is scrubbed at ingest.

Sources: [Mask sensitive data with the OTel Collector \(DT docs\)](#) — `redaction`
(`blocked_values` / `allow_all_keys` / `summary`) and `transform`
(`replace_all_patterns` , per-signal statements); processor-order and span-name caveats.

5. Exporters

Exporters send processed telemetry to backends.

Common Exporters

Exporter	Target	Signals	Notes
<code>otlphttp</code>	OTLP over HTTP	All	Dynatrace, many backends
<code>otlp</code>	OTLP over gRPC	All	Lower overhead
<code>debug</code>	Console/logs	All	Development
<code>prometheus</code>	Prometheus endpoint	Metrics	Pull model
<code>file</code>	Local files	All	Debugging

OTLP HTTP Exporter (Dynatrace)

```
exporters:
  otlphttp:
    endpoint: https://{your-env}.live.dynatrace.com/api/v2/otlp
    headers:
      Authorization: Api-Token ${DT_API_TOKEN}
```

Debug Exporter

```
exporters:
  debug:
    verbosity: detailed
    sampling_initial: 5
    sampling_thereafter: 200
```

Multiple Exporters (Fan-Out)

```
exporters:
  otlphttp/dynatrace:
    endpoint: https://${DT_ENV_ID}.live.dynatrace.com/api/v2/otlp
    headers:
      Authorization: Api-Token ${DT_API_TOKEN}
  otlphttp/backup:
    endpoint: https://backup-backend.example.com

service:
  pipelines:
    traces:
      receivers: [otlp]
      processors: [batch]
      exporters: [otlphttp/dynatrace, otlphttp/backup] # Both!
```

6. Extensions

Extensions provide operational capabilities for the Collector itself.

Common Extensions

Extension	Purpose
health_check	Liveness/readiness endpoints
pprof	Go profiling data
zpages	Debug pages for pipelines

Note: The `memory_ballast` extension is **deprecated** and was removed from default Helm chart configurations. Use the `GOMEMLIMIT` environment variable instead (set to ~80% of the container memory limit). See [Collector discussion #9264](#).

Configuration

```
extensions:
  health_check:
    endpoint: 0.0.0.0:13133
  pprof:
    endpoint: 0.0.0.0:1777
  zpages:
    endpoint: 0.0.0.0:55679

service:
  extensions: [health_check, pprof, zpages]
```

7. Connectors

Connectors link pipelines, acting as both exporter and receiver.

Use Cases

Connector	Function
<code>span_metrics</code> (<code>spanmetrics</code> before contrib v0.151.0)	Generate metrics from spans
<code>servicegraph</code>	Build service dependency graph
<code>count</code>	Count data points

Span Metrics Connector

```
connectors:
  span_metrics:
    dimensions:
      - name: http.request.method
      - name: http.response.status_code
    histogram:
      explicit:
        buckets: [100ms, 500ms, 1s, 5s]

service:
  pipelines:
    traces:
      receivers: [otlp]
      exporters: [span_metrics, otlphttp]
    metrics:
      receivers: [span_metrics] # Receives from connector
      exporters: [otlphttp]
```

8. Configuration Basics

Minimal Configuration

```
receivers:
  otlp:
    protocols:
      grpc:
        endpoint: 0.0.0.0:4317
      http:
        endpoint: 0.0.0.0:4318

processors:
  batch:
    timeout: 10s

exporters:
  otlphttp:
    endpoint: https://${DT_ENV_ID}.live.dynatrace.com/api/v2/otlp
    headers:
      Authorization: Api-Token ${DT_API_TOKEN}

service:
  pipelines:
    traces:
      receivers: [otlp]
      processors: [batch]
      exporters: [otlphttp]
    metrics:
      receivers: [otlp]
      processors: [batch]
      exporters: [otlphttp]
    logs:
      receivers: [otlp]
      processors: [batch]
      exporters: [otlphttp]
```

Configuration Best Practices

Practice	Reason
Always use <code>batch</code> processor	Efficiency

Practice	Reason
Always use <code>memory_limiter</code>	Stability
Put <code>memory_limiter</code> first	Fail fast
Use environment variables for secrets	Security
Validate config before deploy	Catch errors early

Summary

In this notebook, you learned:

- Collector overview and distributions
- Pipeline architecture: receivers → processors → exporters
- Common receivers (OTLP, Prometheus, Jaeger)
- Essential processors (batch, `memory_limiter`, filter)
- Exporters for Dynatrace and other backends
- Extensions for operational capabilities
- Connectors for pipeline linking
- Configuration fundamentals

Next Steps

Next Notebook	Topic
OTEL-03: Collector Deployment	Deployment patterns
OTEL-04: Trace Instrumentation	Instrumenting code

References

- [OTel Collector Documentation](#)

- [Collector Configuration](#)
- [Collector Releases](#)

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.