

S2D-01: Splunk to Dynatrace Migration

- Getting Started

Series: S2D — Splunk to Dynatrace Migration | **Notebook:** 1 of 9 | **Created:** January 2026 | **Last Updated:** 06/23/2026

Overview

This notebook series provides comprehensive guidance for migrating monitoring capabilities from Splunk to Dynatrace. Whether you're moving dashboards, alerts, reports, or log queries, these notebooks will help you understand the key differences between platforms and make informed translation decisions.




Table of Contents

- [1. Key Platform Differences](#)

- 2. [Migration Planning](#)
- 3. [Series Roadmap](#)
- 4. [Quick Start: Verify Your First Log Source](#)

Sprint 1.337 (April 2026) Updates Affecting S2D

Three sprint-1.337 changes simplify the Splunk → Dynatrace migration:

1. **OneAgent primary fields/tags at the source** — hosts emit standardized fields (`dt.security_context` , `dt.cost.costcenter` , `dt.cost.product`) and customer-defined primary tags as top-level attributes. When converting Splunk indexer-level metadata routing to Grail, this lets you skip a class of OpenPipeline parse processors that would otherwise be needed for OneAgent-instrumented sources.
2. **OpenPipeline extraction processor recommended-field suggestions** — when converting Splunk field extractions to OpenPipeline DPL processors, the UI flags permission-relevant fields (`dt.security_context` , `loglevel`) and Smartscape identifiers — preventing accidental promotion of sensitive content.
3. **Configuration API → Settings v2 + Platform tokens** — new automation should target Settings v2 paths and use Platform tokens (`dt0s16` / `dt0s01` , `Authorization: Bearer`).

These changes reinforce the Gen3-first migration baseline.

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Grail enabled
Permissions	<code>logs.read</code> , <code>settings.read</code>
Splunk Access	Ability to view existing queries, alerts, dashboards
Knowledge	Basic familiarity with both Splunk SPL and Dynatrace DQL

Learning Objectives

By the end of this series, you will be able to:

- 1. Validate that required log data is available in Dynatrace
- 2. Translate SPL queries to DQL
- 3. Migrate Splunk alerts to Anomaly Detectors or Workflows
- 4. Convert Splunk dashboards to Dynatrace format
- 5. Apply consistent naming standards to migrated assets
- 6. Request metric extraction for performance-critical queries

Key Platform Differences

Understanding the fundamental differences between Splunk and Dynatrace is critical for a successful migration.

Data Model

Aspect	Splunk	Dynatrace
Storage	Indexes	Buckets (Grail)
Query Language	SPL (Search Processing Language)	DQL (Dynatrace Query Language)
Data Types	Events (logs)	Logs, Metrics, Traces, Events, Entities
Schema	Schema-on-read	Schema-on-read with semantic model
Retention	Index-based	Bucket-based with tiered storage

Alerting Model

Aspect	Splunk	Dynatrace
Execution	Scheduled queries	Continuous monitoring

Aspect	Splunk	Dynatrace
Frequency	User-defined schedule	Every minute
Evaluation	Single point-in-time	Sliding window
Threshold	Total over period	Per-minute samples
Intelligence	Rule-based	Dynatrace Intelligence powered

Migration Planning

A successful migration follows this general sequence:

Phase 1: Discovery

1. **Inventory existing assets** - Document all dashboards, alerts, reports, and saved searches
2. **Identify data sources** - Map log sources to hosts, clusters, or namespaces
3. **Prioritize by criticality** - Focus on production monitoring first

Phase 2: Data Validation

1. **Verify log ingestion** - Confirm logs are flowing into Dynatrace
2. **Check field mappings** - Ensure required fields are available
3. **Configure ingest rules** - Add any missing log sources

Phase 3: Query Translation

1. **Convert SPL to DQL** - Translate query syntax
2. **Validate results** - Compare output with Splunk
3. **Optimize performance** - Apply DQL best practices

Phase 4: Alert Migration

1. **Choose alert type** - Anomaly Detector or Workflow
2. **Translate thresholds** - Apply conversion formula
3. **Configure notifications** - Set up alerting profiles

Phase 5: Dashboard Migration

1. **Recreate visualizations** - Build equivalent charts
2. **Apply filters** - Configure time and segment filters
3. **Validate accuracy** - Compare with Splunk dashboards

Series Roadmap

Notebook	Topic	Description
S2D-01	Getting Started	This notebook - overview and planning
S2D-02	Locating Logs	Validating log data availability
S2D-03	SPL to DQL	Query translation fundamentals
S2D-04	Anomaly Detectors	Translating alerts to continuous monitoring
S2D-05	Workflow Alerts	When and how to use workflow-based alerting
S2D-06	ArrayMovingSum	Handling extended timeframes (>1 hour)
S2D-07	Metric Creation	Creating log metrics via OpenPipeline
S2D-08	Dashboard Migration	Converting Splunk dashboards
S2D-09	Naming Standards	Organization and naming conventions

Quick Start: Verify Your First Log Source

Before diving deeper, let's verify that logs are available in Dynatrace. This query shows log sources and their record counts.

```
// View available log sources and counts
fetch logs, from:-1h
| summarize count = count(), by:{log.source}
| sort count desc
| limit 20
```

Check logs by host name to match Splunk source mapping.

```
// View log counts by host
fetch logs, from:-1h
| summarize count = count(), by:{host.name}
| sort count desc
| limit 20
```

For Kubernetes workloads, check by cluster and namespace.

```
// View log counts by Kubernetes attributes
fetch logs, from:-1h
| filter isNotNull(k8s.cluster.name)
| summarize count = count(), by:{k8s.cluster.name, k8s.namespace.name}
| sort count desc
| limit 20
```

Next Steps

Continue to **S2D-02: Locating Logs in Dynatrace** to learn the complete process for validating log data availability before beginning your migration.

References

- [Dynatrace DQL Reference](#)
 - [anomaly-detection \(DT docs\)](#)
 - [Dynatrace Workflows](#)
 - [Log Management](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.