

S2D-07: Metric Creation from Logs

Series: S2D — Splunk to Dynatrace Migration | **Notebook:** 7 of 9 | **Created:** January 2026 | **Last Updated:** 10/05/2026

Overview

Log-based metrics can be generated by OpenPipeline during log ingest. This provides significant benefits for frequently-executed queries used in dashboards and alerts.

This notebook explains when to request metric extraction and how to provide the necessary information.

Table of Contents

1. [Why Create Metrics from Logs?](#)
 2. [When to Request Metric Extraction](#)
 3. [Metric Types](#)
 4. [Information Required for Metric Request](#)
 5. [Example Metric Request](#)
 6. [Best Practices](#)
 7. [Using Log-Based Metrics](#)
 8. [Request Process](#)
 9. [Metric Extraction Checklist](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with OpenPipeline
Permissions	<code>logs.read</code> , OpenPipeline access (typically platform team)
Knowledge	Completed SPL to DQL translation

Learning Objectives

By the end of this notebook, you will be able to:

1. Understand the benefits of metric extraction
2. Identify candidates for metric creation
3. Prepare the required information for a metric request
4. Follow best practices for metric design

Why Create Metrics from Logs?

Log-based metrics provide several advantages:

Benefit	Description
Decreased query cost	Metrics are pre-aggregated, reducing DQL processing
Improved query speed	Metrics are optimized for time-series queries
Increased retention	Metrics have longer default retention than logs
Reduced license usage	Metric queries are more efficient than log queries

When to Request Metric Extraction

Request metric creation after SPL to DQL translation is complete.

Good Candidates

Queries that:

- Generate timeseries results (`makeTimeseries` or `summarize`)
- Run frequently (dashboards, alerts)
- Have stable filter criteria
- Aggregate simple counts or numeric values

Poor Candidates

Queries that:

- Need full log content
- Use complex parsing logic
- Have frequently changing filters
- Are run infrequently (ad-hoc analysis)

Metric Types

Counter Metrics

Count the number of log records matching certain criteria.

Use when: You want to count occurrences (errors, requests, events)

```
// Counter metric candidate: Error log count
fetch logs, from:-24h
| filter status == "ERROR"
| filter matchesPhrase(k8s.deployment.name, "checkout-service")
| makeTimeseries count = count(), by:{k8s.deployment.name}, interval:1m
```

Value Metrics

Track a numeric value extracted from log records.

Use when: You want to track response times, sizes, durations, etc.

```
// Value metric candidate: Response time from logs
// response_time is parsed here at query time; for a metric it must exist at
// ingest (see Consider Custom Fields)
fetch logs, from:-24h
| parse content, "DATA? 'response_time=' DOUBLE:response_time"
| filter isNotNull(response_time)
| filter matchesPhrase(k8s.deployment.name, "api-service")
| makeTimeseries avg_response = avg(response_time), by:{k8s.deployment.name},
interval:1m
```

Information Required for Metric Request

1. Metric Type

Type	Description
Counter	Count of matching logs
Value	Numeric field from log records

For **Value** metrics, specify the field containing the value (e.g., `response_time`).

2. Filter Statement

The filter portion of your query.

Limitations:

- Must be a single continuous filter
- No custom parsing (`parse`) or field calculations (`fieldsAdd`)
- Only these functions are supported:
 - `matchesPhrase`
 - `matchesValue`
 - `isNull`
 - `isNotNull`
- Basic operators and equality

3. Dimensions (Split By)

The `by:{}` parameter from your `summarize` or `makeTimeseries` command.

This defines how the metric will be segmented.

Example Metric Request

Original DQL Query

```
// Query to convert to metric
fetch logs, from:-24h
| filter status == "ERROR"
| filter matchesPhrase(k8s.cluster.name, "production")
| filter matchesPhrase(k8s.namespace.name, "ecommerce")
| makeTimeseries error_count = count(), by:{k8s.deployment.name,
k8s.namespace.name}, interval:1m
```

Metric Request Details

Field	Value
Metric Type	Counter
Filter	status == "ERROR" AND matchesPhrase(k8s.cluster.name, "production") AND matchesPhrase(k8s.namespace.name, "ecommerce")
Dimensions	k8s.deployment.name , k8s.namespace.name
Metric Name	log.ecommerce.error_count

status == "ERROR" covers every error-or-worse level (SEVERE, ERROR, CRITICAL, ALERT, EMERGENCY). Confirm with your platform team that status is populated before the metric-extraction stage of your pipeline; if it is not, list the levels explicitly instead.

Best Practices

Keep Metrics General

Don't create multiple similar metrics for similar alerts. Consolidate by using dimensions.

Instead of:

- `checkout_errors`
- `payment_errors`
- `cart_errors`

Create:

- `service_errors` with dimension `k8s.deployment.name`

Avoid High-Cardinality Dimensions

Don't use unique values as dimensions:

- Transaction IDs
- Session IDs
- Request IDs
- Timestamps

These create too many metric series and impact performance.

Consider Custom Fields

If your metric requires custom fields not available at ingest time:

1. Request a log processing rule to extract the field first
2. Then request the metric extraction using that field

Using Log-Based Metrics

Once created, metrics are queried with `timeseries` instead of `fetch logs`. A counter metric is read with `sum()`; `avg()` returns the average per-data-point value, not the count.

```
// Query the extracted metric (after it's created)
// Note: This is an example - metric name will depend on your configuration
// Counter metric - use sum(); avg() returns the per-data-point value
timeseries sum(log.ecommerce.error_count), from:-1h, by:{k8s.deployment.name}
```

Request Process

Metric extraction is typically managed by a platform or DevOps team through OpenPipeline configuration.

Steps

1. Complete SPL to DQL translation
2. Identify queries suitable for metric extraction
3. Document metric type, filter, and dimensions
4. Submit request to platform team
5. Wait for Terraform deployment (timing varies)
6. Update dashboards/alerts to use new metric

Metric Extraction Checklist

Item	Status
DQL query validated	☐
Metric type identified (Counter/Value)	☐
Filter uses only supported functions	☐
Dimensions defined (no high-cardinality)	☐
Metric name follows naming convention	☐
Request submitted	☐

Next Steps

- **S2D-08: Dashboard Migration** - Converting Splunk dashboards
- **S2D-09: Naming Standards** - Organizational conventions

References

- [tutorial-log-processing-pipeline \(DT docs\)](#)
- [Log metrics \(DT docs\)](#)
- [DQL Matcher Functions](#)

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.