# S2S-09: Step 9 — Optimize: Cutover Validation and Decommission

> **Series:** S2S — SaaS to SaaS Migration | **Notebook:** 9 of 9 | **Phase:** Run | **Step:** Optimize | **Created:** March 2026 | **Last Updated:** 07/24/2026

## Overview

This is the final step of the SaaS-to-SaaS migration. With parallel operation validated (Step 8), Dynatrace Intelligence baselines established, and stakeholders trained, you are ready to execute the cutover. This notebook covers the go/no-go decision, cutover execution steps, post-cutover validation, rollback procedures, source tenant decommission, and a lessons learned framework.

> **S2S Migration Journey — 3 Phases / 9 Steps**
>
> **Plan:** 1. Discover | 2. Strategize | 3. Design
>
> **Upgrade:** 4. Prepare | 5. Execute | 6. Integrate
>
> **Run:** 7. Expand | 8. Enable | **9. Optimize**

### Order of Operations (Steps 9–11)

Within the 9-step framework, Step 9 maps to three operational actions:

| Operation | Action | Description |
|---|---|---|
| **9** | Validate | Final data flow and performance validation |
| **10** | Cutover | Full switch to target tenant |
| **11** | Decommission | Source tenant deactivation and deprovisioning |

---

# Table of Contents



---

# Prerequisites

| Requirement | Details |
|---|---|
| **Step 8 Complete** | Parallel operation validated, stakeholders trained, baselines established |
| **Go/No-Go Authority** | Decision-maker identified and available for cutover approval |
| **Maintenance Window** | Scheduled and communicated to all stakeholders |
| **Rollback Plan** | Documented and tested — source tenant still active and receiving data |
| **Communication Ready** | Cutover announcement drafted for all stakeholder audiences |

# 1. Go/No-Go Checklist

The go/no-go decision should be made in a structured meeting with all stakeholders present. Every item below must be **Go** for cutover to proceed. Any **No-Go** item requires documented remediation and a reschedule.

## Infrastructure Readiness

| Checkpoint | Status | Verified By |
|---|---|---|
| All hosts reporting to target tenant | [ ] | Platform team |
| All K8s clusters monitored via DynaKube in target | [ ] | Platform team |
| Cloud integrations (AWS/Azure/GCP) active in target | [ ] | Cloud team |
| ActiveGates routing to target tenant | [ ] | Platform team |

## Configuration Readiness

| Checkpoint | Status | Verified By |
|---|---|---|
| All Settings 2.0 configuration deployed | [ ] | Platform team |
| OpenPipeline rules validated (enrichment, routing, masking) | [ ] | Platform team |
| SLO definitions deployed and evaluating | [ ] | SRE team |
| Alerting profiles and notification rules active | [ ] | SRE team |
| Dashboards and notebooks accessible | [ ] | All teams |
| Workflows and automations tested | [ ] | Platform team |

## Security Readiness

| Checkpoint | Status | Verified By |
|---|---|---|
| IAM policies deployed (groups, bindings) | [ ] | Security team |
| SSO/SAML configured and tested | [ ] | Identity team |
| API tokens generated and distributed | [ ] | Platform team |
| Credential vault populated | [ ] | Security team |

### Data Readiness

| Checkpoint | Status | Verified By |
|---|---|---|
| Log volumes match between source and target (±10%) | [ ] | Platform team |
| Span volumes match between source and target (±10%) | [ ] | Platform team |
| Dynatrace Intelligence baselines stable (2+ weeks of data) | [ ] | SRE team |
| SLOs evaluating within expected range | [ ] | SRE team |

## 2. Cutover Execution Steps

Execute these steps in order during the scheduled maintenance window. Each step includes estimated duration and rollback instructions.

| Step | Action | Duration | Rollback |
|---|---|---|---|
| 1 | **Announce** — Send cutover start notification to all stakeholders | 5 min | N/A |
| 2 | **Suppress** — Enable maintenance window in source tenant | 5 min | Disable maintenance window |
| 3 | **Redirect** — Update DNS/proxy for API endpoints to target tenant | 15 min | Revert DNS/proxy |
| 4 | **Verify** — Confirm agents are reporting to target tenant only | 15 min | Re-enable dual reporting |
| 5 | **Enable** — Activate alerting in target tenant (remove staging routing) | 10 min | Re-route to staging |
| 6 | **Validate** — Run post-cutover validation queries (Section 3) | 30 min | If validation fails, rollback |
| 7 | **Disable** — Stop data collection in source tenant (set to read-only) | 10 min | Re-enable collection |
| 8 | **Communicate** — Send cutover complete notification | 5 min | N/A |

| Step | Action | Duration | Rollback |
|---|---|---|---|
| 9 | **Monitor** — Active monitoring for 4–24 hours post-cutover | Ongoing | Full rollback if critical issues |

> **Total estimated cutover window: 2–4 hours** (including monitoring buffer)

### Cutover Timing Recommendations

| Timing | Pros | Cons |
|---|---|---|
| **Weekend morning (low traffic)** | Minimal user impact, low alert volume | Less staff available |
| **Mid-week morning** | Full team available | Higher traffic, more visible |
| **End of business Friday** | Weekend to stabilize | Fatigue risk, limited weekend support |

## 3. Post-Cutover Validation Queries

Run these queries in the **target tenant** immediately after cutover to confirm that all data sources are flowing correctly.

## Host Count Validation

```
// Post-cutover: Total host count by cloud provider
fetch dt.entity.host
| fieldsAdd provider = if(isNotNull(awsNameTag), then: "AWS",
    else: if(isNotNull(azureResourceGroupName), then: "Azure", else: "On-Premises"))
| summarize count = count(), by:{provider}
| sort count desc

// Smartscape note (dt.entity.* is deprecated but still functional): the classic cloud-tag
// fields (awsNameTag / azureResourceGroupName / gcpProjectId) are not Smartscape node fields.
// On Smartscape, smartscapeNodes "HOST" exposes cloud.provider directly — e.g.
//   smartscapeNodes "HOST" | summarize count = count(), by:{cloud.provider}
// (aws / azure / gcp; null = on-premises) — which replaces the tag-presence if-chain.
// Keep the classic query above; the live-topology count caveat also applies.
```

## Service Count Validation

```
// Post-cutover: Service inventory by technology
fetch dt.entity.service
| summarize count = count(), by:{serviceType}
| sort count desc

// Smartscape equivalent (dt.entity.* is deprecated but still functional):
//   smartscapeNodes "SERVICE"
//   | summarize count = count(), by:{dt.service.sdv1_type}
//   | sort count desc
// Caveat: Smartscape reflects CURRENT live topology and can report fewer entities
// than the classic entity store; for a pre-migration discovery inventory keep the
// classic query above.
// Field maps: serviceType -&gt; dt.service.sdv1_type.
```

## Log Flow Validation

```
// Post-cutover: Log volume by loglevel (last 1 hour)
// Compare with pre-cutover baseline from source tenant
fetch logs, from:-1h
| summarize record_count = count(), by:{loglevel}
| sort record_count desc
```

## Span Flow Validation

```
// Post-cutover: Span volume by span.kind (last 1 hour)
fetch spans, from:-1h
| summarize span_count = count(), by:{span.kind}
| sort span_count desc
```

## Enrichment Validation

```
// Post-cutover: Verify OpenPipeline enrichment is active
fetch logs, from:-1h
| fieldsAdd has_security_context = isNotNull(dt.security_context)
| summarize
    total = count(),
    enriched = countIf(has_security_context == true)
| fieldsAdd enrichment_pct = round(toDouble(enriched) / toDouble(total) *
100.0, decimals: 1)
| fieldsAdd status = if(enrichment_pct &gt; 50.0, then: "Enrichment healthy",
else: "Check OpenPipeline config")
```

## Detected Problem Validation

```
// Post-cutover: Active detected problems (should be normal operational level)
fetch dt.davis.problems, from:-24h
| filter event.status == "ACTIVE"
| summarize active_problems = count()
| fieldsAdd assessment = if(active_problems &lt; 10, then: "Normal range",
    else: if(active_problems &lt; 50, then: "Elevated — review problems",
else: "High — investigate immediately"))
```

### MTTR Baseline

```
// Post-cutover: Mean time to resolve (MTTR) for closed problems
// This establishes the MTTR baseline in the target tenant
fetch dt.davis.problems, from:-7d
| filter event.status == "CLOSED"
| filter dt.davis.is_frequent_event == false and dt.davis.is_duplicate ==
false
| fieldsAdd duration_hours = resolved_problem_duration / 1h
| summarize
    avg_mttr_hours = avg(duration_hours),
    median_mttr_hours = median(duration_hours),
    problems_resolved = count()
| fieldsAdd avg_mttr_hours = round(avg_mttr_hours, decimals: 2),
    median_mttr_hours = round(median_mttr_hours, decimals: 2)
```

## 4. Rollback Procedures

If critical issues are discovered during or after cutover, follow these rollback procedures by component.

### Rollback Decision Criteria

| Severity | Criteria | Action |
|---|---|---|
| **Critical** | >25% of hosts not reporting, or no log/span flow | Full rollback immediately |
| **High** | Key SLOs evaluating at 0%, or critical alerts not firing | Partial rollback, keep agents dual-reporting |
| **Medium** | Minor data gaps, non-critical dashboards missing data | Continue with remediation plan |
| **Low** | Cosmetic issues, minor configuration differences | Fix forward |

### Rollback by Component

| Component | Rollback Action | Time to Rollback |
| --- | --- | --- |
| **OneAgent** | Update `host_group` or connection endpoint back to source tenant | 5–15 min per host (automated via deployment tool) |
| **DynaKube** | Update DynaKube CR `apiUrl` to source tenant | 5 min + pod restart |
| **Cloud integrations** | Re-enable in source tenant, disable in target | 15 min |
| **Alerting** | Disable target alerting, re-enable source | 10 min |
| **SLOs** | No rollback needed — source SLOs still evaluating during parallel | N/A |
| **DNS/proxy** | Revert API endpoint routing to source tenant | 15 min |

## 5. Source Tenant Decommission

After a successful cutover and monitoring buffer period, the source tenant can be decommissioned. This is a gradual process — do not rush it.

### Decommission Timeline

| Week | Action | Notes |
| --- | --- | --- |
| **Week 0** | Cutover complete — source set to read-only | No new data ingested |
| **Week 1–2** | Active monitoring buffer — source available for historical queries | Keep source accessible |
| **Week 3–4** | Export any remaining artifacts (problem reports, session replay screenshots) | Document what was exported |
| **Week 5–6** | Disable source tenant monitoring agents (remove OneAgents, delete DynaKube CRs) | Confirm no data flowing to source |
| **Week 7** | Revoke all API tokens and OAuth clients in source tenant | Security cleanup |

| Week | Action | Notes |
|---|---|---|
| **Week 8** | Request tenant deprovisioning from Dynatrace account team | Final step |

### Pre-Decommission Checklist

| Item | Action | Status |
|---|---|---|
| Historical problem reports exported | Export as PDF or CSV for compliance records | [ ] |
| Key dashboards screenshotted | Capture for reference | [ ] |
| Audit log exported | Required for compliance in some industries | [ ] |
| All users notified | Source tenant access will be revoked | [ ] |
| API tokens revoked | Remove all active tokens | [ ] |
| OAuth clients deleted | Remove all client credentials | [ ] |
| SSO/SAML disconnected | Remove identity provider integration | [ ] |
| Account team notified | Request formal deprovisioning | [ ] |

## 6. Documentation and Handover

After migration, update all documentation to reflect the new environment.

### Documents to Update

| Document | Updates Required |
|---|---|
| **Architecture diagrams** | New tenant URLs, updated network flows |
| **Runbooks** | API endpoints, token references, escalation paths |
| **Disaster recovery plan** | New tenant in DR procedures |
| **Change management records** | Migration completion, lessons learned |
| **Vendor management** | Updated contract references, licensing terms |

| Document | Updates Required |
| --- | --- |
| **Security documentation** | New IAM architecture, token inventory |
| **Onboarding guides** | New tenant URLs, access request procedures |

# 7. Lessons Learned Framework

Capture lessons learned while the migration is still fresh. Organize by category for maximum reusability.

## Capture Template

| Category | What Went Well | What Could Improve | Action Items |
| --- | --- | --- | --- |
| **Planning** | ___ | ___ | ___ |
| **Configuration** | ___ | ___ | ___ |
| **Agent Migration** | ___ | ___ | ___ |
| **Data Validation** | ___ | ___ | ___ |
| **Stakeholder Communication** | ___ | ___ | ___ |
| **Cutover Execution** | ___ | ___ | ___ |
| **Tooling (Monaco/Terraform)** | ___ | ___ | ___ |
| **Dynatrace Intelligence / Baselines** | ___ | ___ | ___ |

## Common Lessons from S2S Migrations

| Lesson | Detail |
| --- | --- |
| Entity ID remapping takes longer than expected | Audit dashboards and SLOs **before** migration, not during |
| baselines need at least 2 full weeks | Plan parallel operation accordingly |
| Credential recreation is a bottleneck | Involve cloud and security teams early |

| Lesson | Detail |
| --- | --- |
| Stakeholder fatigue is real | Keep communication concise and predictable |
| Extensions 2.0 must be reinstalled from Hub | Neither Monaco nor Terraform can export them |
| Webhook URLs are easily missed | Audit all notification integrations systematically |

## 8. Series Conclusion

This completes the 9-step SaaS-to-SaaS migration framework.

### Journey Summary

| Phase | Step | Notebook | Focus |
| --- | --- | --- | --- |
| **Plan** | 1 | S2S-01 | Discover — Inventory source tenant, identify migration scenario |
| **Plan** | 2 | S2S-02 | Strategize — Select approach, timeline, risk assessment |
| **Plan** | 3 | S2S-03 | Design — Target tenant architecture (IAM, buckets, pipelines) |
| **Upgrade** | 4 | S2S-04 | Prepare — Build target tenant, deploy configuration |
| **Upgrade** | 5 | S2S-05 | Execute — Redirect agents, migrate OneAgent and DynaKube |
| **Upgrade** | 6 | S2S-06 | Integrate — Reconnect cloud integrations and extensions |
| **Run** | 7 | S2S-07 | Expand — OpenPipeline, SLOs, alerting migration |
| **Run** | 8 | S2S-08 | Enable — Parallel operation, stakeholder handover |
| **Run** | 9 | S2S-09 | Optimize — Cutover, validation, decommission |

### Key Principles Reinforced Throughout

| Principle | Why It Matters |
| --- | --- |
| **The 90/10 Rule** | 90% of config migrates automatically; 10% requires 90% of the effort |
| **Tags over entity IDs** | Tag-based selectors survive migration; hardcoded entity IDs do not |
| **Monaco for config, Terraform for IAM** | Use each tool where it excels |
| **Parallel operation is mandatory** | Historical data does not migrate — plan for dual-tenant operation |
| **Dynatrace Intelligence needs time** | Baselines require 2–4 weeks; communicate this to stakeholders |
| **Communicate early and often** | Migration success is as much about people as technology |

## 9. Step Completion Checklist

| Checkpoint | Status |
| --- | --- |
| Go/no-go meeting completed — all four categories green | [ ] |
| Cutover executed during scheduled maintenance window | [ ] |
| Post-cutover validation queries all pass (hosts, services, logs, spans, enrichment) | [ ] |
| detected problem volume within normal range | [ ] |
| MTTR baseline established in target tenant | [ ] |
| Rollback plan confirmed not needed (or rollback executed and re-cutover scheduled) | [ ] |
| Source tenant set to read-only | [ ] |
| Source tenant decommission timeline agreed | [ ] |
| All documentation updated (runbooks, architecture, DR, security) | [ ] |
| Lessons learned captured and shared | [ ] |

| Checkpoint | Status |
| --- | --- |
| Migration formally closed with stakeholder sign-off | [ ] |

---

## Summary

In Step 9 — the final step — you:

- Conducted a structured go/no-go assessment across infrastructure, configuration, security, and data readiness
- Executed the cutover following a 9-step sequence with defined rollback procedures for each step
- Validated the target tenant with DQL queries confirming host counts, service counts, log flow, span flow, and enrichment
- Established MTTR and detected problem baselines in the target tenant
- Planned source tenant decommission on an 8-week timeline
- Captured lessons learned for future migrations

**The SaaS-to-SaaS migration is complete.** The target tenant is now the production monitoring environment.

### Additional Resources

- Dynatrace Grail Data Lakehouse
- Dynatrace Intelligence Anomaly Detection
- Monaco Configuration as Code
- Dynatrace Terraform Provider
- DQL Reference

---