

S2S-94 LAB: Retiring AWS for Azure

— Moving the Environment and the Workloads

Series: S2S — SaaS to SaaS Migration | **Reference:** 94 — AWS-to-Azure Cloud Retirement LAB | **Created:** September 2026 | **Last Updated:** 10/06/2026

Overview

An appendix LAB for an organization that is **completely retiring AWS in favor of Azure**. Two moves happen at once, and they are easy to conflate:

1. **The Dynatrace environment moves** from an AWS-hosted SaaS cluster to an Azure-hosted one. That is a tenant move: a new target environment, and everything in S2S-01 to S2S-09 applies.
2. **The monitored workloads move** from AWS to Azure. That is a monitoring change: new cloud connections, new log ingest paths, agents on new hosts.

Because workloads move in waves, there is a **dual-cloud window**: the new Azure-hosted environment monitors workloads still running on AWS *and* workloads already on Azure. The window ends when the last AWS workload is gone, and only then are the AWS connection and the AWS-hosted source environment retired.

This LAB is the ordered runbook for that case. It does not repeat the nine steps; it sequences them for this scenario and adds what is specific to it. Two FAQ entries frame it: **FAQ-25** (what carries over to a new tenant) and **FAQ-17** (the cutover invariants — Go/No-Go gate, parallel-run window, rollback triggers, decommission).

Table of Contents

1. [The Scenario in One Table](#)

2. [Phase 0 — Decide and Provision the Azure-Hosted Target](#)
3. [Phase 1 — Identity, Access, and Configuration](#)
4. [Phase 2 — Connect Both Clouds to the Target](#)
5. [Phase 3 — Agent Waves](#)
6. [Phase 4 — Log and Event Ingest Cut-over](#)
7. [Phase 5 — Validate the Dual-Cloud Window](#)
8. [Phase 6 — Decommission AWS and the Source Environment](#)
9. [Open Questions and Unverified Items](#)
10. [Summary](#)

Prerequisites

Requirement	Details
S2S series	S2S-01 to S2S-05 read; S2S-06 §1–2 for cloud connections
Source environment	The AWS-hosted Dynatrace SaaS environment, admin access
Commercial agreement	An Azure-hosted target agreed with your Dynatrace account team (Azure regions are available on request)
AWS access	Rights to deploy CloudFormation stacks in every monitored AWS account
Azure access	Rights to create a service principal and assign it on each monitored subscription
Tooling	Monaco v2.x, Terraform (IAM), <code>oneagentctl</code> on hosts, <code>kubectl</code> for DynaKube

1. The Scenario in One Table

	Before	Dual-cloud window	After
Dynatrace environment	AWS-hosted source	Azure-hosted target is primary; source readable, no longer ingesting from migrated waves	Azure-hosted target only
AWS workloads	Monitored by the source	Re-pointed to the target, wave by wave, while they still run	Gone
Azure workloads	— (or monitored by the source)	Monitored by the target from day one	Monitored by the target
AWS connection	On the source	On the target (and removed from the source)	Removed
Azure connection	— (or classic integration on the source)	On the target, one subscription at a time	On the target

Sequencing recommendation. Stand up the Azure-hosted target **before** the workload migration starts, so that new Azure workloads install OneAgent against the target directly. A workload that moves to Azure while still reporting to the AWS-hosted source has to be re-pointed a second time later. This follows from the one-destination rule for OneAgent (FAQ-25 §6) rather than from a Dynatrace document, so treat it as planning guidance.

2. Phase 0 — Decide and Provision the Azure-Hosted Target

Dynatrace documents that *"Data is stored in Amazon Web Services (AWS), Microsoft Azure, or Google Cloud data centers"*, and lists Azure regions (East US, West US 3, West Europe, Canada Central, UAE, Switzerland North, Australia East) as *"Available on request. Talk to your Dynatrace sales contact."* It also states *"Environments hosted on Azure use dedicated Azure storage accounts"* — useful for the data-residency conversation.

Choose between the two ways of getting an Azure-hosted environment (the full comparison is in **S2S-04** §1):

Decision	Standard SaaS on an Azure region	Azure Native Dynatrace Service
Provisioned by	Your Dynatrace account team	You, from the Azure portal — <i>"available via a private offer"</i>
Account	Settled with the account team	<i>"The integration will create a new Dynatrace environment and account; it can't run on an existing Dynatrace SaaS environment."</i>
Billing	Your Dynatrace agreement	<i>"your Dynatrace license consumption becomes a part of your regular Azure bill"</i>
SSO	SAML — Microsoft Entra ID is a documented IdP	Entra ID; <i>"The integration works across a single Entra ID environment."</i>
Azure Cloud Platform Monitoring on it	Documented	Not stated — see §9

Runbook:

#	Action	Done when
0.1	Agree region and path with the account team and your Azure commercial owner	Written decision, including how the source commitment and target consumption overlap
0.2	Provision the target	Environment reachable; admin can sign in
0.3	Record both environments' platform state (Classic vs Latest surfaces, versions)	Recorded in the migration plan (S2S-01 §8)
0.4	Freeze the source configuration once the export is taken	Freeze communicated (S2S-04 §6)

Sources:

- [Data security controls \(DT docs\)](#) — *"Environments hosted on Azure use dedicated Azure storage accounts."*
- [Azure Native Dynatrace Service \(DT docs\)](#) — *"When you purchase Dynatrace SaaS through the Azure Marketplace, your Dynatrace license consumption becomes a part of your regular Azure bill."*

3. Phase 1 — Identity, Access, and Configuration

#	Action	Reference	Done when
1.1	Configure SAML SSO with Microsoft Entra ID as a new enterprise application for the target	S2S-04 §2	Two users in different groups sign in with the expected permissions
1.2	Keep group claims under Entra ID's limit — <i>"The number of groups emitted in a token is limited to 150 for SAML assertions"</i>	S2S-99 #43	Group claims arrive for your largest user
1.3	Rebuild the IP allowlist on the target (per environment, inbound)	S2S-04 §2	CI/CD runners and admins reach the target API
1.4	Deploy IAM with Terraform	S2S-04 §2, S2S-05 §3	Groups, policies, bindings applied
1.5	Export with <code>monaco download</code> , deploy with <code>monaco deploy</code> — the SaaS Upgrade Assistant is documented for a Managed source only, so import through it only after rehearsing that field practice	S2S-10	Dry-run clean; deploy to the target succeeds
1.6	Deploy ActiveGates the target needs (routing, extensions, private synthetic locations) — on Azure where the workloads are going	S2S-04 §4	<code>smartscapeNodes</code> "ACTIVEGATE" lists them in the target

Unverified — egress addresses. If your firewalls, webhook receivers or ITSM tools admit Dynatrace by source IP, the Azure-hosted cluster's egress addresses will differ from the AWS-hosted source's. No primary source for them was found for this LAB; request them from Dynatrace before cutover rather than reusing the old list.

Sources: [Group claims \(Microsoft Learn\)](#) — *"The number of groups emitted in a token is limited to 150 for SAML assertions and 200 for JWT, including nested groups."*, [IP allowlist \(DT docs\)](#).

4. Phase 2 — Connect Both Clouds to the Target

The target needs **both** connections during the window: AWS for the workloads that have not moved yet, Azure for the ones that have. Both are Latest Dynatrace Cloud Platform Monitoring connections, and neither needs an ActiveGate for polling (**S2S-06** §1).

Azure — switch per subscription, never overlap. The docs: *"Do not onboard Azure subscriptions already monitored by the classic Azure integration, and avoid monitoring the same subscription across multiple Azure connections—both increase the risk of API throttling and service interruptions."* If the source already monitors some Azure subscriptions, move each one in a single change window:

#	Action	Done when
2.1	Create a dedicated service principal for the target — <i>"Do not share it across Dynatrace environments"</i>	Principal exists, assigned on the subscription
2.2	Remove the subscription from the source's integration (classic or Latest)	Source no longer polls it
2.3	Add the subscription to the target's Azure connection	Its resources appear in the target (query in §7)

AWS — a new connection from the target. *"All AWS connection creation methods are powered by CloudFormation"*: create the connection in the target and deploy its stack into each account. The AWS onboarding page read for this LAB carries no double-monitoring warning; keep any period in which both source and target poll the same account short, and confirm cost and API-quota implications with your account team.

#	Action	Done when
2.4	Create the AWS connection in the target; deploy its CloudFormation stack per account	AWS resources appear in the target
2.5	Remove the AWS integration from the source once the target shows the same accounts	Source stops polling AWS

Sources:

~ [Create your first Azure connection \(DT docs\)](#) — *"Use a dedicated service principal exclusively for Dynatrace. Do not share it across Dynatrace environments or use it for other non-Dynatrace workloads."*

– [AWS Cloud Platform Monitoring \(DT docs\)](#) — "All AWS connection creation methods are powered by CloudFormation as Infrastructure-as-Code (IaC) engine."

5. Phase 3 — Agent Waves

Two kinds of wave run side by side:

Wave type	What happens	Mechanism
Re-point — AWS workloads still running	The host stays on AWS; its OneAgent moves to the target	<code>oneagentctl --set-server=... --set-tenant=... --set-tenant-token=... --restart-service</code> , then restart monitored applications (S2S-05 §5)
Re-platform — workloads rebuilt on Azure	New hosts on Azure; OneAgent installed from the target	Standard OneAgent installation against the target — no re-point
Kubernetes — EKS clusters still running / new AKS clusters	EKS: delete the DynaKube and apply one for the target. AKS: apply a DynaKube for the target	"starting with Dynatrace Operator version 1.3.0, editing <code>spec.apiUrl</code> is not allowed" (S2S-05 §6)

A host reports to one environment at a time, so the source/target overlap is **per wave** (**FAQ-25 §6**). After each wave, check where hosts are — by cloud and region — in the target:

```
// Target: hosts by cloud provider and region – run after every wave
// cloud.provider is aws / azure / gcp on cloud-hosted hosts; the region field
differs per cloud
// (aws.region vs azure.location), so coalesce them into one column.
smartscapeNodes "HOST", from:-2h
| fieldsAdd provider = coalesce(cloud.provider, "none (on-premises or
undetected)"),
    region = coalesce(aws.region, azure.location)
| summarize hosts = count(), by:{provider, region}
| sort provider asc, hosts desc
```

Over the window, the `aws` rows should shrink and the `azure` rows grow; the sum across the source and the target should match your discovery baseline (allowing for hosts that

were retired rather than migrated).

6. Phase 4 — Log and Event Ingest Cut-over

Source of data	On the source	On the target	Cut-over rule
OneAgent log collection	Follows the agent	Follows the agent	Moves with each wave automatically
AWS CloudWatch logs	Source's forwarder or Firehose streams	Firehose streams from the target's AWS connection	Switch per log group with its workload's wave; retire at AWS decommission
Azure resource / activity logs	Classic log forwarder, if any	Event Hubs ingest on the target	Switch per subscription with 2.2–2.3
Azure resource events	—	Event Grid system topics	Enable per subscription
Classic Azure log forwarder (if kept as fallback)	Direct via Cluster API, or through an ActiveGate	Redeployed against the target	<i>"Logs older than 24 hours are rejected"</i> — switch within the change window; a forwarder reointed after a longer gap cannot backfill

OpenPipeline rules and bucket routing that match on AWS log sources need Azure equivalents before the Azure ingest starts, or those records land in the default buckets (S2S-07).

Sources:

- [AWS Cloud Platform Monitoring \(DT docs\)](#) — *"Subscribe CloudWatch log groups to auto-generated Firehose streams for immediate ingestion and analysis."*
- [Azure Cloud Platform Monitoring \(DT docs\)](#) — *"Subscribe Azure resources to Event Grid system topics and forward resource lifecycle events such as blob creation and deletion, resource group changes, and service health alerts to Dynatrace."*
- [Set up the Azure log forwarder \(DT docs, Dynatrace Classic\)](#) — *"Logs older than 24 hours are rejected (considered too old by the Dynatrace log ingest endpoint)."*

7. Phase 5 — Validate the Dual-Cloud Window

Run these in the **target** throughout the window. First, connection coverage — cloud resources per provider:

```
// Target: cloud resources per provider, from the AWS and Azure connections
smartscapeNodes "*", from:-2h
| filter startsWith(type, "AWS_") or startsWith(type, "AZURE_")
| fieldsAdd provider = if(startsWith(type, "AWS_"), then: "aws", else:
"azure")
| summarize {resources = count(), resource_types = countDistinct(type)}, by:
{provider}
```

A provider with no row means its connection is missing or not yet polling. Next, compare what the connections see with what OneAgent sees — virtual machines from each connection against OneAgent hosts on each cloud:

```
// Target: VM coverage – connection-discovered instances vs OneAgent-monitored
hosts
smartscapeNodes "AWS_EC2_INSTANCE", from:-2h
| summarize instances = count()
| fieldsAdd resource = "AWS EC2 instances (AWS connection)"
| append [smartscapeNodes "AZURE_MICROSOFT_COMPUTE_VIRTUALMACHINES", from:-2h
    | summarize instances = count()
    | fieldsAdd resource = "Azure virtual machines (Azure connection)"]
| append [smartscapeNodes "HOST", from:-2h
    | filter in(cloud.provider, {"aws", "azure"})
    | summarize instances = count(), by:{cloud.provider}
    | fieldsAdd resource = concat("OneAgent hosts on ", cloud.provider)
    | fieldsRemove cloud.provider]
```

The two numbers per cloud are not expected to match — a connection sees every instance in its scope, OneAgent only the hosts you installed it on. What matters is the trend: on Azure, OneAgent hosts should grow with each re-platform wave; on AWS, both should fall toward zero.

Finally, the **baseline clock** from **FAQ-25 §5**: how much history each host has in the target. Migrated and newly built hosts start at zero, and their detectors are only as good as that history:

```
// Target: days of history per host (the FAQ-25 baseline clock)
// interval:24h rather than 1d – calendar durations are not supported here.
timeseries cpu = avg(dt.host.cpu.usage), from:-30d, interval:24h, by:
{dt.entity.host}
| fieldsAdd days_in_window = arraySize(cpu)
| fieldsAdd days_with_data = arraySize(arrayRemoveNulls(cpu))
| fields dt.entity.host, days_in_window, days_with_data
| sort days_with_data asc
| limit 20
```

Route the noise from low-history hosts to a staging channel rather than disabling alerting (FAQ-25 §5). Log ingest per cloud is a quick cross-check that Phase 4 is complete:

```
// Target: log records per cloud provider (last hour)
// cloud.provider is populated only on records that carry cloud context; a
large null row is
// normal (OneAgent-collected logs from on-premises hosts, API ingest, and
similar).
fetch logs, from:-1h
| summarize records = count(), by:{cloud.provider}
| sort records desc
```

8. Phase 6 — Decommission AWS and the Source Environment

AWS goes last, and in this order. First confirm, in the target, that no AWS host is still reporting — widening the window so a host that stopped recently is still listed:

```
// Target: AWS hosts seen in the last 7 days, split by whether they still
report
// still_reporting = false means the host was seen in the window but not in
the last 2 hours.
smartscapeNodes "HOST", from:-7d
| filter cloud.provider == "aws"
| fieldsAdd still_reporting = lifetime[end] > now() - 2h
| summarize hosts = count(), by:{still_reporting, aws.region}
| sort hosts desc
```

Without `from:-7d`, only nodes seen in the default window are returned, and a host that went quiet a day ago would not appear at all. The decommission gate is **no**

`still_reporting = true` rows — and a zero here is only meaningful if the same query returned AWS hosts earlier in the window (never trust a zero on its own).

#	Action	Done when
6.1	Last AWS workload retired or moved	Query above shows no <code>still_reporting = true</code> rows
6.2	CloudWatch log groups unsubscribed; Firehose streams idle	No AWS rows in the log query (§7)
6.3	Remove the AWS connection from the target, and remove what its CloudFormation stacks created in each account (deleting a stack removes the resources it created — confirm against your AWS change process)	No AWS resources in the connection-coverage query
6.4	Stop ingest into the source: revoke its ingest tokens, disable its remaining integrations; its data stays readable for its bucket retention	S2S-09 §5
6.5	Export what must outlive the source (audit logs, key problem reports)	S2S-09 §5 pre-decommission checklist
6.6	Request deprovisioning of the AWS-hosted source environment from the account team	Formal confirmation received
6.7	Remove the source's SAML application from Entra ID, revoke its tokens and OAuth clients	S2S-09 §5

FAQ-17 covers the invariants that apply to this final step as to any cutover: the rollback triggers are decided before the decommission starts, and the decommission is not reversible.

9. Open Questions and Unverified Items

Item	Status	What to do
Azure Cloud Platform Monitoring / Clouds app on an Azure-Native-created environment	The Azure Native Dynatrace Service page (labelled Dynatrace Classic) does not mention either; read 09/28/2026	Confirm with Dynatrace before choosing the Azure Native path

Item	Status	What to do
Egress IP addresses of the Azure-hosted cluster	No primary source found	Request them from Dynatrace; update firewall and webhook allowlists before cutover
Parallel AWS polling from two environments	No documented limit or warning found on the AWS onboarding page (the Azure page does warn)	Keep it short; confirm cost and API-quota implications with the account team
Commercial overlap (source commitment vs target consumption; Azure Marketplace billing for Azure Native)	Not a documentation question	Settle with the account team in Phase 0

10. Summary

Phase	Outcome
0	Azure-hosted target agreed and provisioned; path (standard vs Azure Native) chosen with the open question settled
1	Entra ID SSO, IP allowlist, IAM and configuration on the target; egress addresses requested
2	Target holds both an AWS and an Azure connection; no Azure subscription on two connections
3	AWS hosts re-pointed and Azure hosts installed against the target, wave by wave
4	Log and event ingest follows the waves; classic forwarders switched inside the 24-hour window
5	Dual-cloud coverage and per-host history measured in the target
6	No AWS host reporting; AWS connection removed; source environment deprovisioned

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.