

S2S-99: Best Practice Summary

Series: S2S | Notebook: 99 | Created: March 2026 | Last Updated: 03/26/2026

Overview

This notebook consolidates every actionable best practice from the S2S series (notebooks 01-12) into a single reference. Each practice is definitive: it tells you exactly what to set, when, and why. Use this as a pre-flight checklist and ongoing reference throughout your SaaS-to-SaaS migration.

Table of Contents

- 1. [Planning and Readiness](#)
- 2. [Configuration Export and Tooling](#)
- 3. [IAM, SSO, and Identity](#)
- 4. [Agent and Operator Migration](#)
- 5. [Configuration Import and Deployment Order](#)
- 6. [Cloud Integration](#)
- 7. [Dashboards, Workflows, and Integrations](#)
- 8. [OpenPipeline and Grail Buckets](#)
- 9. [Parallel Operation and Data Continuity](#)
- 10. [SLOs and Alerting](#)
- 11. [Cutover Execution](#)
- 12. [Post-Migration and Decommission](#)

Prerequisites

Requirement	Details
S2S Series	Read S2S-01 through S2S-12 for full context
Source Tenant	Active Dynatrace SaaS with admin access
Target Tenant	Provisioned Dynatrace SaaS with admin access
CLI Tools	Monaco CLI v2.x, Terraform v1.5+ with provider ~> 1.91

1. Planning and Readiness

Source: S2S-01

#	Best Practice	Setting / Value	Priority	Category
1	Use phased-by-environment migration strategy	Migrate dev first, then staging, then production over 6-12 weeks	Critical	Strategy
2	Expect the 90/10 rule	90% of config migrates automatically; budget 90% of effort for the remaining 10% (entity ID remapping, integration repointing, IAM redesign)	Critical	Planning
3	Run readiness assessment before starting	Complete every item on the Pre-Migration Checklist: licensing, network, auth, IAM, SSO, cloud creds, naming convention, parallel period, rollback plan, stakeholder alignment	Critical	Planning
4	Inventory source tenant with DQL before export	Run entity count queries (<code>fetch dt.entity.host \ summarize count() , same for service, process_group</code>) to establish baseline numbers	Critical	Discovery
5	Prefix resources when consolidating multiple tenants	Use source tenant identifier as prefix (e.g., <code>us-east - Critical Alerts</code>) on all dashboards, rules, and SLOs to prevent naming collisions	Critical	Naming
6	Tag everything by business unit before splitting a tenant	Apply auto-tags identifying which business unit owns each resource before splitting one tenant into many	Recommended	Naming
7	Build a comprehensive endpoint inventory for regional relocation	Document every agent URL, API call, webhook, and integration endpoint that references the source tenant URL	Critical	Discovery
8	Align migration timing with stable traffic periods	Avoid migration during holidays, sales events, or other atypical traffic patterns	Recommended	Scheduling

2. Configuration Export and Tooling

Source: S2S-02

#	Best Practice	Setting / Value	Priority	Category
9	Use <code>monaco download</code> for bulk export	Run <code>monaco download manifest.yaml --environment source-tenant</code> to export all configuration in one operation	Critical	Tooling
10	Store exported configuration in Git	Commit the Monaco export directory to a Git repository immediately after download	Critical	Version Control
11	Use Monaco for everything except IAM	Monaco v2 covers all 8 config types: settings, document, automation, bucket, segment, slo-v2, openpipeline, and classic api	Critical	Tooling
12	Use Terraform exclusively for IAM policies, groups, and bindings	IAM is the only resource type that requires Terraform; everything else works with Monaco	Critical	Tooling
13	Validate export contains no secrets	Run <code>grep -r "dt0c01" projects/</code> after export and confirm 0 matches	Critical	Security
14	Verify export completeness by comparing schema counts	Compare <code>ls projects/full-export/settings/ wc -l</code> against Settings API <code>totalCount</code> per schema	Recommended	Validation
15	Run <code>monaco validate manifest.yaml</code> before any deploy	Catches JSON validity errors, missing references, and dependency issues before they reach the target	Critical	Validation
16	Audit configuration changes from last 30 days before export	Query audit logs (<code>fetch logs, from:-30d filter matchesPhrase(log.source, "audit")</code>) to identify recently changed settings	Recommended	Discovery

3. IAM, SSO, and Identity

Source: S2S-03

#	Best Practice	Setting / Value	Priority	Category
17	Redesign IAM during consolidation instead of lift-and-shift	Use the migration as the opportunity to fix group sprawl and overprivileged policies; redesign from scratch for the consolidated tenant	Critical	IAM
18	Create a new SAML application in your IdP for the target tenant	New Entity ID, new ACS URL; do not reuse the source tenant's SAML app	Critical	SSO
19	Test SSO with a pilot user before cutover	SAML configuration issues are the #1 day-of-cutover blocker	Critical	SSO
20	Configure SCIM provisioning for the target tenant	Set new SCIM URL and token; users sync automatically on first login	Recommended	User Migration
21	Create new OAuth clients in target tenant	Source tenant OAuth client secrets cannot be exported; create fresh clients with matching scopes	Critical	Credentials
22	Create new API tokens in target tenant	Source tenant API tokens cannot be exported; create new tokens and update all consumers before cutover	Critical	Credentials
23	Use Terraform with <code>account-idm-read</code> , <code>account-idm-write</code> , and <code>iam-policies-management</code> OAuth scopes for IAM	These three scopes are required to manage groups, policies, and bindings	Critical	IAM
24	Map cloud provider identity constructs when changing providers	AWS IAM roles to Azure Service Principals, Azure AD Groups to GCP Google Groups, etc.	Recommended	Cloud Identity
25	Limit Azure AD group claims to 150 groups per SAML assertion	Azure AD has a hard limit; use group filtering if your org exceeds this	Recommended	SSO

4. Agent and Operator Migration

Source: S2S-04

#	Best Practice	Setting / Value	Priority	Category
26	Migrate agents phased by environment (dev, staging, production)	Dev first (weeks 1-2), staging (weeks 3-4), production (weeks 5-6); gives Davis AI time to establish baselines	Critical	Agent Migration
27	Reconfigure OneAgent with <code>oneagentctl</code>	Run <code>oneagentctl --set-server=<target-url>/communication --set-tenant=<target-id> --set-tenant-token=<target-token></code> ; no application restart required	Critical	Agent Migration
28	Reinstall ActiveGates in the target tenant	ActiveGates cannot be reconfigured like OneAgents; install fresh from the target tenant UI	Critical	Agent Migration
29	Use Dynatrace Operator v1.8.1+ with DynaKube API <code>v1beta5</code> or <code>v1beta6</code>	Operator 1.8.x removes <code>v1beta3</code> and auto-converts to <code>v1beta6</code> ; update manifests from <code>v1beta3</code> before upgrading	Critical	Kubernetes
30	Install Operator via Helm from <code>oci://public.ecr.aws/dynatrace/dynatrace-operator</code> version 1.8.1	Use <code>--atomic</code> flag for rollback safety	Critical	Kubernetes
31	Budget 540m CPU and 872Mi memory per node for Dynatrace overhead	OneAgent (100m/512Mi) + CSI Driver components (440m/360Mi) per node; control plane adds 650m/320Mi per cluster	Recommended	Capacity
32	Recreate network zones in target before migrating agents	Export with <code>monaco download --specific-settings builtin:networkzones</code> and deploy to target	Critical	Network
33	Keep source tenant running throughout migration window	Do not decommission source ActiveGates or revoke source tokens until all agents are confirmed in the target	Critical	Rollback
34	Automate fleet-wide <code>oneagentctl</code> with config management tools	Use Ansible, Chef, Puppet, AWS SSM, or	Recommended	Agent Migration

#	Best Practice	Setting / Value	Priority	Category
		Azure Automation for large fleets		
35	Update IRSA (AWS), Managed Identity (Azure), or Workload Identity (GCP) to reference target tenant	Cloud-specific K8s identity must point to the new tenant	Critical	Kubernetes

5. Configuration Import and Deployment Order

Source: S2S-05

#	Best Practice	Setting / Value	Priority	Category
36	Deploy configuration in dependency order	Buckets (1st) -> Enrichment rules (2nd) -> OpenPipeline (3rd) -> Segments (4th) -> Gen2 config (5th-13th) -> Alerting/SLOs (14th-18th) -> Dashboards/Workflows (19th-21st) -> IAM last (22nd-24th)	Critical	Deployment
37	Deploy IAM groups, policies, and bindings last	IAM WHERE clauses reference schemas, buckets, and security contexts that must already exist	Critical	IAM
38	Use <code>monaco deploy manifest.yaml --environment target</code> for the entire import (except IAM)	Monaco resolves dependencies automatically within a single run	Critical	Tooling
39	Run <code>monaco deploy --dry-run</code> before every real deploy	Preview what will change before committing to the target tenant	Critical	Validation
40	Replace all hardcoded entity IDs with tag-based filters before import	Convert <code>entityId("HOST-xxxx")</code> to <code>tag(app:checkout), tag(env:production)</code> in dashboards, SLOs, and notification rules	Critical	Entity Remapping
41	Look up new entity IDs in target tenant using <code>fetch dt.entity.<type> \ filter contains(entity.name, "<name>")</code>	Use DQL to find new entity IDs after agents report to target	Recommended	Entity Remapping
42	Prefix resources with source tenant identifier during consolidation	Set <code>name = "\${var.source_prefix} - \${resource_name}"</code> in Terraform variables	Recommended	Naming

6. Cloud Integration

Source: S2S-06

#	Best Practice	Setting / Value	Priority	Category
43	Recreate cloud provider credentials manually in the target tenant	<code>monaco download</code> cannot export AWS, Azure, or K8s credentials; recreate IAM roles, service principals, and service accounts for the target	Critical	Cloud
44	Create new AWS IAM role with trust policy referencing the target tenant's Dynatrace account ID and external ID	Get these values from the target tenant's AWS integration setup wizard	Critical	AWS
45	Create new Azure App Registration with <code>Monitoring Reader</code> role on target subscriptions	New Tenant ID, Client ID, Client Secret for the target tenant	Critical	Azure
46	Create new GCP Service Account with <code>Monitoring Viewer</code> and <code>Compute Viewer</code> roles	Generate new JSON key for the target tenant integration	Critical	GCP
47	Deploy new AWS Log Forwarder Lambda pointing to the target tenant	Update S3 event notifications to trigger the new Lambda function	Critical	AWS
48	Update Azure Event Hub consumer group for the target tenant	Configure new ActiveGate or function to consume from Event Hub	Critical	Azure
49	Standardize tag naming across cloud providers	Use consistent keys (<code>app</code> , <code>env</code> , <code>team</code>) regardless of AWS/Azure/GCP	Recommended	Multi-Cloud
50	Schedule credential rotation for all cloud provider integrations post-migration	Rotate within 30 days of cutover	Recommended	Security

7. Dashboards, Workflows, and Integrations

Source: S2S-07

#	Best Practice	Setting / Value	Priority	Category
51	Audit all dashboards for hardcoded entity IDs before import	Search dashboard JSON for <code>entityId(, HOST- , SERVICE- , PROCESS_GROUP-</code> patterns and replace with tag-based filters	Critical	Dashboards
52	Reassign dashboard ownership to a target tenant user or service account	Source user accounts may not exist in target; set <code>owner</code> field to valid target identity	Critical	Dashboards
53	Create new workflow service user (actor) in the target tenant	Source service users cannot be migrated; create new service user and set as workflow actor	Critical	Workflows
54	Update all webhook URLs that reference source-tenant-specific endpoints	Verify custom webhooks accept traffic from the target tenant's IP range	Critical	Notifications
55	Recreate Synthetic ActiveGate locations in the target tenant	Private Synthetic locations are tenant-specific; update location references in all monitor configs	Critical	Synthetic
56	Use a classic API token (not OAuth) for Synthetic monitor migration	Synthetic monitors require classic token as of provider v1.88.0+	Critical	Synthetic
57	Reinstall Extensions 2.0 from the Dynatrace Hub in the target tenant	Neither Monaco nor Terraform supports Extensions 2.0 export/import; manual reinstall required	Critical	Extensions
58	Extensions 2.0 require a host-based ActiveGate, not K8s-based	Ensure the target tenant has an appropriate host-based AG deployment for extensions	Critical	Extensions
59	Webhook integrations (Teams, Slack, PagerDuty, ServiceNow, Jira) use the same URLs across tenants	No URL changes needed for these standard integrations; just recreate the notification rule in the target	Recommended	Notifications

8. OpenPipeline and Grail Buckets

Source: S2S-08

#	Best Practice	Setting / Value	Priority	Category
60	Deploy Grail buckets before OpenPipeline routing rules	Routing rules that target nonexistent buckets cause data loss	Critical	Deployment Order
61	Deploy in order: buckets -> enrichment rules -> OpenPipeline rules	This three-step sequence ensures data flows correctly from the first byte	Critical	Deployment Order
62	Prefix bucket names with source tenant identifier when consolidating	E.g., <code>us-east_app_logs</code> , <code>eu-west_app_logs</code> to prevent collisions	Recommended	Naming
63	Match retention policies exactly for regulated environments	Production: set to same days as source (e.g., 360 days for PCI); verify compliance requirements before cutover	Critical	Compliance
64	Update bucket references in OpenPipeline routing rules to target bucket names	Routing rules from export contain source bucket names; replace with target bucket names	Critical	Data Routing
65	Use a single Monaco deploy for buckets + enrichment + OpenPipeline together	<code>monaco deploy</code> resolves the dependency order automatically	Recommended	Tooling

9. Parallel Operation and Data Continuity

Source: S2S-09

#	Best Practice	Setting / Value	Priority	Category
66	Run parallel operation for a minimum of 30 days	Historical data does not migrate; parallel is the only path to continuity	Critical	Parallel
67	Use phased parallel model to control cost	Migrate agents in waves (dev -> staging -> prod) so you never run 2x host units simultaneously; expect 1.1-1.5x host unit cost	Critical	Cost
68	Allow 2-4 weeks for Davis AI baselines to stabilize	Availability baselines: 2-3 days; response time and error rate: 1-2 weeks; resource utilization: 2-4 weeks (full business cycle)	Critical	Baselines
69	Migrate dev/staging first to give Davis AI a head start	Dev agents in target from week 1 provides 4+ weeks of baseline data before production cutover	Recommended	Baselines
70	Communicate the SLO baseline gap to stakeholders	SLOs start at 0% in the target and accumulate data over 7-30 days depending on evaluation window	Critical	Communication
71	Notify engineering teams 4 weeks before cutover	Include migration timeline, agent cutover schedule, and expected impacts	Critical	Communication
72	Notify SRE/on-call 2 weeks before cutover	Cover alert routing changes and dual-tenant monitoring procedures	Critical	Communication
73	Warn about Davis AI false positives during first 2-4 weeks post-cutover	Baselines are relearning; expect noisy alerting until stabilized	Recommended	Baselines
74	Negotiate parallel license period with Dynatrace	Contact your Dynatrace account team to discuss temporary parallel licensing to reduce cost	Recommended	Cost
75	Export key problem history as documentation before decommissioning source	Problem history cannot migrate; export critical incidents for post-mortem reference	Recommended	Data

10. SLOs and Alerting

Source: S2S-10

#	Best Practice	Setting / Value	Priority	Category
76	Convert all SLO metric expressions from entity IDs to tag-based filters	Replace <code>type(SERVICE),entityId(SERVICE-xxx)</code> with <code>type(SERVICE),tag(app:checkout),tag(env:production)</code>	Critical	SLOs
77	Replace management zone ID references with management zone names	Convert <code>mzId(123)</code> to <code>mzName("Production")</code> because zone IDs change between tenants	Critical	SLOs
78	Set initial SLO evaluation window to match available data	Start with <code>ROLLING_3_DAYS</code> , expand to <code>ROLLING_WEEK</code> after 7 days, then to full window after 30 days	Critical	SLOs
79	Align SLO cutover with calendar boundaries	If using calendar-based SLOs, switch at the start of a month or quarter to avoid partial evaluation	Recommended	SLOs
80	Port alerting profile severity filters and tag filters as-is	Ensure the referenced tags exist in the target tenant before deploying alerting profiles	Critical	Alerting
81	Verify maintenance window timezone settings match the target tenant	Cron expressions are timezone-sensitive; confirm schedules are correct after import	Recommended	Alerting
82	Update maintenance window entity scopes to reference target tenant entities	Source entity references in scope definitions must be remapped	Critical	Alerting

11. Cutover Execution

Source: S2S-11

#	Best Practice	Setting / Value	Priority	Category
83	Freeze configuration changes in the source tenant before final sync	No settings modifications during the cutover window	Critical	Change Control
84	Run a final <code>monaco download -> monaco deploy sync</code> immediately before cutover	Captures any configuration changes made since initial export	Critical	Sync
85	Verify host count in target matches source	Run <code>fetch dt.entity.host \ summarize count() </code> in both tenants and compare	Critical	Validation
86	Verify service count in target matches source	Run <code>fetch dt.entity.service \ summarize count() </code> in both tenants and compare	Critical	Validation
87	Verify logs are flowing in target	Run <code>fetch logs, from:-1h \ summarize log_count = count() </code> and confirm <code>log_count > 0</code>	Critical	Validation
88	Verify spans are flowing in target	Run <code>fetch spans, from:-1h \ summarize span_count = count() </code> and confirm <code>span_count > 0</code>	Critical	Validation
89	Verify enrichment tags are populating	Run <code>fetch logs, from:-1h \ filter isNotNull(dt.security_context) \ summarize count() </code> and confirm <code>count > 0</code>	Critical	Validation
90	Disable alerting in source tenant after cutover	Prevents duplicate alerts during the transition tail	Critical	Alerting
91	Set source tenant to read-only retention after cutover	Preserves historical data for reference without active monitoring cost	Critical	Decommission
92	Keep a fallback local admin account for SSO failures	If SAML is misconfigured on day of cutover, local admin bypasses SSO	Critical	SSO
93	Obtain stakeholder sign-off from SRE, dev, security, management, and compliance	Each team validates their domain: alerting, dashboards, IAM, SLO reporting, retention	Critical	Governance

12. Post-Migration and Decommission

Source: S2S-12

#	Best Practice	Setting / Value	Priority	Category
94	Keep source tenant running for 30 days after cutover before decommissioning	Read-only mode for reference; do not deprovision until all stakeholders confirm	Critical	Decommission
95	Tune Davis AI anomaly detection baselines during weeks 1-2 post-cutover	Suppress known false positives; adjust sensitivity for noisy services	Critical	Optimization
96	Expand SLO evaluation windows to full duration by week 4	Move from <code>ROLLING_3_DAYS</code> to <code>ROLLING_WEEK</code> to <code>ROLLING_MONTH</code> as data accumulates	Recommended	SLOs
97	Rotate all credentials (API tokens, OAuth clients, cloud provider keys) within 30 days of cutover	Post-migration credential hygiene; revoke any temporary migration credentials	Critical	Security
98	Export final audit logs from source tenant before decommission	Compliance requires audit trail retention; export before access is lost	Critical	Compliance
99	Revoke all API tokens and deactivate all OAuth clients in source	Prevents unauthorized access to the deprecated tenant	Critical	Security
100	Remove SAML/SSO application from IdP for source tenant	Eliminates stale IdP configuration; prevents confusion	Recommended	SSO
101	Update all runbooks, CI/CD pipelines, and architecture diagrams to reference target tenant	Replace source tenant URLs, API endpoints, and dashboard links everywhere	Critical	Documentation
102	Consolidate redundant Synthetic monitors post-migration	If consolidating tenants, deduplicate monitors that tested the same endpoints from different tenants	Optional	Optimization
103	Conduct a lessons-learned review within 2 weeks of cutover	Capture planning gaps, tool issues, communication effectiveness, timeline accuracy, and surprise entity ID references	Recommended	Process
104	Replace all remaining entity ID references in dashboards and SLOs with tag-based filters	Post-migration cleanup; ensures configuration is fully portable for future migrations	Recommended	Optimization

Summary

This notebook contains **104 best practices** across 12 categories. The breakdown by priority:

Priority	Count	Guidance
Critical	73	Must implement. Skipping any of these risks data loss, access failures, or migration rollback.
Recommended	27	Implement unless there is a documented reason not to. These reduce risk and cost.
Optional	4	Implement if resources allow. These improve efficiency but are not blockers.

The Five Principles of S2S Migration

- 1. **Export early, validate often** -- do not wait for cutover to discover missing config
- 2. **Tags over entity IDs** -- make all configuration portable before migration
- 3. **Parallel is not optional** -- historical data does not migrate; 30 days minimum
- 4. **IAM last** -- deploy the resources first, then lock them down with policies
- 5. **Communicate relentlessly** -- migration is as much about people as technology

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.