

SL2DT-01: Overview & Migration Strategy

Series: SL2DT — Sumo Logic to Dynatrace | **Notebook:** 1 of 10 | **Created:** April 2026 | **Last Updated:** 06/23/2026

Overview

Goal of this notebook: establish the strategy, mental model, and vocabulary for migrating from Sumo Logic to Dynatrace. Every decision in the following nine notebooks ladders up to a choice made here.

This is not a how-to. It is the *why* — why Gen3, why Anomaly Detection over lift-and-shift thresholds, why train-the-trainer, why a bucket-plus-attribute `_sourceCategory` strategy, why OpenPipeline at ingest rather than parsing at query time.

Scope: logs, dashboards, monitors, scheduled searches, Field Extraction Rules. Metrics and traces are addressed where they intersect log migration. Cloud SIEM migration is explicitly out of scope — it is a separate workstream.

Sprint 1.337 (April 2026) Updates Affecting SL2DT

Three sprint-1.337 changes simplify the Sumo Logic → Dynatrace migration:

- OneAgent primary fields/tags at the source.** Hosts now emit standardized fields (`dt.security_context` , `dt.cost.costcenter` , `dt.cost.product`) and customer-defined primary tags as top-level attributes on every signal — eliminating a class of OpenPipeline parse processors needed during the Sumo→DT cut. See SL2DT-03 (Log Ingest Architecture) for the ingest-time configuration.
- OpenPipeline extraction processor recommended-field suggestions.** When converting Sumo Field Extraction Rules (FERs) to OpenPipeline DPL processors (SL2DT-03), the UI now flags permission-relevant fields and Smartscape identifiers — preventing the most common misconfiguration: promoting sensitive content into a permission-relevant position.

3. **Configuration API → Settings v2 acceleration + Platform tokens.** New automation (SL2DT-08) should target Settings v2 paths and use Platform tokens (`dt0s16 / dt0s01` , `Authorization: Bearer`). Classic `dt0c01` still works for legacy paths but is not the default for new pipelines.

These changes reinforce — they don't change — the Gen3-first migration baseline documented in this series.

Table of Contents

1. [Why This Series Exists](#)
 2. [Platform Mental Model — Sumo vs Dynatrace](#)
 3. [The Five-Wave Migration Pattern](#)
 4. [Train-the-Trainer Delivery Model](#)
 5. [MVP vs Full Fidelity — The Cut-Scope Decision](#)
 6. [Anti-Patterns to Avoid](#)
 7. [Success Criteria](#)
 8. [What to Read Next](#)
-

Prerequisites


Requirement	Details
Audience	Migration lead + architects scoping the engagement
Format	Orientation — read before any hands-on work
Prior knowledge	Sumo Logic search experience (helpful); Dynatrace fundamentals (ONBRD-01)
Tenant	Not required for this notebook

1. Why This Series Exists

Sumo-to-Dynatrace migrations fail for predictable reasons:

Failure mode	Why it happens	Mitigation (which notebook addresses it)
1:1 lift-and-shift of static thresholds	Teams treat Dynatrace as "Sumo with a different UI"	SL2DT-05 anomaly detection framework
<code>_sourceCategory</code> chaos post-migration	No upfront taxonomy decision	SL2DT-02 inventory, SL2DT-03 ingest design
Dashboards migrate but nobody uses them	Dashboard inventory wasn't audit-log-pruned	SL2DT-02 cut-scope analysis
Monitors fire too much	DQL anomaly baseline needs different tuning than Sumo thresholds	SL2DT-05 tuning playbook
Governance gap at cutover	Sumo role model doesn't map 1:1 to IAM	SL2DT-07 role mapping
Parallel-run runs too long	No validation exit criteria defined	SL2DT-09 cutover runbook

This series is the procedural spine that prevents those failures.



2. Platform Mental Model — Sumo vs Dynatrace

The platforms solve the same problem differently. Mapping the concepts is the first job.

Sumo Logic Concept	Dynatrace Equivalent	Notes
<code>_sourceCategory</code> (taxonomy)	Grail bucket + optional attribute	Plan in SL2DT-03
Hosted Collector	OneAgent / OTel Collector / OpenPipeline endpoint	Per-source choice
Installed Collector	OneAgent on host	Auto-instrumented
Source (file/stream config)	OneAgent log source config or OTel receiver	Ingest-side
Partition	Bucket (Grail)	Retention + IAM boundary
Index (field index)	— (Grail indexes all fields)	No manual index management
Field Extraction Rule	OpenPipeline processor	DPL pattern
Saved Search	Notebook section or Workflow DQL task	
Scheduled Search	Workflow with cron trigger + DQL + action	
Monitor (threshold)	Workflow with DQL + threshold, or anomaly detection	Prefer Dynatrace Intelligence
Monitor (outlier/anomaly)	Anomaly Detection	Native
Dashboard	Notebook or Dashboard (DT)	1:1 in most cases
Dashboard Panel	DQL section within notebook/dashboard	
Lookup Table	Grail Lookup Table (Resource Store API)	
Role (RBAC)	Platform IAM Group + Policy	

Sumo Logic Concept	Dynatrace Equivalent	Notes
Role Search Filter	Policy with bucket filter	
Data Volume Alert	Bucket retention + cost per bucket	

Load-bearing concept: `_sourceCategory` . Everything downstream depends on how it's mapped. SL2DT-03 walks through the three strategies (bucket / attribute / entity tag) with a decision framework.

3. The Five-Wave Migration Pattern

Migrate in waves, not a single big-bang. Each wave has entry criteria, exit criteria, and owners.

Wave	Asset Class	Entry Criteria	Exit Criteria	Typical Duration
Wave 0	Inventory + cut scope	Sumo audit-log pulled	Cut scope signed off by owners	2 weeks
Wave 1	Ingest + OpenPipeline + buckets	Wave 0 complete; taxonomy chosen	All high-volume sources dual-writing	3–4 weeks
Wave 2	Core monitors (top-10 business-critical)	Wave 1 complete; dual ingest running	First 10 monitors validated	2 weeks
Wave 3	Dashboards (high-usage only)	Wave 2 complete	Top-N dashboards live + users trained	3–4 weeks
Wave 4	Remaining monitors + dashboards (app teams)	Wave 3 complete; train-the-trainer done	Signoff per team	6–8 weeks
Wave 5	Cutover + Sumo decommission	All waves validated; parallel run window elapsed	Sumo contract terminated	2–3 weeks

Critical waves: Wave 0 (inventory) and Wave 1 (ingest parity). These are gating — nothing downstream works without them.

Parallelizable: Waves 2 and 3 can overlap with Wave 4 for app teams who are ready early. Coordinate through the migration PM.

4. Train-the-Trainer Delivery Model

For engagements with 50+ teams and 100+ power users, direct delivery doesn't scale. Use train-the-trainer:

1. **Core migration team** (typically SREs, 4–8 engineers) learns Dynatrace deeply — SL2DT-04 through SL2DT-08.
2. **Core team demonstrates** by migrating 5 representative real-world assets per domain:
 - 1–2 dashboards
 - 2–3 monitors (mix of static + anomaly)
 - 1 scheduled search
 - 1 FER → OpenPipeline processor
3. **Core team documents** the translation as a worked example, commits to a shared repo.
4. **App teams convert their own assets** using the examples as templates. Each team delivers a PR for core-team review.
5. **Core team reviews + approves** — becomes the gate for decommissioning the Sumo-side asset.

This pattern is proven in the NR2DT migrations (NRLC-01 through NRLC-09 apply the same model). Don't skip the review step — it is where best-practice violations get caught before they proliferate.

Why App Teams Resist (and How to Handle It)

Ownership matters. App teams see monitor conversion as "more work without new value." Common patterns:

- **"Just recreate my static threshold"** — App team wants to copy the number from Sumo without evaluating if anomaly detection would reduce noise. Push back with data (SL2DT-05 provides the framework).

- **"I'll do it later"** — App team defers, risking cutover date. Migration PM tracks team-by-team progress in a shared burn-down.
- **"The translation tool got it wrong"** — App team dismisses low-confidence output without engaging. Core team's review role includes translation debugging (SL2DT-04).

5. MVP vs Full Fidelity — The Cut-Scope Decision

Migrations fail when the scope is "everything in Sumo." A typical Sumo deployment has 5,000–10,000 dashboards and monitors; audit logs routinely show that 70–80% of them haven't been accessed in 90+ days.

The cut-scope decision is a project-plan make-or-break. Get it wrong in either direction:

Direction	Symptom	Consequence
Too aggressive	Cut assets that are actually used quarterly	Escalations post-cutover; trust damaged
Too permissive	Everything migrates	Budget overrun; team burnout; low-value assets dominate review cycles

Decision Framework

Pull the Sumo audit log. For each asset, record:

- **Last accessed date** — if > 90 days, candidate for cut
- **Unique users** — if 1 user, candidate for "personal" (migrate only if owner requests)
- **Team ownership** — if team is dissolved / reorg'd, re-assign or retire
- **Complexity** — scripts/schedules add migration cost; weigh against usage

Present the cut list to each team owner. Require explicit signoff on cut vs migrate vs retire.

Output: `cut-scope.md` — your audit trail for every retired asset. Commit it.

6. Anti-Patterns to Avoid

Six anti-patterns recur in every Sumo-to-Dynatrace migration. Watch for them from day one.

6.1 Lift-and-Shift Static Thresholds

App teams' first instinct is to copy the `> 100` threshold from Sumo to Dynatrace. This is wrong in two ways:

1. **Baseline differs.** Dynatrace's ingest path (OneAgent + OpenPipeline) may produce slightly different metric values than Sumo. 100 in Sumo $\neq$ 100 in DT.
2. **Dynatrace Intelligence is better.** For 80%+ of static-threshold monitors, Anomaly Detection produces fewer false positives and catches real anomalies that static thresholds miss.

Fix: SL2DT-05 has a decision framework — "when does static win vs Dynatrace Intelligence."

6.2 `_sourceCategory` Flat-Mapped to a Single Field

Tempting: "just put `_sourceCategory` as a custom attribute, queries work the same."
Problems:

- No bucket isolation → one noisy team's logs bloat your high-retention bucket
- IAM can't use the attribute for policy filtering (some versions)
- Cost allocation is per-bucket, not per-attribute

Fix: SL2DT-03 covers the mixed bucket + attribute strategy.

6.3 Parsing at Query Time Instead of Ingest Time

Sumo's FERs parse at ingest. Teams sometimes skip this step in DT and use `parse content, ...` in every dashboard query.

- Ingest-time parsing = parsed fields indexed and queryable
- Query-time parsing = re-parsed on every execution, slower, no structured index

Fix: SL2DT-03 + SL2DT-04 — FER equivalents go into OpenPipeline.

6.4 Dashboard Sprawl

Cutting 8,000 dashboards to 2,000 is still too many. Target is under 500 for most organizations, with clear ownership per dashboard.

Fix: SL2DT-02 cut scope + SL2DT-06 dashboard consolidation.

6.5 "We'll Figure Out IAM at the End"

Bucket and policy decisions made late force re-ingest. IAM is Wave 1, not Wave 5.

Fix: SL2DT-07 is intentionally placed mid-series to cue you: don't defer governance.

6.6 Running Sumo + DT in Parallel Forever

"Parallel run" is a validation window, not an architecture. Cost doubles. Teams bifurcate. The cutover date must be a commitment, not aspirational.

Fix: SL2DT-09 defines explicit exit criteria for the parallel-run window.

7. Success Criteria for the Migration

Dimension	Metric	Target
Monitor fidelity	% of Sumo monitors with Dynatrace equivalent (validated)	≥ 95%
Dashboard fidelity	% of in-scope dashboards recreated and validated	≥ 90%
Query translation confidence	% of SumoQL translated at HIGH confidence	≥ 75%
Dynatrace Intelligence adoption	% of new monitors that use anomaly detection (vs static)	≥ 40%
Alert noise reduction	DT alert count during parallel-run vs Sumo baseline	≤ Sumo baseline
Team signoff	% of app teams signed off by cutover date	100%

Dimension	Metric	Target
Cost target	DT + parallel Sumo cost tracking vs budget	within 110%
Cutover date	Sumo contract termination	on or before target date

Track these on a weekly dashboard during Waves 2–4. Use Anomaly Detection on the alert-count metric itself to catch regressions.

8. What to Read Next

- **SL2DT-02 — Assessment & Inventory** — pull the Sumo audit log, produce cut scope
- **SL2DT-03 — Log Ingest Architecture** — `_sourceCategory` mapping, bucket design, OneAgent + OTel + OpenPipeline

If you're just here for a specific question:

- Query translation? → SL2DT-04 and the `sumoql-to-dql` skill
- Monitor conversion? → SL2DT-05
- Automation? → SL2DT-08

Companion assets

- [sumoql-to-dql skill](#) — translation tables
 - [SL2DT AGENT-TASKS.md](#) — brief for building the `Dynatrace-SumoLogic` migration tool
 - [NR2DT series](#) — parallel pattern for New Relic migrations
 - [S2D series](#) — parallel pattern for Splunk migrations
-

11. References

Dynatrace platform overview

- [Grail \(DT docs\)](#)
- [DQL Reference \(DT docs\)](#)
- [OpenPipeline \(DT docs\)](#)
- [Workflows \(DT docs\)](#)
- [Anomaly detection \(DT docs\)](#)
- [Identity and access management \(DT docs\)](#)

Sumo Logic platform reference

- [Sumo Logic search operators \(Sumo Logic docs\)](#)
- [Sumo Logic monitors \(Sumo Logic docs\)](#)
- [Sumo Logic dashboards \(Sumo Logic docs\)](#)
- [Sumo Logic users and roles \(Sumo Logic docs\)](#)

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace or Sumo Logic. Always verify information against the official [Dynatrace documentation](#) and [Sumo Logic documentation](#).