

SL2DT-08: Automation & GitOps

Series: SL2DT — Sumo Logic to Dynatrace | **Notebook:** 8 of 11 | **Created:** April 2026 | **Last Updated:** 07/30/2026

Overview

Goal of this step: automate the promotion path for dashboards, monitors, workflows, IAM, OpenPipeline, and bucket configs. Migration is only sustainable if post-cutover changes go through the same CI/CD discipline — not via UI clickops that drift out of source control.

Automation applies both to the migration itself (programmatically import 8,000 dashboards faster than clicking) and to ongoing operations (every config change is a PR).

Table of Contents

1. [What You'll Produce](#)
 2. [Tool Selection — Terraform vs Monaco vs the Future Migration Tool](#)
 3. [Terraform — Bucket, IAM, Settings 2.0 Config](#)
 4. [Monaco — Bulk Config Promotion](#)
 5. [CI/CD Pipeline Patterns](#)
 6. [Automating the Sumo Extract → DT Import Flow](#)
 7. [Ansible / CloudFormation / Cloud-Specific Handoffs](#)
 8. [Post-Migration: Change Management](#)
 9. [Step Exit Criteria](#)
-

Prerequisites

Requirement	Details
Audience	Platform engineering, DevOps, SRE
Tools	Git, Terraform ≥ 1.6 , Monaco $\geq 2.x$, Python 3, cloud CLIs (AWS/GCP/Azure)
Dynatrace access	Platform Token with broad write scopes; separate service account for CI
Prior reading	AUTOM-01 (Configuration automation overview), K8S-07 (GitOps for Dynatrace), <code>Dynatrace-NewRelic</code> for migration-tool patterns

1. What You'll Produce

Artifact	Purpose
<code>terraform/</code>	All IaC for buckets, IAM, OpenPipeline, Workflows
<code>monaco/</code>	Bulk configs (dashboards, monitors, settings)
<code>.github/workflows/</code> (or equivalent CI)	Pipelines for <code>validate</code> → <code>plan</code> → <code>apply</code>
<code>scripts/extract-and-import/</code>	<code>Sumo</code> → DT migration pipeline scripts
<code>runbooks/rollback.md</code>	How to revert a bad config change

2. Tool Selection — Terraform vs Monaco vs the Future Migration Tool

Tool	Use for	Avoid for
Terraform (Dynatrace provider)	Buckets, IAM groups/policies, Settings 2.0 schemas, some Workflows	Bulk dashboards (slow), complex dashboard JSON

Tool	Use for	Avoid for
Monaco	Dashboards, monitors, settings in bulk; promotion across environments	IAM (use Terraform)
Dynatrace-SumoLogic migration tool (future — see AGENT-TASKS.md)	Sumo extract + translate + import	Ongoing change management
Direct API (scripts)	One-offs, custom orchestration, Document API for notebooks	Anything that needs state tracking
Dynatrace UI	Investigation, exploration — never for config changes post-cutover	Migration execution

Principles

- **One tool per concern.** Don't manage the same asset in both Terraform and Monaco.
- **State reviewable.** Every config change goes through PR with `terraform plan` or `monaco deploy --dry-run` output.
- **Idempotency.** Every pipeline must be safe to re-run.

3. Terraform — Bucket, IAM, Settings 2.0 Config

Provider Setup

```

terraform {
  required_providers {
    dynatrace = {
      source  = "dynatrace-oss/dynatrace"
      version = "~> 1.60"
    }
  }
}

provider "dynatrace" {
  dt_env_url      = var.dt_env_url
  dt_platform_token = var.dt_platform_token # Bearer token
}

```

Organizing Modules

```
terraform/
├─ modules/
│   ├─ bucket/           # Grail buckets
│   ├─ iam-group-policy/ # Group + policy + binding
│   ├─ openpipeline/     # Pipelines + processors
│   └─ workflow/         # Workflow configs
├─ envs/
│   ├─ prod/
│   │   └─ main.tf
│   ├─ preprod/
│   └─ dev/
└─ variables.tf
```

Example — Bucket Module Call

```
module "bucket_prod_logs" {
  source      = "../../modules/bucket"
  name        = "custom_logs_prod"
  display_name = "Production Logs"
  retention   = 90
}

module "iam_prod_readers" {
  source      = "../../modules/iam-group-policy"
  group_name  = "g_prod_readers"
  description = "Read-only prod logs"
  policy_statements = [
    "ALLOW storage:logs:read WHERE storage:bucket-name = \"custom_logs_prod\";"
  ]
}
```

State Management

Use a remote backend (S3 + DynamoDB lock, or GCS, or Terraform Cloud):

```
terraform {  
  backend "s3" {  
    bucket      = "tfstate-dynatrace-migration"  
    key         = "prod/dynatrace.tfstate"  
    region      = "us-east-1"  
    dynamodb_table = "tfstate-lock"  
  }  
}
```

Never commit `.tfstate` files to git.

4. Monaco — Bulk Config Promotion

Monaco is the right tool for dashboards/monitors in bulk. It excels at multi-tenant promotion (dev → preprod → prod).

Project Structure

```
monaco/  
├─ environments.yaml      # Tenant list + auth  
├─ projects/  
│   ├─ dashboards/  
│   │   ├─ config.yaml  
│   │   └─ definitions/  # One JSON/YAML per dashboard  
│   ├─ monitors/  
│   └─ openpipeline/
```

environments.yaml (non-secret)

```
environmentGroups:
  - name: production
    environments:
      - name: prod
        url:
          type: value
          value: https://&lt;env-id&gt;.apps.dynatrace.com
        auth:
          token:
            type: environment
            name: DT_PROD_TOKEN
      - name: preprod
        environments:
          - name: preprod
            ...
```

Deploying Dashboards

```
monaco deploy -e environments.yaml projects/dashboards --environment-group
production --dry-run
monaco deploy -e environments.yaml projects/dashboards --environment-group
production
```

Template Variables

Monaco supports per-environment templating:


```
configs:
  - id: dashboard_payments_api
    type: dashboard
    template: templates/payments-api.json
    parameters:
      tenant_name: "{{{ .name }}}"
```

Migration Use Case — Bulk Dashboard Import

If you have 2,000 translated dashboards, iterating via Terraform's `dynatrace_document` is slow (one API call per resource, sequential apply). Monaco batches and deploys in parallel — typical 5–10x faster for bulk work.

5. CI/CD Pipeline Patterns

Typical Pipeline Stages



GitHub Actions Example

```
name: dynatrace-config
on:
  pull_request:
    branches: [main]
  push:
    branches: [main]

jobs:
  validate:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - run: python3 scripts/validate-dql.py terraform/
      - run: monaco lint monaco/
      - uses: hashicorp/setup-terraform@v3
      - run: |
          cd terraform/envs/prod
          terraform init -backend=false
          terraform validate

  plan:
    if: github.event_name == 'pull_request'
    needs: validate
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: hashicorp/setup-terraform@v3
      - run: |
          cd terraform/envs/prod
          terraform init
          terraform plan -no-color > plan.txt
      - uses: actions/github-script@v7
        with:
          script: |
            const fs = require('fs');
            const plan =
fs.readFileSync('terraform/envs/prod/plan.txt', 'utf8');
            github.rest.issues.createComment({
              issue_number: context.issue.number,
              owner: context.repo.owner, repo: context.repo.repo,
              body: ``\n' + plan + '\n``
            });

  apply:
    if: github.event_name == 'push'
```

```
needs: validate
runs-on: ubuntu-latest
environment: production
steps:
  - uses: actions/checkout@v4
  - uses: hashicorp/setup-terraform@v3
  - run: |
      cd terraform/envs/prod
      terraform init
      terraform apply -auto-approve
  - run: monaco deploy -e environments.yaml monaco/projects --environment-
group production
```

Secrets Management

Platform Tokens and API keys live in:

- **GitHub Actions:** repository secrets
- **GitLab CI:** protected variables
- **Jenkins:** Credentials plugin
- **CircleCI:** contexts

Never commit tokens. Every `*.env` file goes in `.gitignore`.

6. Automating the Sumo Extract → DT Import Flow

For bulk migration work (Waves 3 and 4), scripts automate the extract + translate + import loop. The eventual `Dynatrace-SumoLogic` migration tool (see `docs/AGENT-TASKS.md`) is the destination; until built, scripts.

Minimal Python Pipeline

```
#!/usr/bin/env python3
"""Extract Sumo dashboards, translate, import to Dynatrace."""
import os, json, requests, base64
from pathlib import Path

SUMO_AUTH = base64.b64encode(
    f"{os.environ['SUMO_ACCESS_ID']}:{os.environ['SUMO_ACCESS_KEY']}".encode()
).decode()
DT_TOKEN = os.environ['DT_TOKEN'] # dt0s16.XXX
DT_URL = os.environ['DT_TENANT']
SUMO_URL = f"https://api.{os.environ['SUMO_REGION']}.sumologic.com"

def sumo_get(path):
    r = requests.get(f"{SUMO_URL}{path}",
                    headers={"Authorization": f"Basic {SUMO_AUTH}"})
    r.raise_for_status()
    return r.json()

def dt_post(path, body):
    r = requests.post(f"{DT_URL}{path}", json=body,
                    headers={"Authorization": f"Bearer {DT_TOKEN}"})
    r.raise_for_status()
    return r.json()

def translate_dashboard(sumo_dashboard):
    # Uses sumoql-to-dql skill mapping tables
    # Returns DT notebook/dashboard JSON
    ...

def validate_dashboard(dt_config) -> list[str]:
    """Local structural preflight on a translated dashboard document.

    Returns a list of problem strings; empty list means the document passed.
    Cheap, offline, and catches the faults a machine translator actually
    makes.
    """
    problems = []

    if dt_config.get("version") != 13:
        problems.append(f"version is {dt_config.get('version')}!r, expected 13")

    tiles = dt_config.get("tiles")
    layouts = dt_config.get("layouts")
    if not isinstance(tiles, dict):
```

```

        problems.append("tiles is missing or not an object")
    if not isinstance(layouts, dict):
        problems.append("layouts is missing or not an object")
    if not isinstance(tiles, dict) or not isinstance(layouts, dict):
        return problems          # nothing further is checkable

    for tile_id, tile in tiles.items():
        for key in ("type", "query", "visualization"):
            if not tile.get(key):
                problems.append(f"tile {tile_id}: missing {key}")
        if tile_id not in layouts:
            problems.append(f"tile {tile_id}: no matching layouts entry")

    for layout_id in layouts:
        if layout_id not in tiles:
            problems.append(f"layouts {layout_id}: no matching tile")

    return problems

def main():
    out = Path("./translated")
    out.mkdir(exist_ok=True)
    skipped = []
    for dash in sumo_get("/api/v2/dashboards")["dashboards"]:
        dt_config =
    translate_dashboard(sumo_get(f"/api/v2/dashboards/{dash['id']}"))
        (out / f"{dash['id']}.json").write_text(json.dumps(dt_config,
    indent=2))

    problems = validate_dashboard(dt_config)
    if problems:
        skipped.append((dash["id"], dash.get("title"), problems))
        continue          # never POST a document that failed preflight

    # Dry-run by default; --apply flag pushes to DT
    # dt_post("/platform/document/v1/documents", dt_config)

    for dash_id, title, problems in skipped:
        print(f"SKIPPED {dash_id} ({title}):")
        for p in problems:
            print(f"  - {p}")
    print(f"{len(skipped)} dashboard(s) skipped on preflight failure")

main()

```

Hook `validate_dashboard` into the **PreflightCheck** structure below as a per-document check, so a translation defect is reported in the same shape as a token or scope problem rather than as a stack trace mid-loop.

Preflight Check (before bulk import)

Adopt the **PreflightCheck** pattern from the future migration tool (see `docs/AGENT-TASKS.md §3`) for script-level migrations:

```
@dataclass
class PreflightCheck:
    name: str
    status: str          # pass / warn / fail
    scopes_min: list[str]
    scopes_recommended: list[str]
    diagnosis: str
    remediation: str

def preflight():
    checks = []
    # Check token prefix determines auth scheme
    prefix = DT_TOKEN[:7]
    if prefix.startswith("dt0s16"):
        auth_scheme = "Bearer"
    elif prefix.startswith("dt0c01"):
        auth_scheme = "Api-Token"
    else:
        checks.append(PreflightCheck(
            "token_prefix", "fail",
            [], [],
            f"Unrecognized token prefix: {prefix}",
            "Generate a Platform Token (prefix dt0s16) or Classic API Token (prefix dt0c01)"
        ))
    # ... more checks
    return checks
```

Auth Scheme by Token Prefix

Critical pattern — every HTTP call must pick the right header:

Token Prefix	Scheme
<code>dt0s16.</code>	<code>Authorization: Bearer <token></code>

Token Prefix	Scheme
dt0s01.	Authorization: Bearer <token> (OAuth)
dt0c01.	Authorization: Api-Token <token>

Wrong scheme → 401 Unsupported authorization scheme , even with every scope.

What the Local Gate Does and Does Not Cover

`validate_dashboard` is a cheap offline filter, not a substitute for the platform's own validation. It catches the structural faults a machine translator actually produces — a wrong `version` , a tile the translator emitted without a `visualization` , a `tiles` entry with no matching `layouts` entry — at zero API cost and before anything reaches a tenant. It cannot tell you whether the document *renders*.

Two limits to be explicit about:

- **A 2xx from the Documents API does not mean the dashboard renders.** The API confirms the document was accepted for storage. Whether the Dashboards app can draw every tile is a separate question, answered only by opening it.
- **The local gate has no model of the dashboard schema.** It knows the handful of invariants written into it. A tile with a valid-looking but wrong `visualization` value, or a `visualizationSettings` block shaped for a different chart type, passes the local check and fails in the app.

So keep the ladder: local preflight on every document → open a **sample** in the Dashboards app on a non-production tenant → then bulk-apply. The sample is what closes the gap, and on a bulk run it is the only thing that catches a systematic translator defect before it multiplies across the whole wave.

Forthcoming/rolling out (SaaS 1.344): a dashboard that fails validation will no longer load. Published 07/27/2026, staged tenant rollout from 07/29/2026 — verify whether it has reached your tenant. Before 1.344 a failing dashboard still loaded and surfaced validation warnings; after, it does not load at all, and Dynatrace names **API- and AI-authored** dashboards as the most affected population — which is exactly what this pipeline emits. This raises the stakes on a bulk loop specifically: a translator defect that used to ship N dashboards with a warning banner now ships N dashboards that do not open. Treat the warning state as a grace period, not a supported state. Full gate: **DASH-07 §5**.

Keep the dry-run default. `--apply` should be an explicit, deliberate act on a wave that has already passed preflight and had a sample opened.

7. Ansible / CloudFormation / Cloud-Specific Handoffs

Many customers run existing automation for infrastructure. Integrate rather than replace.

Ansible (host-based automation)

OneAgent installation via Ansible playbook:

```
- name: Install Dynatrace OneAgent
  hosts: all
  become: true
  tasks:
    - name: Download OneAgent installer
      get_url:
        url: "{{ dt_installer_url }}"
        dest: /tmp/oneagent.sh
        headers:
          Authorization: "Api-Token {{ dt_pass_token }}"
    - name: Install OneAgent
      shell: /bin/sh /tmp/oneagent.sh --set-host-group="{{ dt_host_group }}"
      args: { creates:
        /opt/dynatrace/oneagent/agent/lib64/liboneagentloader.so }
```

CloudFormation / AWS

Use the Dynatrace AWS Clouds app (native integration) rather than custom lambdas. Clouds app push path is settled practice — no ActiveGate needed for most AWS metrics.

Terraform for AWS role creation:

```
resource "aws_iam_role" "dynatrace_monitoring" {
  name = "DynatraceMonitoringRole"
  assume_role_policy = jsonencode({
    Version = "2012-10-17"
    Statement = [{
      Effect    = "Allow"
      Principal = { AWS = "arn:aws:iam::509560245411:role/dynatrace-
monitoring-role" }
      Action    = "sts:AssumeRole"
      Condition = { StringEquals = { "sts:ExternalId" = var.external_id } }
    }]
  })
}
```

Azure / GCP

Native cloud integrations via Dynatrace tenant UI or API. Avoid custom exporters unless absolutely needed.

Decision Matrix

Source Platform	Ingest Path
On-prem Linux/Windows hosts	OneAgent via Ansible
AWS EC2 + services	OneAgent (from host) + AWS Clouds app (for services)
AWS Lambda	Dynatrace Lambda extension
Azure VM + services	OneAgent + Azure integration
GCP GCE + services	OneAgent + GCP integration
Kubernetes	DynaKube operator
CloudWatch / EventBridge	AWS Clouds app

8. Post-Migration: Change Management

Once Sumo is decommissioned (SL2DT-09), every Dynatrace config change must go through the same pipeline.

The New Baseline

- All monitors, dashboards, buckets, IAM, OpenPipeline configs in git
- Every change is a PR
- No UI clickops for production changes (UI for exploration only)

Enforcement

- CI blocks any PR with invalid DQL, schema errors, or bucket misconfig
- Required reviews: 1 engineer + 1 platform-team approver for prod
- Settings 2.0 schema drift detection: weekly automated diff between tenant state and git state

Drift Detection Workflow

```
# .github/workflows/drift-detect.yml
name: drift-detect
on:
  schedule:
    - cron: "0 6 * * 1" # Weekly Monday
jobs:
  drift:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - run: terraform plan -detailed-exitcode
      - run: monaco deploy --dry-run --report drift-report.md
      - if: failure()
        uses: actions/github-script@v7
        with:
          script: |
            github.rest.issues.create({
              owner: context.repo.owner, repo: context.repo.repo,
              title: 'Dynatrace config drift detected',
              body: 'See attached diff'
            });
```

9. Step Exit Criteria

- [] All buckets + IAM + OpenPipeline managed via Terraform
- [] Dashboards + monitors managed via Monaco
- [] CI pipeline runs on every PR: validate → plan → apply
- [] Secrets managed outside git (CI vault, not `.env` in repo)
- [] Preflight-check pattern used in all migration scripts
- [] Ansible / CloudFormation integration documented per source platform
- [] Drift detection scheduled + alerting on failure
- [] Rollback runbook in `runbooks/rollback.md`

Next step: SL2DT-09 — Cutover, Validation & Decommission (parallel run, validation gates, Sumo teardown).

11. References

Dynatrace automation and configuration-as-code

- [Configuration as code \(DT docs\)](#)
- [Terraform configuration \(DT docs\)](#)
- [Monaco configuration \(DT docs\)](#)
- [Dynatrace Terraform provider \(Terraform Registry\)](#)
- [Monaco \(Dynatrace GitHub\)](#)
- [API authentication \(DT docs\)](#)
- [Platform tokens \(DT docs\)](#)

Sumo Logic automation (source)

- [Sumo Logic Terraform provider \(Terraform Registry\)](#)
- [Sumo Logic API \(Sumo Logic docs\)](#)

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace or Sumo Logic. Always verify information against the official [Dynatrace documentation](#) and [Sumo Logic documentation](#).