

SL2DT-10: Migrating Telegraf-Collected Metrics

Series: SL2DT — Sumo Logic to Dynatrace | **Notebook:** 10 of 11 | **Created:** July 2026 | **Last Updated:** 07/20/2026

Overview

Goal of this step: bring a Telegraf-collected metric estate across to Dynatrace — inventory it, choose a target ingest path per source, reshape the metric keys, and prove parity during the parallel run.

Steps 03 through 06 cover the **log** path end to end. Metrics are a separate workstream, and a Sumo account whose operational visibility is largely Telegraf-collected metrics has no chapter in that sequence. The metric keys also change shape at the boundary **regardless of which path you choose**, so every metric query, dashboard panel, and monitor threshold needs rewriting.

Where this fits — read this as a parallel track, not a post-cutover step. This notebook is numbered 10 because it was added after the original nine steps, not because metric migration happens last. Run it alongside the log path: inventory with **SL2DT-02**, choose ingest paths alongside **SL2DT-03**, rewrite queries alongside **SL2DT-04**, and rebuild monitors and dashboards alongside **SL2DT-05** and **SL2DT-06**. Deferring it to after cutover means running two cutovers.

Table of Contents

1. [What You'll Produce](#)
 2. [Inventory the Metric Estate](#)
 3. [Choose the Target Ingest Path](#)
 4. [Telegraf to Dynatrace — Plugin Mechanics](#)
 5. [Metric-Key Reshaping at the Boundary](#)
 6. [Parity Validation During the Parallel Run](#)
 7. [Step Exit Criteria](#)
 8. [References](#)
-

Prerequisites

Requirement	Details
Audience	Migration lead + the team that owns the Telegraf configurations
Format	Procedural — produce four artifacts checked into the migration repo
Sumo access	Metrics sources readable; Telegraf <code>.conf</code> files retrievable from the hosts
Dynatrace access	Tenant with metric ingest available; for the remote route, a token with <code>metrics.ingest</code> scope
Prior reading	SL2DT-02 (inventory and cut scope), FAQ-11 (how metrics work in Dynatrace), OPIPE-04 (cardinality budgets)

1. What You'll Produce

Artifact	Format	Purpose
<code>inventory/metrics.json</code>	JSON	Every Sumo metrics source and Telegraf input: measurement, fields, tag keys, volume
<code>metric-key-map.md</code>	Markdown	Source metric name → Dynatrace metric key, per chosen path
<code>metric-cut-scope.md</code>	Markdown	Migrate / retire / defer per metric, with owner signoff
<code>parity-report.md</code>	Markdown	Sumo vs Dynatrace comparison across the parallel window

These gate the metric half of cutover in **SL2DT-09**. `metric-key-map.md` is the artifact that unblocks query rewriting in **SL2DT-04** — produce it early, because every downstream rewrite depends on it.

2. Inventory the Metric Estate

SL2DT-02 § 6 inventories collectors for the log path. Metric sources need the same treatment, but the questions differ: what matters is not only where the data comes from, but what each series will cost and what it will *key on* after the boundary.

Metric Source Types to Enumerate

Sumo-side source	What to capture
Telegraf via Installed Collector	Enabled <code>[[inputs.*]]</code> plugins, measurement names, field names, tag keys

Sumo-side source	What to capture
Hosted metrics source (HTTP)	Sending system, payload format, metric naming scheme
StatsD / Graphite source	Metric naming scheme, aggregation interval
Prometheus scrape	Scrape targets, job labels, relabeling rules
Host metrics from the Installed Collector	Whether OneAgent already covers the same signal — frequently it does

Per-Metric Record

For every metric series that survives cut scope, record:

- **Input plugin** and **measurement** name
- **Field names** — each field becomes its own Dynatrace metric key (see § 5)
- **Tag keys** — these become dimensions; count the distinct values per key
- **Cardinality estimate** — distinct tag-value combinations per measurement
- **Volume** — data points per minute
- **Consumers** — which Sumo dashboards, monitors, and scheduled searches read it

Cardinality is the cost driver. A measurement with only a handful of fields but one high-cardinality tag — a container ID, a request ID, a user ID — multiplies into a large number of series. Apply the cardinality budgets from **OPIPE-04** during inventory rather than after ingest; § 5 explains why correcting this later is a break rather than an edit.

The cheapest metric is the one you never migrate. Apply the same usage-based cut scope here that SL2DT-02 applies to dashboards and monitors: a metric that no dashboard or monitor has read in 90 days is a retire candidate, not a migration candidate.

3. Choose the Target Ingest Path

Four viable targets. They are ordered not by effort but by how much of the platform you still have afterwards.

Path	When it wins	What you get	What it costs
1. OneAgent native	The signal is one OneAgent already instruments — hosts, processes, common runtimes, many databases	Entities and topology, and Davis baselining, without ingesting anything custom	Deploy OneAgent; accept its metric names rather than your own

Path	When it wins	What you get	What it costs
2. Hub extension	A pre-built Dynatrace extension exists for the technology	A curated metric set with a supported configuration surface	Extension setup; confirm its metric set covers what your dashboards actually read
3. Keep Telegraf → <code>outputs.dynatrace</code>	You need the fastest cutover, or the source has no OneAgent or extension equivalent	Existing collection keeps working; ingest becomes a config change	Metric data points only — and the keys still change shape (§ 5)
4. OTel Collector → OTLP	You are standardizing on OpenTelemetry across signals anyway	One collection tier for metrics, traces, and logs	Delta temporality is mandatory (§ 4); naming differs again (§ 5)

Choosing a Target Path for Telegraf-Collected Metrics

Ordered by how much of the platform you keep afterwards — not by effort

Most platform value

Fastest to ship

1 OneAgent native
native instrumentation

WHEN IT WINS

OneAgent already instruments the technology

WHAT YOU GET

Entities, topology and Davis baselining

WHAT IT COSTS

Deploy OneAgent; accept its metric names

Prefer where it applies

2 Hub extension
pre-built package

WHEN IT WINS

A pre-built Dynatrace extension exists

WHAT YOU GET

Curated metric set on a supported config surface

WHAT IT COSTS

Setup; verify it covers what dashboards read

Check the Hub first

3 Keep Telegraf
`outputs.dynatrace`

WHEN IT WINS

Fastest cutover, or no OneAgent/extension path

WHAT YOU GET

Existing collection keeps working — a config change

WHAT IT COSTS

Metric data points only — no entity topology

Transition, not destination

4 OTel Collector
via OTLP

WHEN IT WINS

Standardizing on OpenTelemetry anyway

WHAT YOU GET

One collection tier for metrics, traces and logs

WHAT IT COSTS

Delta temporality is mandatory; keys differ

Verify keys per route

No path preserves your source metric names. Any route reshapes the key — map old to new before cutover (§ 5). Most estates end up mixing paths. That is a legitimate end state — but record a revisit date for anything parked on path 3.

Decide Per Source, Not Per Estate

Most migrations end up mixing paths — OneAgent for the host and runtime signals it already covers, an extension where one exists, and Telegraf for the long tail. That is a legitimate end state, not a failure to converge.

Path 3 is a transition state, not a destination — write that down explicitly. In community practice the most consistent trap is that "keep Telegraf" ships first, works, and is then never revisited. Ingested Telegraf metrics arrive as metric data points carrying dimensions; they do not by themselves produce the entity topology that OneAgent-detected technologies do. Before choosing a metrics-only path for a technology OneAgent could instrument, confirm what topology

and baselining you are giving up — and if you choose it anyway, record a revisit date in `metric-cut-scope.md`.

Related: [FAQ-14](#) makes the same argument for SQL Server — caller-side and server-side collection are complements, and a script-only posture (here, a Telegraf-only posture) is a transition rather than an end state.

4. Telegraf to Dynatrace — Plugin Mechanics

The `outputs.dynatrace` plugin ships with Telegraf. There are two endpoint modes, and the choice affects authentication **and** enrichment.

Endpoint Modes

Mode	url	Token	Automatic enrichment
Local OneAgent	empty, or <code>http://127.0.0.1:14499/metrics/ingest</code> (the default)	None required	Host ID and host name added as dimensions
Remote environment	<code>https://{env-id}.live.dynatrace.com/api/v2/metrics/ingest</code>	<code>metrics.ingest</code> scope	None — you supply dimensions yourself

Telegraf metric ingestion requires **OneAgent 1.201+**, and the default metric ingestion port is **14499**.

```
## Local OneAgent endpoint — no API token required
[[outputs.dynatrace]]
  ## Leave empty to use the local OneAgent endpoint
  ## (default: http://127.0.0.1:14499/metrics/ingest)
  url = ""
  prefix = "telegraf"

## Remote environment endpoint
[[outputs.dynatrace]]
  url = "https://{your-environment-id}.live.dynatrace.com/api/v2/metrics/ingest"
  api_token = "dt0c01.*****" # metrics.ingest scope only
  prefix = "telegraf"
```

Behaviours That Surprise People

Behaviour	Detail
Everything is a gauge by default	"All metrics are reported as gauges, unless they are specified to be delta counters using the <code>additional_counters</code> or <code>additional_counters_patterns</code> config option." A Telegraf counter you do not list lands as a gauge.
Batching is 1,000 lines per request	This is a client-side batching constant in the metric library the plugin uses — <i>not</i> an API ceiling. The ingest API limits the payload to 1 MB and states there is no limit on the number of metrics. Do not size your pipeline against 1,000.
<code>prefix</code> supplies its own separator	The plugin documents <code>prefix</code> as "prepended to all metric names (will be separated with a <code>.</code>)" — you do not add the dot.

`prefix` **trailing-dot trap**. The Dynatrace documentation page shows `prefix = "telegraf."` with a trailing dot, while the plugin's own sample configuration uses `prefix = "telegraf"` without one and states the separator is added automatically. Use the no-dot form, then **verify the resulting keys in your tenant** before building dashboards on them — a doubled dot is easy to ship and awkward to unpick after ingest.

5. Metric-Key Reshaping at the Boundary

No path preserves your source metric names. This is the most disruptive fact in a metric migration, and it is why `metric-key-map.md` gates the query rewrite in SL2DT-04.

How `outputs.dynatrace` Builds the Key

The plugin composes the Dynatrace metric key from the Telegraf **measurement** and **field**, joined with a dot:

```
metricName := tm.Name() + "." + field.Key
```

So a Telegraf `apache` input producing a `BusyWorkers` field arrives as:

Stage	Name
Telegraf measurement + field	<code>apache</code> + <code>BusyWorkers</code>
Dynatrace metric key	<code>apache.BusyWorkers</code>
With <code>prefix = "telegraf"</code>	<code>telegraf.apache.BusyWorkers</code>

Whatever name your Sumo dashboards referenced, compare it against this form and record **both** in `metric-key-map.md`. They will not match, and every panel and monitor threshold reading the old name needs rewriting.

Sourcing note: this composition is verified from the plugin **source**, not its README — the README documents the `prefix` option but never states the `measurement.field` rule. Cited accordingly in § 8.

Key Syntax Rules That Will Reject Your Data

Rule	Constraint
Allowed characters	Lowercase and uppercase letters, numbers, hyphens (-), underscores (_)
Non-Latin letters	Not allowed
Leading character	A key cannot start with a number or a hyphen; sections cannot start with a hyphen
Length	3–255 characters

The `dt.` prefix is reserved, and violations are dropped silently. "Metric data ingested with the `dt.` prefix will be dropped," and dimensions ingested with the `dt.` prefix "are ignored." If your measurement or your configured `prefix` produces a `dt.`-prefixed key, the data never arrives and nothing raises an error. Check this before the parallel run, not during it.

The OTLP Route Names Differently Again

If you choose path 4, do not assume its keys match path 3. OTLP ingest applies its own transformations: "a metric key may be suffixed automatically depending on the payload (for example, `.count` for counters and `.gauge` for gauges)," and invalid characters "will be replaced with underscores."

Verify, don't assume. The exact key a given Telegraf input produces *through an OpenTelemetry Collector* depends on the receiver in that collector, which sits outside Dynatrace's documentation. Send one representative metric through your chosen route into a non-production tenant and read the actual key back (§ 6) before committing a naming convention to `metric-key-map.md`.

6. Parity Validation During the Parallel Run

Metric parity is a different test from log parity. You are not diffing record counts — you are confirming that every key in `metric-key-map.md` exists, carries the dimensions you expected, and is populated

across the same window Sumo covers.

Step 1 — Confirm the Keys Arrived Under the Names You Expect

```
// List the metric keys that actually landed under your Telegraf prefix.  
// Compare against metric-key-map.md: anything missing never arrived, and  
// anything unexpected is a naming surprise worth catching before cutover.  
metrics  
| filter startsWith(metric.key, "telegraf.")  
| fields metric.key  
| sort metric.key asc  
| limit 200
```

Step 2 — Confirm a Migrated Key Is Actually Populated

Existence is not continuity. For every key a monitor or dashboard depends on, confirm data points arrive across the whole parallel window — not merely from the moment you enabled the output.

```
// Replace the metric key with one from metric-key-map.md.  
// A flat or gappy result means the Telegraf output is not running everywhere  
// you assume it is – check per host before declaring parity.  
timeseries datapoints = count(telegraf.apache.BusyWorkers), from:-24h, interval:1h  
| fieldsAdd total = arraySum(datapoints)  
| fields total
```

What "Parity" Means for Metrics

Check	Passing looks like
Key presence	Every key in <code>metric-key-map.md</code> returns from the query above
Dimension fidelity	The tag keys inventoried in § 2 are present as dimensions
Continuity	Data points across the full parallel window, not only since switch-on
Host coverage	The metric arrives from every host running the input, not a subset
Consumer rewrite	Every dashboard panel and monitor that read the old name now reads the new key

Record the outcome in `parity-report.md`. Metric parity failures are usually **coverage** failures — a Telegraf config rolled to some hosts but not all — and those stay invisible if you only check that the key exists.

7. Step Exit Criteria

Do not include metrics in the SL2DT-09 cutover until every box is checked:

Checkpoint	Status
<code>inventory/metrics.json</code> complete — every metric source, measurement, field, and tag key	[]
Cut scope applied — retire candidates identified and signed off by their owners	[]
Target path chosen per source, with every path-3 choice carrying a revisit date	[]
<code>metric-key-map.md</code> complete — old name → new key for every surviving metric	[]
No resulting key collides with the reserved <code>dt.</code> prefix	[]
Cardinality estimated against the OPIPE-04 budgets	[]
Every dependent dashboard panel and monitor threshold rewritten to the new keys	[]
<code>parity-report.md</code> shows key presence, dimension fidelity, continuity, and host coverage	[]
Rollback understood — Sumo metric collection still running until cutover is declared	[]

8. References

Dynatrace — metric ingestion

- [Send Telegraf metrics to Dynatrace \(DT docs\)](#) — requires OneAgent 1.201+; default ingestion port 14499
- [OneAgent metric API \(DT docs\)](#) — host ID and host name added automatically on the local endpoint
- [Metric ingestion protocol \(DT docs\)](#) — metric key character, section, and length rules
- [POST ingest metrics \(DT docs\)](#) — `metrics.ingest` scope; 1 MB payload limit
- [Metrics limits \(DT docs\)](#) — reserved `dt.` prefix; data ingested with it is dropped
- [About OTLP metrics ingest \(DT docs\)](#) — delta temporality requirement; automatic key suffixing

Telegraf — the `outputs.dynatrace` plugin

- [Dynatrace output plugin README \(InfluxData GitHub\)](#) — endpoint modes, `prefix`, gauge-by-default behaviour
- [dynatrace.go \(InfluxData GitHub\)](#) — the `measurement + "." + field` key composition

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace or Sumo Logic. Always verify information against the official [Dynatrace documentation](#) and [Sumo Logic documentation](#).