

SL2DT-99: Summary & Runbook Index

Series: SL2DT — Sumo Logic to Dynatrace | **Notebook:** 11 of 11 | **Created:** April 2026 | **Last Updated:** 10/05/2026

Overview

Goal of this notebook: single-page index for the entire SL2DT series. Use as a reference card during migration work, a hand-off document for incoming teams, and a quick-start for future Sumo-to-Dynatrace engagements.

This notebook is intentionally terse. Each section points to the authoritative detail in SL2DT-01 through SL2DT-10.

Table of Contents

1. [Migration Framework Summary](#)
 2. [Series Index — Which Notebook Does What](#)
 3. [Critical Decisions Checklist](#)
 4. [Artifacts Produced Across All Steps](#)
 5. [Failure-Mode Reference](#)
 6. [Quick Reference — Common DQL Patterns](#)
 7. [Engagement Timeline Template](#)
 8. [Companion Assets](#)
 9. [References](#)
-

Prerequisites

Requirement	Details
Audience	Anyone onboarding to a Sumo-to-Dynatrace engagement
Prior reading	None required — this notebook is the entry point

1. Migration Framework Summary

Sumo-to-Dynatrace in 9 steps + summary:

```
01 Overview      → strategy, mental model, anti-patterns
02 Inventory     → audit-log-driven cut scope
03 Ingest        → OneAgent + OTel + OpenPipeline + buckets
04 Translate     → SumoQL → DQL, confidence-scored
05 Monitors      → Dynatrace Intelligence / Workflow / Metric Event
decisions
06 Dashboards    → Notebooks (analytical) + Dashboards (operational)
07 Governance    → Platform IAM, bucket-scoped policies
08 Automation    → Terraform + Monaco + CI/CD
09 Cutover       → 3-tier validation, decommission
99 Summary       → this notebook
```

Typical duration: **3–6 months** for a mid-size deployment (100–500 active users, 2k–8k dashboards/monitors). Large enterprises (10k+ assets, multi-tenant) may need 9–12 months.

Target outcomes:

- Sumo contract terminated by deadline
- $\geq 95\%$ monitor + dashboard fidelity
- $\geq 40\%$ new monitors using anomaly detection (not static thresholds)
- Alert noise $\leq$ Sumo baseline

2. Series Index — Which Notebook Does What

#	Notebook	When to read	Primary artifacts
01	Overview & Migration Strategy	First, always	Strategy doc, risk register
02	Assessment & Inventory	Before any hands-on work	<code>inventory/*.json</code> , <code>cut-scope.md</code> , <code>taxonomy-map.md</code>
03	Log Ingest Architecture	Wave 1	Bucket configs, OpenPipeline pipelines, FER conversions
04	SumoQL → DQL Translation	Wave 1 / 2	<code>translations/all-queries.csv</code>
05	Monitor/Alert Conversion	Wave 2	Dynatrace Intelligence detectors, Workflows, Metric Events
06	Dashboard Conversion	Wave 3 / 4	Notebook + Dashboard JSON
07	User Governance & Access	Wave 1 (do early!)	IAM groups/policies/bindings
08	Automation & GitOps	Wave 1 ongoing	Terraform + Monaco configs, CI pipelines
09	Cutover, Validation & Decommission	Wave 5	Cutover runbook, validation report, retro
10	Telegraf Metric Migration	Parallel track — run alongside 02-06	<code>inventory/metrics.json</code> , <code>metric-key-map.md</code> , <code>parity-report.md</code>
99	Summary (this notebook)	Any time — reference card	—

3. Critical Decisions Checklist

Six decisions shape the entire migration. Make them early, document them in

`DECISIONS.md` :

Decision	Default	When to deviate
<code>_sourceCategory</code> strategy	Mixed: bucket for coarse isolation + attribute for query grouping	Bucket-only if retention tiers dominate; attribute-only if uniform retention
Ingest path per source	OneAgent for hosts, OTel Collector for cloud/HTTP	Sumo Forwarding if OneAgent can't be deployed
Monitor conversion	anomaly detection for metric-backed alerts	Static threshold only when truly static
Dashboard target	Notebooks for analytical, Dashboards for operational	Hybrid for SRE pages
Automation tool split	Terraform for IAM/buckets/pipelines, Monaco for dashboards/monitors	Single-tool if team mandates
Parallel-run duration	2–4 weeks	Longer only with documented cause (rarely a good sign)

4. Artifacts Produced Across All Steps

Cumulative list. All committed to the migration repo.

Artifact	Step	Purpose
<code>inventory/*.json</code>	SL2DT-02	Asset inventory
<code>inventory/audit-usage.csv</code>	SL2DT-02	Usage data for cut scope
<code>taxonomy-map.md</code>	SL2DT-02	<code>_sourceCategory</code> → bucket/attribute
<code>cut-scope.md</code>	SL2DT-02	Per-asset decision + owner signoff
<code>terraform/buckets.tf</code>	SL2DT-03	Bucket configs

Artifact	Step	Purpose
<code>openpipeline/*.json</code>	SL2DT-03	Pipeline + processor configs
<code>translations/all-queries.csv</code>	SL2DT-04	Per-query DQL + confidence
<code>monitors/decision-matrix.csv</code>	SL2DT-05	Dynatrace Intelligence/Workflow/Metric-event decision
<code>dashboards/*</code>	SL2DT-06	Notebook/Dashboard JSON
<code>iam/groups.tf</code> , <code>policies.tf</code>	SL2DT-07	IAM as code
<code>.github/workflows/*.yaml</code>	SL2DT-08	CI pipelines
<code>cutover/runbook.md</code>	SL2DT-09	Cutover step-by-step
<code>cutover/validation-report.md</code>	SL2DT-09	3-tier validation results
<code>cutover/post-cutover-retro.md</code>	SL2DT-09	Retro after 30-day stable window

5. Failure-Mode Reference

When things go wrong, consult this table first:

Symptom	Root cause category	Fix (notebook)
App teams recreating static thresholds in DT	Dynatrace Intelligence not explained	SL2DT-05 §3
Bucket writes to wrong bucket	OpenPipeline matcher wrong	SL2DT-03 §4
Queries time out	No <code>from:</code> or too wide range	SL2DT-04 §3 Rule 1

Symptom	Root cause category	Fix (notebook)
Sort/filter on timeseries fails	Array not reduced	SL2DT-04 §3 Rule 3
Parse returns null	DPL syntax wrong	SL2DT-03 §8, sumoql-to-dql skill
Users can't see their data	Bucket policy wrong	SL2DT-07 §3
ServiceNow incidents not created	Webhook payload not rebuilt	SL2DT-05 §6
Alert volume exploded	baselines still learning	SL2DT-05 §8
Dashboard blank after migration	Time range or bucket mismatch	SL2DT-06 §3
SSO assigns no permissions	Claim → group mapping wrong	SL2DT-07 §4

6. Quick Reference — Common DQL Patterns

Count logs per scope

```
// 6. Quick Reference – Common DQL Patterns
fetch logs, from:-1h
| summarize c = count(), by:{sumo.source_category}
| sort c desc
```

Top errors by host

```
// Top errors by host
// status == "ERROR" covers ERROR, SEVERE, CRITICAL, ALERT, FATAL and
// EMERGENCY; loglevel == "ERROR" misses SEVERE.
fetch logs, from:-1h
| filter sumo.source_category == "prod/api"
| filter status == "ERROR"
| summarize c = count(), by:{host.name}
| sort c desc
| limit 10
```

Error rate timeseries

```
// Error rate timeseries
fetch logs, from:-1h
| filter sumo.source_category == "prod/api"
| fieldsAdd is_error = if(status == "ERROR", 1, else:0)
| makeTimeseries {total = count(), errors = sum(is_error)}, interval:1m
| fieldsAdd error_pct = 100.0 * errors[] / total[]
```

Latency percentiles

```
// Latency percentiles
fetch logs, from:-1h
| filter sumo.source_category == "prod/api"
| parse content, "LD? 'latency=' INT:latency"
| summarize {
    p50 = percentile(latency, 50),
    p95 = percentile(latency, 95),
    p99 = percentile(latency, 99)
}, by:{http.path}
```

Recent detected problems

```
// Recent detected problems (dt.davis.problems: one record per problem)
fetch dt.davis.problems, from:-24h
| fields event.start, display_id, event.name, event.category, event.status
| sort event.start desc
| limit 20
```

Timeseries with array reduction


```
// Pattern when you need to filter/sort after makeTimeseries
timeseries avg_cpu = avg(dt.host.cpu.usage), from:-1h, by:{dt.entity.host}
| fieldsAdd max_cpu = arrayMax(avg_cpu)
| filter max_cpu > 80
| sort max_cpu desc
```

Lookup for enrichment

```
// Lookup for enrichment
//
// Corrected 08/12/2026: `dt.tags.<name>` is not a field. Entity tags
are a single `tags` field
// holding an ARRAY of "key:value" strings, so a tag is matched with an
iterative expression -
// iAny(startsWith(tags[], "team:")) - not addressed as a nested field. The
old cell failed with
// FIELD_DOES_NOT_EXIST. Pass the whole `tags` array through the lookup and
read it downstream.
fetch logs, from:-1h
| filter sumo.source_category == "prod/api"
| summarize c = count(), by:{host.name}
| lookup [fetch dt.entity.host, from:-7d
        | fieldsAdd hostname = entity.name
        | fields hostname, tags],
        sourceField:host.name, lookupField:hostname,
        fields:{tags}
```

7. Engagement Timeline Template

Scaled to a 6-month mid-size migration. Adjust for your scope.



Parallel tracks that can overlap

- IAM (SL2DT-07) runs in parallel with ingest (SL2DT-03)
- Translation (SL2DT-04) can start once taxonomy is mapped, before ingest is complete
- Automation (SL2DT-08) runs continuously from Month 2 onward

8. Companion Assets

Skill

- `sumoql-to-dql` skill — translation tables, confidence scoring (internal authoring aid, not a published asset)
- `references/mapping-tables.md` — complete operator/function/field maps
- `references/examples.md` — 15 worked translations
- `references/out-of-scope.md` — features not covered in v1

Parallel series

- **NR2DT** — New Relic → Dynatrace procedural
- **NRLC** — New Relic component deep dives
- **S2D** — Splunk → Dynatrace
- **M2S** — Managed → SaaS
- **S2S** — SaaS → SaaS

Official documentation

- [Sumo Logic Docs](#)
 - [Dynatrace Docs](#)
-

9. References

Dynatrace — by topic

Platform / DQL / Grail

- [DQL Reference \(DT docs\)](#)
- [DPL Pattern Reference \(DT docs\)](#)
- [Grail \(DT docs\)](#)
- [OpenPipeline \(DT docs\)](#)

Ingest

- [Ingest from \(DT docs\)](#)
- [Dynatrace OneAgent \(DT docs\)](#)
- [OpenTelemetry ingest \(DT docs\)](#)
- [Logs \(DT docs\)](#)
- [Dynatrace API \(DT docs\)](#)

Intelligence and workflows

- [Anomaly detection \(DT docs\)](#)
- [Problems app \(DT docs\)](#)
- [Workflows \(DT docs\)](#)
- [Notifications and alerting \(DT docs\)](#)
- [Maintenance windows \(DT docs\)](#)

Visualization

- [Dashboards and notebooks \(DT docs\)](#)
- [Notebooks \(DT docs\)](#)

Access and tokens

- [Identity and access management \(DT docs\)](#)
- [Manage user permissions and policies \(DT docs\)](#)
- [Access tokens \(DT docs\)](#)
- [Platform tokens \(DT docs\)](#)
- [API authentication \(DT docs\)](#)

Automation

- [Configuration as code \(DT docs\)](#)
- [Terraform configuration \(DT docs\)](#)
- [Monaco configuration \(DT docs\)](#)
- [Dynatrace Terraform provider \(Terraform Registry\)](#)
- [Monaco \(Dynatrace GitHub\)](#)

Sumo Logic — by topic

- [Sumo Logic search operators \(Sumo Logic docs\)](#)
- [Sumo Logic parse operators \(Sumo Logic docs\)](#)
- [Sumo Logic monitors \(Sumo Logic docs\)](#)
- [Sumo Logic dashboards \(Sumo Logic docs\)](#)
- [Sumo Logic users and roles \(Sumo Logic docs\)](#)
- [Sumo Logic API \(Sumo Logic docs\)](#)
- [Sumo Logic Terraform provider \(Terraform Registry\)](#)

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace or Sumo Logic. Always verify information against the official [Dynatrace documentation](#) and [Sumo Logic documentation](#).