

SLO-01: SLO and SLI Fundamentals

Series: SLO — Service Level Objectives | **Notebook:** 1 of 5 | **Created:** June 2026 |
Last Updated: 07/01/2026

Overview

Service Level Objectives turn "is the service healthy?" from a matter of opinion into a measured, agreed number. This notebook defines the three building blocks — **SLI**, **SLO**, **error budget** — explains how Dynatrace models them in the modern SLO app, and helps you choose your first handful of SLOs.

This is the conceptual foundation for the series. SLO-02 builds the SLI queries, SLO-03 covers composition and error budgets in depth, SLO-04 covers alerting, and SLO-05 covers provisioning SLOs as code.

Table of Contents

- [1. The Three Building Blocks](#)
 - [2. Why SLOs Matter](#)
 - [3. The Dynatrace SLO Model](#)
 - [4. Anatomy of an SLO](#)
 - [5. Choosing Your First SLOs](#)
 - [6. Cross-Series Pointers](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS Gen3 with Grail and the SLO app
IAM permissions (SLO app itself)	<code>slo:slos:read</code> , <code>slo:slos:write</code> , <code>slo:objective-templates:read</code> — the gating permissions for the app; a minimal read/write grant is <code>ALLOW slo:slos:read, slo:objective-templates:read; ALLOW slo:slos:write;</code>
Data-source permissions	<code>storage:metrics:read</code> , <code>storage:spans:read</code> , plus whichever of <code>storage:logs:read</code> / <code>storage:bizevents:read</code> / <code>storage:events:read</code> / <code>storage:buckets:read</code> your SLI queries touch; <code>settings:objects:read/write</code> to create SLOs
Concepts	Basic DQL (see the SPANS and DBMON series); familiarity with your service topology

1. The Three Building Blocks

Term	What it is	Example
SLI — Service Level Indicator	A measured ratio of <i>good</i> events to <i>total</i> events, expressed as a percentage	99.61% of requests succeeded
SLO — Service Level Objective	A target for an SLI over a defined window	availability $\geq$ 99.5% over 30 days
Error budget	The allowed shortfall: <code>100% - target</code>	0.5% of requests may fail before the SLO is breached

An **SLI** answers *how are we doing right now*. An **SLO** sets *how good is good enough*. The **error budget** converts the gap into a spendable quantity — once it is gone, the SLO is breached, and that is the signal teams act on.

A note on SLAs: an **SLA** (Service Level Agreement) is a contract with consequences (credits, penalties). SLOs are your internal targets, and you almost always set them *stricter* than any SLA so you find out before the customer does.

2. Why SLOs Matter

- **They translate technical health into business language.** "p95 checkout latency under 1s, 99.9% of the month" is something a product owner can reason about.
- **They give you an error budget.** A budget reframes reliability from "zero errors" (impossible, paralysing) to "errors are fine until the budget runs low" — which then drives concrete decisions: keep shipping, or slow down and harden.
- **They are the highest-signal alert source.** Alerting on *burn rate* against the budget catches both the acute outage and the slow leak, with far less noise than threshold alerts on raw metrics (SLO-04).

3. The Dynatrace SLO Model

Dynatrace has two generations of SLO tooling:

	SLO Classic	Modern SLO app (use this)
SLI definition	Metric expressions, fixed indicator types	DQL-based — any <code>good ÷ total</code> query over metrics, spans, logs, or bizevents
API	Service-level Objectives API classic (<code>/api/v2/slo</code>)	SLO Service Public API
Backing store	classic config; Settings 2.0 schema <code>builtin:monitoring.slo</code>	the SLO service itself, not a Settings schema
As code	<code>dynatrace_slo_v2</code> Terraform resource (metric-selector SLIs only)	<code>dynatrace_platform_slo</code> Terraform resource (provider v1.78.0+) / Monaco (v2.22+)
Data scope	service metrics	all Grail data types

A correction worth reading if you have codified SLOs already (07/31/2026). This notebook and SLO-05 previously placed `builtin:monitoring.slo` and `dynatrace_slo_v2` in the modern column. They belong to the classic generation: `dynatrace_slo_v2` authenticates with classic `slo.read / slo.write` plus `settings.*` scopes and writes through `builtin:monitoring.slo`, whereas the

modern app has its own **SLO Service Public API** and a separate `dynatrace_platform_slo` resource using OAuth `slo:slos:*` scopes. The readiness scan additionally flags both `/api/v2/slo` and `builtin:monitoring.slo` as blocked at upgrade. If your SLOs are in Terraform as `dynatrace_slo_v2`, they are on the classic path — see SLO-05 §2.

New work should use the **modern SLO app** with DQL-based SLIs — it covers every data type and is the path that is actively developed. Dynatrace's own guidance is to still prefer a metric-based SLI over a DQL one where a suitable pre-aggregated metric exists, since metric queries are faster and cheaper to evaluate; use `makeTimeseries` for event/log-derived SLIs where no such metric exists (SLO-02). If you have SLO Classic definitions, Dynatrace publishes an upgrade path.



4. Anatomy of an SLO

Every modern Dynatrace SLO is four decisions:

1. **The SLI query** — a DQL query that returns the good-versus-total ratio. This is the bulk of the work and is covered in SLO-02.
2. **The target** — the percentage you commit to (e.g. `99.5`). Optionally a **warning** level above the target that surfaces "getting close" before an actual breach.
3. **The evaluation window** — *rolling* (e.g. last 30 days, always moving) or *calendar* (this month). Rolling windows give smoother trend signals; calendar windows align to business reporting. SLO-03 covers the trade-off.

4. **What happens on breach** — burn-rate alerting and routing, covered in SLO-04.

The error budget is not a separate decision — it falls out of the target ($100\% - \text{target}$).
The art is in the first three choices.

5. Choosing Your First SLOs

The most common failure mode is too many SLOs, none trusted. Start narrow:

- **Pick 2–4 business-critical services** — the user journeys that matter (checkout, login, search), not every microservice.
- **1–2 SLIs per service** — availability plus either latency or error rate. Resist the urge to measure everything.
- **Set achievable targets first.** Measure the current SLI for a week, then set the target slightly below observed performance. A target you breach on day one teaches the team to ignore SLOs.
- **Tighten over time.** SLOs are living guidance — review quarterly, ratchet up as the service matures.

Anti-pattern: copying "99.9%" onto every service because it sounds rigorous. A target must reflect what the service can actually deliver and what its users actually need.

6. Cross-Series Pointers

- **SLO-02** — building the SLI queries (availability, latency, error rate, custom).
- **SLO-03** — error budgets, burn rate, composite and weighted SLOs.
- **SLO-04** — burn-rate alerting and routing SLO breaches.
- **SLO-05** — provisioning SLOs as code (`dynatrace_slo_v2` , Monaco).
- **AIOPS-02** — anomaly detection; SLO breaches and anomalies are complementary detection inputs.
- **WFLOW** — routing the problems an SLO breach raises.

Sources: [Service-Level Objectives \(DT docs\)](#), [SLOs for all data types \(Dynatrace News\)](#).

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.