WEBRUM-06: Performance Analysis

> **Series:** WEBRUM — Web Real User Monitoring | **Notebook:** 6 of 10 | **Created:** March 2026 | **Last Updated:** 10/31/2026

Overview

Web performance directly impacts user experience, conversion rates, and SEO rankings. While Core Web Vitals (covered in WEBRUM-03) provide high-level scoring, a deeper performance analysis requires understanding the full page load waterfall — from DNS lookup to load complete — and how performance varies across geographies, network conditions, and device types.

Table of Contents

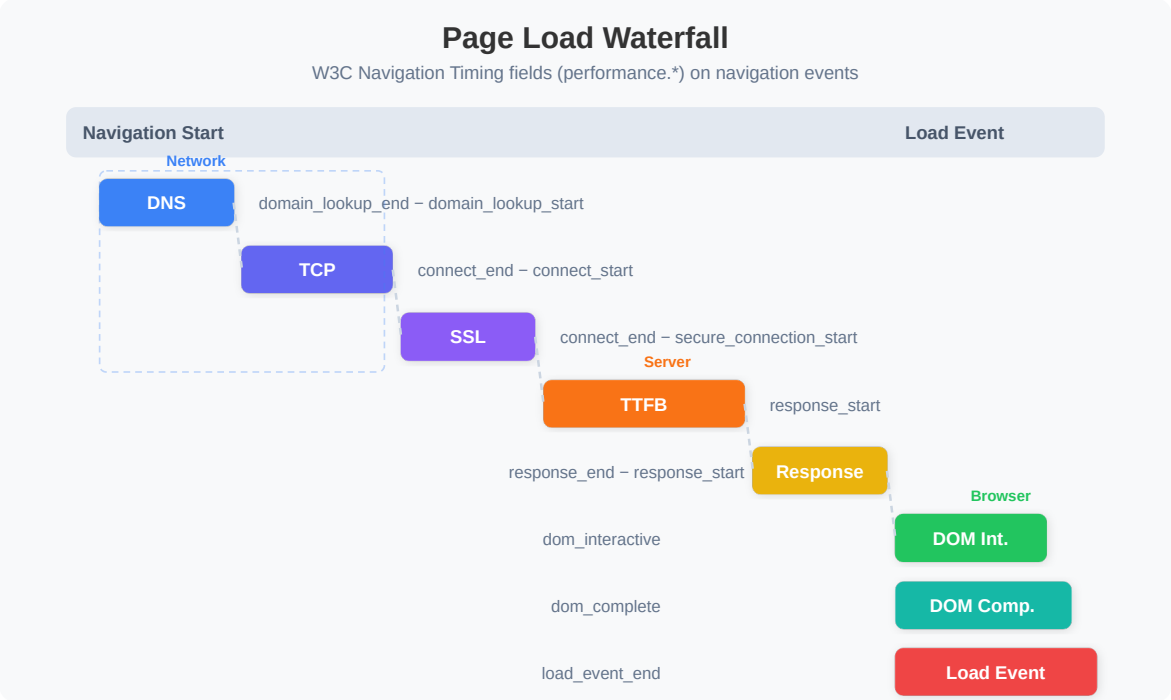
1. Page Load Waterfall
2. Time to First Byte (TTFB)
3. DOM Interactive and Load Event
4. Performance by Geography
5. Performance by Network and Device
6. Identifying Slow Pages
7. Performance Trends
8. Summary and Next Steps

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Grail enabled
RUM Enabled	Web applications with detailed timing capture
Permissions	`storage:user.events:read`, `storage:user.sessions:read`
Previous Notebooks	WEBRUM-01: RUM Fundamentals, WEBRUM-03: Core Web Vitals

1. Page Load Waterfall

Every page load follows a sequence of phases captured by the browser's Navigation Timing API. Dynatrace records these milestones for each load action:



Timing Milestones Explained

Milestone	Description	Optimization Focus
DNS time	Domain name resolution	Use DNS prefetching, reduce DNS lookups
TCP connect	Establishing TCP connection	Enable HTTP/2, use CDN edge servers
SSL time	TLS handshake	Enable TLS 1.3, OCSP stapling
TTFB	Server processing + first byte transit	Optimize server response time
Response time	Full HTML download	Compress, minimize HTML size
DOM interactive	HTML parsed, DOM ready for interaction	Defer non-critical JS, reduce blocking resources
DOM complete	All sub-resources loaded	Lazy load images, async load scripts
Load event	`window.onload` fires	Final milestone; all resources ready

Where the Timings Live

All of these milestones are fields on **navigation events** — select them with `characteristics.has_w3c_navigation_timings == true`. Each `performance.*` value is a duration measured from the start of the navigation, so `performance.load_event_end` is the page-load time. Dynatrace's own page-load query uses exactly this field and filter.

Do **not** measure page load with `duration` on a page summary. A page summary covers the page's whole life — *"A page instance begins with each hard navigation and ends when the next document request starts or when the tab is closed"* — so its `duration` is how long the page stayed open, not how long it took to load. On the synthetic-only validation tenant (24 h, 10/05/2026) the median page-summary `duration` was 4.7 s against a median load time of 0.36 s; on real-user traffic, where tabs stay open for minutes, the gap is far larger.

> **Sources:** Monitor web performance with DQL (DT docs) — *"RUM captures page load timings from the W3C Navigation Timing API as built-in metrics and as fields on navigation events."*; Navigation-

> related events — semantic dictionary (DT docs) — *"A page instance begins with each hard navigation and ends when the next document request starts or when the tab is closed."*; Requests — semantic dictionary (DT docs) — *"The end time of the load event handler phase."* **Dictionary:** model `rum_page_summary` — *"A page instance begins with each hard navigation and ends when the next document request starts or when the tab is closed."*, read 10/05/2026.

```
// Page load waterfall — average phase timings for the top 10 pages, all from
navigation events.
// The server phase is performance.response_start (TTFB); the previous
`avg(ttfb.waiting_duration)`
// column read the pre-request wait/redirect phase, and is empty on navigation
events anyway.
fetch user.events, from:-24h
| filter characteristics.has_w3c_navigation_timings == true
| summarize {page_views = count(),
    avg_dns_ms = avg(performance.domain_lookup_end -
performance.domain_lookup_start) / 1ms,
    avg_connect_ms = avg(performance.connect_end - performance.connect_start)
/ 1ms,
    avg_ttfb_ms = avg(performance.response_start) / 1ms,
    avg_download_ms = avg(performance.response_end -
performance.response_start) / 1ms,
    avg_dom_interactive_ms = avg(performance.dom_interactive) / 1ms,
    avg_load_event_ms = avg(performance.load_event_end) / 1ms},
    by:{page.detected_name}
| sort page_views desc
| limit 10
```

2. Time to First Byte (TTFB)

TTFB measures the time from the browser sending the request to receiving the first byte of the response. It reflects:

- **Server processing time** — How long the backend takes to generate the response
- **Network latency** — Round-trip time between client and server
- **CDN performance** — Cache hit/miss at the edge

Google recommends a TTFB of **≤ 800ms** for a good user experience.

In New RUM the TTFB value is `web_vitals.time_to_first_byte` (stable, a duration). Where it is not populated, `ttfb.value` (experimental, in milliseconds) carries *"The*

responseStart value of the navigation timing" — the same measurement. The queries below read the stable field and fall back to the experimental one. They do **not** use `ttfb.waiting_duration`: that field is *"The total time from when the user initiates loading to when the navigation request is handled"* — the wait before the request, usually redirects — and reading it as TTFB under-reports by an order of magnitude.

> **Sources:** Navigation-related events — semantic dictionary (DT docs) — *"The responseStart value of the navigation timing."*, *"Long waiting durations are usually caused by HTTP redirects."*; Monitor web performance with DQL (DT docs) (computes TTFB from `web_vitals.time_to_first_byte`).

```
// TTFB corrected 10/05/2026: coalesce(web_vitals.time_to_first_byte / 1ms,
ttfb.value) — the
// stable duration field where populated, else the experimental responseStart
value (already ms).
// ttfb.waiting_duration is the pre-request wait/redirect phase, not TTFB.
// TTFB analysis by page — identify pages with slow server response
fetch user.events, from:-24h
| filter characteristics.has_page_summary == true
| fieldsAdd ttfb_ms = coalesce(web_vitals.time_to_first_byte / 1ms,
ttfb.value)
| filter isNotNull(ttfb_ms)
| summarize {page_views = count(), avg_ttfb = avg(ttfb_ms),
    p75_ttfb = percentile(ttfb_ms, 75), p95_ttfb = percentile(ttfb_ms, 95)},
    by:{page.detected_name}
| filter page_views &gt; 10
| sort p75_ttfb desc
| limit 10
```

```
// TTFB corrected 10/05/2026: coalesce(web_vitals.time_to_first_byte / 1ms,
ttfb.value) — the
// stable duration field where populated, else the experimental responseStart
value (already ms).
// ttfb.waiting_duration is the pre-request wait/redirect phase, not TTFB.
// TTFB distribution — classify into Good / Needs Improvement / Poor
fetch user.events, from:-24h
| filter characteristics.has_page_summary == true
| fieldsAdd ttfb_ms = coalesce(web_vitals.time_to_first_byte / 1ms,
ttfb.value)
| filter isNotNull(ttfb_ms)
| fieldsAdd ttfb_category = if(ttfb_ms &lt;= 800, "Good",
    else: if(ttfb_ms &lt;= 1800, "Needs Improvement",
    else: "Poor"))
| summarize action_count = count(), by:{ttfb_category}
| sort action_count desc
```

3. DOM Interactive and Load Event

DOM Interactive is when the HTML has been fully parsed and the DOM tree is ready — but images, stylesheets, and sub-frames may still be loading. This is when users can first interact with the page.

Load Event fires when all resources (images, scripts, stylesheets, iframes) have finished loading. The gap between DOM Interactive and Load Event reveals how much time is spent on sub-resource loading.

```
// DOM Interactive vs Load Event — identify resource-heavy pages
fetch user.events, from:-24h
| filter characteristics.has_w3c_navigation_timings == true
| filter isNotNull(performance.dom_interactive) and
isNotNull(performance.load_event_end)
| summarize {page_views = count(),
    avg_dom_interactive = avg(performance.dom_interactive),
    avg_load_event = avg(performance.load_event_end)},
    by:{page.detected_name}
| fieldsAdd resource_load_gap = avg_load_event - avg_dom_interactive
| filter page_views &gt; 10
| sort resource_load_gap desc
| limit 10
```

A large gap between DOM Interactive and Load Event suggests the page has many sub-resources (large images, heavy scripts, third-party tags) that delay full load completion.

Consider lazy loading and async script loading to reduce this gap.

4. Performance by Geography

Performance varies significantly by user location due to network latency, CDN coverage, and server proximity.

```
// Page-load time corrected 10/05/2026: performance.load_event_end on
navigation events
// (has_w3c_navigation_timings). A page summary's `duration` is how long the
page was OPEN.
// Page load performance by country — identify slow regions (TTFB by region:
next query)
fetch user.events, from:-24h
| filter characteristics.has_w3c_navigation_timings == true
| filter isNotNull(geo.country.name)
| summarize {page_views = count(),
    avg_load_ms = avg(performance.load_event_end) / 1ms,
    p75_load_ms = percentile(performance.load_event_end, 75) / 1ms},
    by:{geo.country.name}
| filter page_views &gt; 20
| sort p75_load_ms desc
| limit 15
```

```
// TTFB corrected 10/05/2026: coalesce(web_vitals.time_to_first_byte / 1ms,
ttfb.value) — the
// stable duration field where populated, else the experimental responseStart
value (already ms).
// ttfb.waiting_duration is the pre-request wait/redirect phase, not TTFB.
// Compare TTFB across regions — CDN effectiveness indicator
fetch user.events, from:-24h
| filter characteristics.has_page_summary == true
| filter isNotNull(geo.continent.name)
| fieldsAdd ttfb_ms = coalesce(web_vitals.time_to_first_byte / 1ms,
ttfb.value)
| summarize {page_views = count(),
    avg_ttfb_ms = avg(ttfb_ms),
    p75_ttfb_ms = percentile(ttfb_ms, 75)},
    by:{geo.continent.name}
| sort p75_ttfb_ms desc
```

> **Tip:** If TTFB is consistently high for specific regions, consider deploying CDN edge servers or regional server instances closer to those users.

5. Performance by Network and Device

Network connection type and device capability significantly impact perceived performance.

Web RUM has no connection-type field to split by: `network.connection.type` is documented as *"Only supported by OneAgent for Mobile"*, and `network.protocol.name` is the OSI application protocol (`http`), not wifi versus cellular. The queries below break page-load time down by browser and operating system instead.

> **Sources:** User events — semantic dictionary (DT docs) — *"The internet connection type. Only supported by OneAgent for Mobile."* **Dictionary:** `network.protocol.name` (`stable`) — *"OSI Application Layer or non-OSI equivalent."*, read 10/05/2026.

```
// Page-load time corrected 10/05/2026: performance.load_event_end on
navigation events
// (has_w3c_navigation_timings). A page summary's `duration` is how long the
page was OPEN.
// Performance by browser — which browsers are slowest?
fetch user.events, from:-24h
| filter characteristics.has_w3c_navigation_timings == true
| filter isNotNull(browser.name)
| summarize {page_views = count(),
    avg_load_ms = avg(performance.load_event_end) / 1ms,
    p75_load_ms = percentile(performance.load_event_end, 75) / 1ms},
    by:{browser.name}
| filter page_views &gt; 20
| sort p75_load_ms desc
| limit 10
```

```
// Page-load time corrected 10/05/2026: performance.load_event_end on
navigation events
// (has_w3c_navigation_timings). A page summary's `duration` is how long the
page was OPEN.
// Performance by OS — desktop vs mobile operating systems (os.name; there is
no os.family field)
fetch user.events, from:-24h
| filter characteristics.has_w3c_navigation_timings == true
| filter isNotNull(os.name)
| summarize {page_views = count(),
    avg_load_ms = avg(performance.load_event_end) / 1ms,
    p75_load_ms = percentile(performance.load_event_end, 75) / 1ms},
    by:{os.name}
| filter page_views &gt; 20
| sort p75_load_ms desc
```

6. Identifying Slow Pages

Find the pages that need optimization attention — ranked by the impact of their slowness (volume x load time).

```
// Page-load time corrected 10/05/2026: performance.load_event_end on
navigation events
// (has_w3c_navigation_timings). A page summary's `duration` is how long the
page was OPEN.
// Slowest pages by p95 load time — worst-case performance
fetch user.events, from:-24h
| filter characteristics.has_w3c_navigation_timings == true
| summarize {page_views = count(),
    avg_load_ms = avg(performance.load_event_end) / 1ms,
    p75_load_ms = percentile(performance.load_event_end, 75) / 1ms,
    p95_load_ms = percentile(performance.load_event_end, 95) / 1ms},
    by:{page.detected_name}
| filter page_views &gt; 20
| sort p95_load_ms desc
| limit 10
```

```
// Page-load time corrected 10/05/2026: performance.load_event_end on
navigation events
// (has_w3c_navigation_timings). A page summary's `duration` is how long the
page was OPEN.
// Weighted impact score — pages with high traffic AND slow loads
fetch user.events, from:-24h
| filter characteristics.has_w3c_navigation_timings == true
| summarize {page_views = count(),
    avg_load_ms = avg(performance.load_event_end) / 1ms},
    by:{page.detected_name}
| fieldsAdd impact_score = page_views * avg_load_ms
| sort impact_score desc
| limit 10
```

> **Tip:** The impact score (page views x average load time) helps prioritize optimization efforts. A moderately slow page with high traffic may be more impactful than a very slow page with few visitors.

7. Performance Trends

Track performance over time to detect regressions and measure the impact of optimizations.

```
// Page-load time corrected 10/05/2026: performance.load_event_end on
navigation events
// (has_w3c_navigation_timings). A page summary's `duration` is how long the
page was OPEN.
// Page load time trend — daily p75 over the last 7 days
fetch user.events, from:-7d
| filter characteristics.has_w3c_navigation_timings == true
| fieldsAdd load_ms = performance.load_event_end / 1ms
| makeTimeseries p75_load_ms = percentile(load_ms, 75), interval:24h
```

```
// TTFB corrected 10/05/2026: coalesce(web_vitals.time_to_first_byte / 1ms,
ttfb.value) — the
// stable duration field where populated, else the experimental responseStart
value (already ms).
// ttfb.waiting_duration is the pre-request wait/redirect phase, not TTFB.
// TTFB trend — hourly p75 over the last 24 hours
fetch user.events, from:-24h
| filter characteristics.has_page_summary == true
| fieldsAdd ttfb_ms = coalesce(web_vitals.time_to_first_byte / 1ms,
ttfb.value)
| filter isNotNull(ttfb_ms)
| makeTimeseries p75_ttfb = percentile(ttfb_ms, 75), interval:1h
```

8. Summary and Next Steps

In this notebook, we covered:

- **Page load waterfall** — Full timing breakdown from DNS to load complete, all on navigation events
- **TTFB analysis** — Server response time measurement and classification
- **DOM timing** — Interactive vs complete timing for resource load gap analysis
- **Geographic performance** — Regional and continental performance differences
- **Device performance** — Impact of browser and OS (web RUM has no connection-type field)
- **Slow page identification** — Ranking pages by p95 load time and impact score
- **Performance trends** — Time-series tracking for regression detection

Next Steps

- **WEBRUM-07: Session Replay** — Visually investigate slow page experiences
- **WEBRUM-08: Dashboards and Alerting** — Build operational RUM dashboards with Apdex

References

- Dynatrace Performance Analysis
- Navigation Timing API
- TTFB Best Practices

- User events — semantic dictionary (DT docs) — *"Used for internal optimization when storing the data and not intended for query usage."*
- Semantic Dictionary changelog 1.349 (DT docs)

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.