

WFLOW-01: Workflow Fundamentals

Series: WFLOW — Workflows and Alert Notifications | **Notebook:** 1 of 10 | **Created:** January 2026 | **Last Updated:** 10/06/2026

Introduction to Dynatrace Workflows

Dynatrace Workflows is the automation engine that enables event-driven automation, scheduled tasks, and integration orchestration. This notebook introduces core concepts, components, and your first workflow.

Table of Contents

- [1. What Are Workflows?](#)
 - [2. Workflows vs Other Automation](#)
 - [3. Workflow Components](#)
 - [4. Accessing Workflows](#)
 - [5. Execution Model](#)
 - [6. Your First Workflow](#)
 - [7. Viewing Execution History](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Platform subscription

Requirement	Details
Permissions	<code>app-engine:apps:run</code> , <code>automation:workflows:read</code> , <code>automation:workflows:write</code> , <code>automation:workflows:run</code> , <code>app-engine:functions:run</code> ; for the §7 queries, <code>storage:system:read</code> and <code>storage:buckets:read</code> on <code>dt.system.events</code> (§4)
Prior Knowledge	Basic Dynatrace navigation

1. What Are Workflows?

Workflows are automated sequences of tasks that execute in response to events, schedules, or manual triggers. They enable:

Capability	Example
Alert Notifications	Send Slack/Teams messages when problems occur
Incident Management	Create PagerDuty/ServiceNow tickets automatically
Auto-Remediation	Restart pods, scale resources, run scripts
Reporting	Generate and email daily reports
Integration	Sync data with external systems
Data Enrichment	Query APIs and enrich events

Key Benefits

- **No Infrastructure** - Runs in Dynatrace, no external servers needed
- **Event-Driven** - Triggers on detected problems, Davis events, other Grail events, schedules
- **Low-Code + Code** - Visual builder with JavaScript option
- **Secure** - Built-in secrets management, RBAC, audit logs
- **Observable** - Execution history, metrics, debugging

2. Workflows vs Other Automation

Dynatrace offers multiple automation approaches. When should you use Workflows?

Approach	Best For	Limitations
Workflows	Event-driven automation, notifications, integrations	Rate limits, execution time limits
Alerting profiles + problem notifications (Classic)	The profile filters problems; the problem notification delivers them (email, webhook, integrations)	One alerting profile per notification, no logic
Apps (AppEngine)	Custom UI, complex applications	Requires development expertise
Extensions	Data collection, custom metrics	Not for automation/notifications
Site Reliability Guardian	Release validation, SLO verification	Specific to deployment validation

Migration from Legacy Alerting

Legacy Alerting	Workflow Equivalent
Problem notification → Email	Detected Problem trigger → Email task
Problem notification → Webhook	Detected Problem trigger → HTTP Request task
Custom integration	Detected Problem trigger → JavaScript + HTTP

Recommendation: New implementations should use Workflows. Alerting profiles are labeled **Dynatrace Classic** and Dynatrace recommends simple workflows for new setups — but no deprecation or end-of-life date has been published, so existing profiles keep working. The one real forcing function is Management Zones: the workflow Problem trigger has no management-zone filter, and Dynatrace's upgrade guidance states *"Management zones are not available in Latest Dynatrace. Replace management zone scoping with Grail record-based field filters."* A profile scoped by an MZ is rebuilt as a problem-triggered workflow scoped by affected-entity tags or a custom DQL matcher (WFLOW-02 §2; MZ2POL-01 §5).

3. Workflow Components

Every workflow consists of these components:

3.1 Trigger

What starts the workflow:

Trigger Type	Starts When	Use Case
Detected Problem	Dynatrace Intelligence detects a problem	Alert notifications
Davis Event	An anomaly detector raised a Davis event (per-alert, before grouping into a problem)	Per-alert automation
Schedule	A fixed time, a time interval, or a cron expression comes due	Daily reports
On-Demand	Manual execution or API call	Testing, ad-hoc runs
Event Trigger	Any event matching a DQL matcher (<code>events</code> , <code>bizevents</code> , <code>security.events</code> , <code>dt.system.events</code>)	Business process automation

3.2 Tasks

Actions the workflow performs:

Task Category	Examples
Notifications	Slack, Teams, Email, PagerDuty
Incident Management	ServiceNow, Jira
Data Queries	DQL queries, entity lookups
HTTP Requests	REST API calls to any endpoint
JavaScript	Custom code execution
Control Flow	Wait, conditions, loops

3.3 Conditions

Logic that controls task execution. Each task carries one `conditions` block: the states its predecessors must end in, an optional custom expression, and what to do when the condition is not met (`SKIP` or `STOP`):

```
conditions:
  states:
    previous_task: OK          # SUCCESS, ERROR, ANY, OK or NOK
    custom: '{{ event()["event.category"] == "AVAILABILITY" }}'
    else: SKIP
```

3.4 Expressions

Dynamic values using Jinja2 syntax:

```
{{ event()["event.name"] }}      # Access trigger data (problem title)
{{ problem_link() }}             # Link to the problem (Problem trigger only)
{{ result("task_name") }}        # Previous task result – shape depends on
task type; see WFLOW-08 §6
{{ input("environment") }}       # Workflow input
```

There is no `env` object in the expression language. Secrets belong in a **connection** or the **Credential Vault**, never in an expression (WFLOW-09 §2).

3.5 Workflow Types, Live vs. Draft, and the Actor


Three more concepts decide whether a workflow runs at all and what it costs.

Simple or standard. *"Simple workflows are limited to a single task, with reduced task options and a restricted set of available actions."* They cannot use Run JavaScript or Run workflow, and *"They carry no workflow hours cost."* Standard workflows support multiple tasks, branches, loops and JavaScript, and *"draw from your workflow hours quota."* Billing counts how long a standard workflow exists, not how often it runs: *"Workflow hours are the number of hours that a workflow has existed in your environment, measured since the point of its creation."* A one-destination notification is a good fit for a simple workflow; ALERT-03 covers the cost trade-off.

Live or draft. *"Triggers only start on live workflows."* Select **Deploy** to make the draft live. A draft can still be run by hand from the **Run** button, and *"Draft-only workflows do not directly consume workflow hours"*.

Actor. "Every task executes under the actor's permissions, regardless of how the workflow was started." By default the actor is the user who created or last updated the workflow, and event triggers only see events that user can read. For production workflows Dynatrace recommends a service user: "We highly recommend using service users as actors for all workflows that are worked on collaboratively and serve a production grade use case."

Visual: Workflow Execution Flow



4. Accessing Workflows

Navigation

1. Open Dynatrace
2. Go to **Apps** in the left navigation
3. Search for and open **Workflows**

Or use the direct URL: `https://<your-environment>/ui/apps/dynatrace.automations/workflows`

Required Permissions

Permission	Allows
<code>app-engine:apps:run</code>	Open the Workflows app (list apps, read app bundles)
<code>automation:workflows:read</code>	View workflows
<code>storage:system:read</code> + <code>storage:buckets:read</code> (scoped to <code>dt.system.events</code>)	View execution history, including the §7 queries
<code>automation:workflows:write</code>	Create, edit, delete workflows (can be limited to <code>automation:workflow-type = "SIMPLE"</code>)
<code>automation:workflows:run</code>	Run workflows manually or via the API
<code>app-engine:functions:run</code>	Use the function executor (needed to write and execute workflows)
<code>automation:workflows:admin</code>	Workflow admin mode: all workflows and executions

These grant access to workflows only: *"To successfully run workflow tasks, the actor might need additional permissions."* Each connector's setup page lists what its actions need (for example `email:emails:send` for Send email).

Workflow Listing

The Workflows app has an **All workflows** tab and an **Executions** tab. A new workflow is private: *"By default, a workflow is only visible to the creator (workflow owner)."* The owner can make it public, or transfer ownership to another user or a group.

5. Execution Model

Understanding how workflows execute is important for designing reliable automation.

Execution Flow

```
Trigger fires
↓
Workflow instance created
↓
Tasks execute (sequential or parallel)
↓
Conditions evaluated at each step
↓
Results captured
↓
Execution complete (success/failure)
```

Platform Limits

Two distinct timeouts apply to every workflow task — confusing them is one of the most common causes of "why did my task fail at exactly 2 minutes?":

Limit	Default	Maximum	Scope
Task timeout	60 minutes	7 days	Wall-clock budget for the whole task, including retries and loops. Configured per task via the <code>timeout</code> field (in seconds).
Dynatrace runtime timeout	120 seconds	—	Per-action execution budget inside the AutomationEngine runtime (applies to DQL queries and individual JavaScript/HTTP calls). Limits a single action and cannot be raised. A timed-out action fails; the task ends in <code>ERROR</code> unless Retry on error or a Loop runs further actions, and the task timeout bounds their total runtime.

Limit	Default	Maximum	Scope
DQL <code>requestTimeoutMilliseconds</code>	No default documented	—	The JavaScript SDK's <code>queryExecute()</code> parameter — how long the call waits for the result before returning a <code>requestToken</code> instead, in milliseconds. It does not stop the query; poll with <code>queryPoll()</code> to collect the result (WFLOW-08 §2).

Why this matters. Raising the task `timeout` does not help a DQL query that's hitting the 120-second runtime budget — you need to narrow the query window, pre-aggregate, or split the work. Conversely, an approval task waiting for a human pager-out doesn't need a runtime budget — it needs a long task timeout (`timeout: 1800` for 30 minutes, `timeout: 86400` for 24 hours).

Concrete handling: see **WFLOW-08 § Configuring Task Timeouts** for the YAML/JSON shape on `execute-dql-query`, `run-javascript`, and approval tasks.

Other documented limits.

- **Event-triggered executions:** *"Each workflow is limited to 1,000 event-triggered executions per hour."* Beyond that, executions are throttled (HTTP 429) for up to one hour. The two Dynatrace pages disagree about repeated breaches: the event-trigger reference says *"If it's exceeded three times within seven days, the trigger is automatically deactivated."*, while the upgrade guide says *"The trigger is not automatically deactivated when the limit is reached repeatedly."* Check the Workflows overview for throttled or deactivated triggers rather than relying on either.
- **Trigger filter size:** the matching expression compiled from all trigger fields is limited to 1,000 characters (WFLOW-02 §2).
- **Workflows per environment:** *"Environment limits are 10,000 workflows per customer environment and 100 per trial environment."*
- **Concurrent executions:** *"There can be as many executions of a workflow running at any given time as requested (within system capacity)."*
- **Input and result size:** *"The workflow default input size is limited to 10 MB."* The execution result has the same 10 MB limit.

Per-workflow task counts are not published. In community practice, keep workflows small (under ~20 tasks) and focused (one trigger → one outcome), and verify behavior under load in your own environment before designing around a specific number.

Execution States

State	Meaning
RUNNING	Currently executing
SUCCESS	Completed successfully
ERROR	A task failed or timed out — there is no separate TIMED_OUT state
CANCELLED	Manually cancelled
SKIPPED / DISCARDED	Task level only — the task's conditions were not met

These are the values `dt.automation_engine.state` carries in `dt.system.events` (the queries in §7 read them); `FAILED`, `SUCCEEDED` and `TIMED_OUT` are not values.

6. Your First Workflow

Let's create a simple workflow that logs a message when manually triggered.

Step 1: Create New Workflow

1. Open **Workflows** app
2. Click **+ Workflow**
3. Name it: `Hello World workflow`

A new workflow starts as a **simple** workflow (shown in the upper-left corner of the editor), and simple workflows exclude Run JavaScript. Dynatrace's own pages differ here: the simple-workflow page excludes the action, while the Get started walkthrough adds it directly. If Run JavaScript is not offered in Step 3, unlock the full functionality, which makes this a standard workflow (§3.5). Leaving it as an undeployed draft keeps it out of workflow-hour billing; delete it when you are done.

Step 2: Configure Trigger

1. Click the trigger box (top of canvas)
2. Select **On demand** trigger
3. This allows manual execution for testing

Step 3: Add a Task

1. Click **+** below the trigger
2. Search for **Run JavaScript**
3. Add the following code:

```
export default async function() {
  const now = new Date().toISOString();
  console.log(`Hello from Dynatrace Workflows at ${now}`);

  return {
    message: "Workflow executed successfully!",
    timestamp: now
  };
}
```

Step 4: Save and Run

1. Click **Save**
2. Click **Run** (play button). *"The first time you run a workflow, you are prompted to authorize the automation service to run a workflow as your user."* Select **Allow and run**.
3. View execution results

Running from the editor works on a draft. Automatic triggers (Schedule, Problem, Event) start only once you select **Deploy** (\$3.5).

Expected Output

The task output should show:

```
{  
  "message": "Workflow executed successfully!",  
  "timestamp": "2026-01-27T12:00:00.000Z"  
}
```

7. Viewing Execution History

Monitor your workflow executions using DQL queries.

```

// Recent workflow executions
// Data object corrected 08/12/2026. Workflow executions are NOT in `events`,
and there is no
// `automation.task.execution` / `automation.workflow.execution` event type in
any spelling – those
// filters matched nothing, silently, in 25 cells. AutomationEngine writes to
`dt.system.events` with
// `event.kind == "WORKFLOW_EVENT"` (5.7M records / 30d here), split by
`event.type`:
//   WORKFLOW_EXECUTION · TASK_EXECUTION · ACTION_EXECUTION ·
WORKFLOW_CREATED/UPDATED/DELETED
// Fields are the `dt.automation_engine.*` family:
//   .workflow.title / .workflow.id      (was workflow.name)
//   .task.name                          (was task.name)
//   .action.function / .action.app      (was task.type – e.g. run-javascript,
send-email,
//                                       snow-search-incidents)
//   .state                             (was task.status / execution.status)
//                                       values RUNNING · SUCCESS · ERROR ·
DISCARDED · SKIPPED
//                                       – note "FAILED" is NOT a value; it
is ERROR
//   .state.is_final                    true only on terminal records –
filter on it, otherwise a
//                                       single execution is counted once as
RUNNING and again as
//                                       SUCCESS/ERROR
//   .state_info                        (was task.error) · duration (was
task.duration)
// Enumerate with:
//   fetch dt.system.events, from:-24h | filter event.kind == "WORKFLOW_EVENT"
| limit 1
fetch dt.system.events, from:-24h
| filter event.kind == "WORKFLOW_EVENT"
| filter event.type == "WORKFLOW_EXECUTION"
| filter dt.automation_engine.state.is_final == true
| fields timestamp, dt.automation_engine.workflow.title,
dt.automation_engine.state, duration
| sort timestamp desc
| limit 25

```

```

// Workflow execution summary
// Data object corrected 08/12/2026. Workflow executions are NOT in `events`,
and there is no
// `automation.task.execution` / `automation.workflow.execution` event type in
any spelling – those
// filters matched nothing, silently, in 25 cells. AutomationEngine writes to
`dt.system.events` with
// `event.kind == "WORKFLOW_EVENT"` (5.7M records / 30d here), split by
`event.type`:
//   WORKFLOW_EXECUTION · TASK_EXECUTION · ACTION_EXECUTION ·
WORKFLOW_CREATED/UPDATED/DELETED
// Fields are the `dt.automation_engine.*` family:
//   .workflow.title / .workflow.id      (was workflow.name)
//   .task.name                          (was task.name)
//   .action.function / .action.app      (was task.type – e.g. run-javascript,
send-email,
//                                       snow-search-incidents)
//   .state                             (was task.status / execution.status)
//                                       values RUNNING · SUCCESS · ERROR ·
DISCARDED · SKIPPED
//                                       – note "FAILED" is NOT a value; it
is ERROR
//   .state.is_final                    true only on terminal records –
filter on it, otherwise a
//                                       single execution is counted once as
RUNNING and again as
//                                       SUCCESS/ERROR
//   .state_info                        (was task.error) · duration (was
task.duration)
// Enumerate with:
//   fetch dt.system.events, from:-24h | filter event.kind == "WORKFLOW_EVENT"
| limit 1
fetch dt.system.events, from:-7d
| filter event.kind == "WORKFLOW_EVENT"
| filter event.type == "WORKFLOW_EXECUTION"
| filter dt.automation_engine.state.is_final == true
| summarize executions = count(), by:{dt.automation_engine.workflow.title,
dt.automation_engine.state}
| sort executions desc
| limit 25

```

```
// Failed workflow executions
// Data object corrected 08/12/2026. Workflow executions are NOT in `events`,
and there is no
// `automation.task.execution` / `automation.workflow.execution` event type in
any spelling – those
// filters matched nothing, silently, in 25 cells. AutomationEngine writes to
`dt.system.events` with
// `event.kind == "WORKFLOW_EVENT"` (5.7M records / 30d here), split by
`event.type`:
//   WORKFLOW_EXECUTION · TASK_EXECUTION · ACTION_EXECUTION ·
WORKFLOW_CREATED/UPDATED/DELETED
// Fields are the `dt.automation_engine.*` family:
//   .workflow.title / .workflow.id      (was workflow.name)
//   .task.name                          (was task.name)
//   .action.function / .action.app      (was task.type – e.g. run-javascript,
send-email,
//                                       snow-search-incidents)
//   .state                              (was task.status / execution.status)
//                                       values RUNNING · SUCCESS · ERROR ·
DISCARDED · SKIPPED
//                                       – note "FAILED" is NOT a value; it
is ERROR
//   .state.is_final                     true only on terminal records –
filter on it, otherwise a
//                                       single execution is counted once as
RUNNING and again as
//                                       SUCCESS/ERROR
//   .state_info                         (was task.error) · duration (was
task.duration)
// Enumerate with:
//   fetch dt.system.events, from:-24h | filter event.kind == "WORKFLOW_EVENT"
| limit 1
fetch dt.system.events, from:-24h
| filter event.kind == "WORKFLOW_EVENT"
| filter event.type == "WORKFLOW_EXECUTION"
| filter dt.automation_engine.state == "ERROR"
| fields timestamp, dt.automation_engine.workflow.title,
dt.automation_engine.state_info
| sort timestamp desc
| limit 25
```

Next Steps

Now that you understand workflow fundamentals, continue with:

Recommended Path

1. **WFLOW-02: Triggers & Event Types** - Configure event-driven triggers
2. **WFLOW-03: Alert Notification Basics** - Send Slack, Teams, email alerts
3. **WFLOW-04: Advanced Notification Routing** - Conditional routing and escalation

Key Takeaways

- Workflows automate responses to events, schedules, and manual triggers
 - Components: Triggers → Tasks → Conditions → Results
 - Use Workflows for notifications, integrations, and auto-remediation
 - Monitor execution history via DQL queries
-

Summary

In this notebook, you learned:

- What Dynatrace Workflows are and when to use them
 - How workflows compare to other automation options
 - Core workflow components (triggers, tasks, conditions)
 - How to access and navigate the Workflows app
 - The execution model and limits
 - How to create your first workflow
 - How to query execution history
-

References

- [Workflows umbrella \(DT docs\)](#)
- [Workflow triggers \(DT docs\)](#)
- [Workflow actions \(DT docs\)](#)
- [Workflow reference / Jinja expressions \(DT docs\)](#)
- [Alerting and notifications umbrella \(DT docs\)](#)

- [Davis Problems app \(DT docs\)](#)
 - [Workflows concepts \(DT docs\)](#)
 - [Build workflows \(DT docs\)](#)
 - [Create a simple workflow \(DT docs\)](#)
 - [Get started with Workflows \(DT docs\)](#)
 - [Manage workflow permissions \(DT docs\)](#)
 - [Monitor workflow executions \(DT docs\)](#)
 - [Automation Workflow consumption \(DT docs\)](#)
 - [Event triggers for workflows \(DT docs\)](#)
 - [Upgrade guide — alerting and notifications \(DT docs\)](#)
 - [Upgrade security notifications — management zones \(DT docs\)](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.